

# SECTION 1

Let's Get Started!

# Demo: Let's See AWS in Action!



# SECTION 2

Create a Free Tier Account on AWS

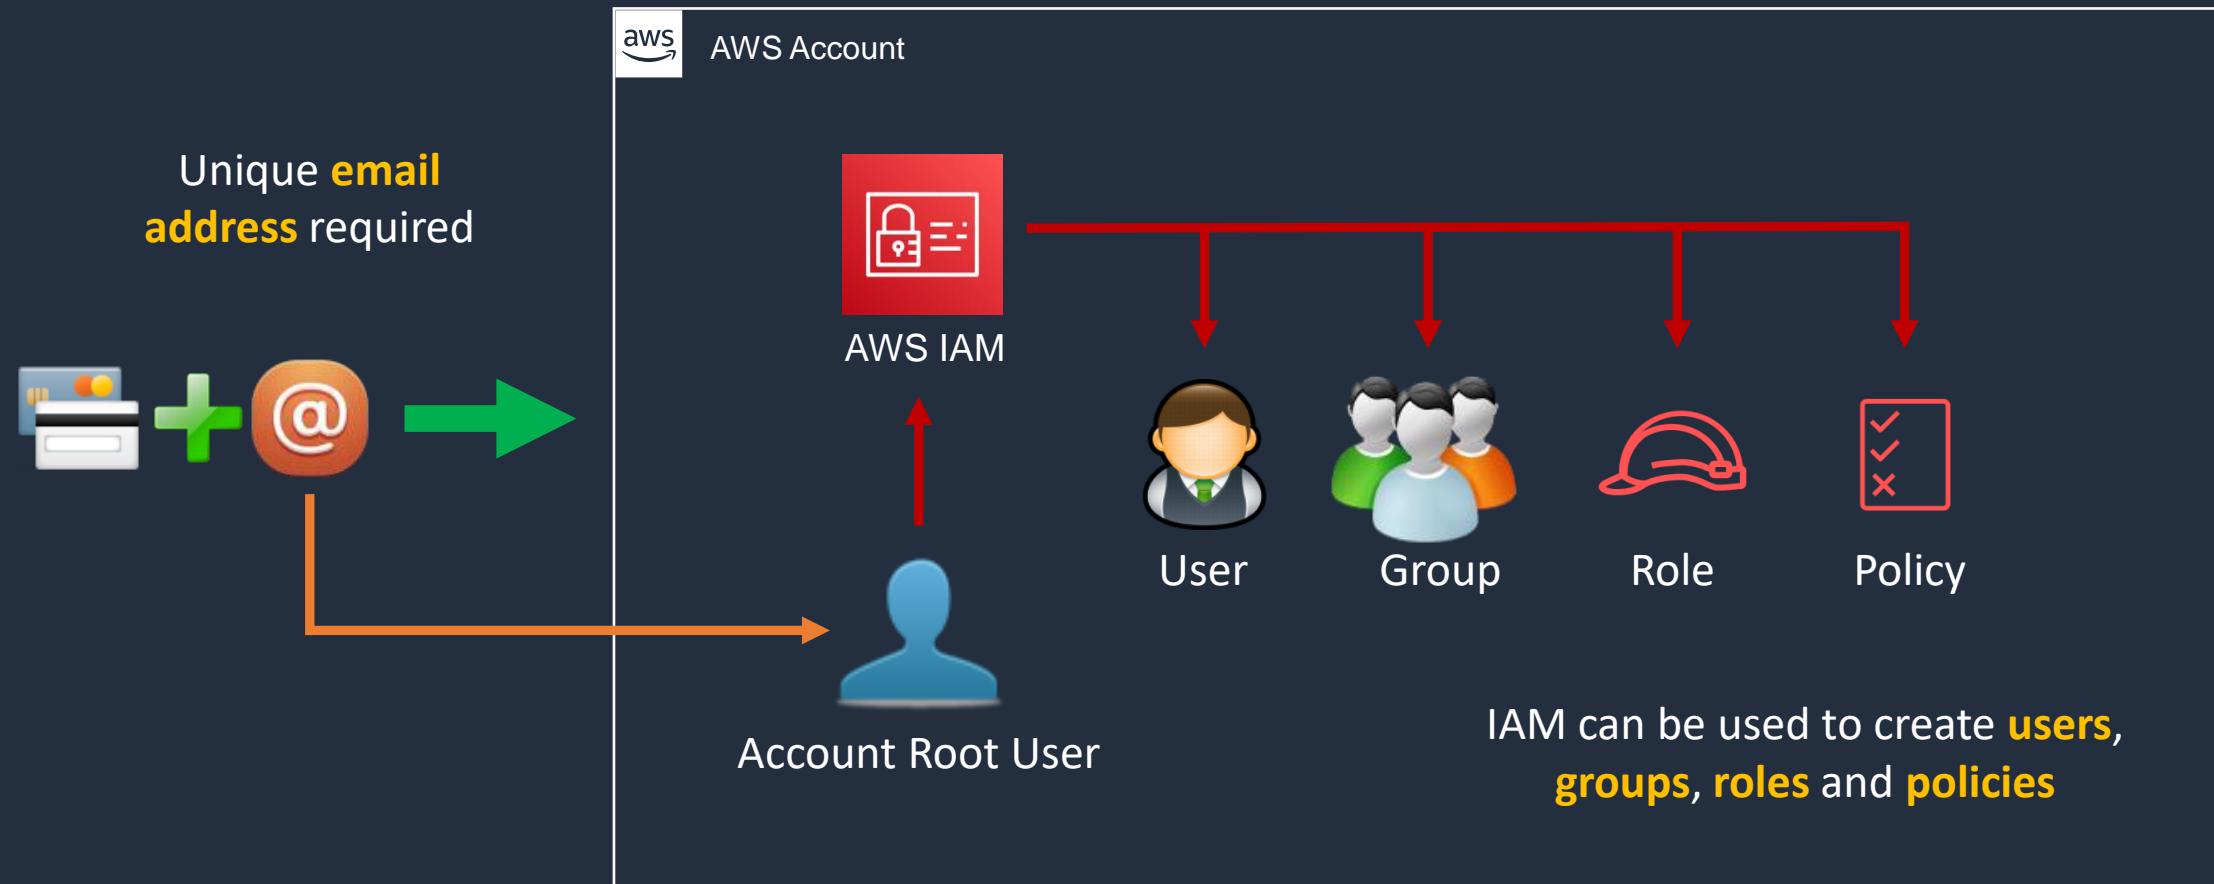
# AWS Account Overview



# AWS Account Overview



It's an IAM best practice to create **individual users** and to avoid using the **Root** account

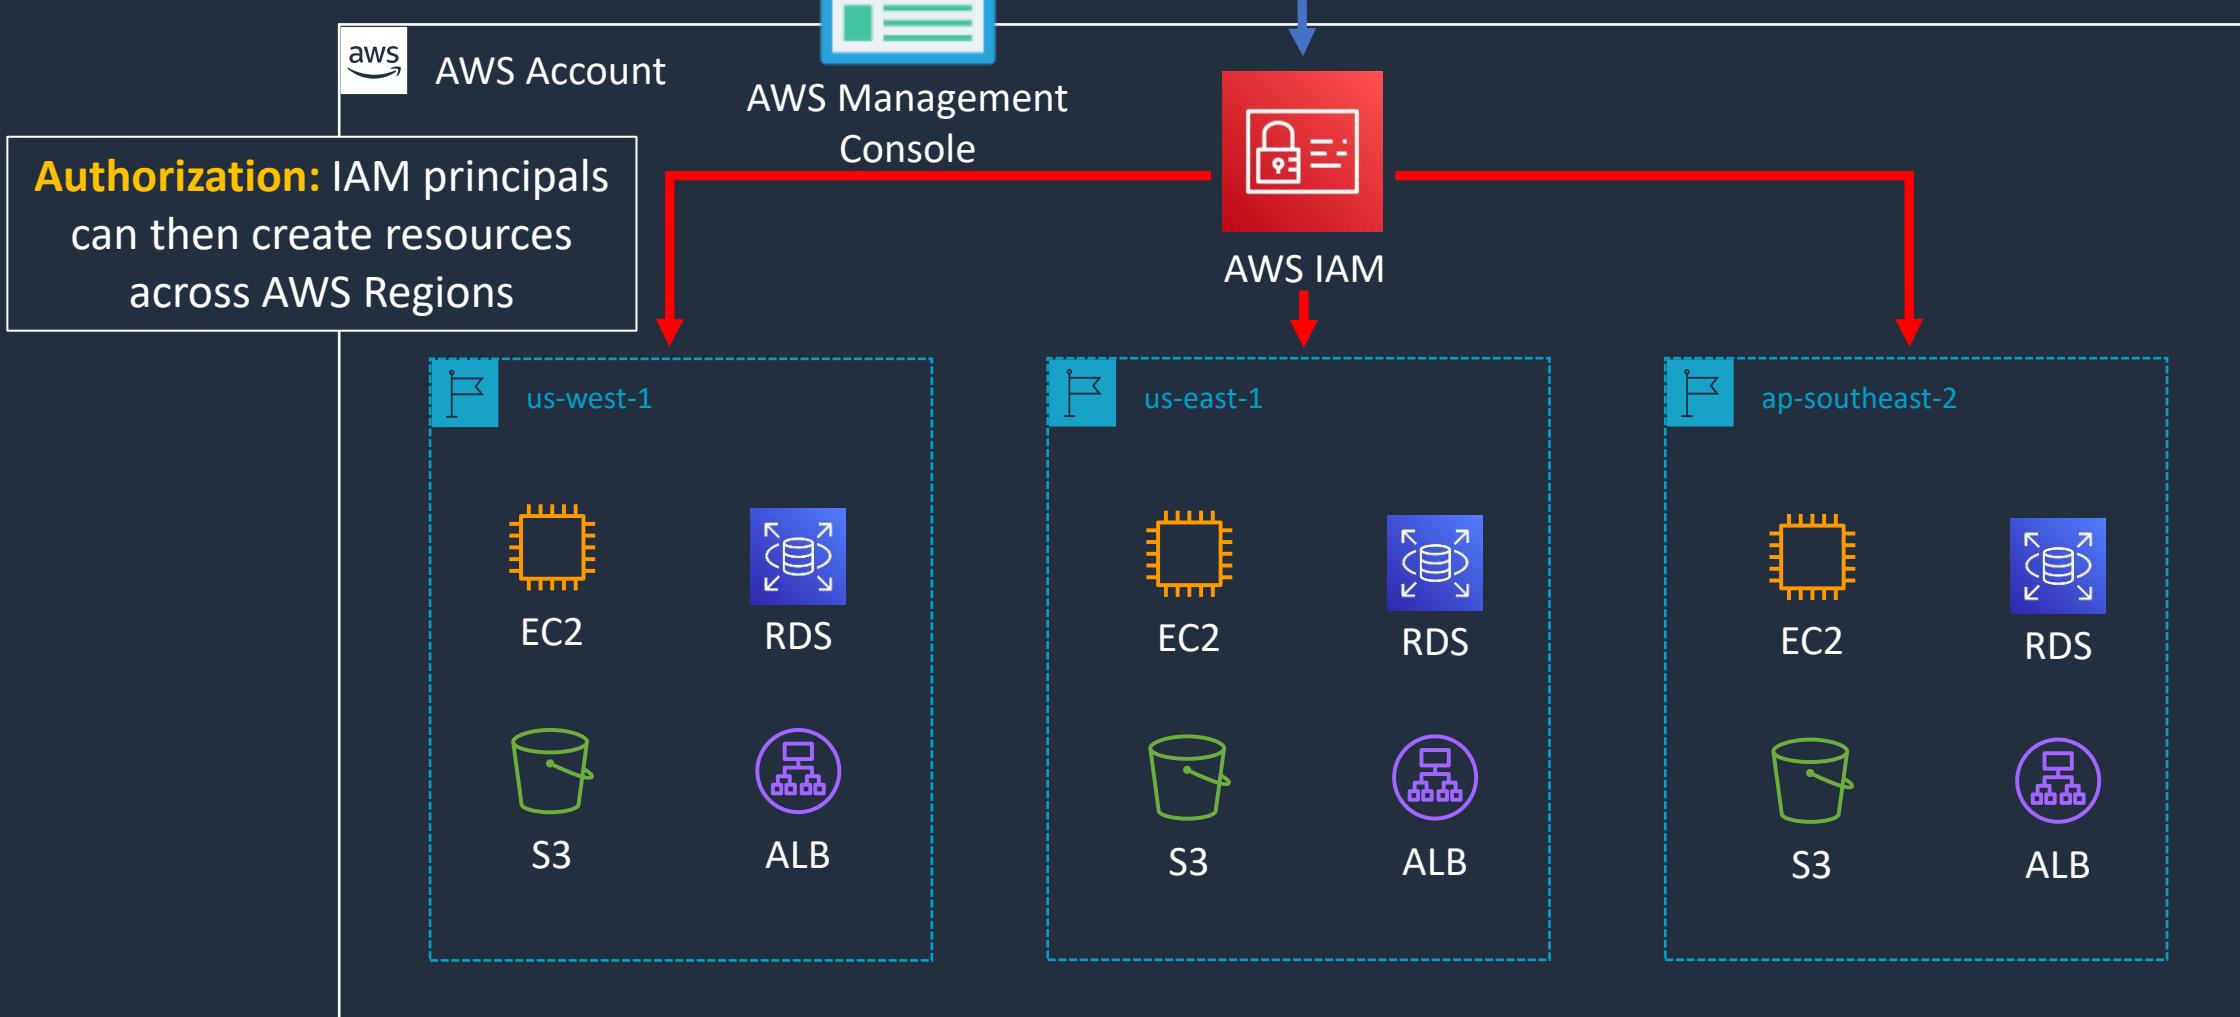


The Root user has **full control** over the account



# AWS Account Overview

**Authentication:** IAM principals authenticate to IAM using the console, API, or CLI



All AWS **identities** and **resources** are created  
within the AWS account

# Create your AWS Free Tier Account



# What you need...



Credit card for setting up the account and paying any bills



Unique email address for this account

john@gmail.com



Check if you can use a **dynamic alias** with an existing email address



john+ACCOUNT-ALIAS-1@gmail.com

john+ACCOUNT-ALIAS-2@gmail.com



AWS account name / alias



Phone to receive an **SMS** verification code

# Configure Account and Create a Budget and Alarm



# Account Configuration

---

- Configure **Account Alias**
- Enable access to billing for **IAM users**
- Update **billing preferences**
- Create a **budget and alarm**

# Install Tools



# SECTION 3

## IT Fundamentals

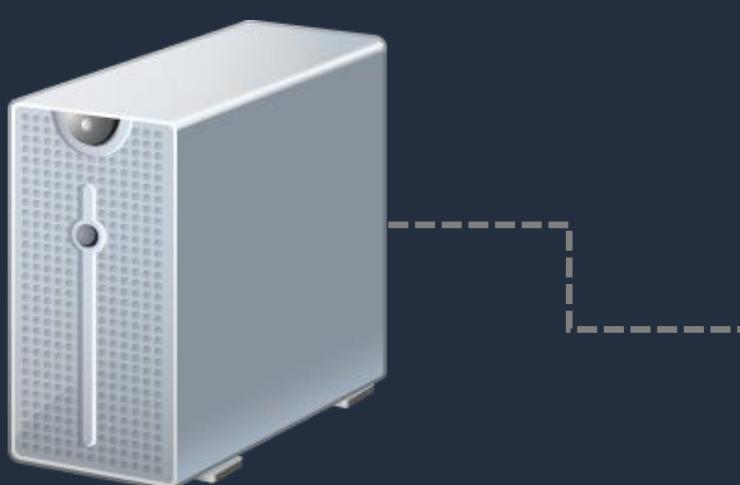
# Client - Server Computing



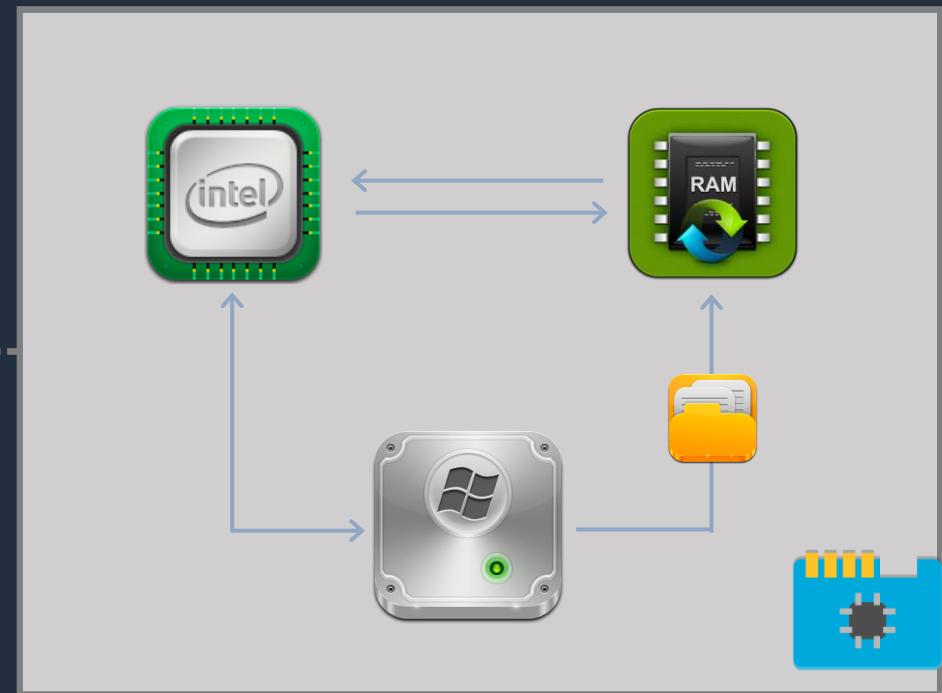


# Basic Architecture of a Computer

---



Central Processing  
Unit (CPU)



RAM is **non-persistent**  
storage or “memory”

**Network Interface**  
**Cards** (NICs) connect  
computers to networks

**Hard Disk Drives**  
(HDDs) offer  
**persistent** storage



# Measurements for Components

---

Component	Unit of Measurement
Central Processing Unit (CPU)	Gigahertz (GHz)
Random Access Memory (RAM)	Gigabytes (GB)
Hard Disk Drive (HDD) / Solid State Drive (SSD)	Gigabytes (GB)
Network Interface Card (NIC)	Megabits per second (Mbps) Gigabits per second (Gbps)

# Servers vs Desktops/Laptops



Server



Servers can be used by many users over a network



Laptop



Desktop

## Server Hardware Build:

- Hardware is more specialized
- Much higher prices compared to desktops / laptops
- Includes redundancy



# Clients and Servers

## Cloud Computing



Servers

Servers running in the cloud offer **services** which include the **application, processing** and **data storage**

**Client devices** are connected via the **Internet**



Cloud Networking



Client Devices

**Client devices** require connectivity via **wired, wireless** or **cellular** networks

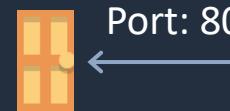


# Connecting to Services

The client application finds the server by **IP address**



Web Server



Port: 80

Protocol: HTTP



File Server



Port: 445

Protocol: SMB



Email Server



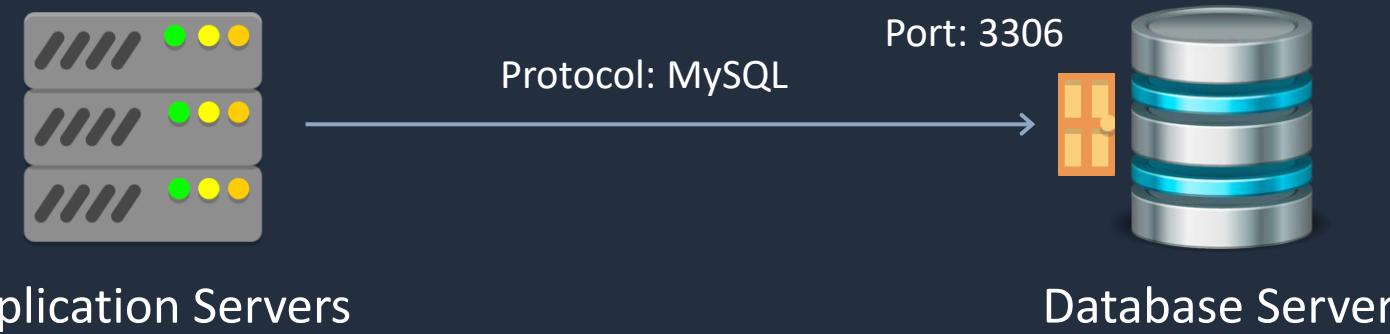
Port: 25

Protocol: SMTP





# Server to Server Connectivity



# Storage - Block vs File vs Object





# Hard Drives

---



Hard Disk Drive (HDD)

- Also known as magnetic drives
- Older technology
- Much slower than SSD
- Much cheaper than SSD



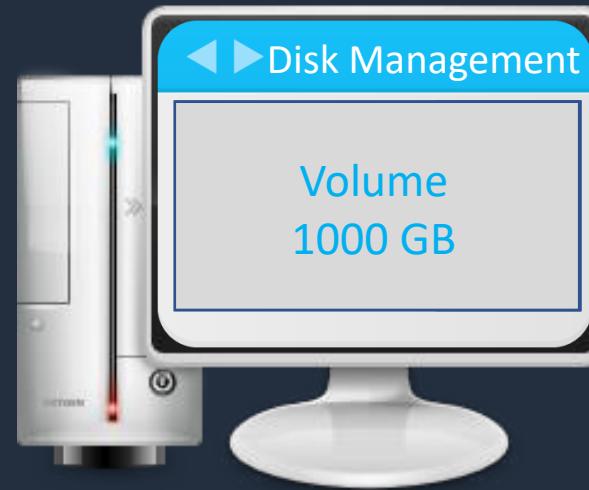
Solid State Drive (SSD)

- Uses flash memory
- Newer technology
- MUCH faster than HDD
- More expensive than HDD



# Hard Drives

Hard drives are  
block-based  
storage systems

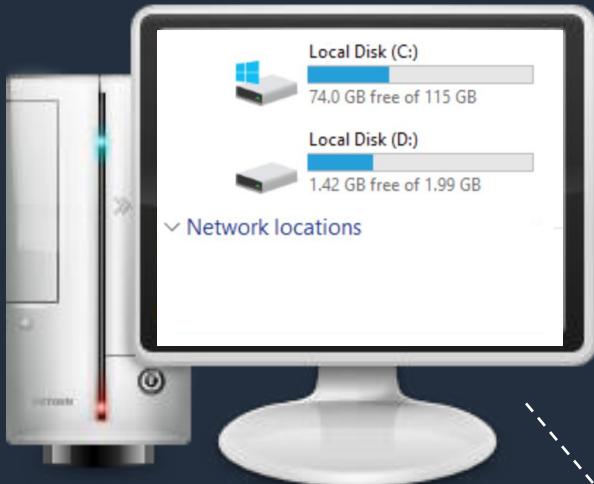


Hard drives are block-based storage systems

The Operating System  
(OS) sees a volume. A  
volume can be  
partitioned and  
formatted

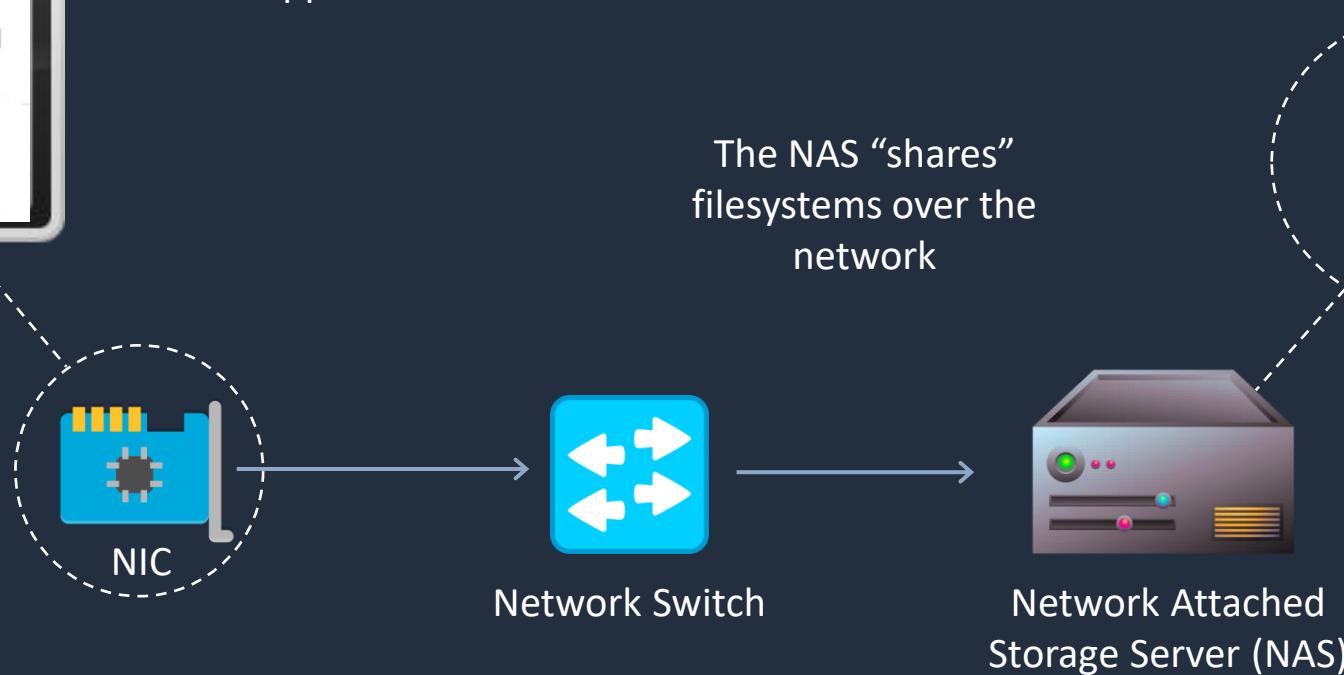


# Network Attached Storage



The Operating System (OS)  
sees a filesystem that is  
mapped to a local drive letter

The NAS “shares”  
filesystems over the  
network



NAS devices are file-based storage systems

# Object Storage Systems

User uploads **objects** using a web browser



The **HTTP protocol** is used with a **REST API** (e.g. GET, PUT, POST, SELECT, DELETE)

There is no hierarchy of objects in the container



Object Storage Container

Objects can be files, videos, images etc.

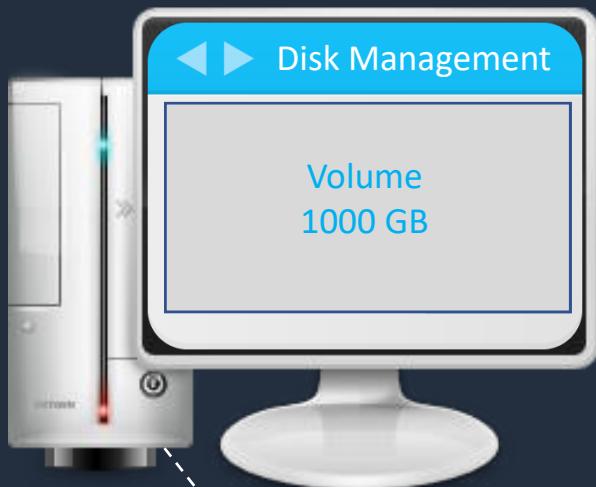
# Block vs File vs Object Storage

The OS sees **volumes** that can be partitioned and formatted

A filesystem can be shared by many users

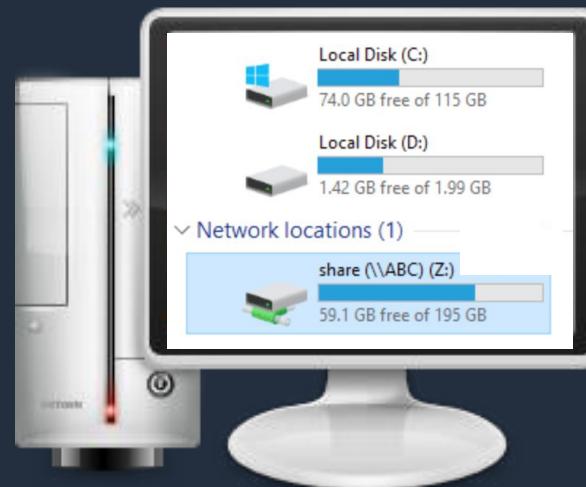
Massively scalable and low cost

## Block Storage



The OS reads/writes at the **block level**. Disks can be internal, or network attached

## File Storage



A filesystem is “**mounted**” to the OS using a network share

## Object Storage



Object Storage Container

There is **no hierarchy** of objects in the container

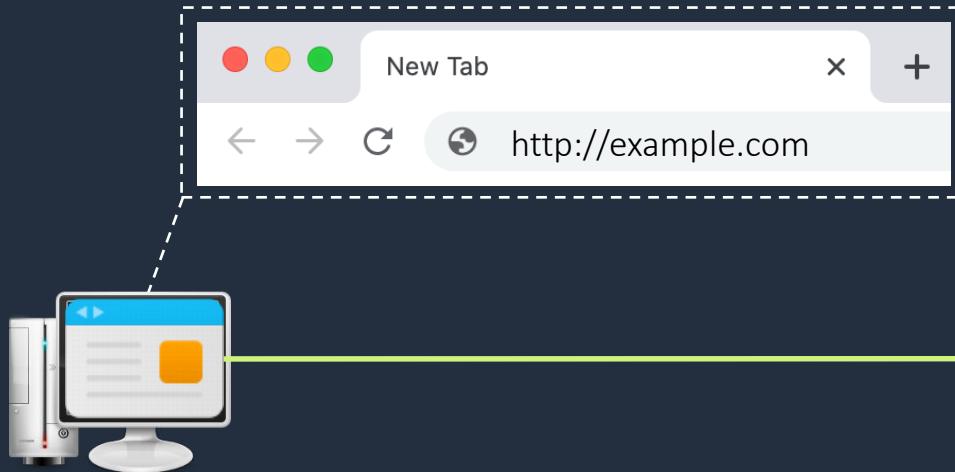
Uses a **REST API**

# IP Addresses and DNS





# What is an IPv4 address?



What's the **IP address** for example.com?



QUESTION: What's the IP address for example.com?

**IP addresses** are the addresses computers use to communicate



Connection is made using **IP address**

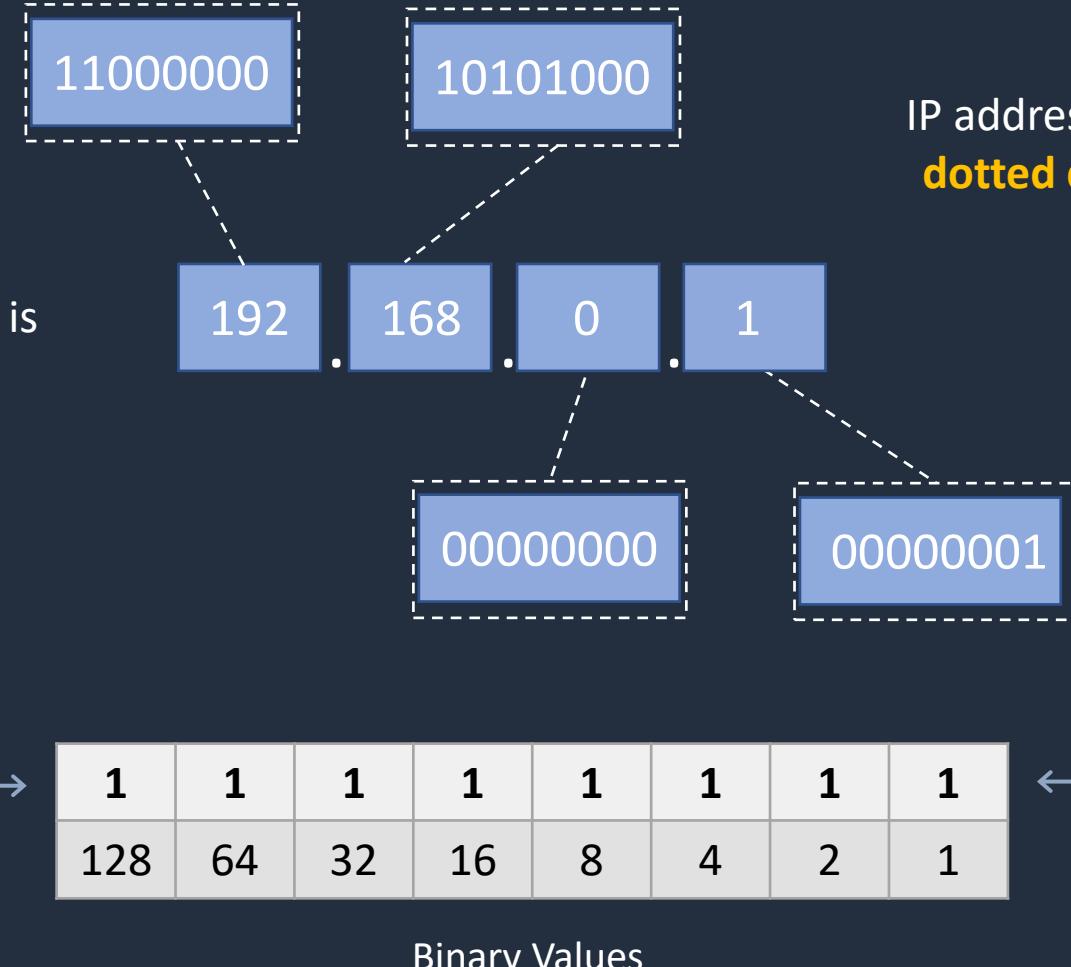
Name	Type	Value
example.com	A	52.134.26.12
test.com	A	137.10.47.51

DNS Zone File



# Structure of an IPv4 Address

Each part of the address is a **binary octet**



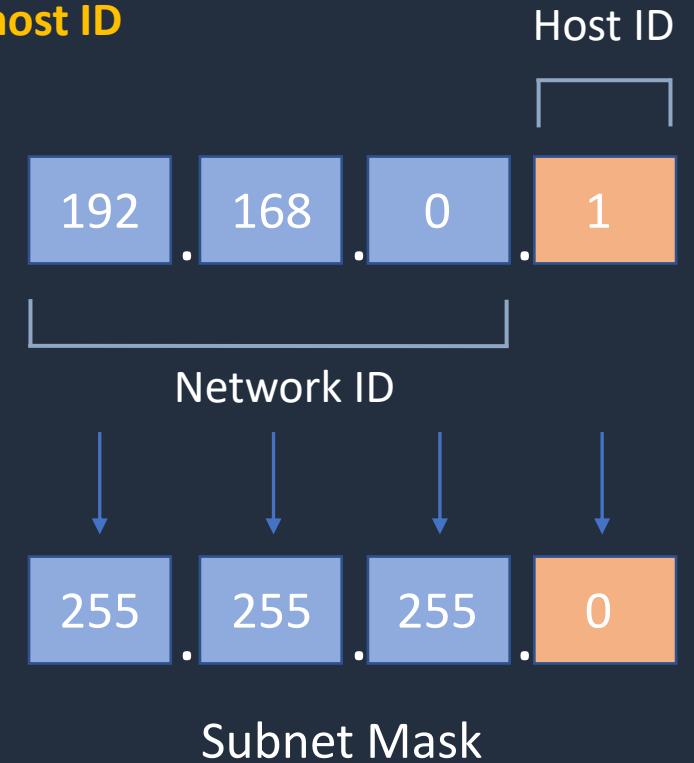
IP addresses are written in  
**dotted decimal notation**

Most significant bit →      ← Least significant bit



# Networks and Hosts

An IPv4 address has a **network** and **host ID**



The **subnet mask** is used to define the **network** and **host ID**



# Networks and Hosts

Network    

192	168	0	0
-----	-----	---	---

Subnet Mask    

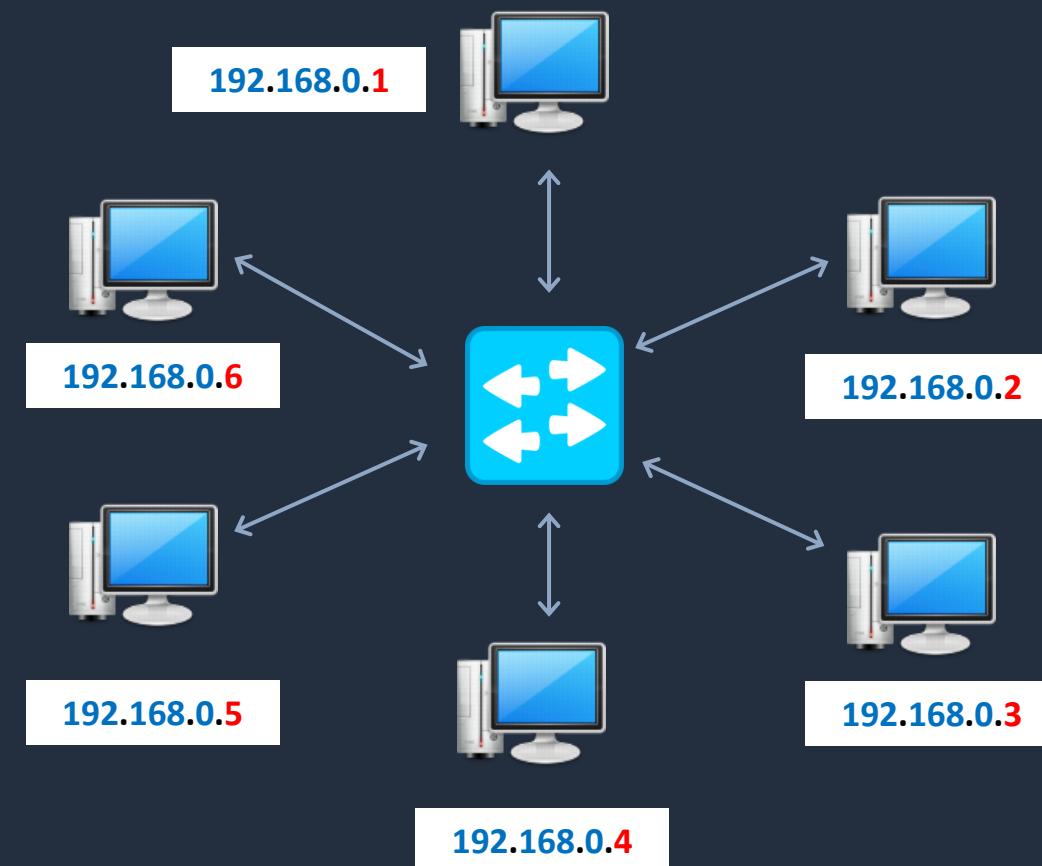
255	255	255	0
-----	-----	-----	---

24 bits

192.168.0.0/24

A **network** and **subnet mask** can also be written in this format

All computers share the same **network ID** and have a unique **host ID**



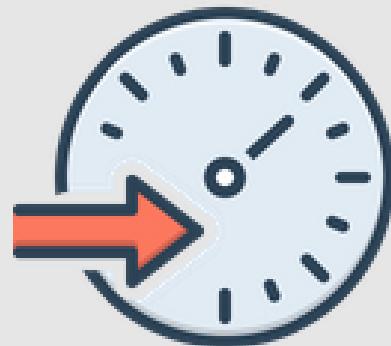


# Private IP Address Ranges



These addresses are reserved  
for **private use** according to  
IETF RFC-1918

# Bandwidth and Latency





# Bandwidth and Latency

---

Bandwidth is the **rate** of data transfer for a **fixed period of time** measured in **Gbps**



Bandwidth can be considered the **width** of the communication band

Latency is the amount of time it takes to send data from one point to another measured in **microseconds** or **milliseconds**



# Bandwidth and Latency

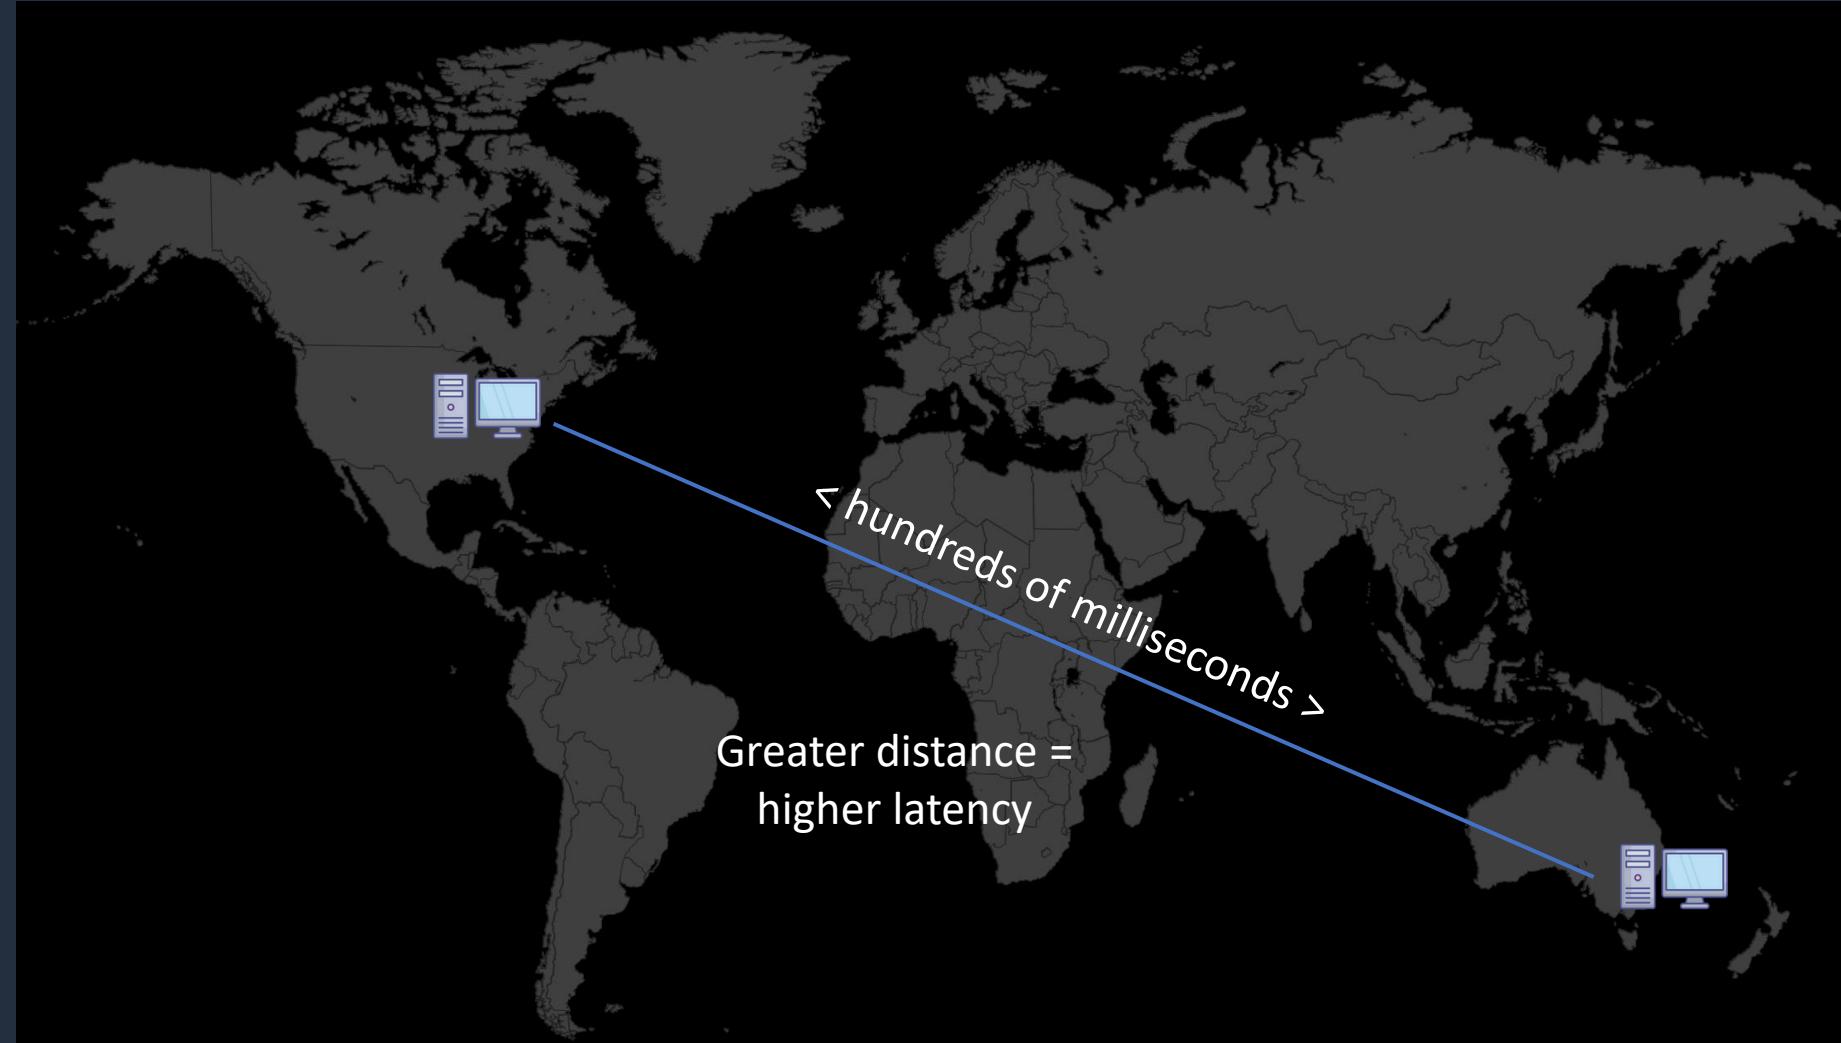
---



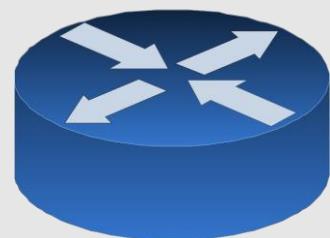


# Bandwidth and Latency

---



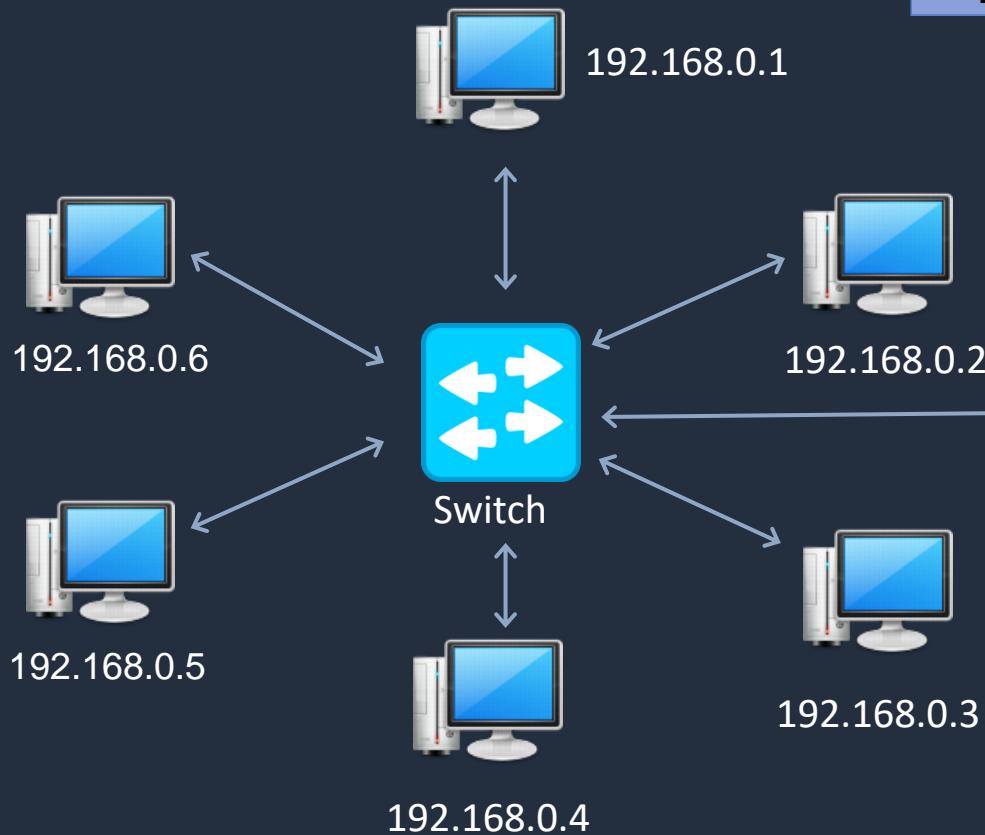
# Networking – Routers, Switches and Firewalls





# Routers and Switches

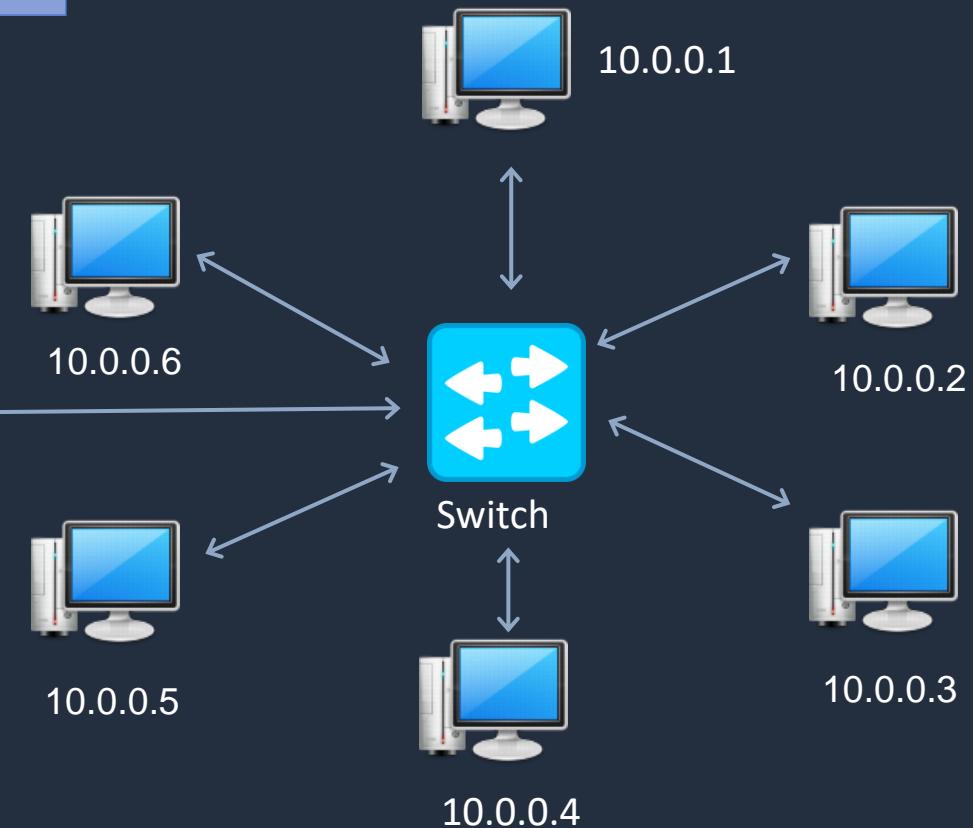
IP Subnet A: 192.168.0.0/24



Destination	Interface
192.168.0.0/24	eth0
10.0.0.0/24	eth1

Route Table

IP Subnet B: 10.0.0.0/24





# Firewalls

POLICY	PROTOCOL	PORT	DESTINATION	SOURCE
ALLOW	HTTP	80	INTERNAL	ANY
ALLOW	HTTPS	443	INTERNAL	ANY
DENY	ANY	ANY	INTERNAL	ANY

Firewall Rules

IP Subnet A



Database Server



Application Server

IP Subnet B



Web Server



Firewall



Firewall



Firewall



The Internet



Database Server



Application Server



DigitalCloud  
TRAINING

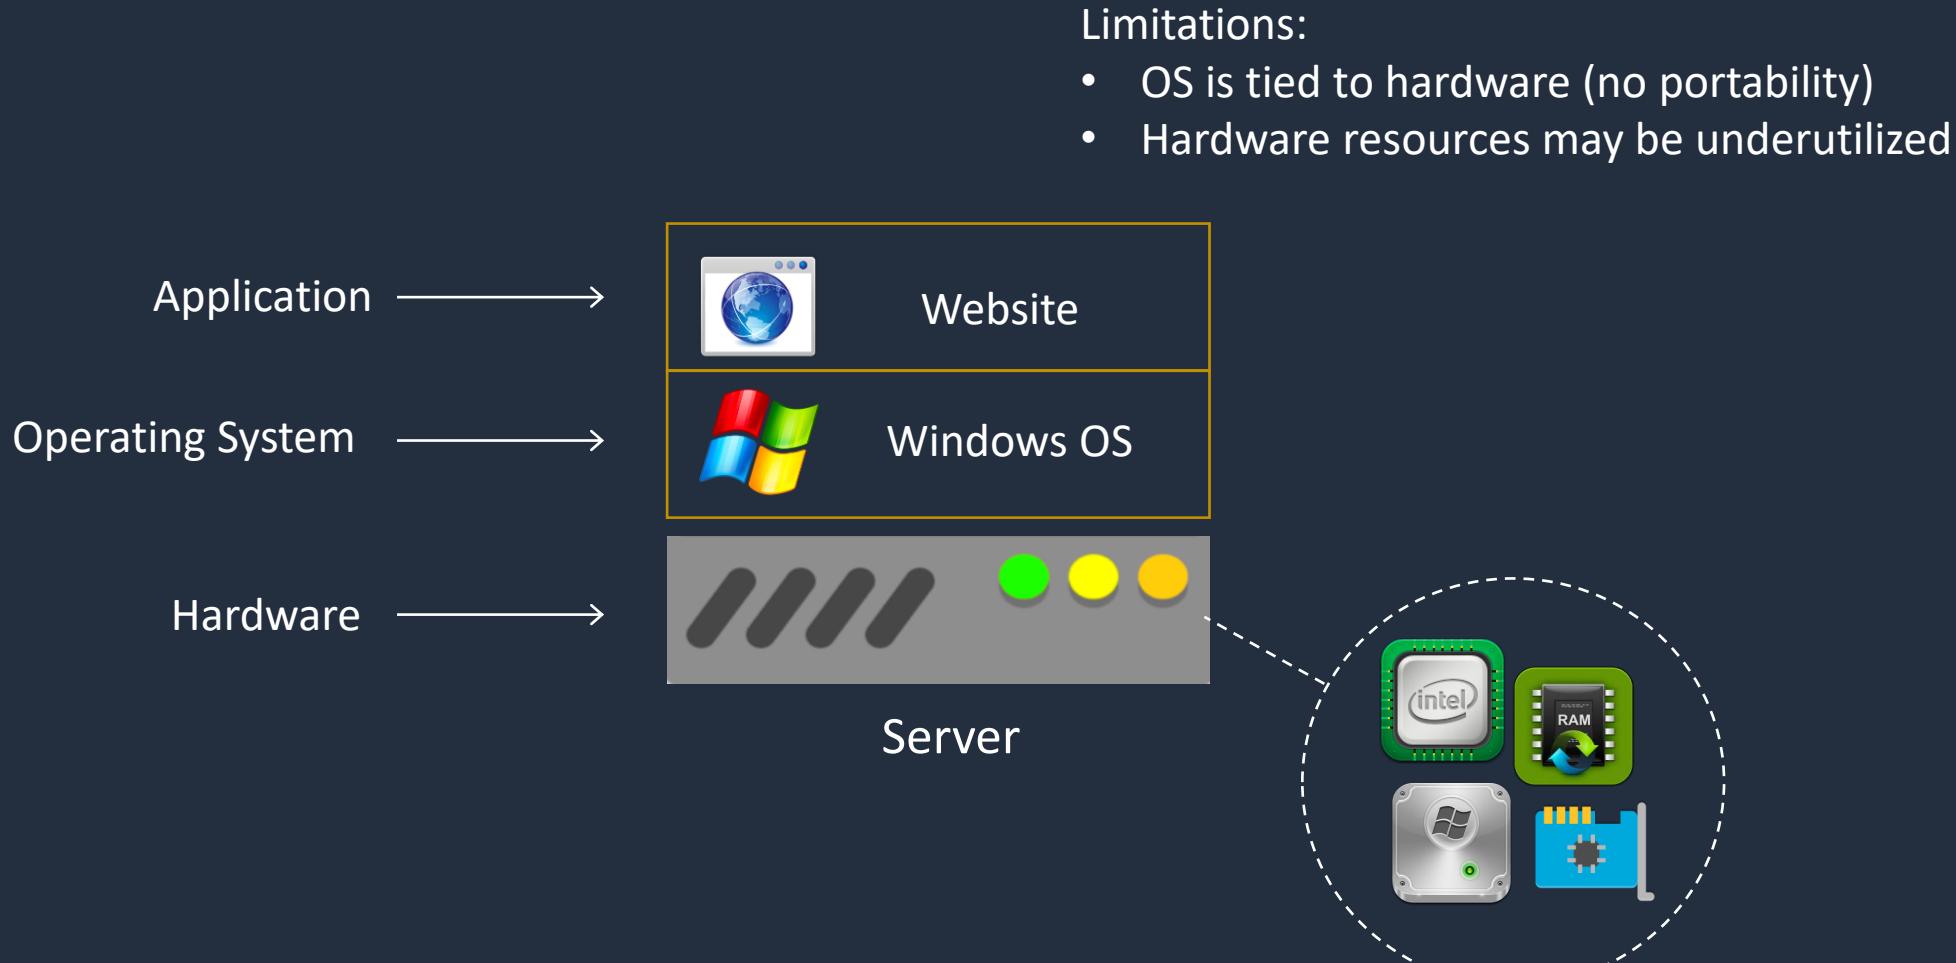
# Server Virtualization and Containers





# Server without Virtualization

---

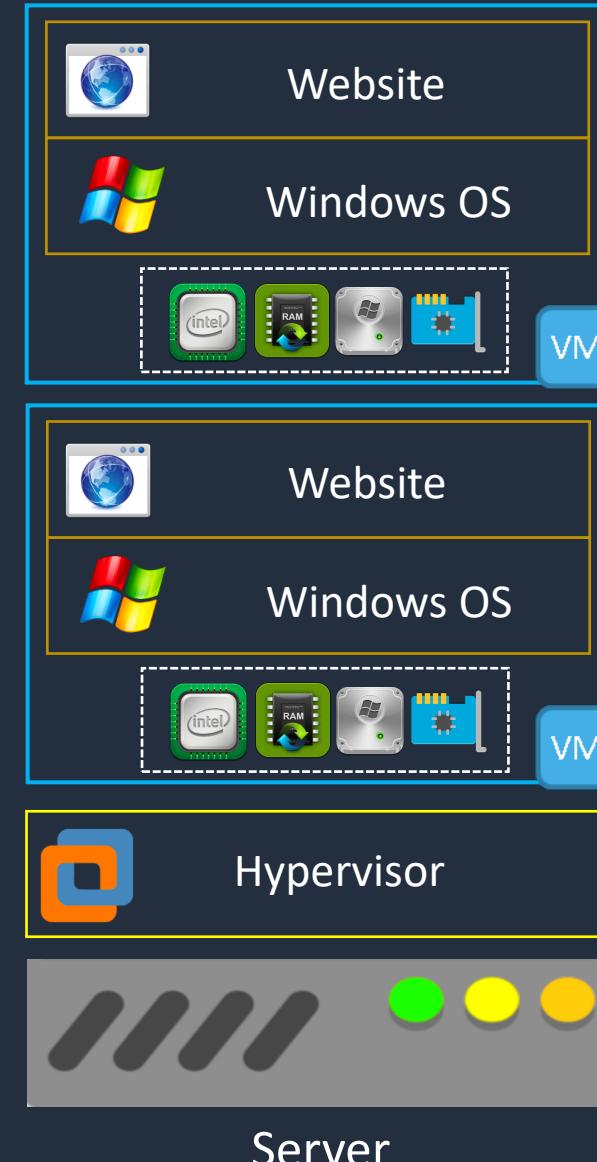




# Server with Virtualization

This is known as a  
**virtual server** or  
"virtual machine"

The **hypervisor** creates a  
layer of abstraction

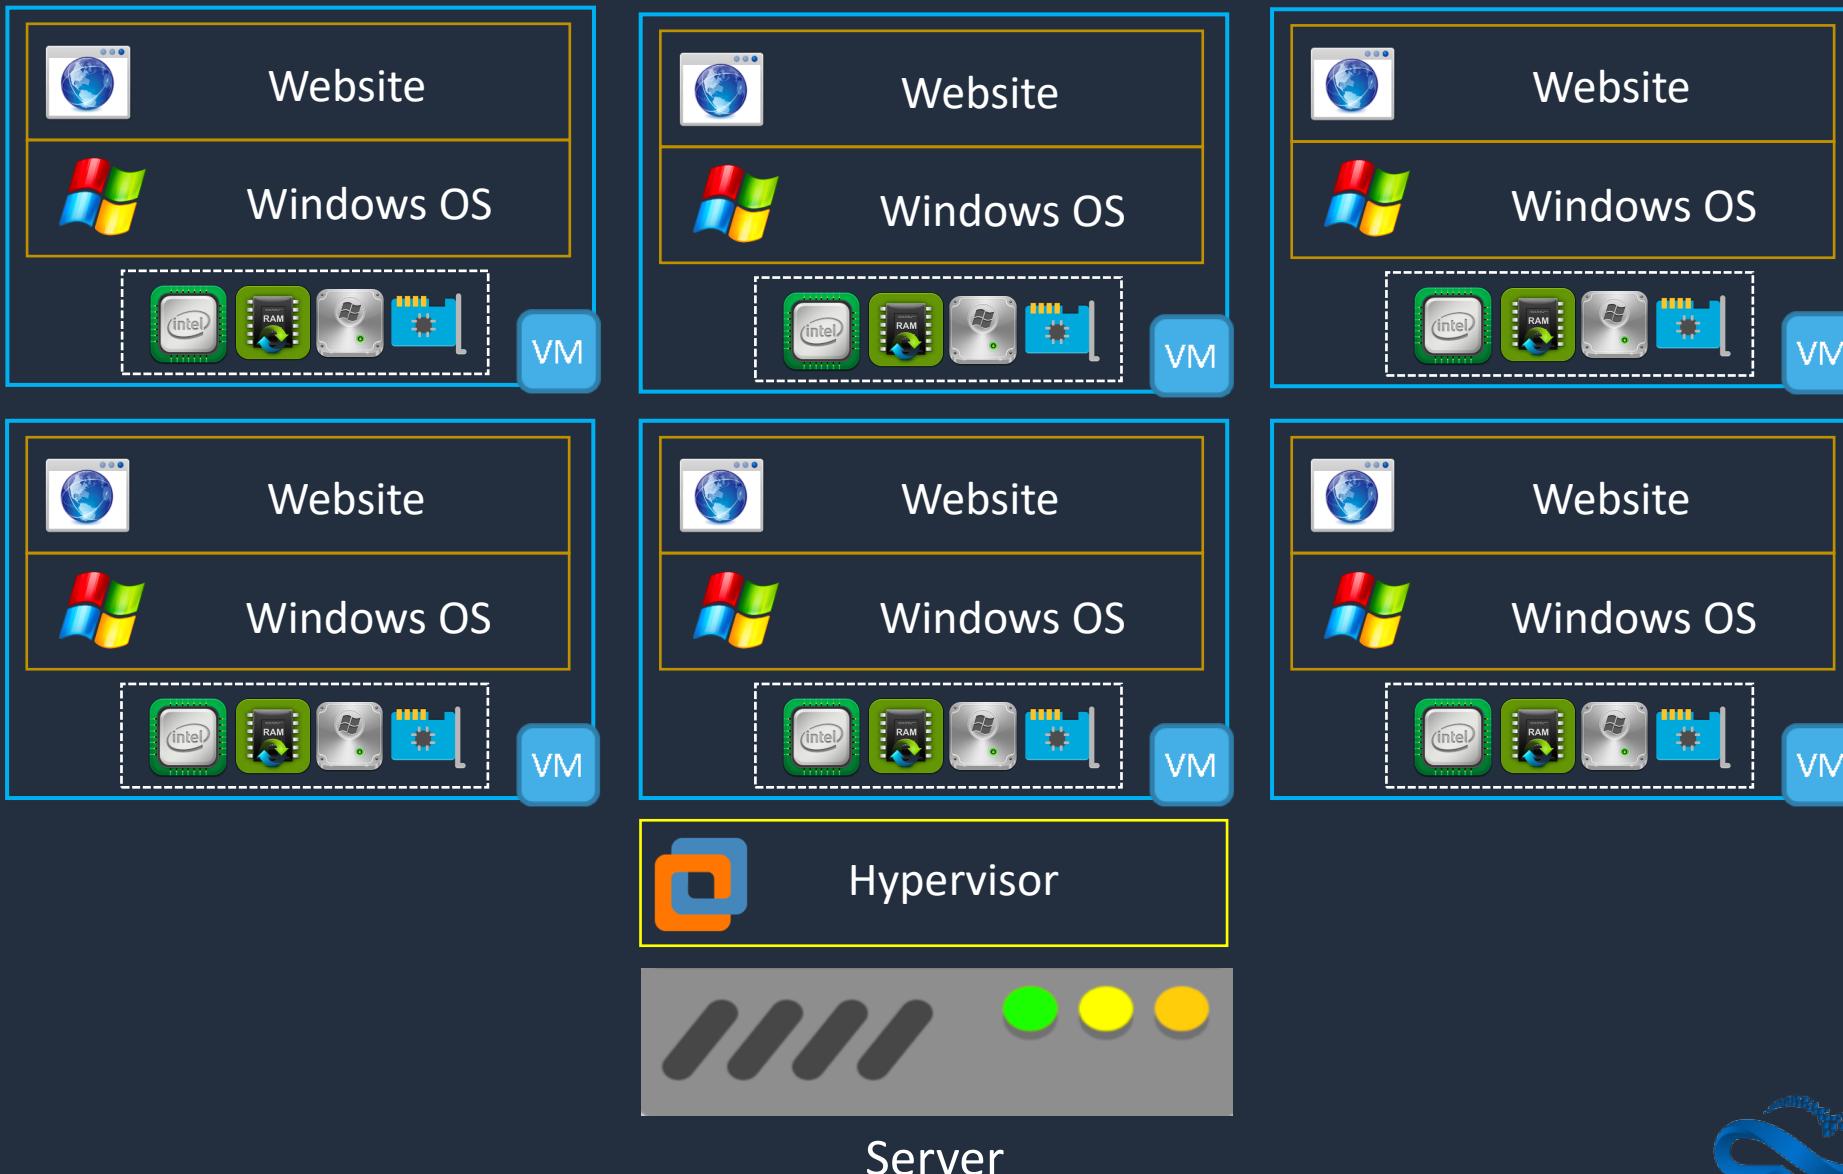


Many VMs can run on the  
same **physical hardware**

**Virtual hardware** is  
presented to the OS



# Server with Virtualization

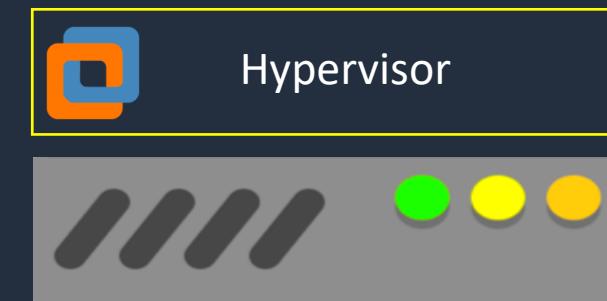




# Server Virtualization: Portability



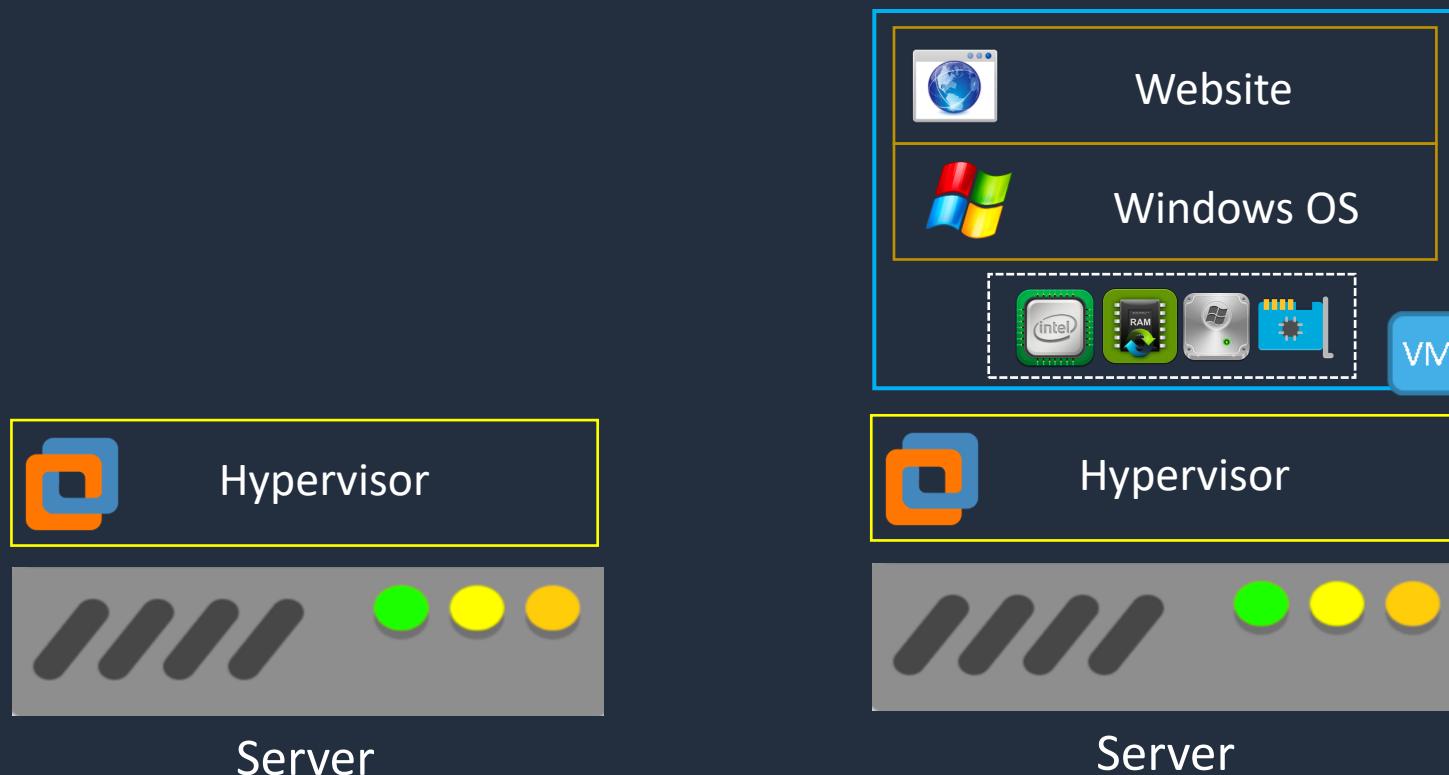
Server



Server



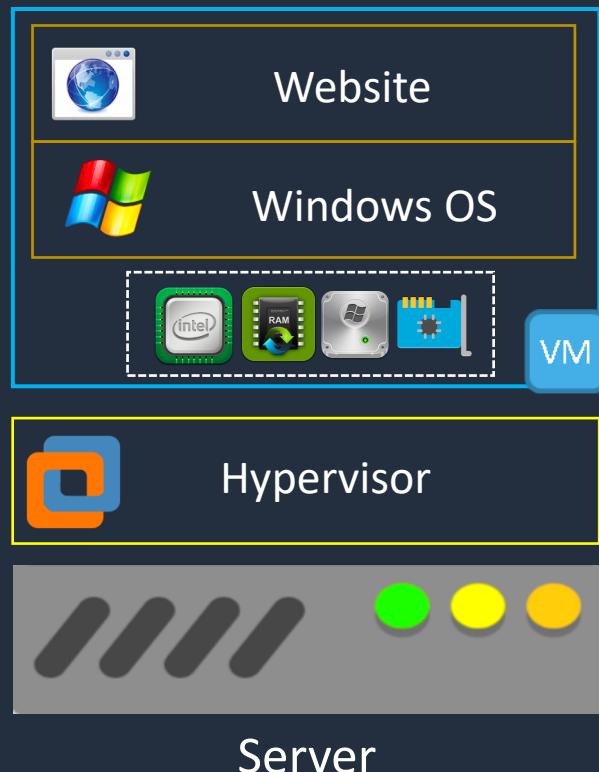
# Server Virtualization: Portability





# Docker Containers

Every VM needs an operating system  
which uses significant resources





# Docker Containers

Containers start up very **quickly**

A container includes all the code, settings, and dependencies for running the application



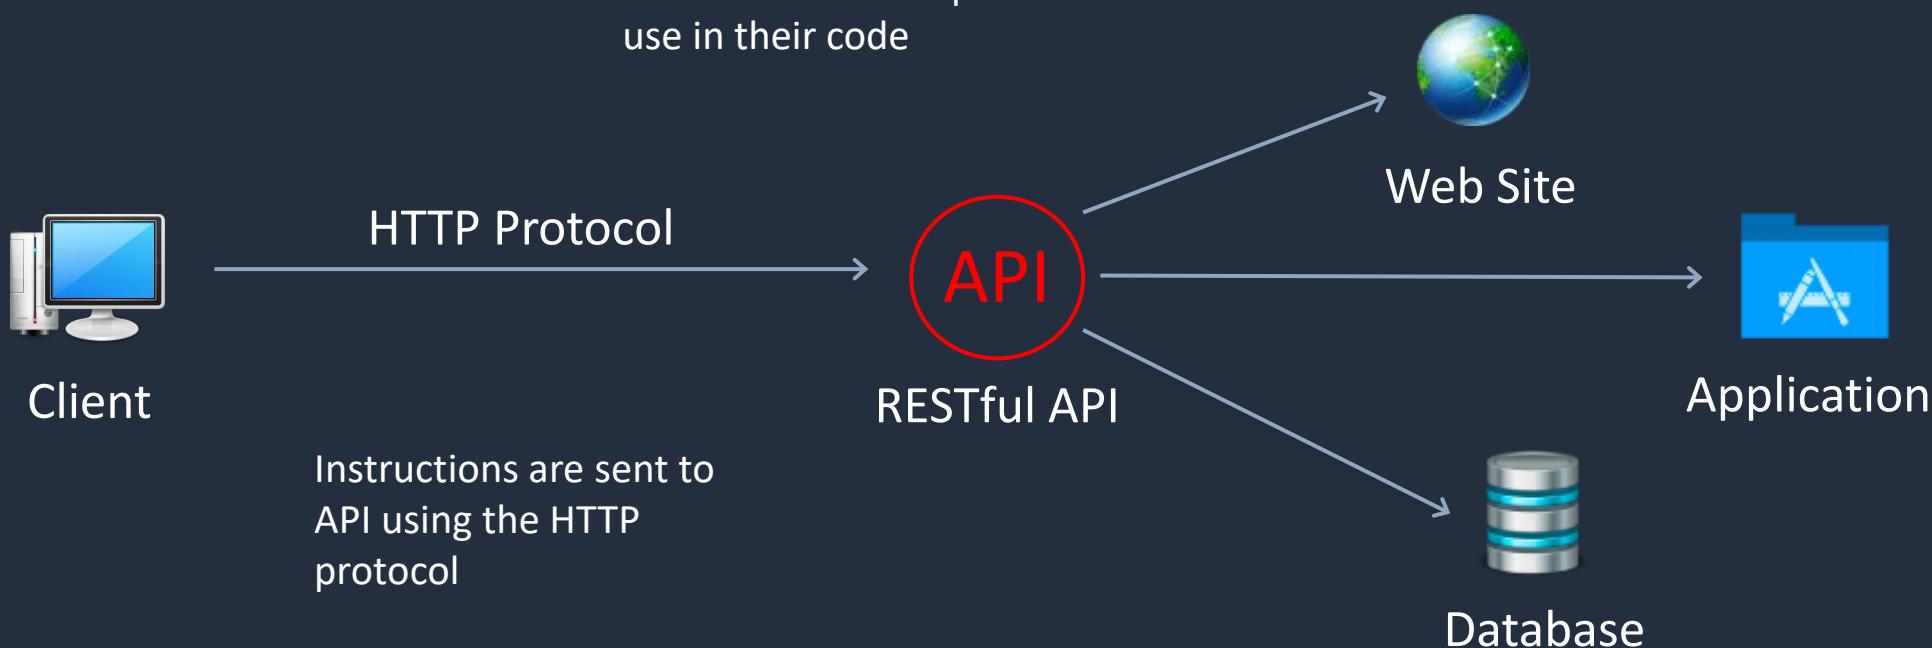
Containers are very **resource efficient**

Each container is **isolated** from other containers

# Application Programming Interfaces (APIs)

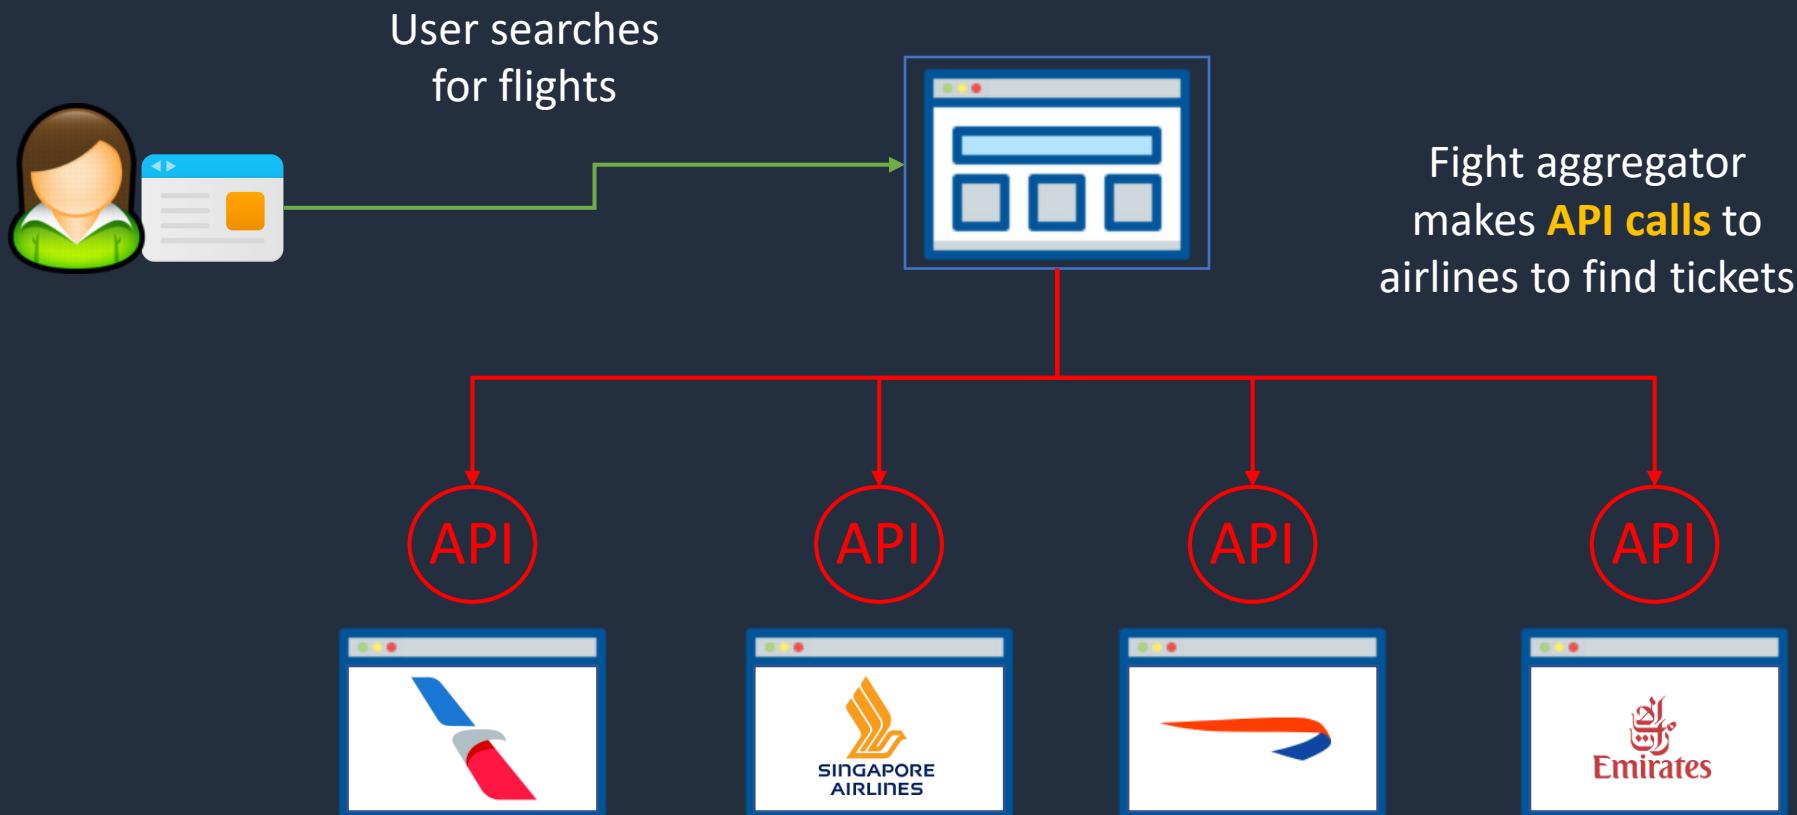


# Application Programming Interfaces (APIs)



# Flight Aggregator Example

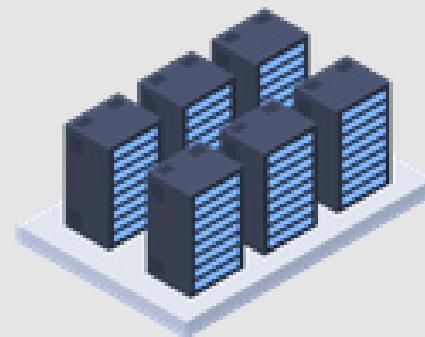
**Flight aggregator** such as  
Momondo or Skyscanner



# SECTION 4

## Cloud Computing Concepts

# Legacy IT / Traditional IT

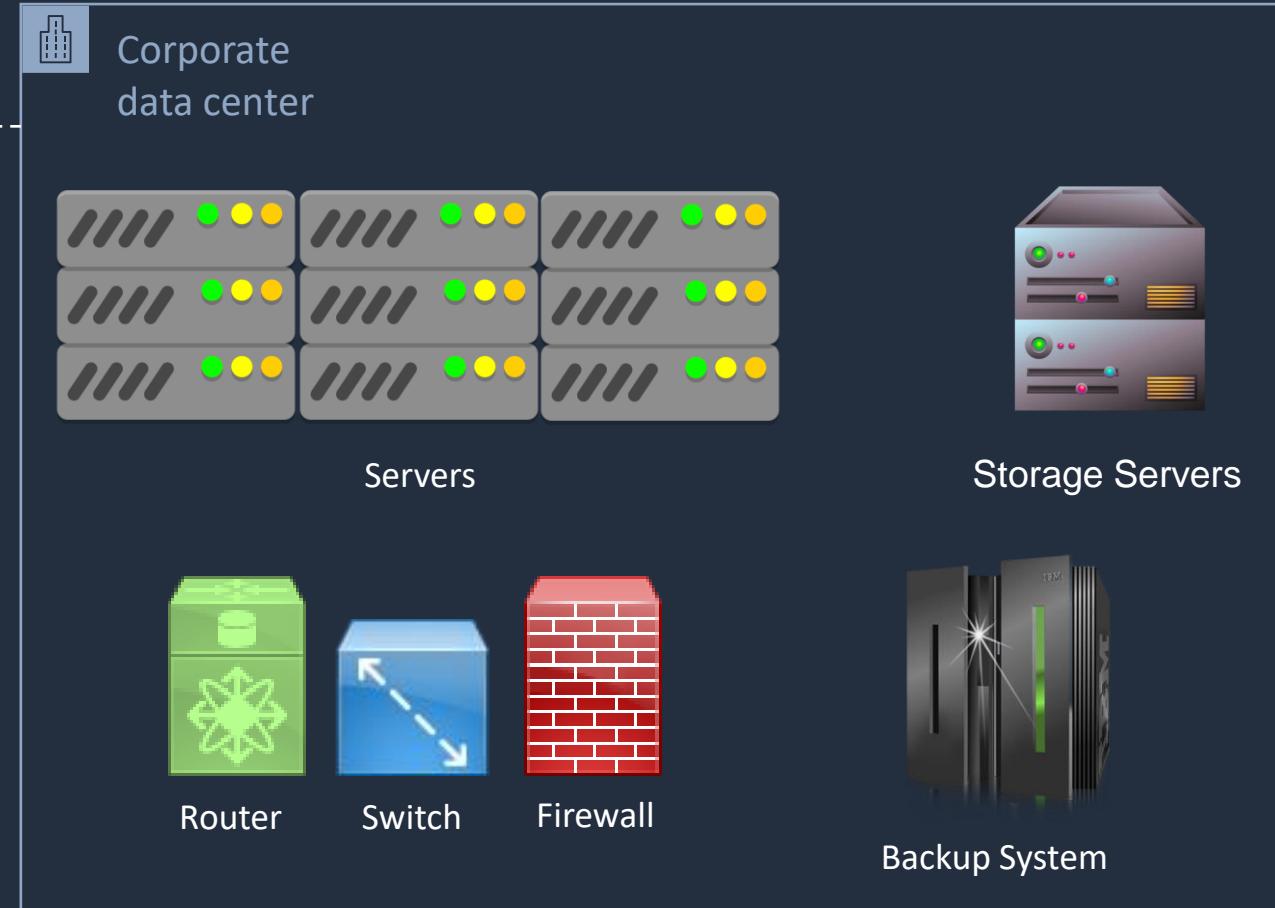




# Traditional IT



This model is very **capital** intensive

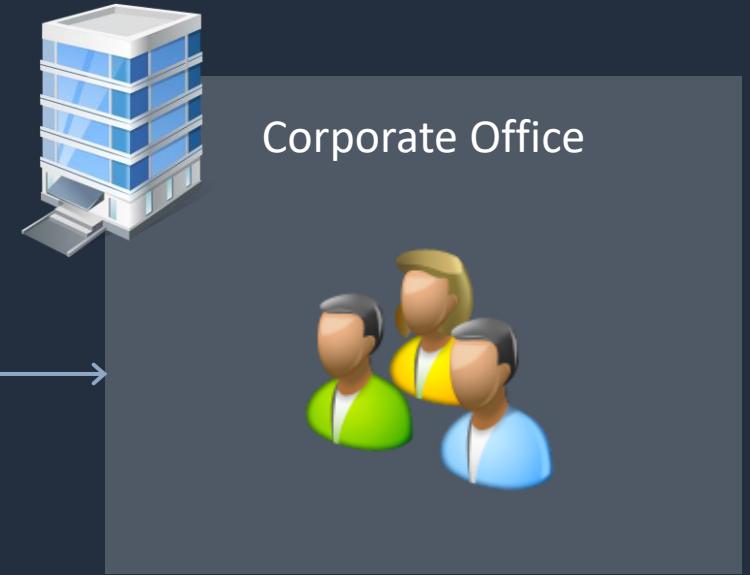
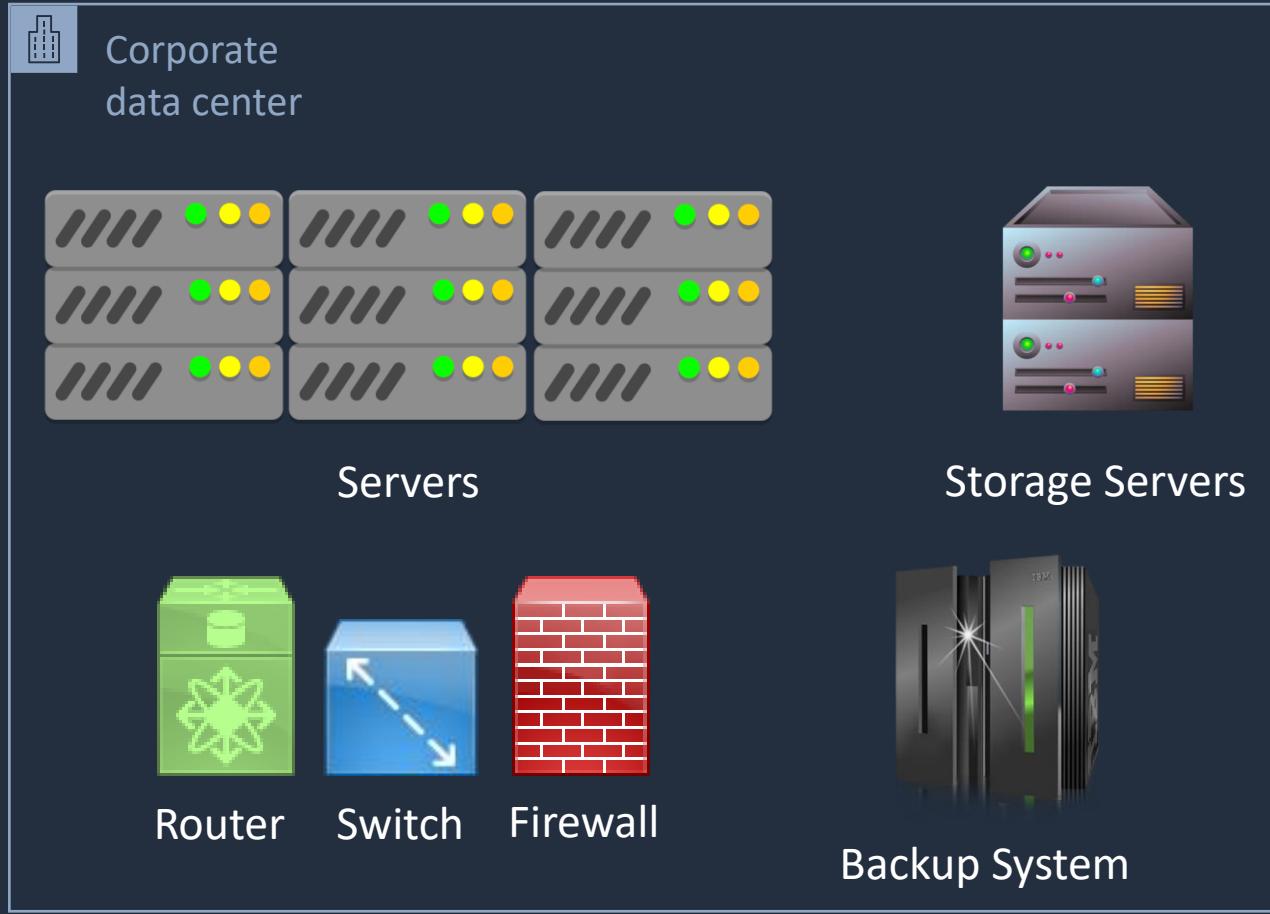


A company typically leases space in a data center, or may own the whole building

The IT equipment is **owned** by the company



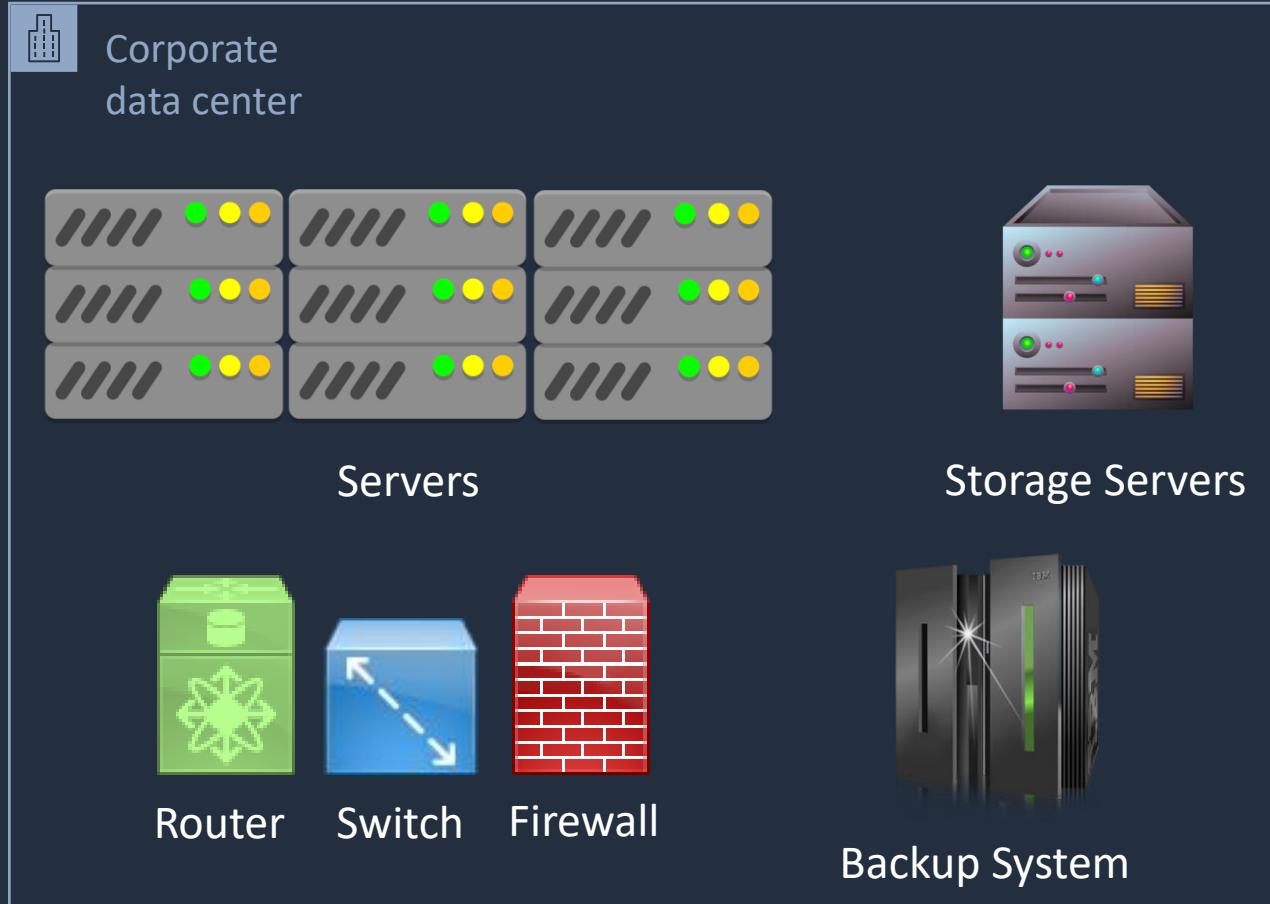
# Traditional IT



IT staff must design,  
build, operate, and  
manage equipment



# Traditional IT



## Costs:

- Data center building
- Data center security
- Physical IT hardware
- Software licensing costs
- Maintenance contracts
- Power
- Internet connectivity
- Staff wages (design, build, operations, maintenance)

# What is Cloud Computing?





# What is Cloud Computing? Well-known examples

## Non-Cloud Services:



Email Server



File Server



Customer Relationship  
Management (CRM)

## Cloud Services:



Gmail



Dropbox



Salesforce

You don't own or manage the infrastructure on which the service runs

Cloud services are offered on a subscription / consumption model

The service scales as demand changes



# What is Cloud Computing? The Key Characteristics

Name	Description
On-demand, self-service	A user can consume cloud resources, as needed, automatically, and without human interaction
Broad network access	Capabilities are available over the network using standard mechanisms. Can be the Internet or a Wide Area Network (WAN)
Resource pooling	The providers resources are pooled and serve multiple consumers using a multi-tenant model
Rapid elasticity	Capabilities can scale “elastically” based on demand
Measured service	Resource usage is monitored and metered



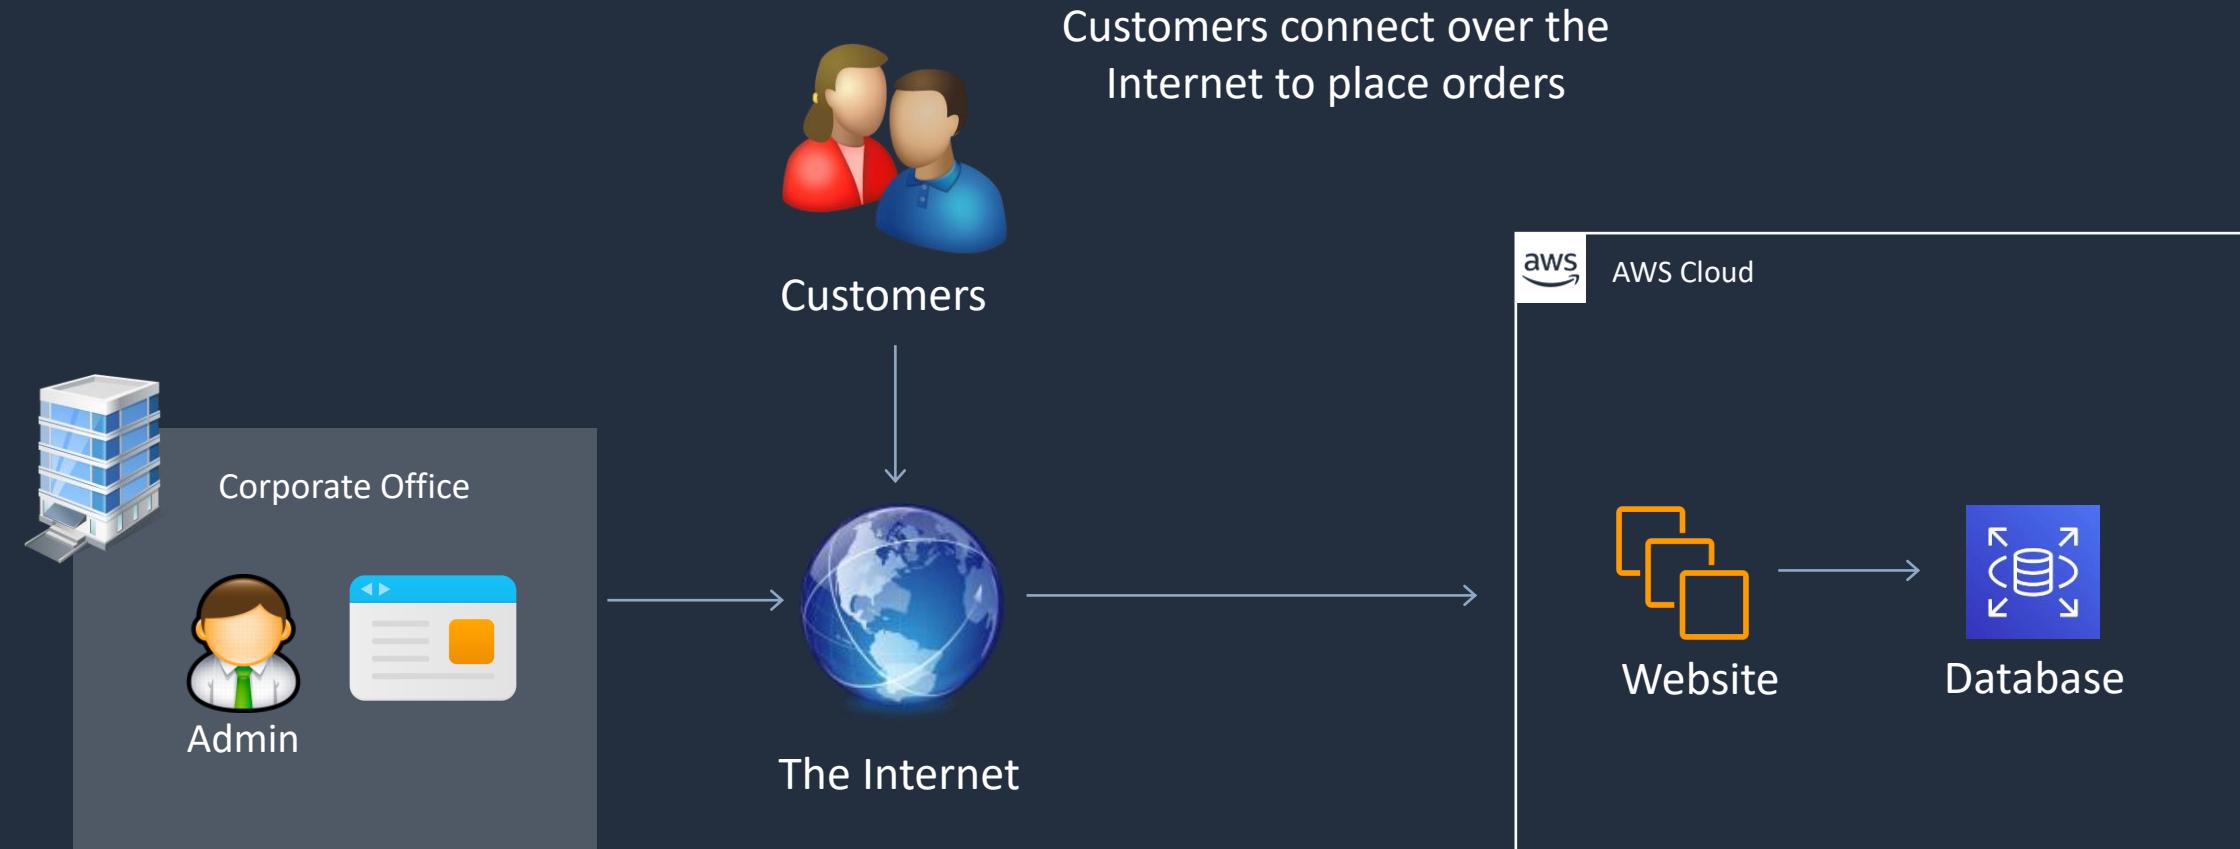
# Deploying an eCommerce Website on-premises (aka the old way)





# Deploying an eCommerce Website in the Cloud

---



The admin deploys an eCommerce website / database to AWS using the AWS Management Console



# Launching Cloud Services: Management Console



## Management Console

A web-based console accessed through a standard web browser

<b>Compute</b> EC2 Lightsail ↗ Lambda Batch Elastic Beanstalk Serverless Application Repository AWS Outposts EC2 Image Builder	<b>Blockchain</b> Amazon Managed Blockchain	<b>Analytics</b> Athena EMR CloudSearch Elasticsearch Service Kinesis QuickSight ↗ Data Pipeline AWS Data Exchange AWS Glue AWS Lake Formation MSK	<b>Business Applications</b> Alexa for Business Amazon Chime ↗ WorkMail Amazon Honeycode
<b>Storage</b> S3 EFS FSx S3 Glacier Storage Gateway AWS Backup	<b>Satellite</b> Ground Station	<b>Quantum Technologies</b> Amazon Braket ↗	<b>End User Computing</b> WorkSpaces AppStream 2.0 WorkDocs WorkLink
<b>Database</b> RDS DynamoDB ElastiCache Neptune Amazon Redshift Amazon QLDB	<b>Management &amp; Governance</b> AWS Organizations CloudWatch AWS Auto Scaling CloudFormation CloudTrail Config OpsWorks	<b>Security, Identity, &amp; Compliance</b> IAM Resource Access Manager Cognito Secrets Manager GuardDuty Inspector Amazon Macie AWS Single Sign-On Certificate Manager Key Management Service CloudHSM Directory Service	<b>Internet Of Things</b> IoT Core FreeRTOS IoT 1-Click IoT Analytics IoT Device Defender IoT Device Management IoT Events IoT Greengrass IoT SiteWise IoT Things Graph



# Launching Cloud Services: Command Line



## Command Line

This command launches a **virtual server** (EC2 instance) on Amazon Web Services



```
aws ec2 run-instances --image-id ami-xxxxxxxx --count 1 --instance-type t2.micro
```



```
aws s3 ls s3://mys3databucket
```

This command lists the contents of a storage container (bucket) on **Amazon S3**



# Launching Cloud Services: Software Development Kit



Code (SDK)

A developer writes code in an integrated development environment (IDE)

The screenshot shows a Visual Studio Code interface with the following details:

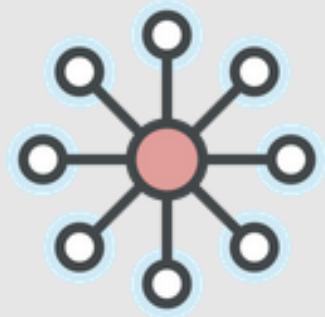
- File Menu:** File, Edit, Selection, View, Go, Debug, Terminal, Help.
- Title Bar:** app.js - MY-SAM-APP (Workspace) - Visual Studio C... -
- Explorer:** Shows the file structure of the MY-SAM-APP workspace, including .aws, templates.json, my-sam-app-nodejs, hello-world, tests, .npmignore, app.js, package.json, .gitignore, event.json, README.md, template.yaml, and MY-SAM-APP.code-workspace.
- Editor:** The app.js file is open, showing a Lambda function handler:

```
my-sam-app-nodejs > hello-world > app.js > lambdaHandler > exports.lambdaHandler
15  *
16  */
Run Locally | Debug Locally | Configure
17 exports.lambdaHandler = async (event, context) => {
18   try {
19     // const ret = await axios(url);
20     response = {
21       'statusCode': 200,
22       'body': JSON.stringify({
23         message: 'hello world',
24         // location: ret.data.trim()
25       })
26     }
27   } catch (err) {
28     console.log(err);
29     return err;
30   }
31
32   return response
33 };
34
```

- Bottom Status Bar:** Ln 29, Col 1 (27 selected) Spaces: 4 UTF-8 CRLF JavaScript AWS:default

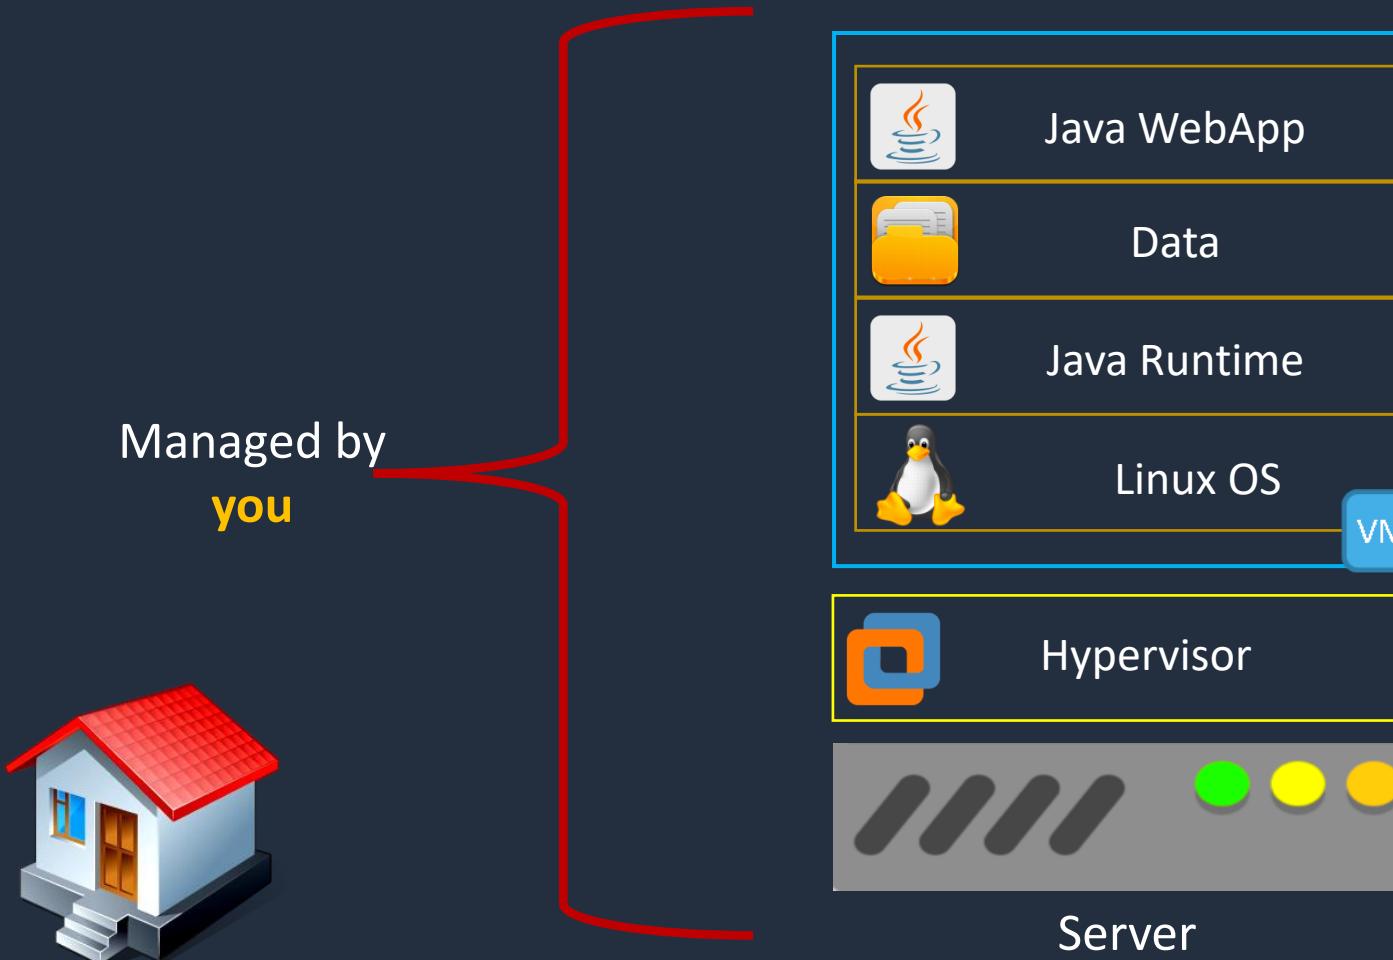
The code leverages the SDK to work with cloud services

# Cloud Computing Service Models





# Cloud Service Models: Private Cloud

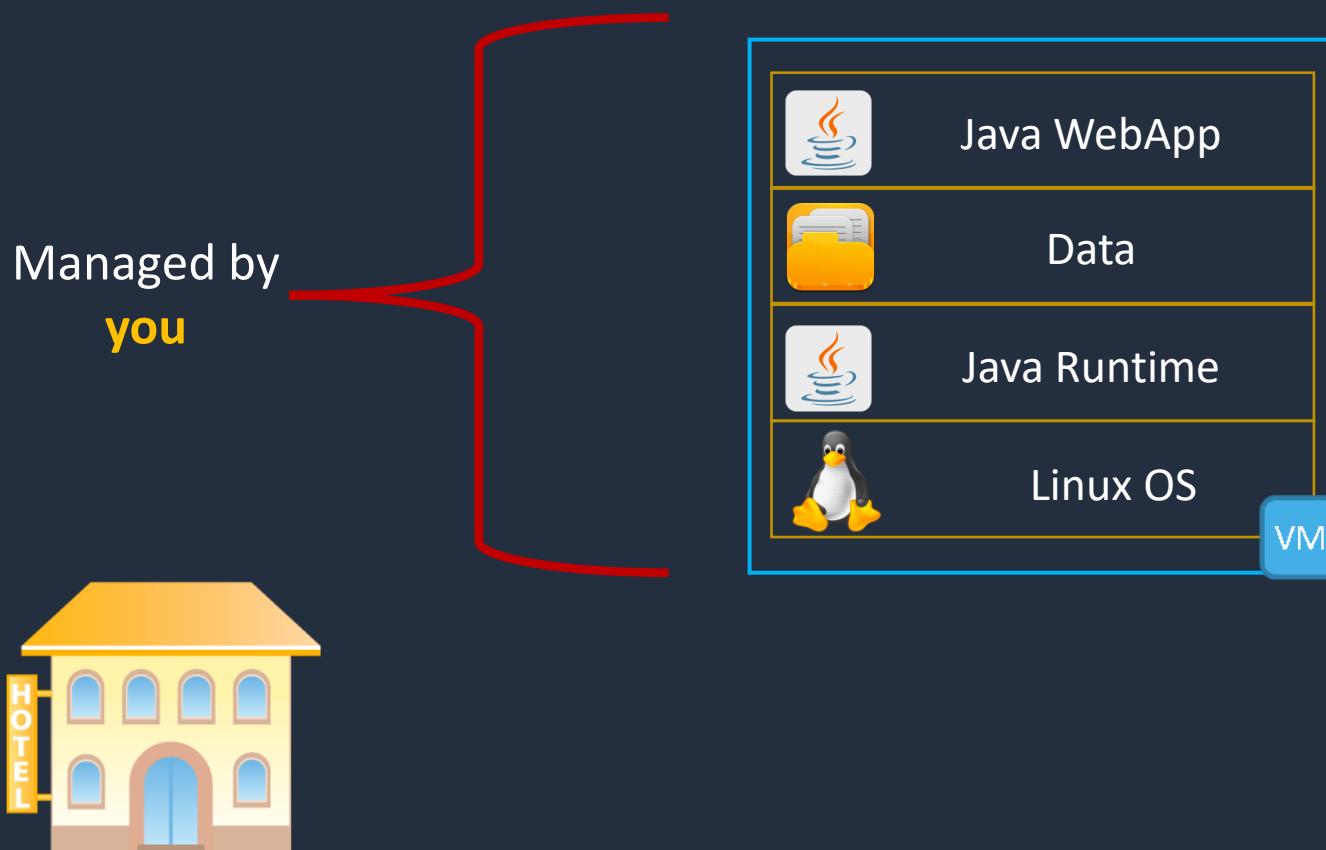


A **private cloud** must also include self-service, multi-tenancy, metering, and elasticity



# Cloud Service Models: Infrastructure as a Service (IaaS)

---



Examples:

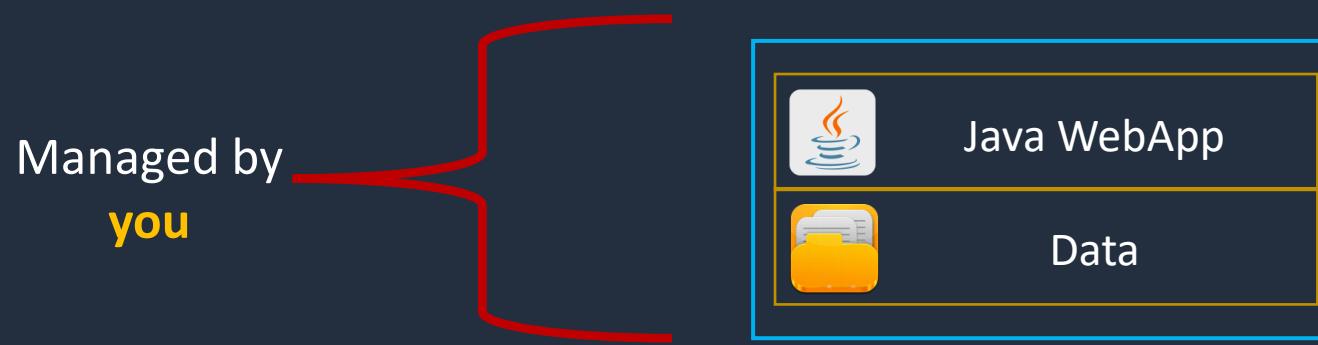
- Amazon Elastic Compute Cloud (EC2)
- Azure Virtual Machines
- Google Compute Engine



# Cloud Service Models: Platform as a Service (PaaS)

---

---



Examples:

- AWS Elastic Beanstalk
- Azure WebApps
- Compute App Engine



# Cloud Service Models: Software as a Service (SaaS)

---

---

Managed by  
you



Java WebApp

Pure consumption  
model

Examples:

- Google Apps
- Salesforce.com
- Zoom



# Cloud Service Models: Comparison

## Private Cloud



You manage **everything** -  
greater responsibility +  
greater control

## IaaS

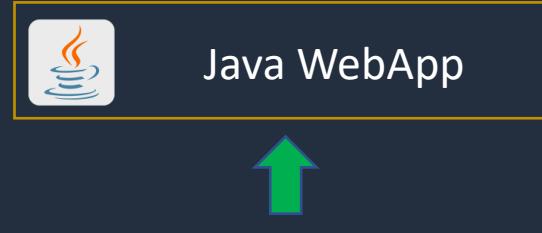


You manage from the  
**virtual server** upwards

## PaaS



## SaaS



You simply **consume** the  
service - little responsibility  
+ little control

# Cloud Computing Deployment Models





# Cloud Computing Deployment Models

---

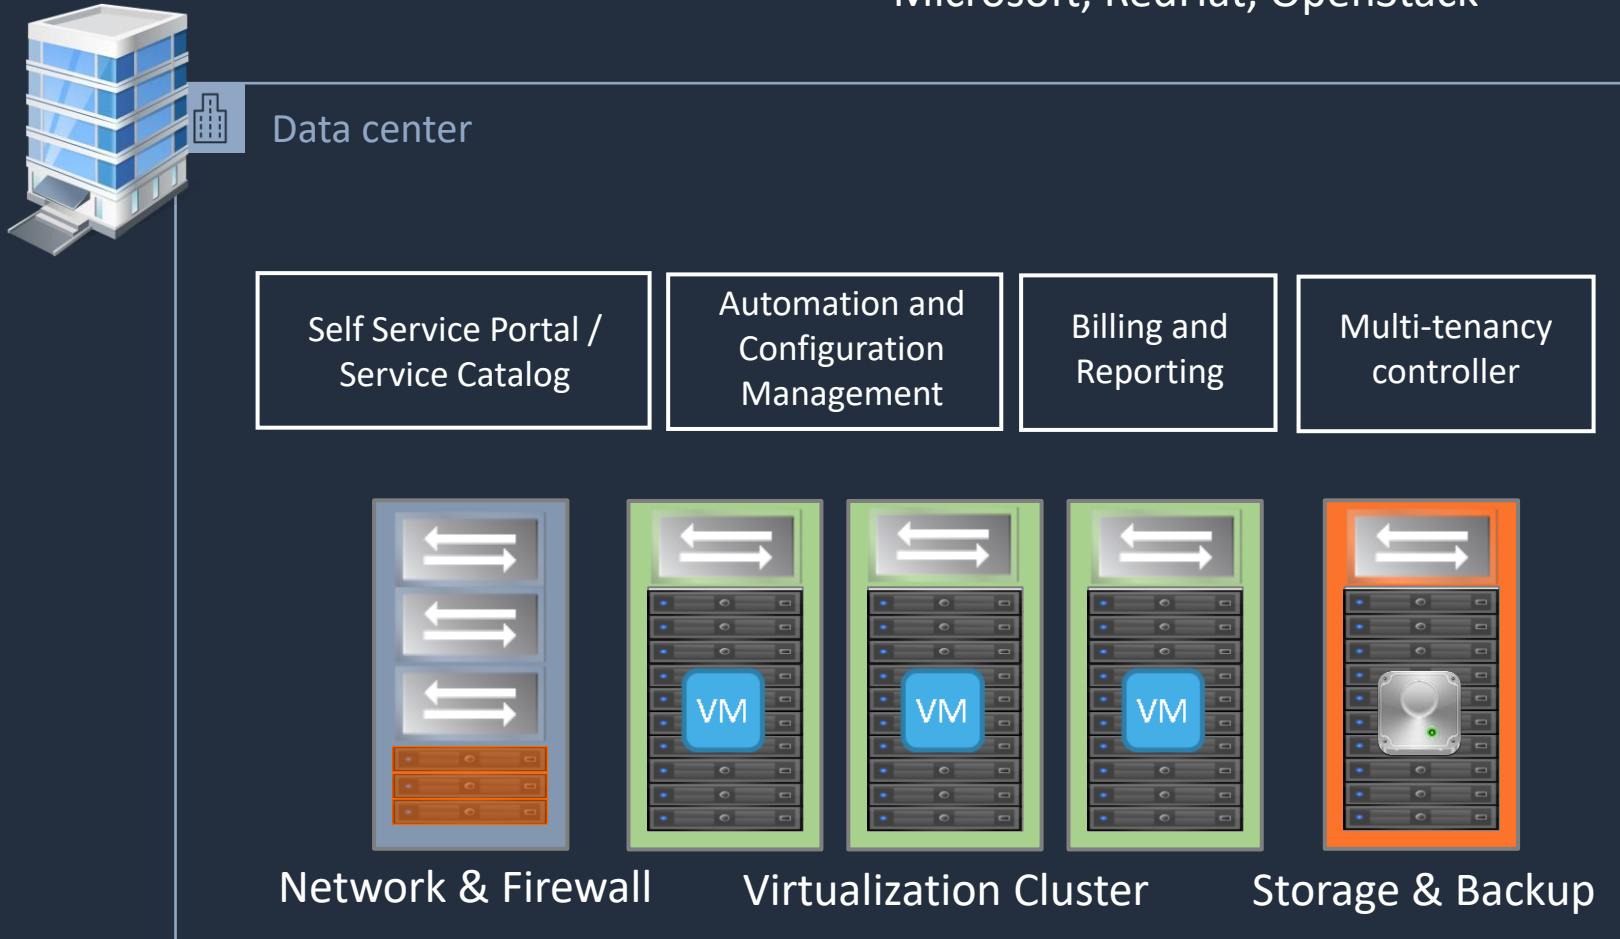
---

Name	Description	Examples
Private Cloud	An enterprise deploys their own infrastructure and applications into their own data center	VMware, Microsoft, RedHat, OpenStack
Public Cloud	The IT services that you consume are hosted and delivered from a third-party and accessed over the Internet	AWS, Microsoft Azure, Google Cloud Platform
Hybrid Cloud	A combination of on-premises, private cloud, and public cloud services are consumed	
Multicloud	Usage of two or more public clouds at a time, and possibly multiple private clouds	



# Private Cloud

Examples are VMware,  
Microsoft, RedHat, OpenStack



Cloud management  
software layer

## Benefits:

- Complete control of the entire stack
- Security – in a few cases, organizations may need to keep all or some of their applications and data in house

You **build** and **manage** the  
cloud deployment

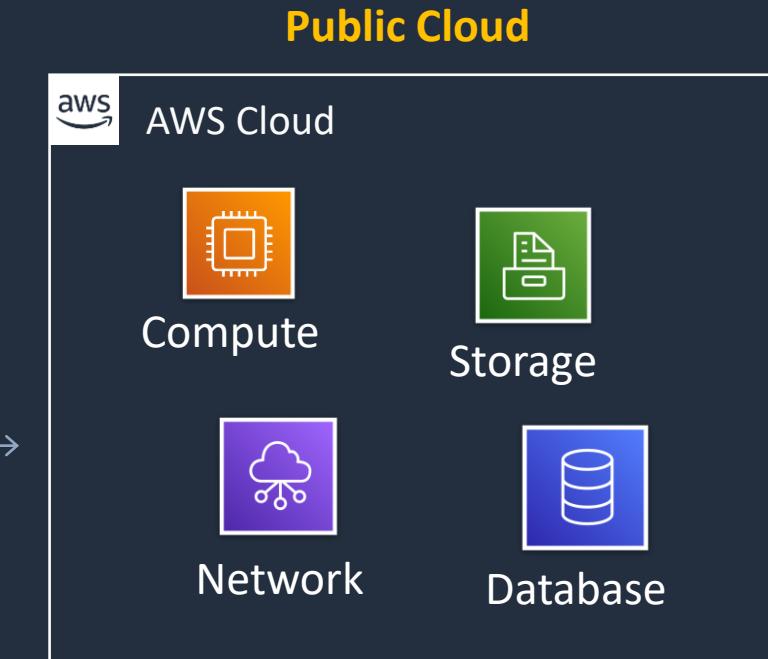


# Cloud Deployment Models - Public Cloud

Examples are AWS, Microsoft Azure, Google Cloud Platform

## Benefits:

- Variable expense, instead of capital expense
- Economies of scale
- Massive elasticity



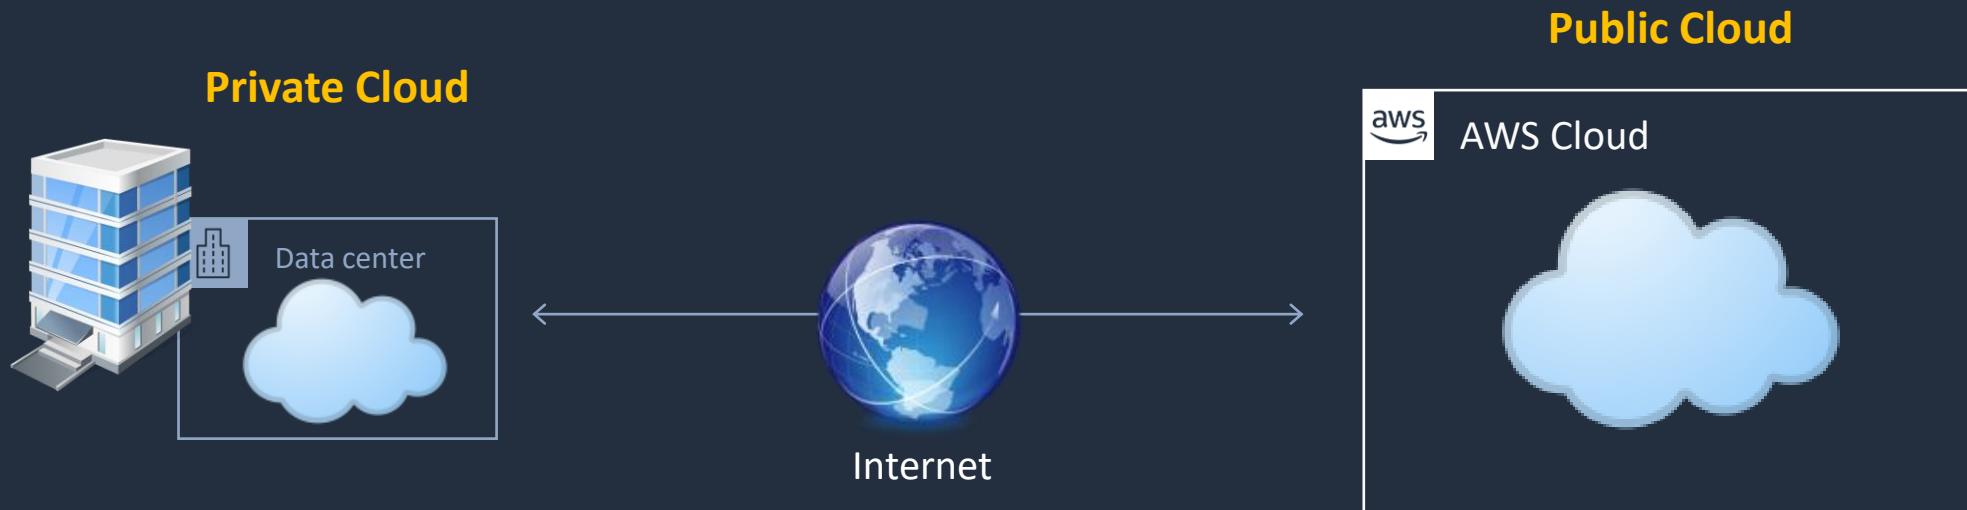
Connected using either the  
**Internet** or a **private link**



# Cloud Deployment Models - Hybrid Cloud

## Benefits:

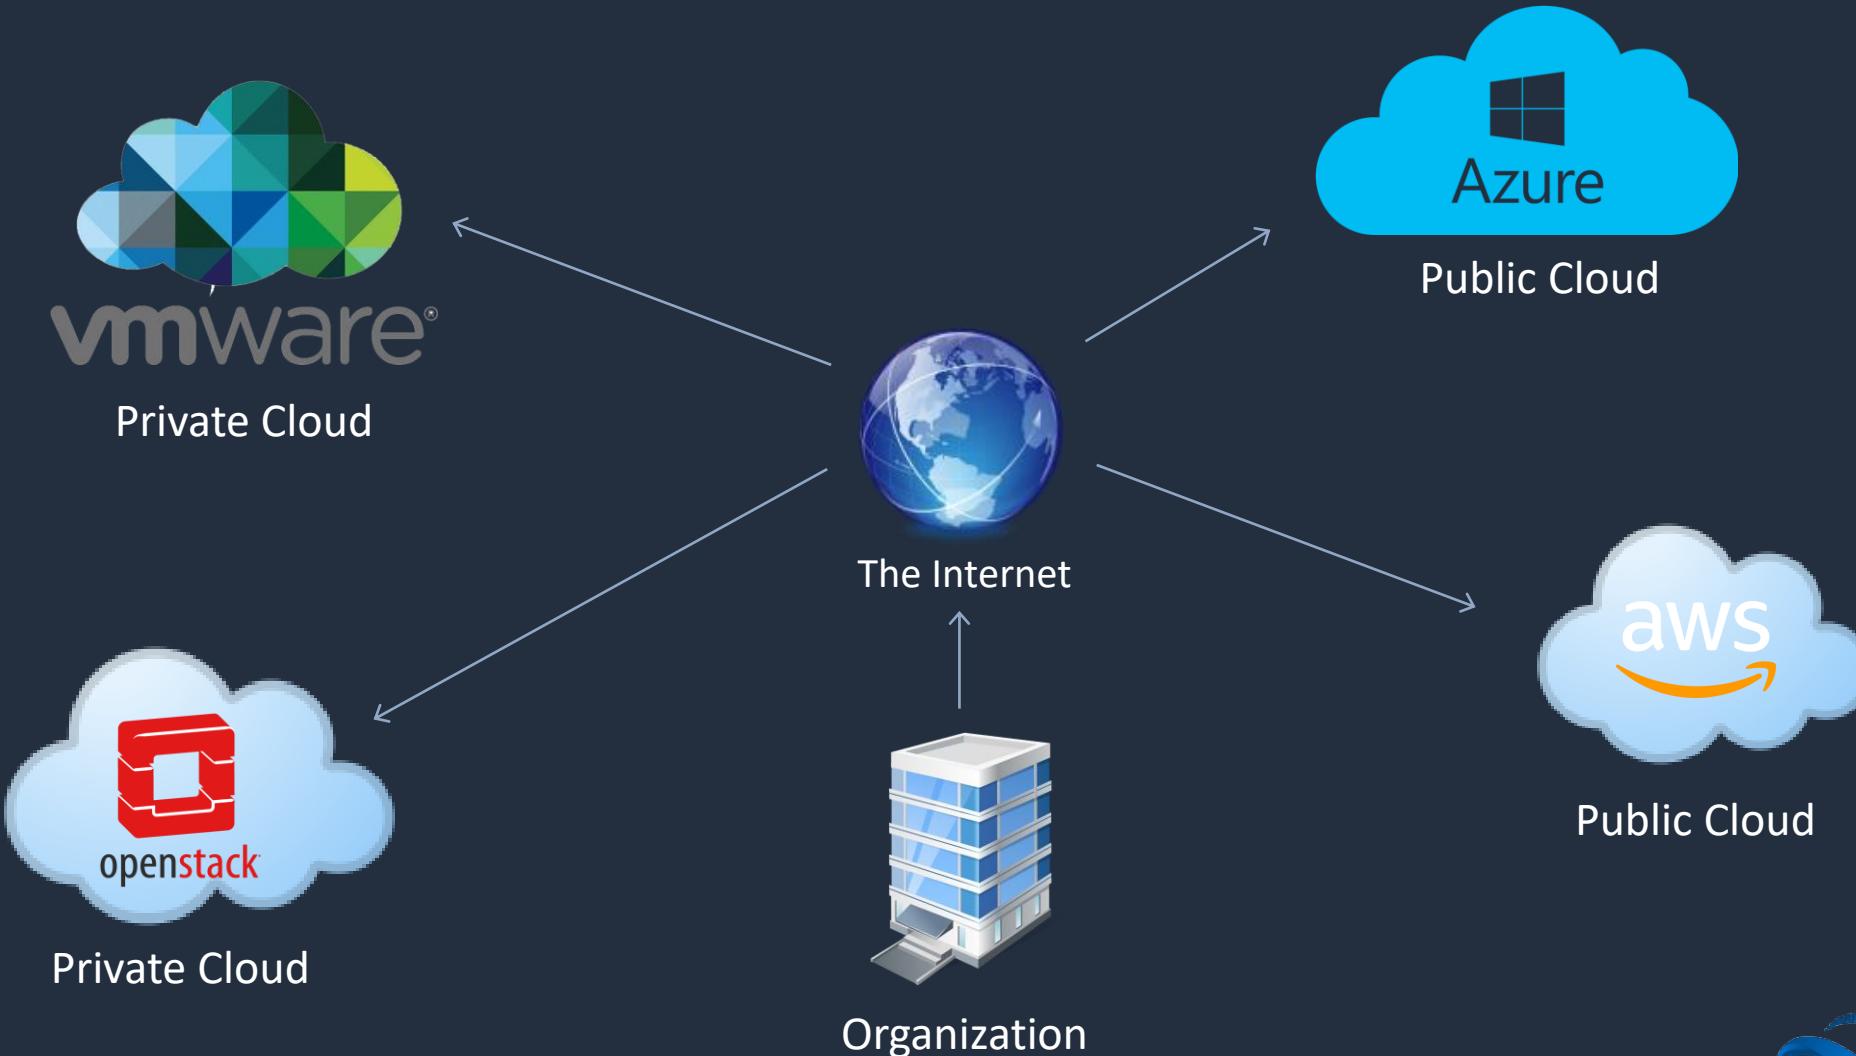
- Allows companies to keep the critical applications and sensitive data in a traditional data center environment or private cloud
- Take advantage of public cloud resources like SaaS, for the latest applications, and IaaS, for elastic virtual resources
- Facilitates portability of data, apps and services and more choices for deployment models



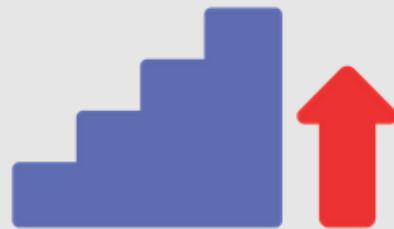
Connected using either the  
**Internet** or a **private** link



# Cloud Deployment Models - Multicloud



# Scaling Up vs Scaling Out





# Stateful vs Stateless Applications

---

Stateless:

No “state” is recorded about the user's session



Person checks a news website



Stateful:

Netflix records what has been watched

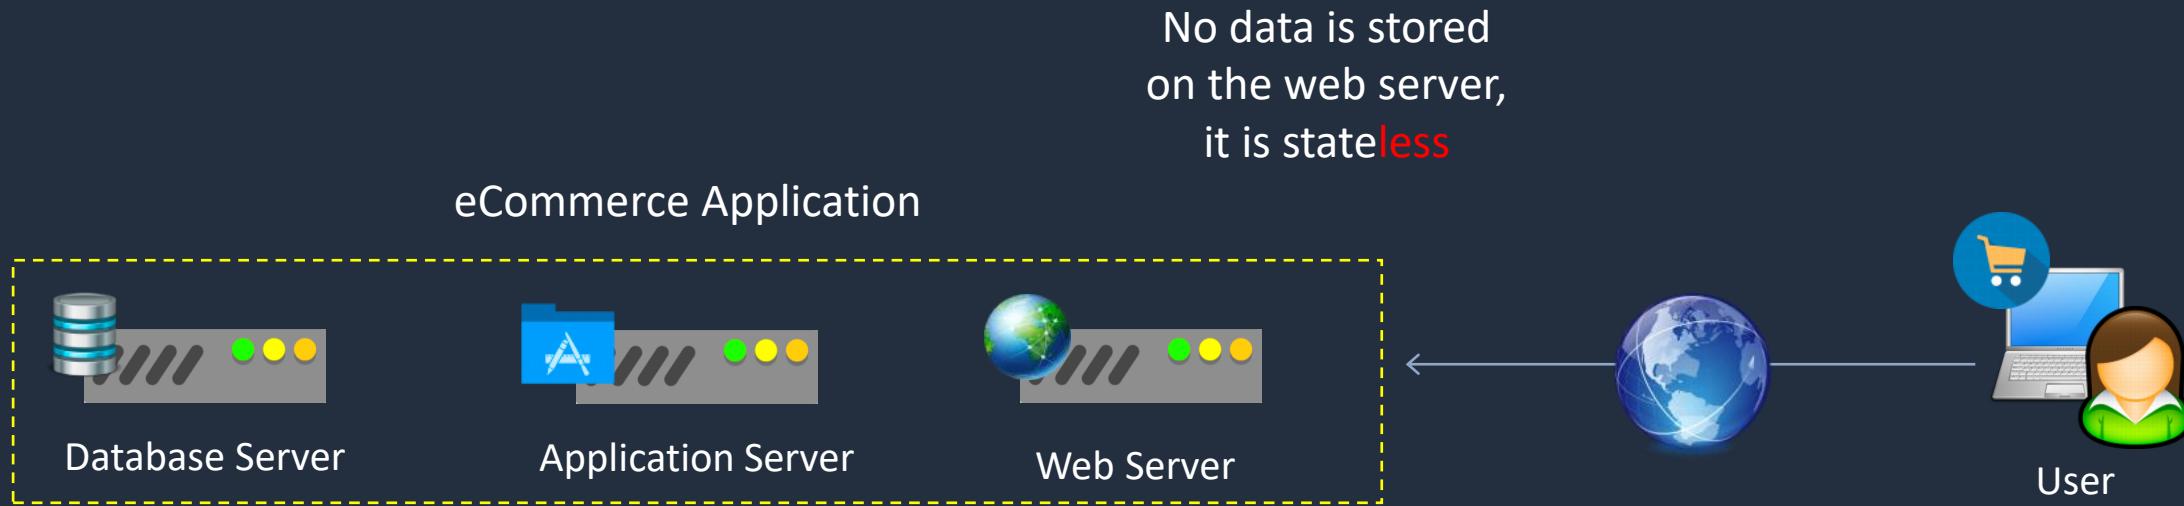


Person logs into Netflix





# Stateful vs Stateless Applications

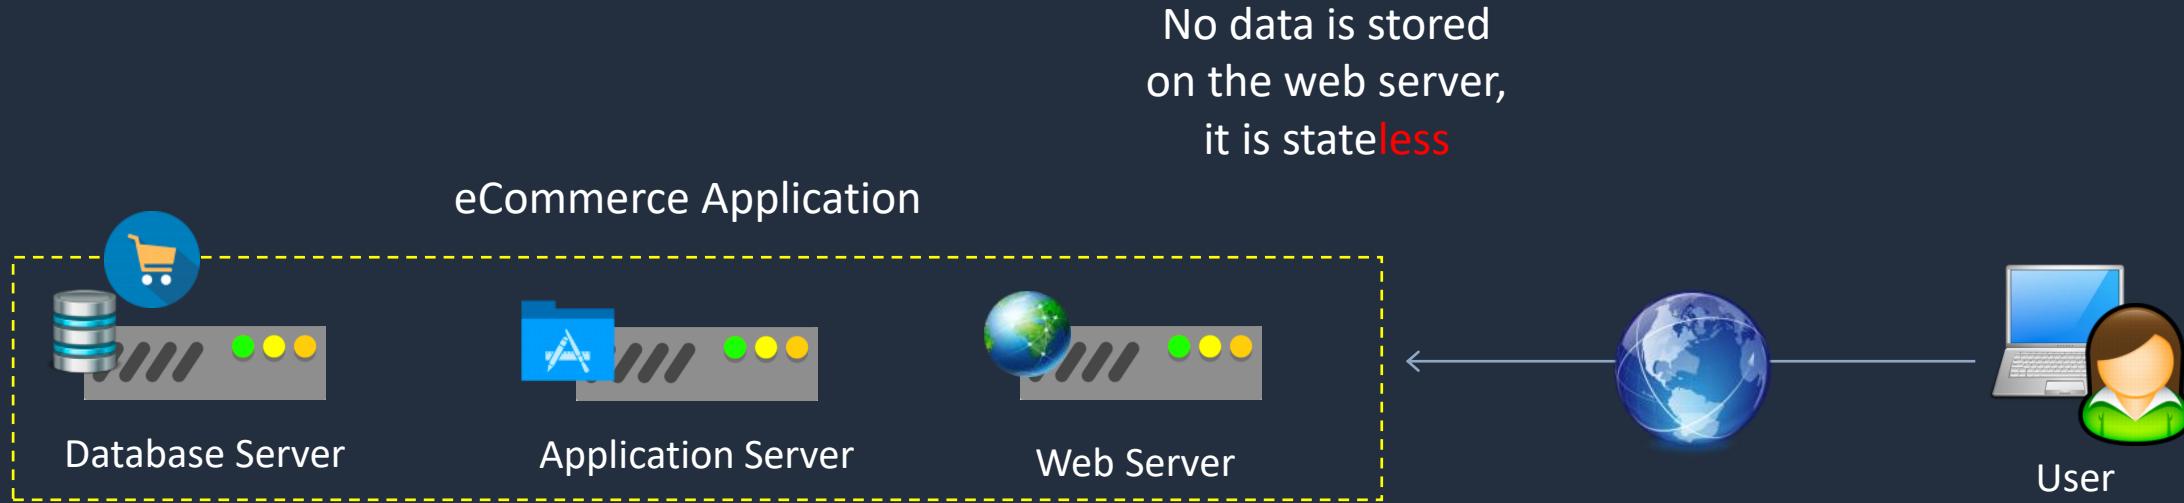


When the user purchases, the application layer processes the order and records the data in the database. This is **stateful**

The cart items are stored in cookies on the computer



# Stateful vs Stateless Applications



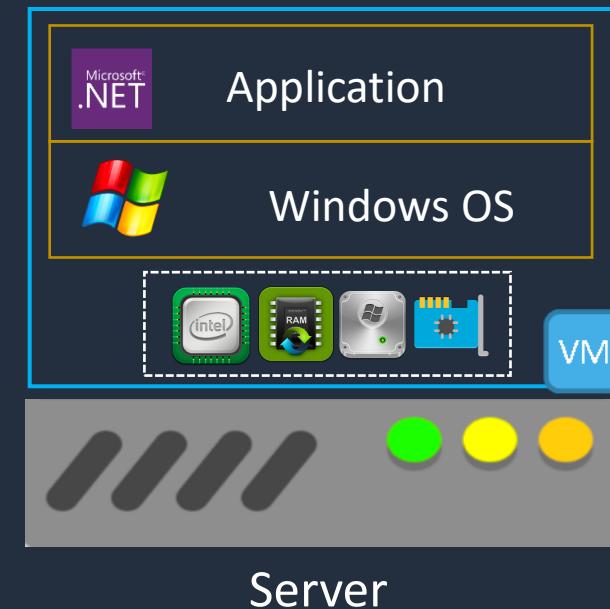
When the user purchases, the application layer processes the order and records the data in the database. This is **stateful**

The cart items are stored in cookies on the computer



# Scalability and Elasticity: Scaling Up

---



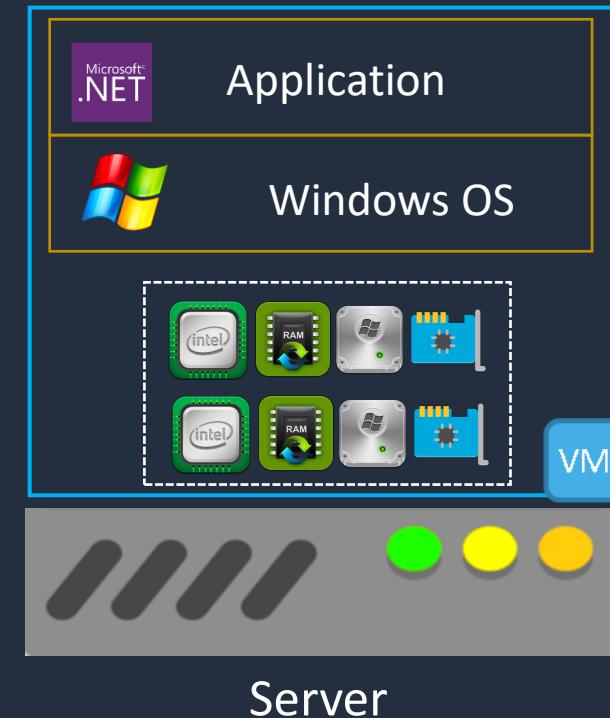
Server



# Scalability and Elasticity: Scaling Up

---

Scaling up means adding resources to the server



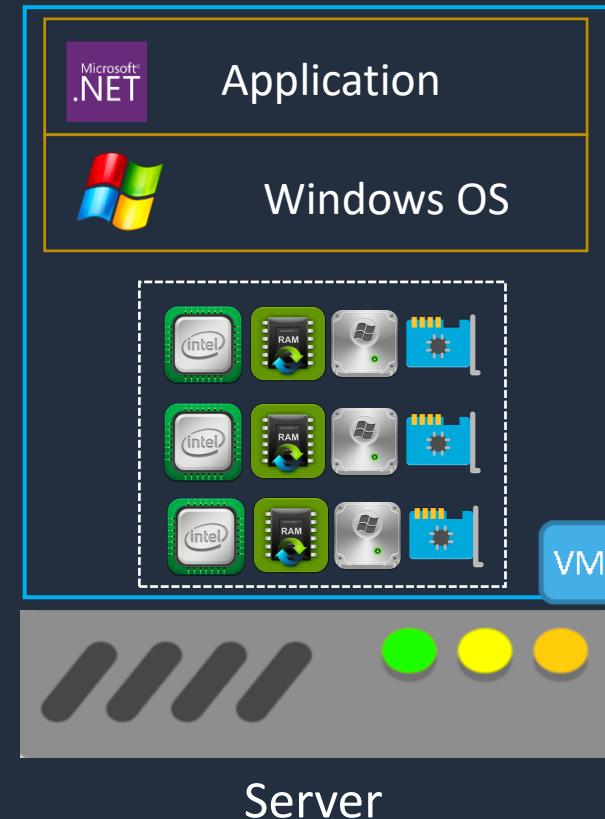
Server



# Scalability and Elasticity: Scaling Up

---

Scaling up means adding resources to the server

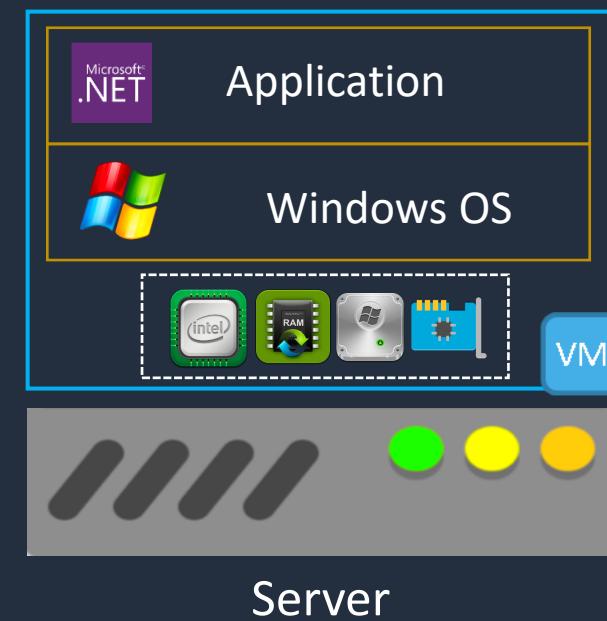


Server



# Scalability and Elasticity: Scaling Out

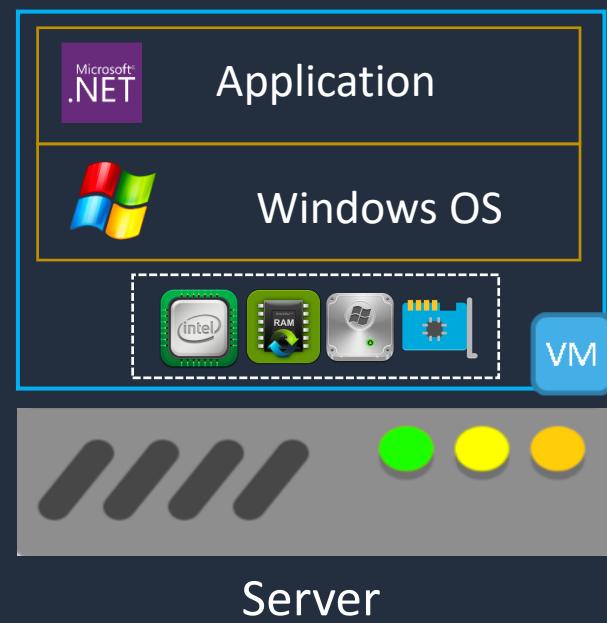
---





# Scalability and Elasticity: Scaling Out

---



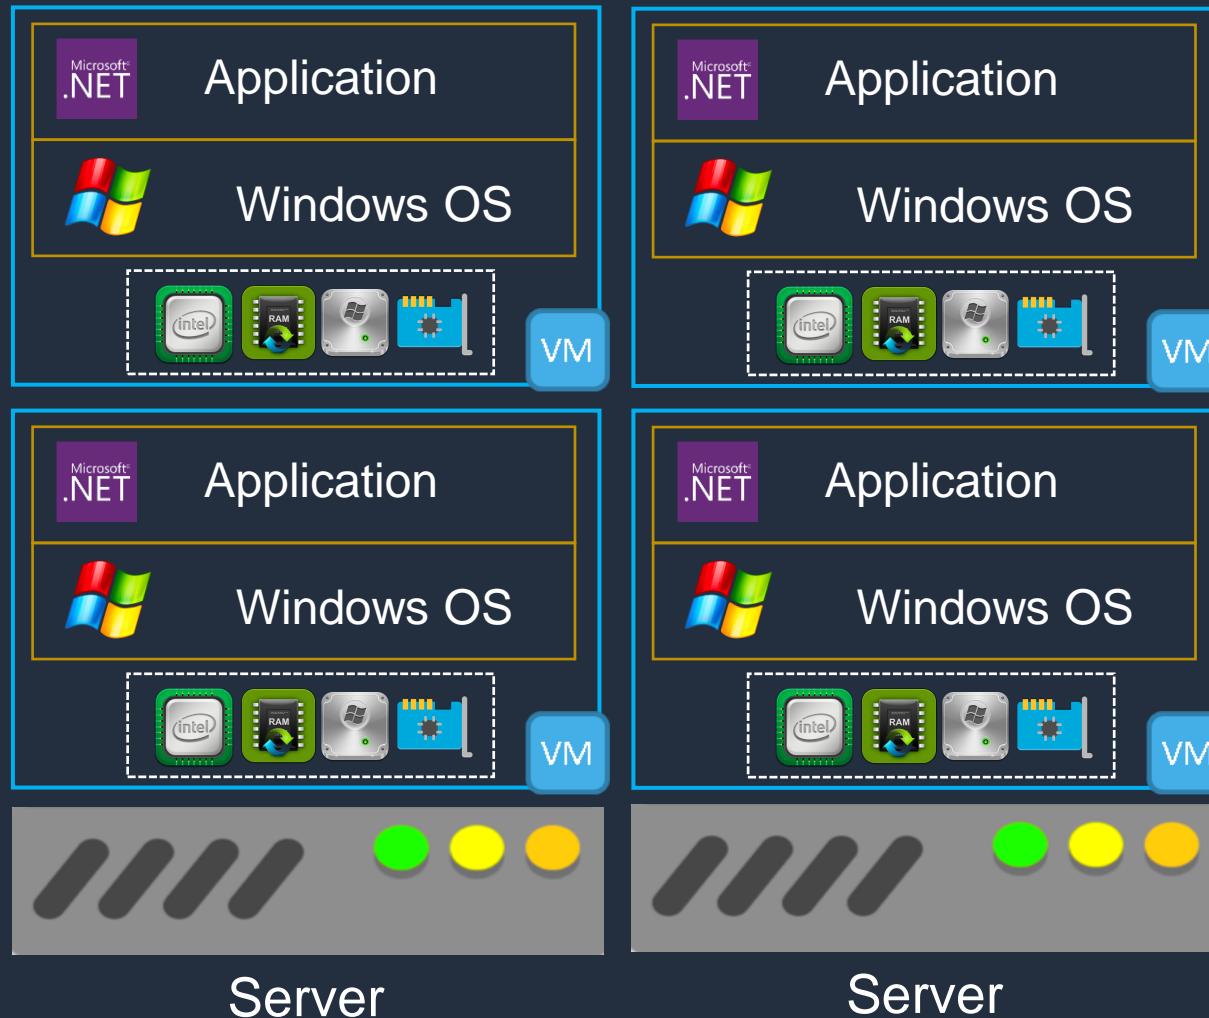
Server



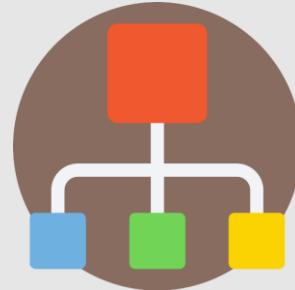
# Scalability and Elasticity: Scaling Out

---

Scaling out means adding additional servers

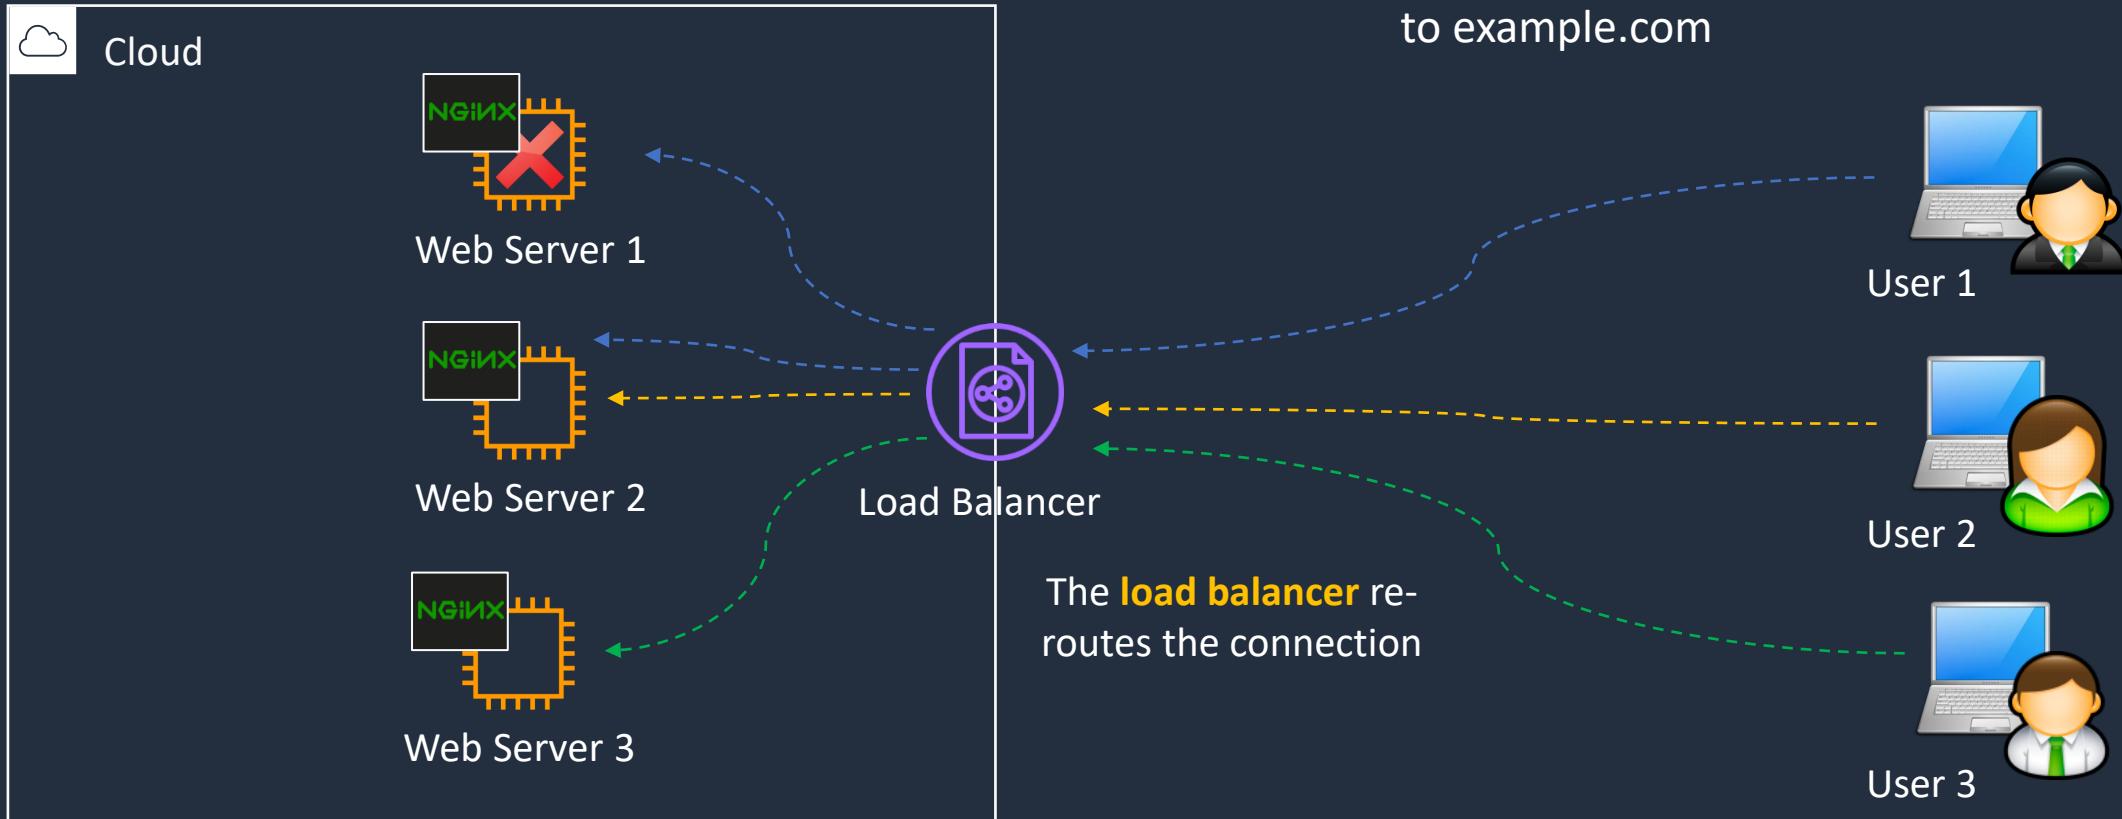


# High Availability and Fault Tolerance





# Load Balancing





# Fault Tolerance

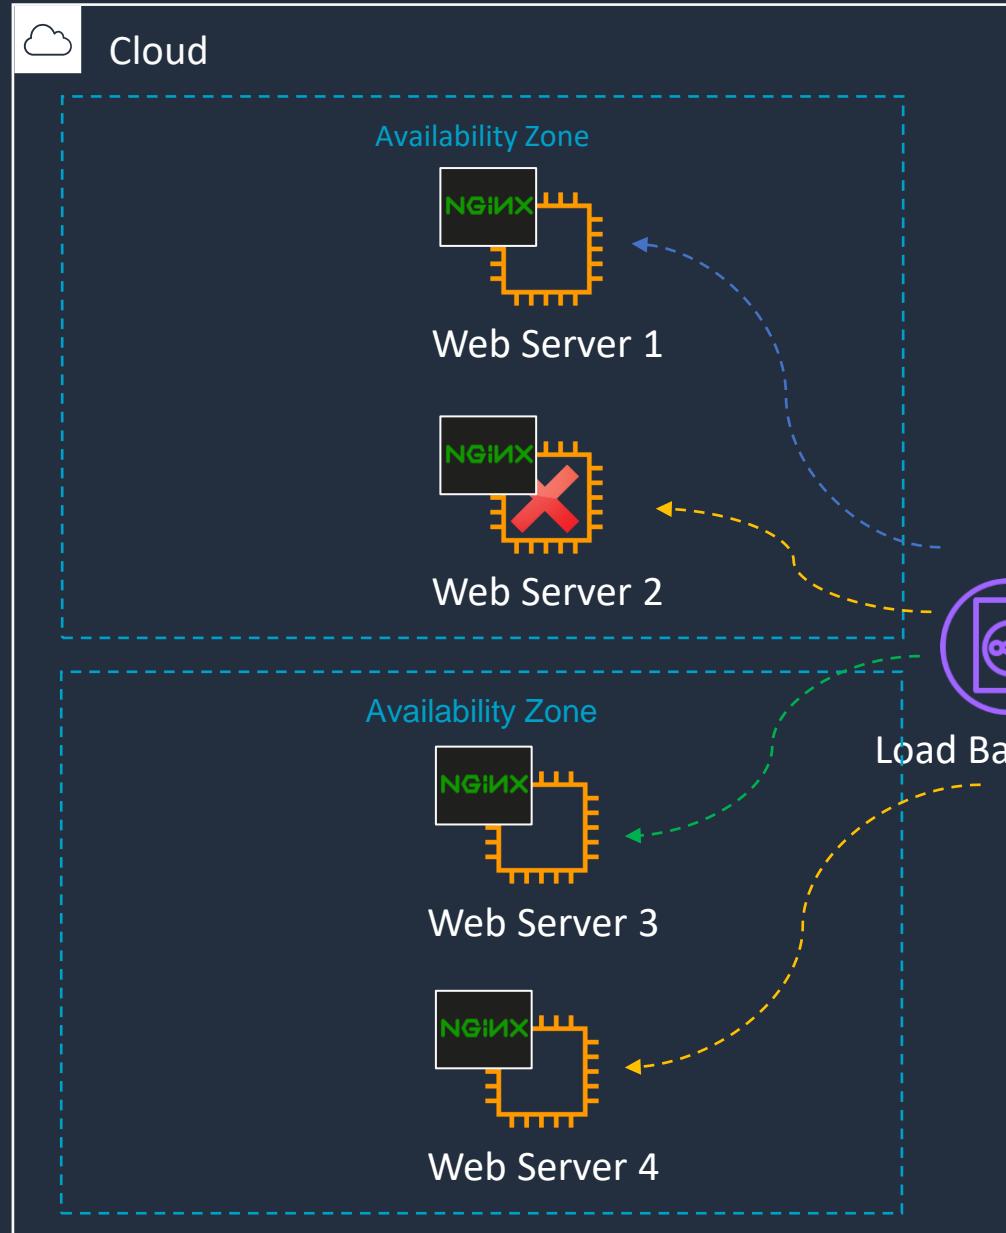
---

Redundant components  
allow the system to  
continue to operate



The system may fail if there  
is no built-in redundancy

# High Availability and Fault Tolerance

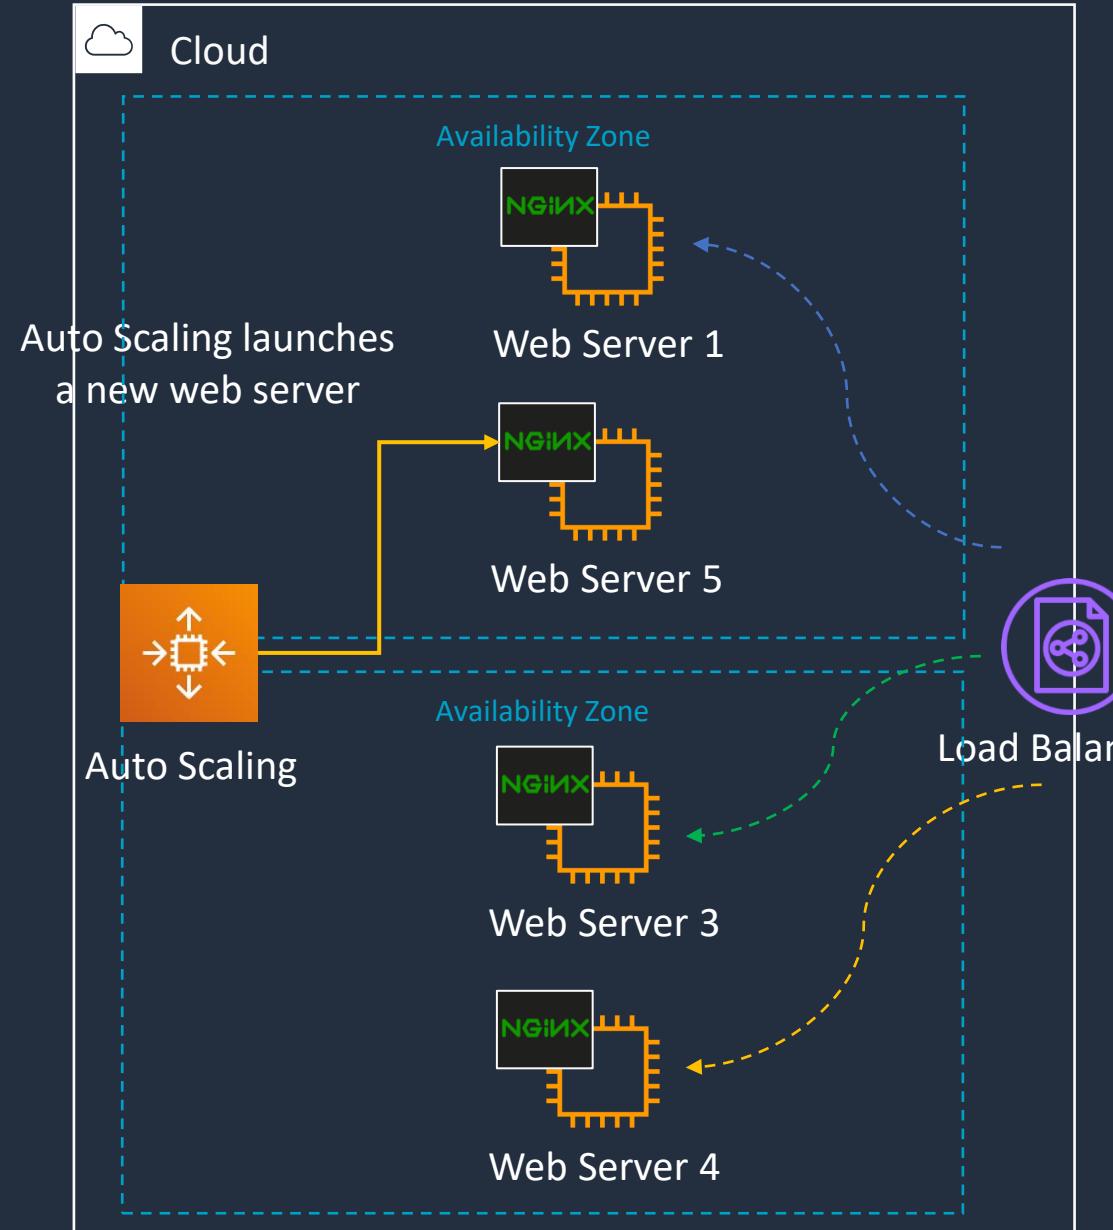


Think of an  
**availability zone** as a  
separate data center





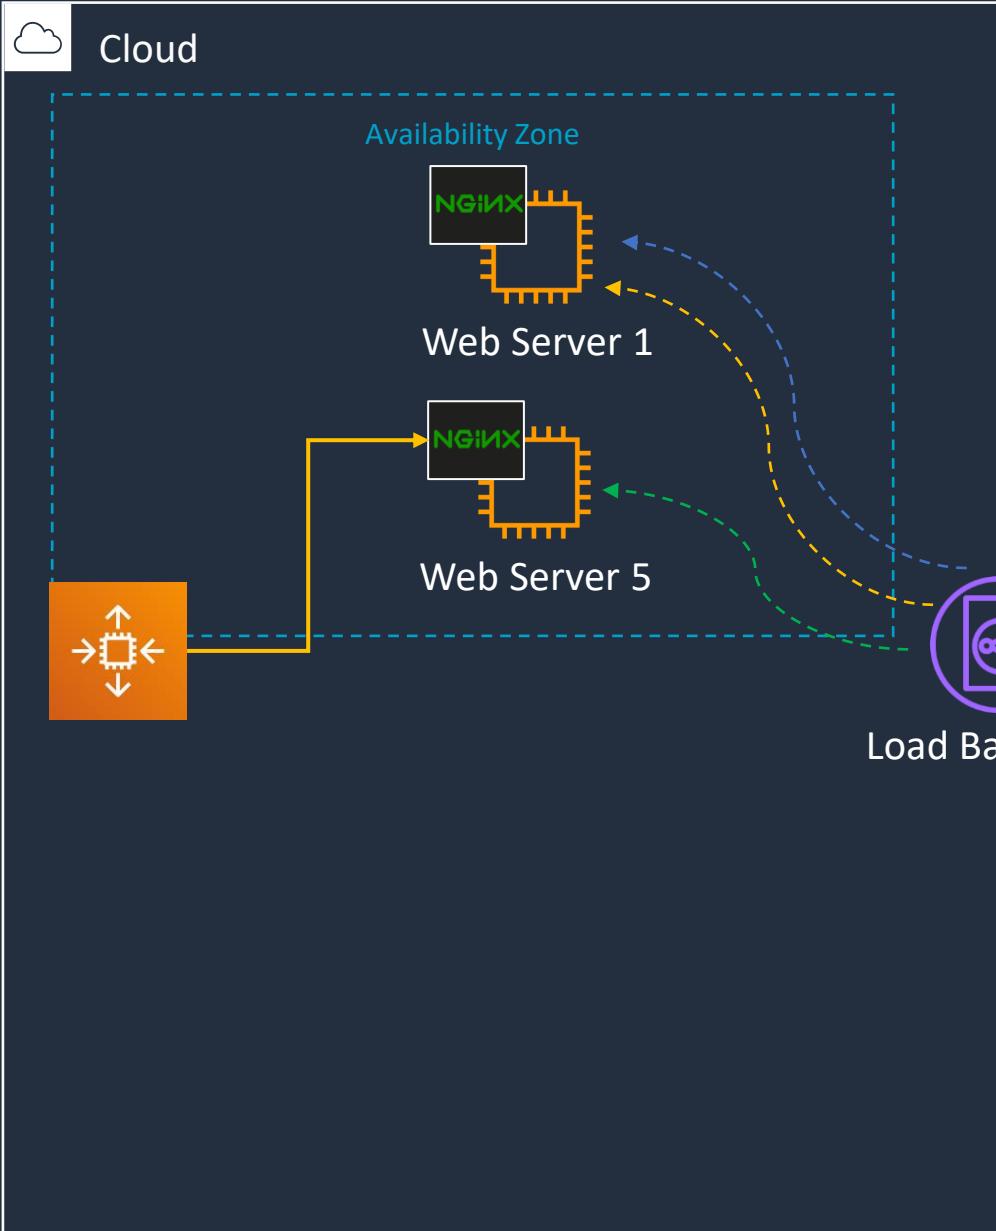
# High Availability and Fault Tolerance



Think of an **availability zone** as a separate data center



# High Availability and Fault Tolerance



Think of an  
**availability zone** as a  
separate data center



User 1



User 2



User 3

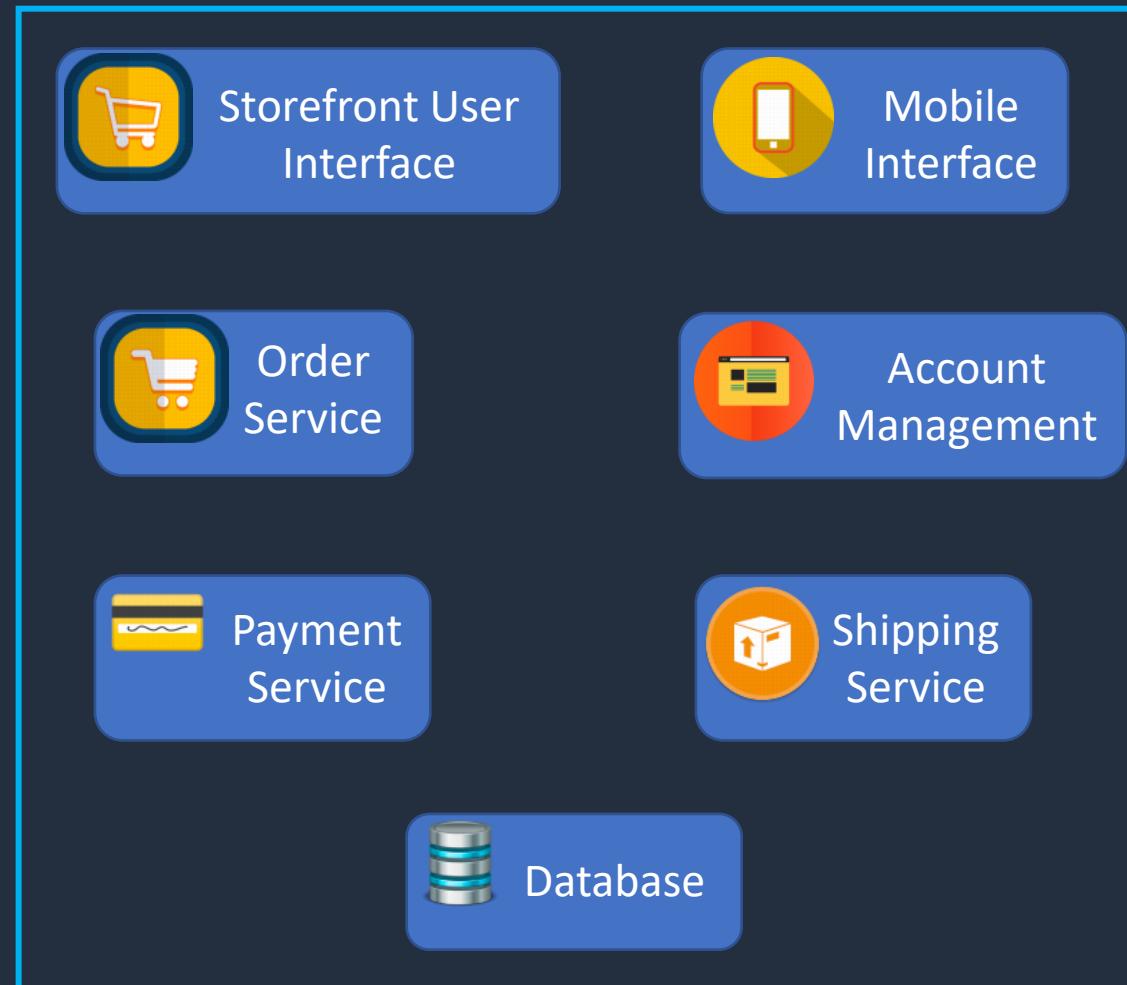
# Monolithic and Microservices Architectures





# Monolithic Application

---





# Monolithic Application

---

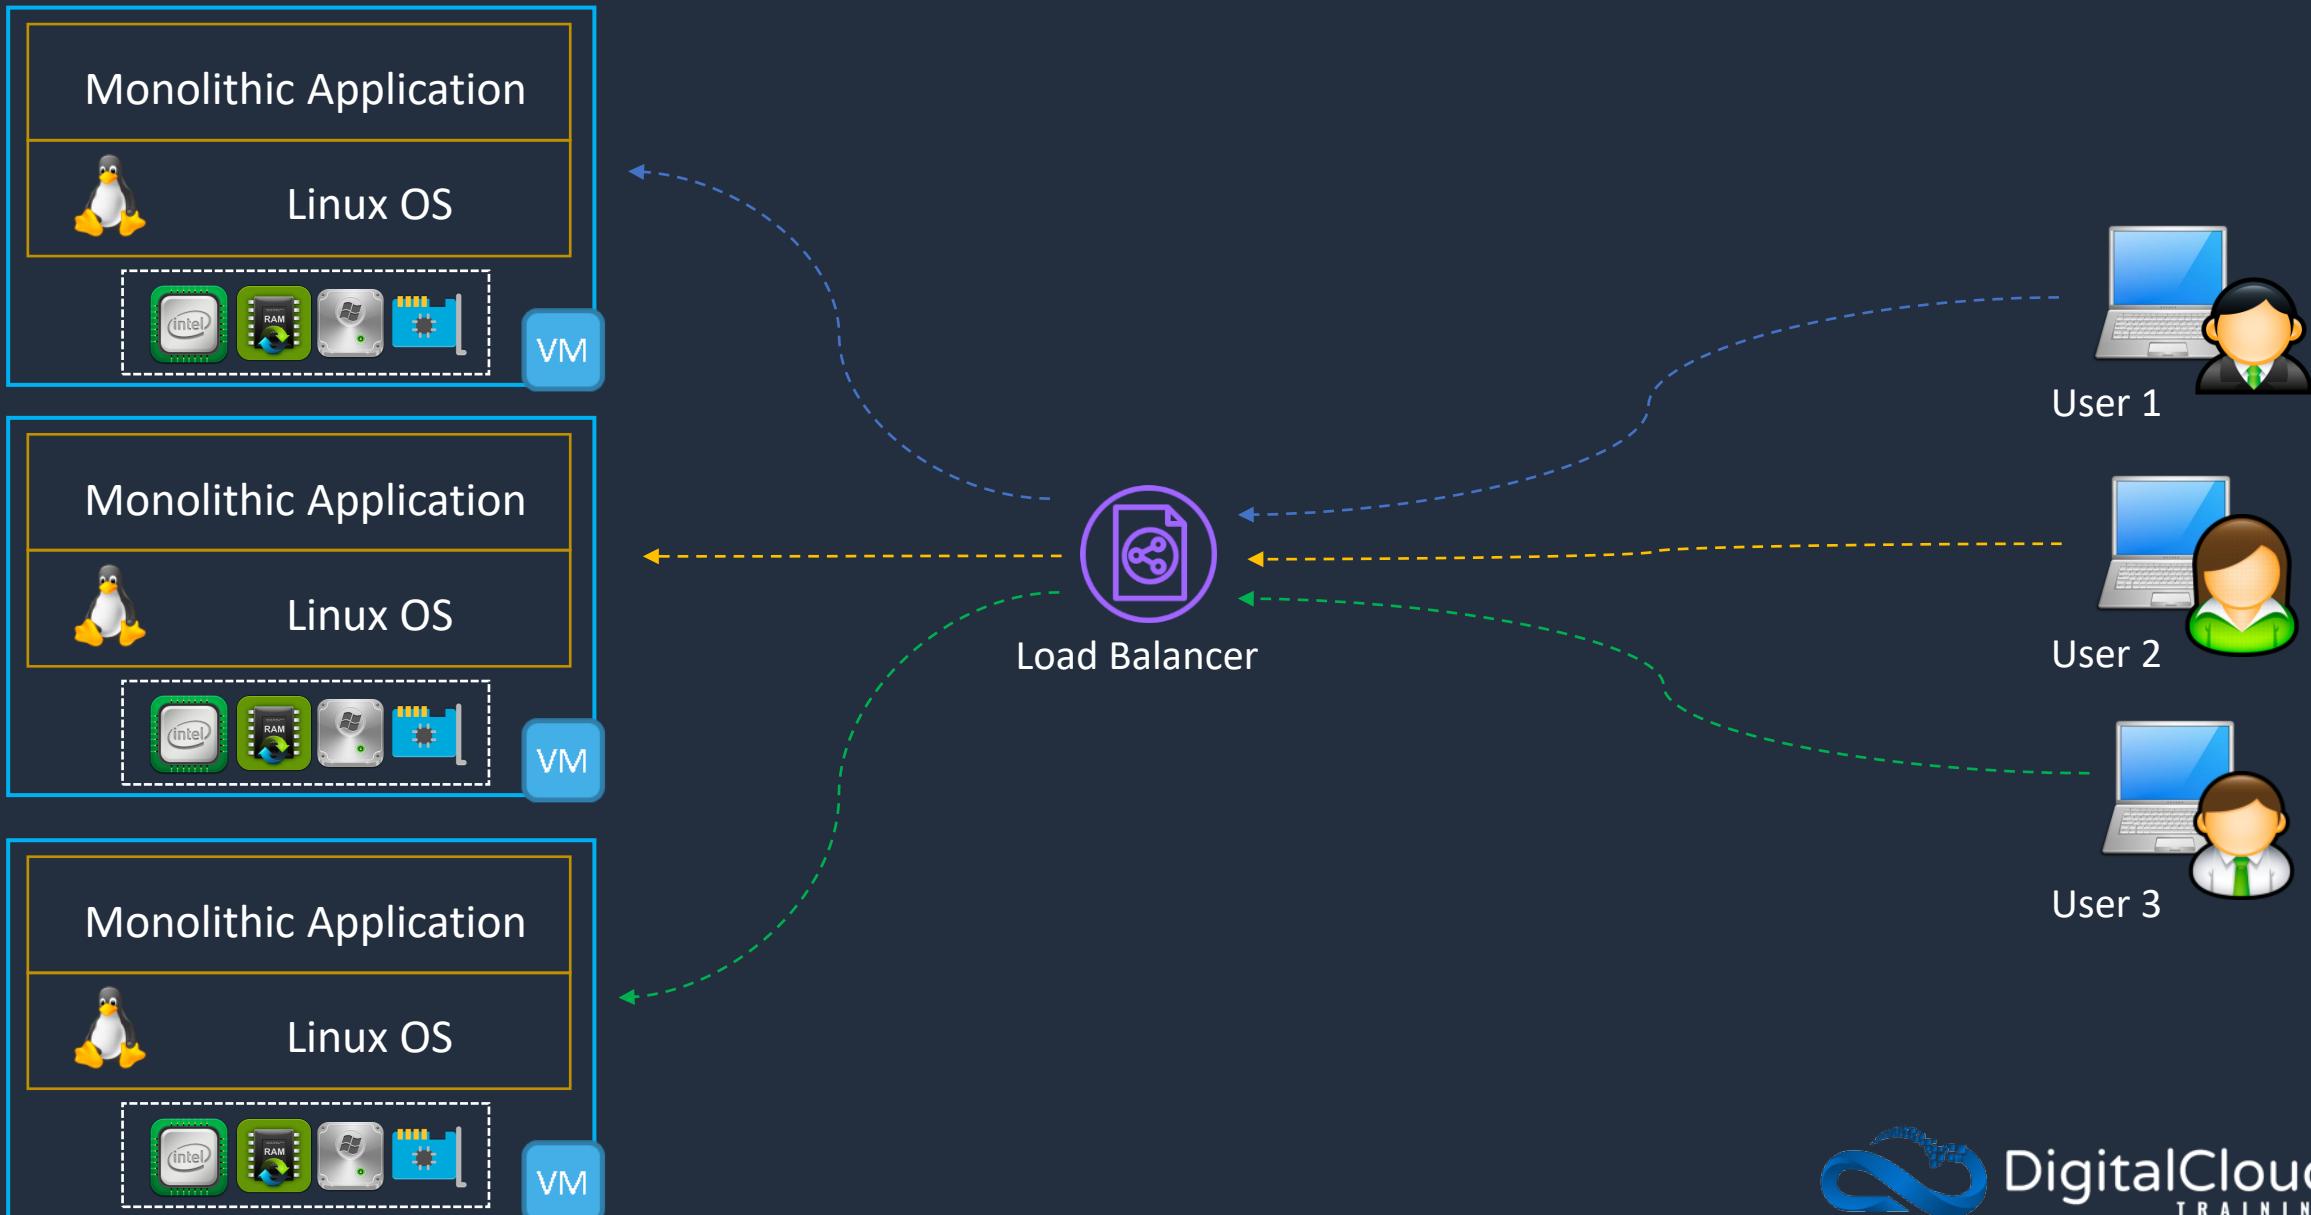
Updates to, or failures  
of, any single  
component can take  
down the whole  
application



The user interface,  
business logic, and  
data access layer are  
combined on a single  
platform



# Monolithic Application





# Monolithic Application

---

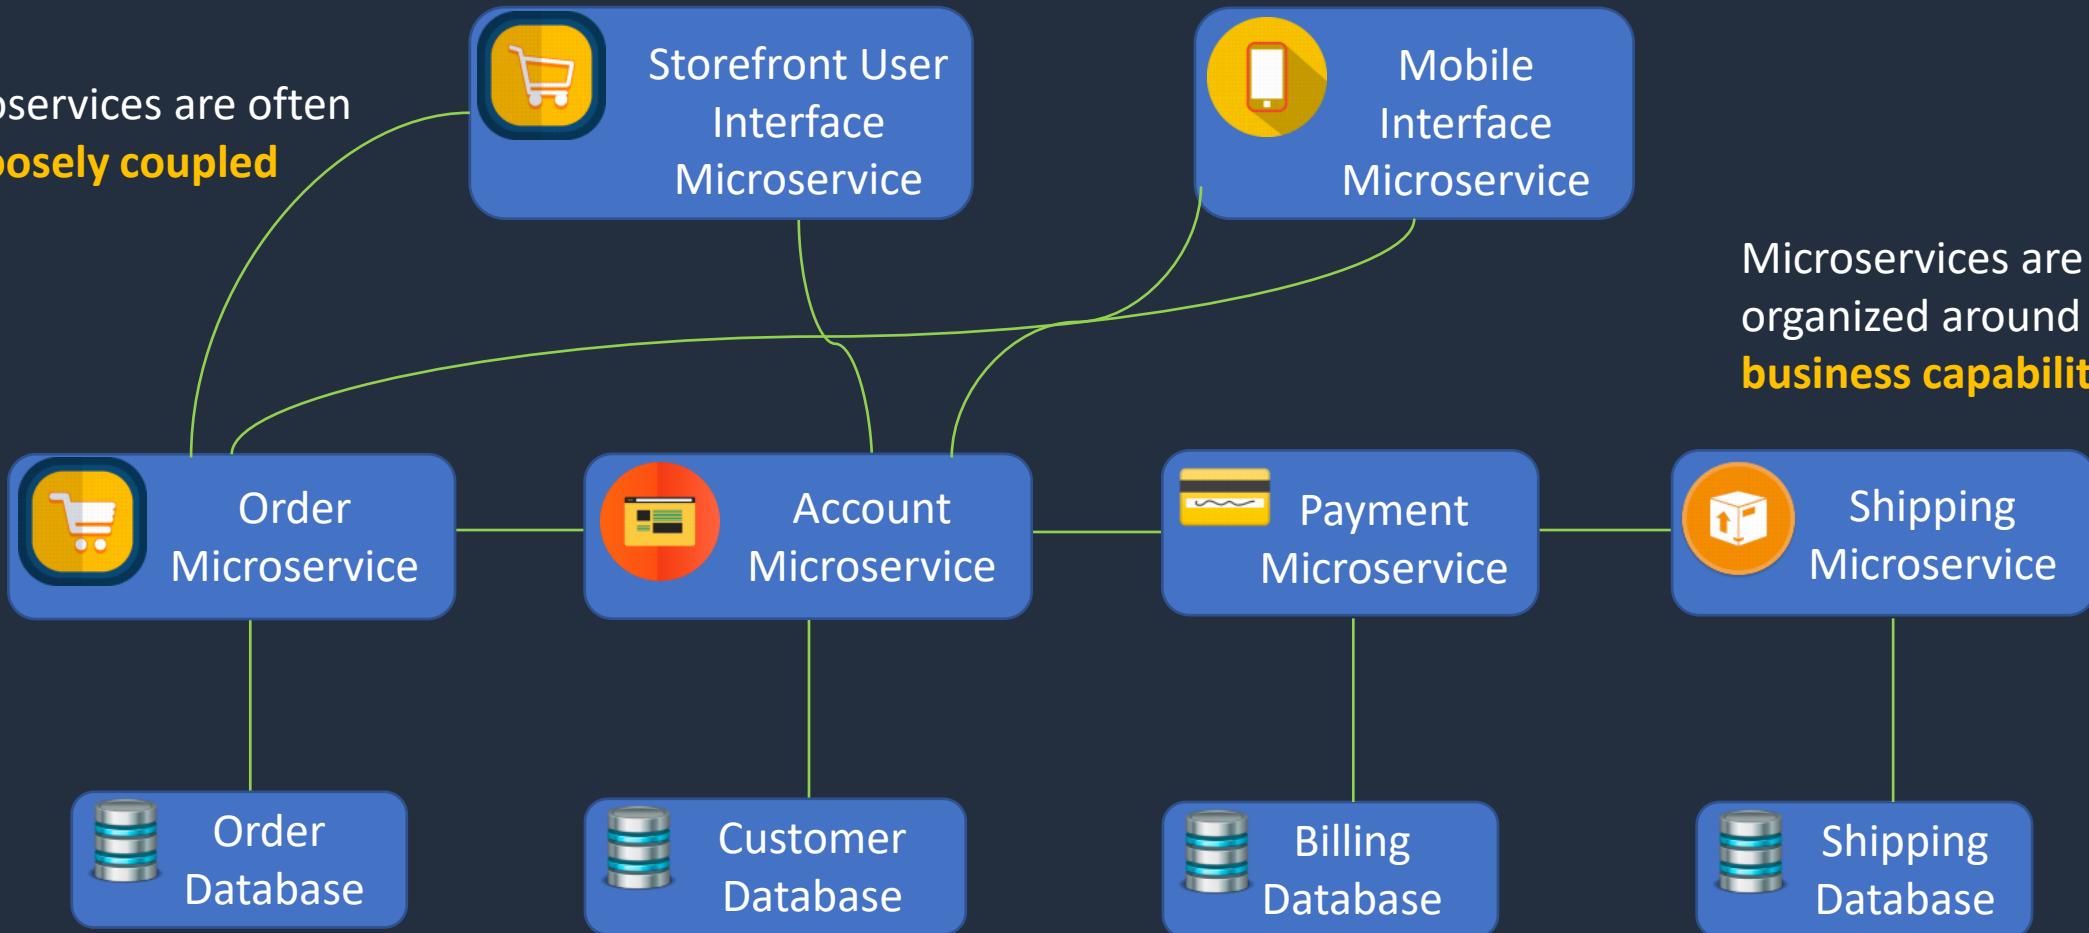




# Microservices Architecture

A **microservice** is an independently deployable unit of code

Microservices are often  
**loosely coupled**





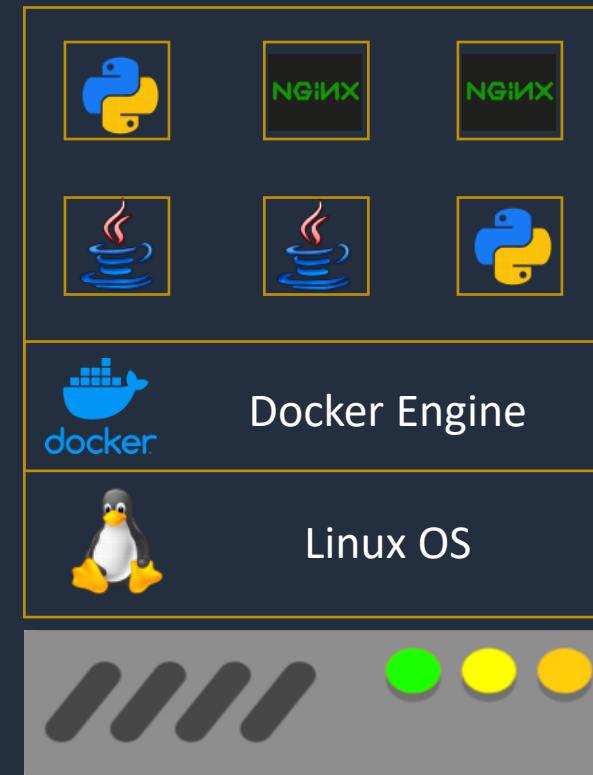
# Microservices using Docker Containers

---

Order  
Microservice

Storefront User  
Interface  
Microservice

Shipping  
Microservice



Server



# Microservices using Docker Containers

Microservices can also be spread across hosts

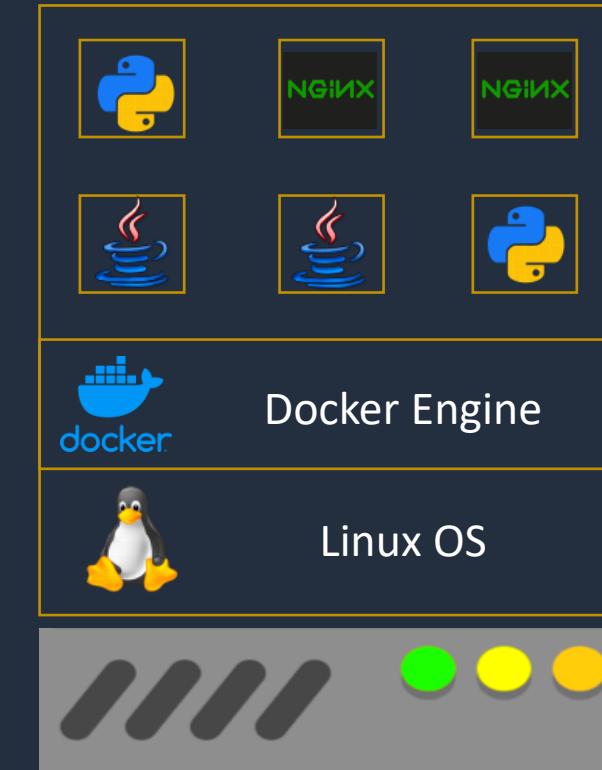
Many instances of each microservice can run on each host



Server



Server



Server



# Microservices: Attributes and Benefits

---

Microservices Attribute	Microservices Benefit
Use of Application Programming Interfaces (APIs)	Easier integrations between application components; assists with loose coupling
Independently deployable blocks of code	Can be scaled and maintained independently
Business-oriented architecture	Development organized around business capabilities; teams may be cross-functional and services may be reused
Flexible use of technologies	Each microservice can be written using different technologies (e.g. programming languages)
Speed and agility	Fast to deploy and update. Easy to include high availability and fault tolerance for each microservice

# SECTION 5

## AWS Access Control and Networking

# Amazon Web Services Overview





# Amazon Web Services (AWS)

Charge for  
services based  
on usage

27 Regions  
around the  
world

Subsidiary  
of Amazon

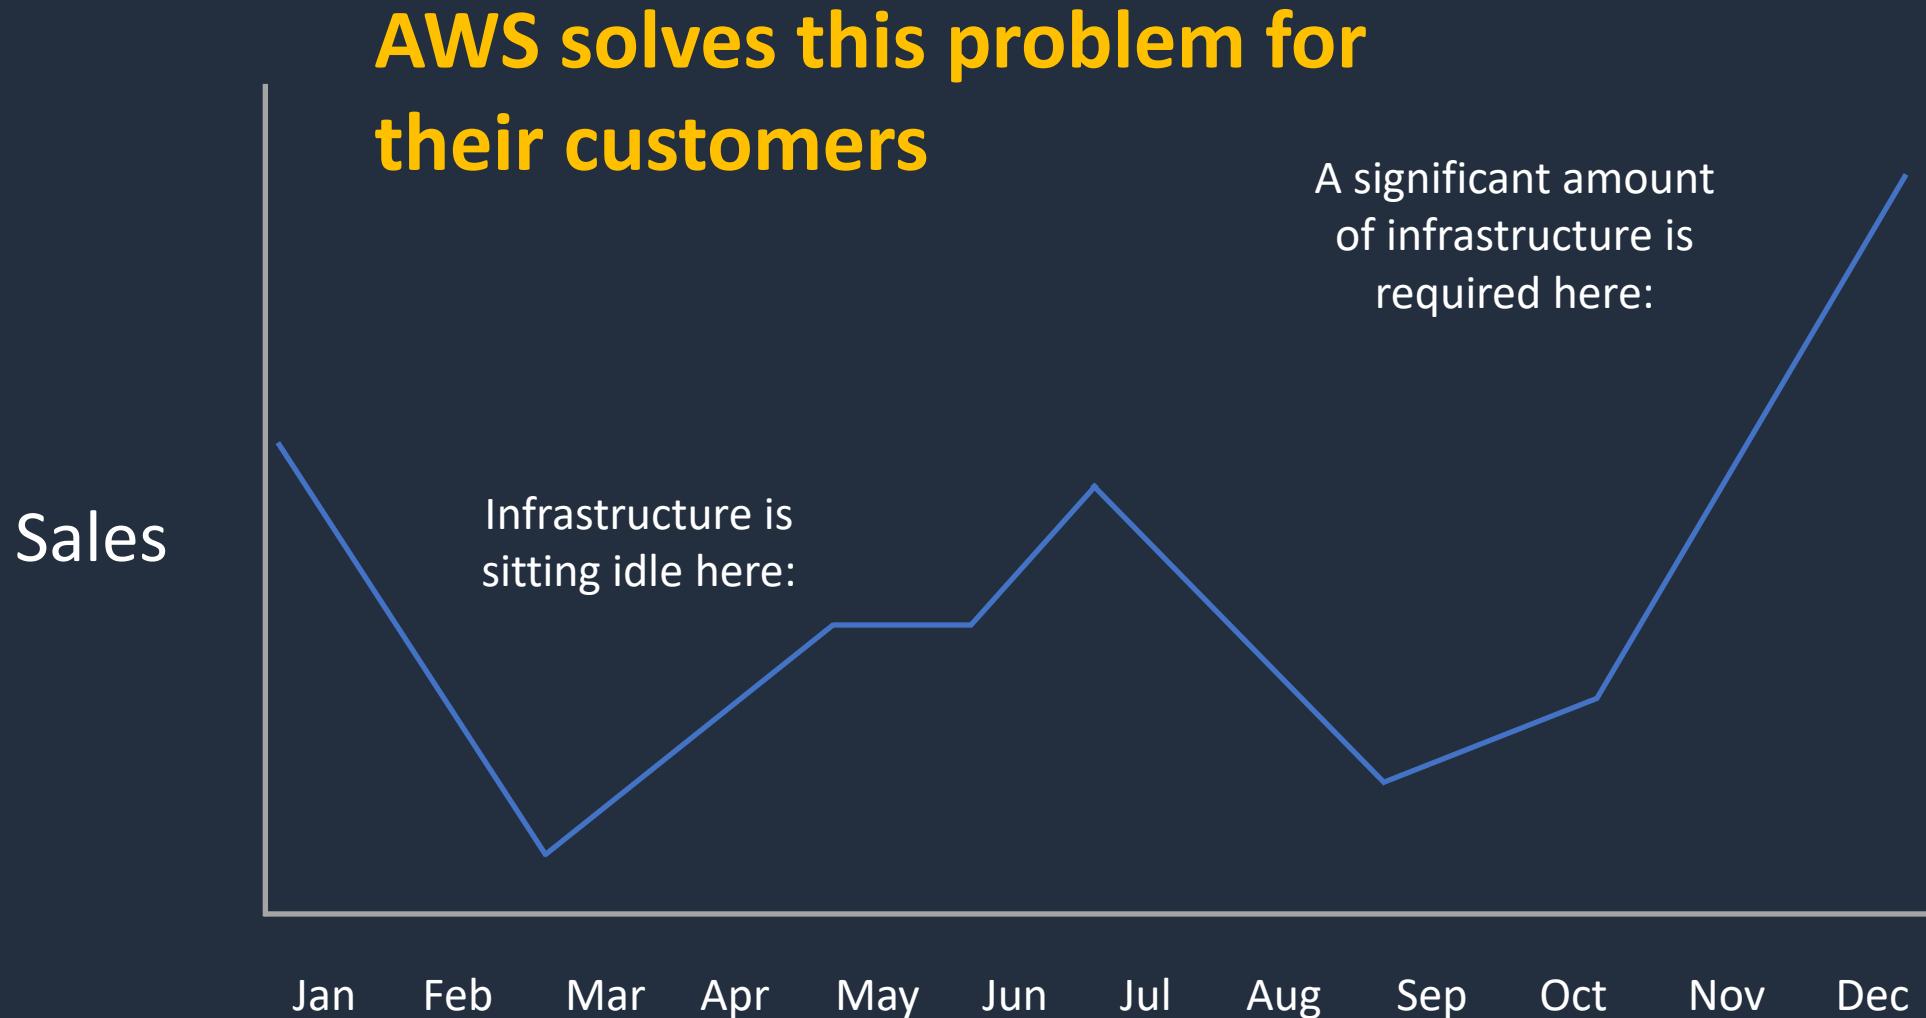


Hyperscale  
Public Cloud  
Provider

Services are  
offered on-  
demand



# Overview of AWS





# AWS Service Categories (a few examples)



Compute



Storage



Database



Analytics



Machine Learning



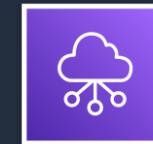
Many more categories  
and over **200** services!



Internet of Things



Media Services



Networking



End User Computing



# Fundamentals of AWS Pricing

## Compute



Amount of resources such as CPU and RAM and duration

## Storage



Quantity of data stored

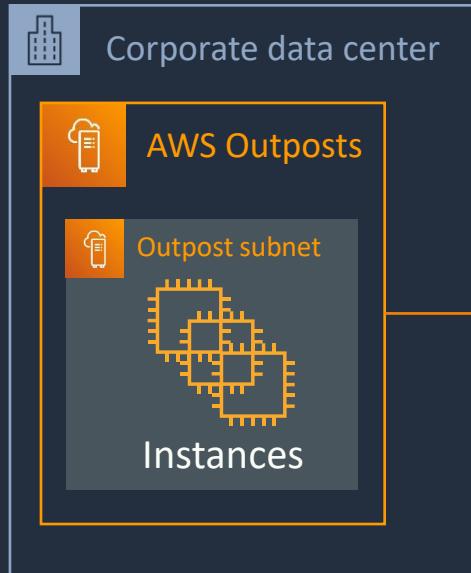
## Outbound Data Transfer



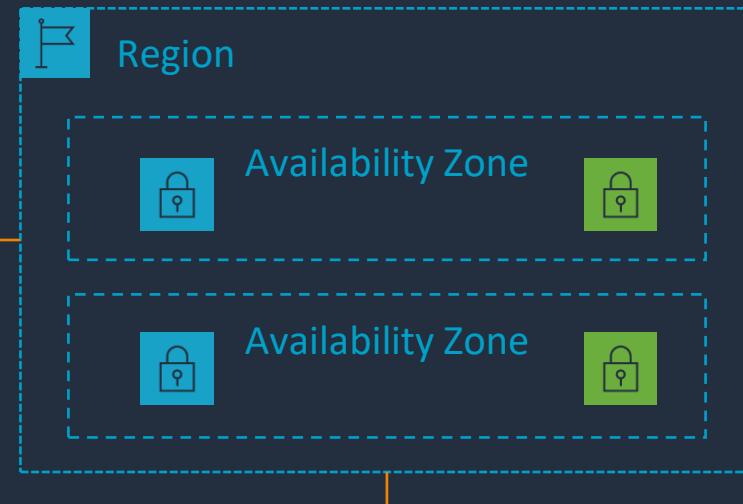
Quantity of data that is transferred out from all services



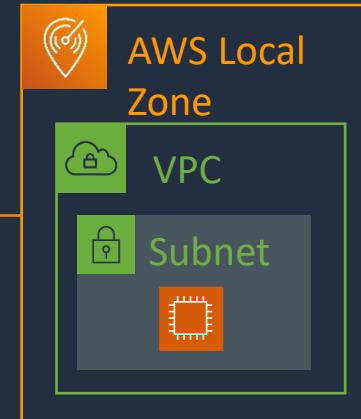
# AWS Global Infrastructure



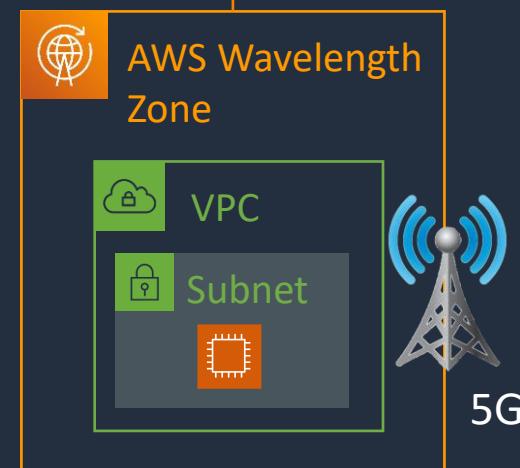
- Use cases:
- Use AWS services on-premises
  - Use AWS APIs on-premises



- Use cases:
- Single-digit ms latency to end users / apps
  - Live video, ML, AR/VR



- Use cases:
- Single-digit ms latency to mobile devices/users
  - Live video, ML, AR/VR



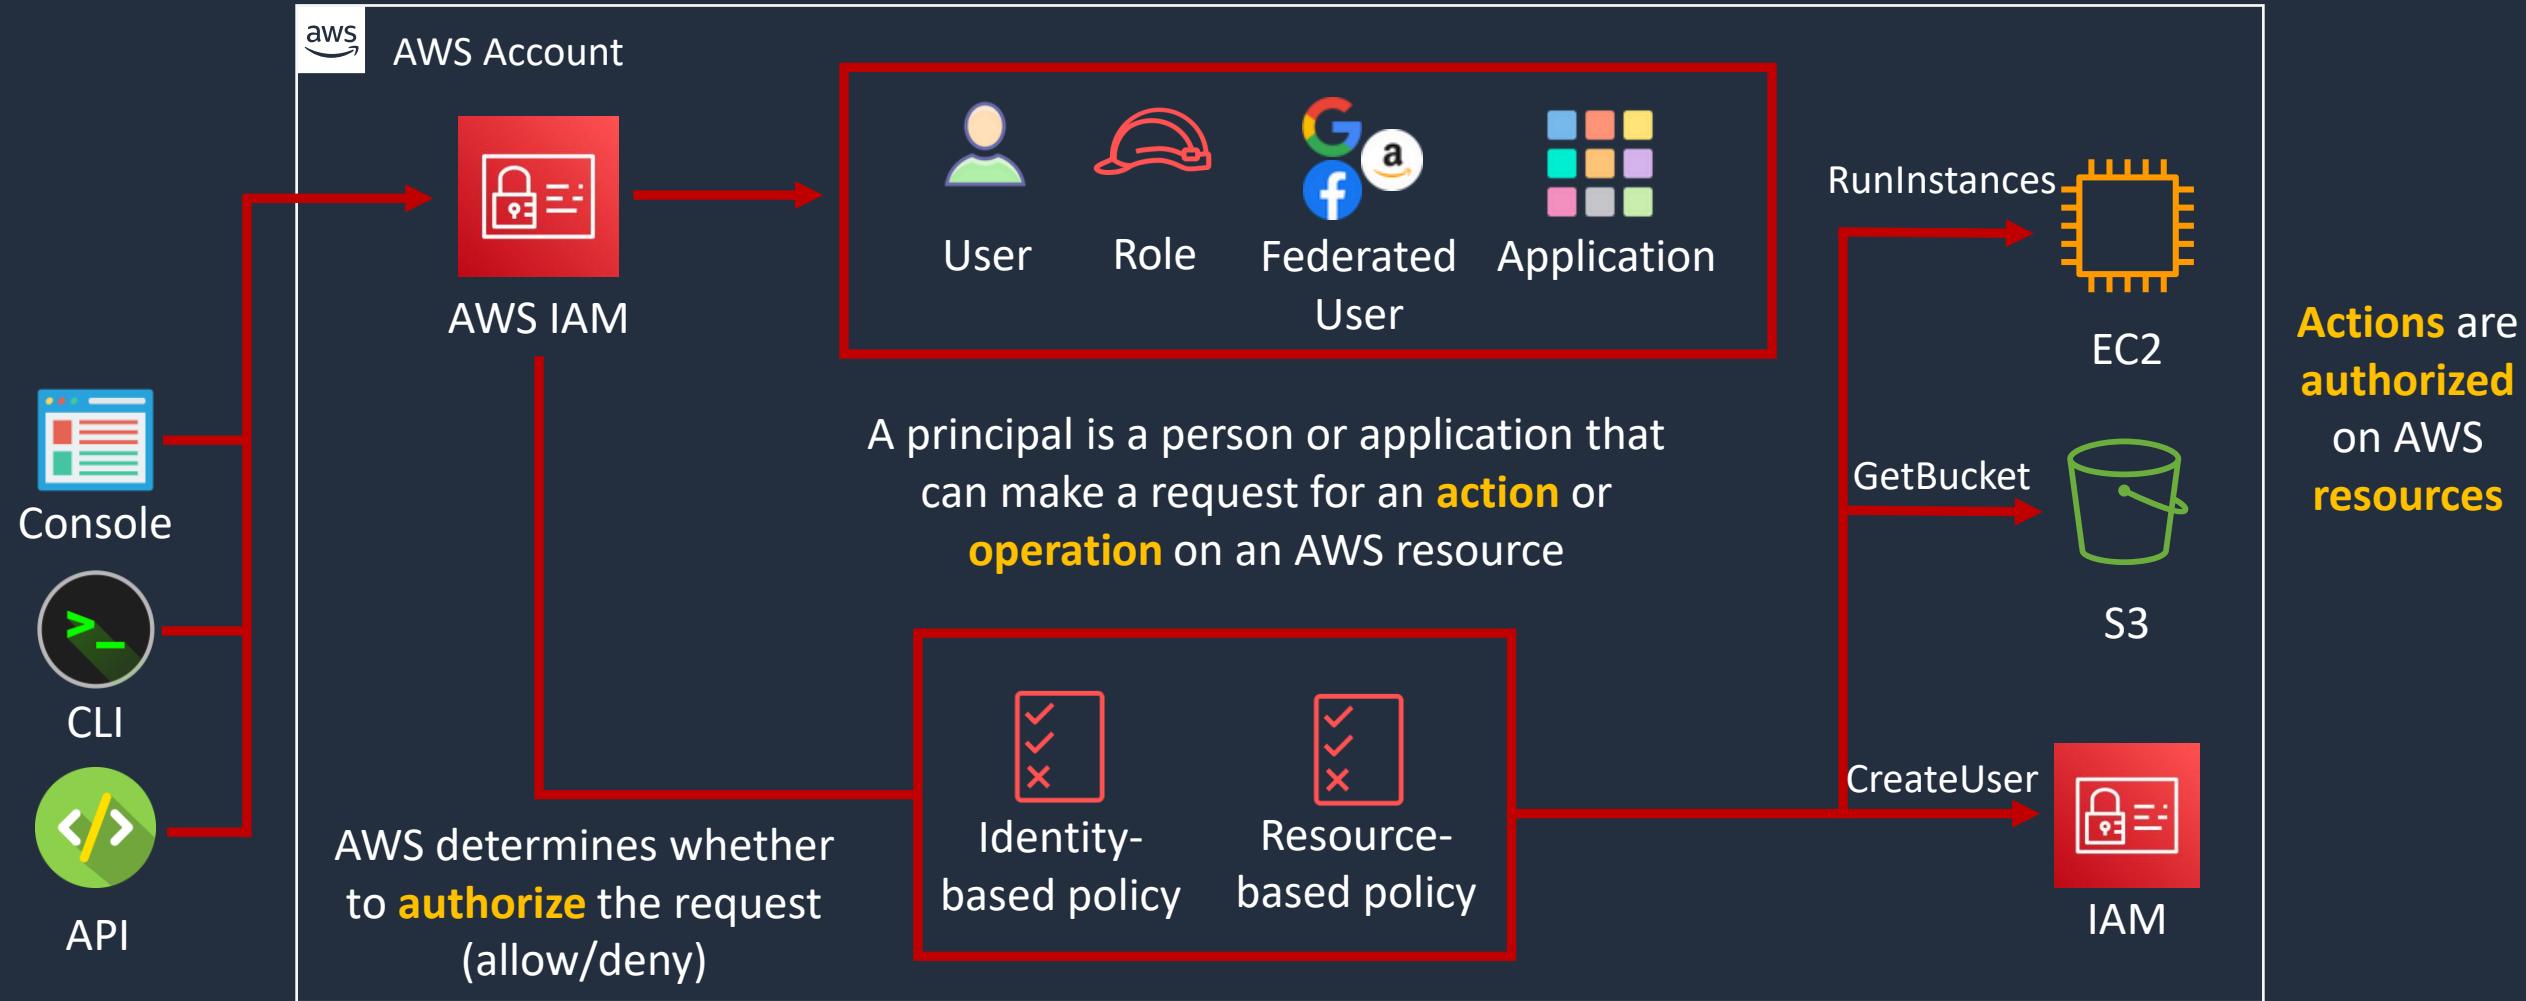
# IAM Overview





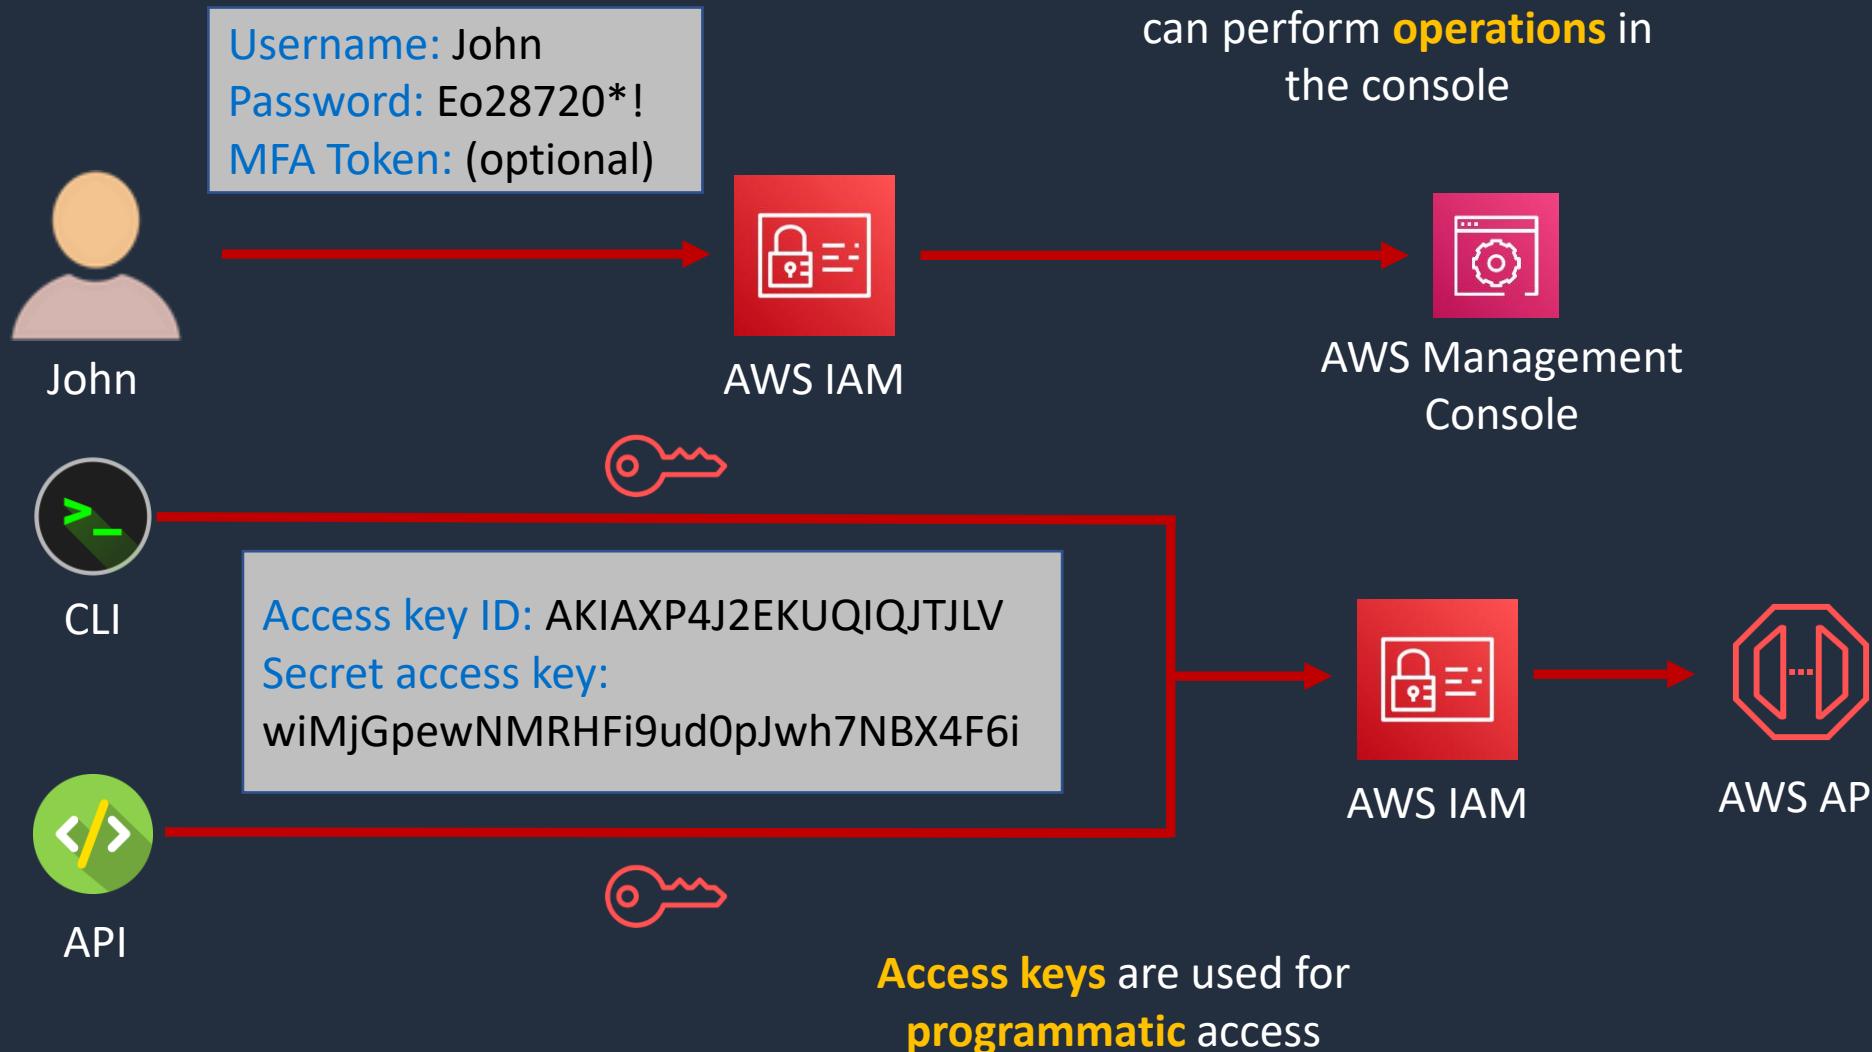
# AWS Identity and Access Management (IAM)

IAM Principals must be **authenticated** to send requests (with a few exceptions)





# IAM Authentication Methods



# Create an IAM User and Group



# Amazon Virtual Private Cloud (VPC)

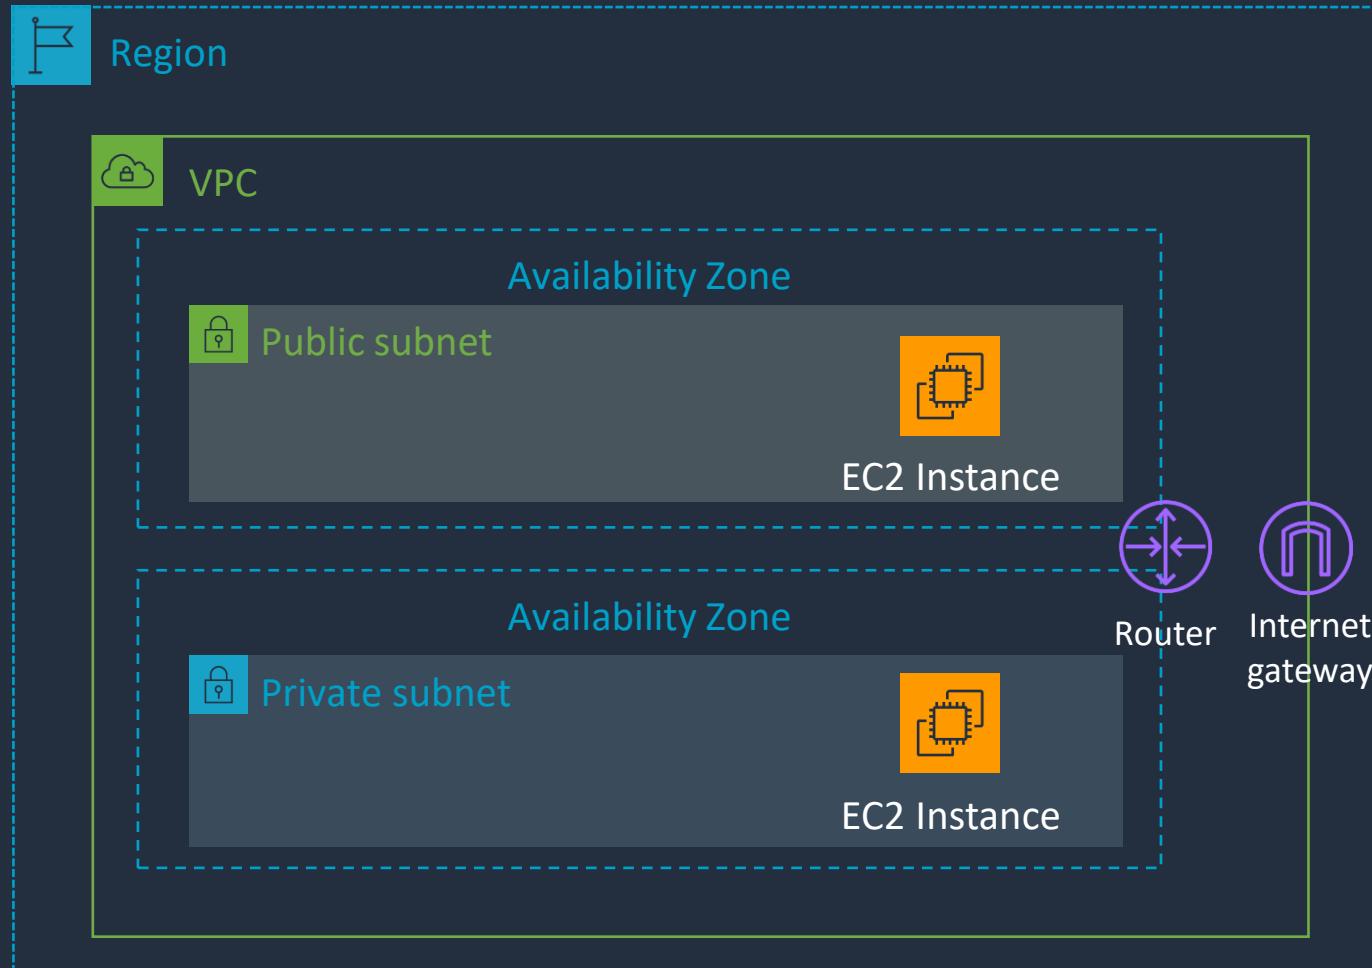




# Amazon VPC

A **VPC** is a logically isolated portion of the AWS cloud within a region

**Subnets** are created within **AZs**



You can launch **EC2 instances** into your VPC subnets

## Main Route Table

Destination	Target
10.0.0.0/16	Local
0.0.0.0/0	igw-id

The **route table** is used to configure the VPC router

An **Internet Gateway** is used to connect to the Internet



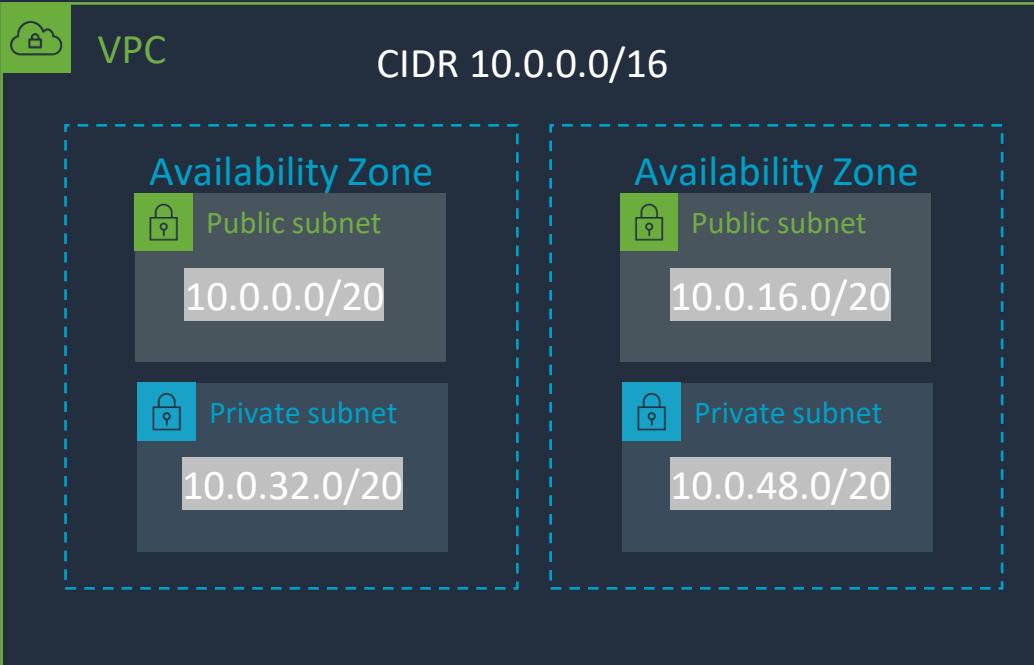
# Amazon VPC

Each **VPC** has a different block of IP addresses

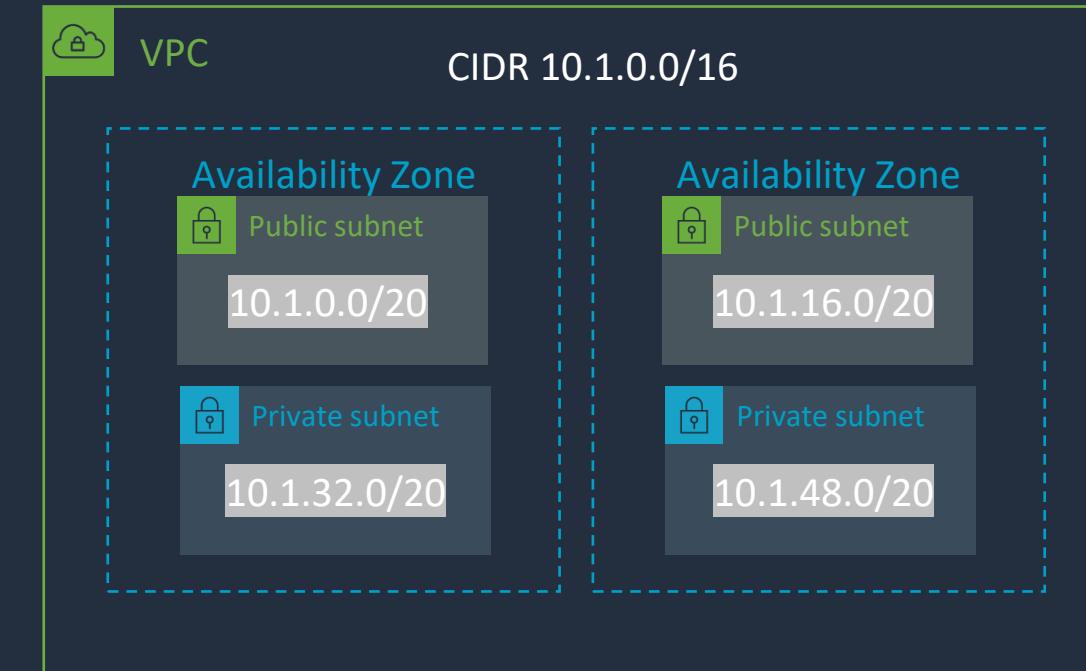
**CIDR** stands for Classless Interdomain Routing



Region



Each subnet has a block of **IP addresses** from the CIDR block

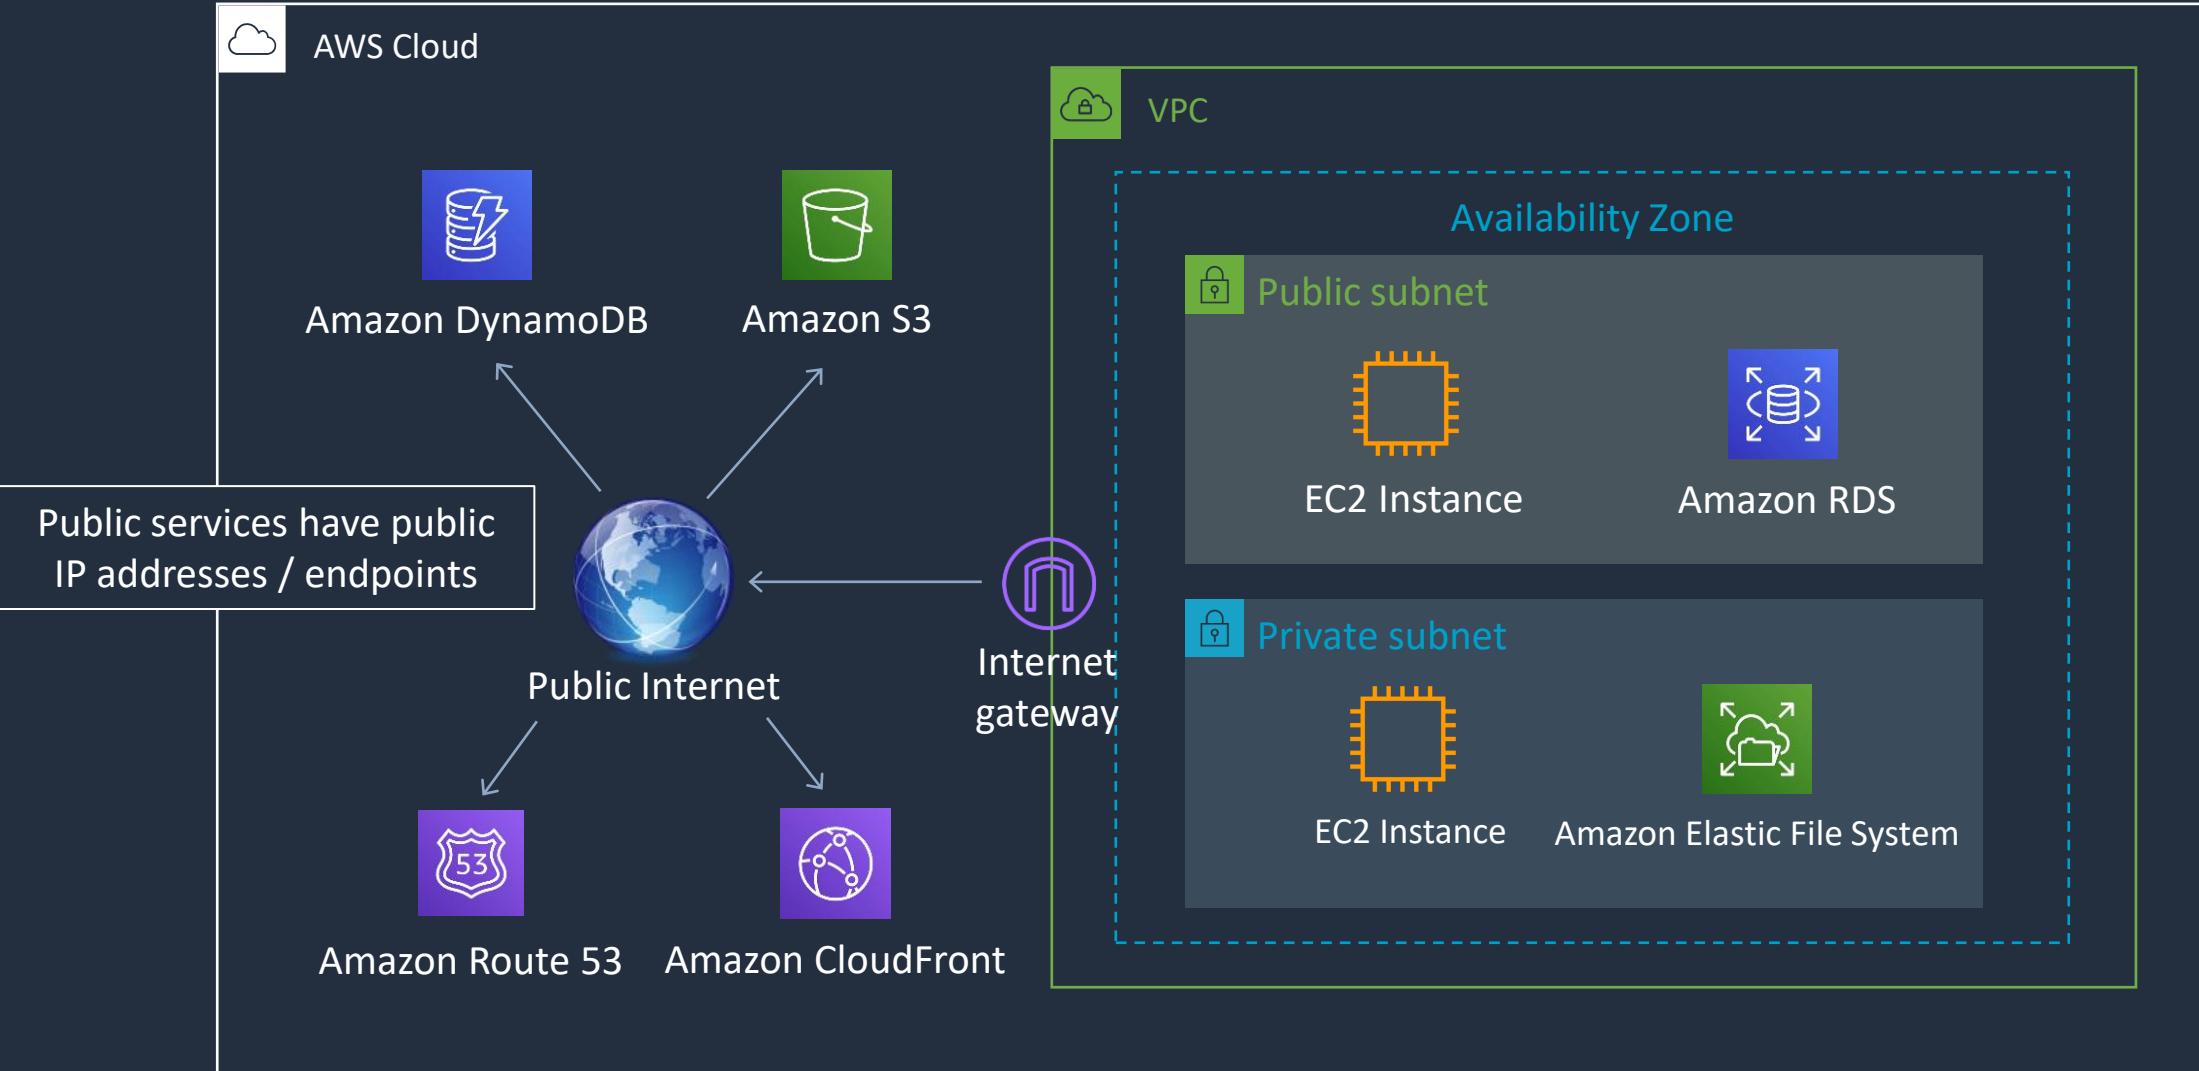


You can create **multiple VPCs** within each region



# AWS Public and Private Services

Private services can have public IP addresses but exist within the VPC

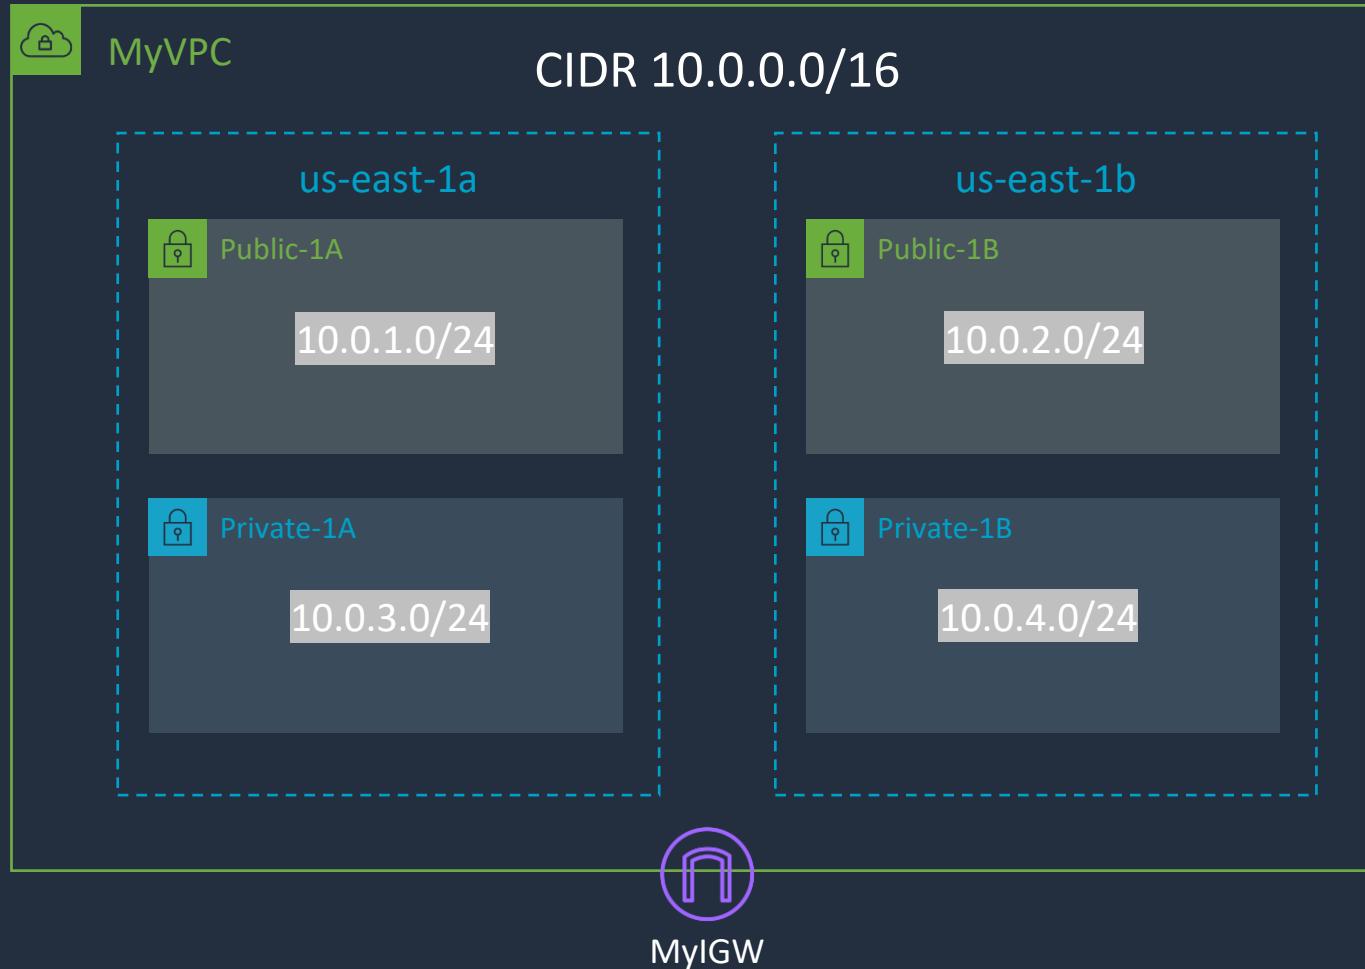


# Create a Custom VPC





# Custom VPC



## Main Route Table

Destination	Target
10.0.0.0/16	Local
0.0.0.0/0	igw-id

## Private-RT Route Table

Destination	Target
10.0.0.0/16	Local

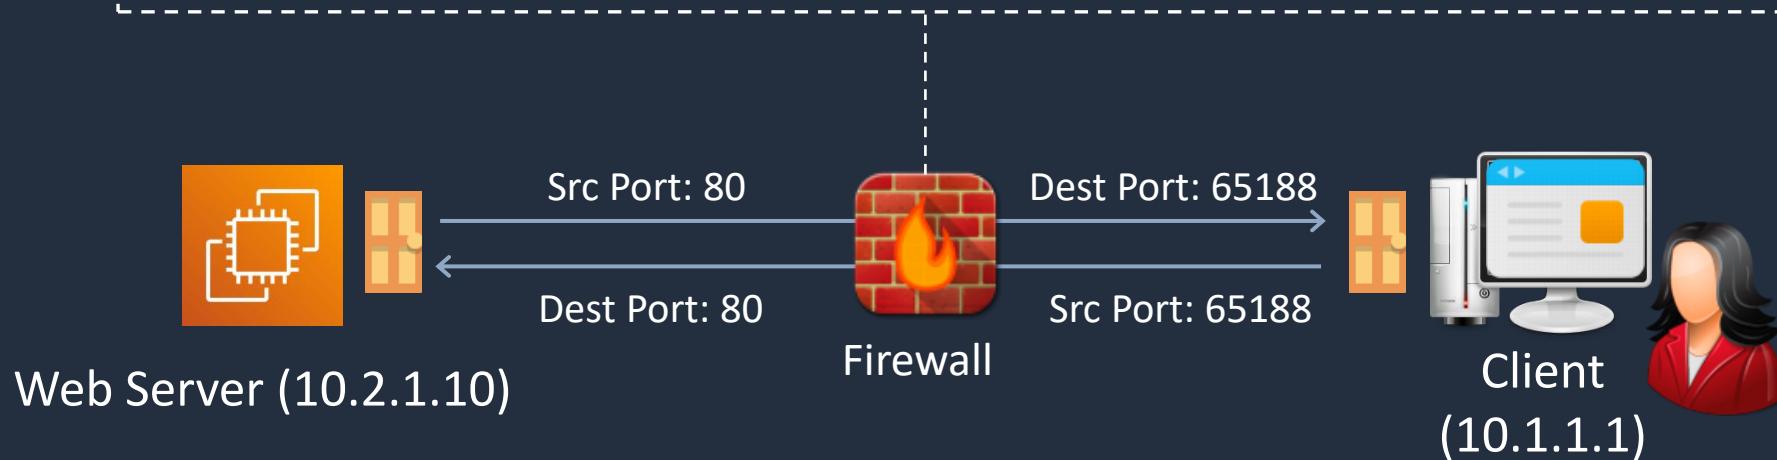
# Security Groups and Network ACLs





# Stateful vs Stateless Firewalls

PROTOCOL	SOURCE IP	DESTINATION IP	SOURCE PORT	DESTINATION PORT
HTTP	10.1.1.1	10.2.1.10	65188	80
HTTP	10.2.1.10	10.1.1.1	80	65188

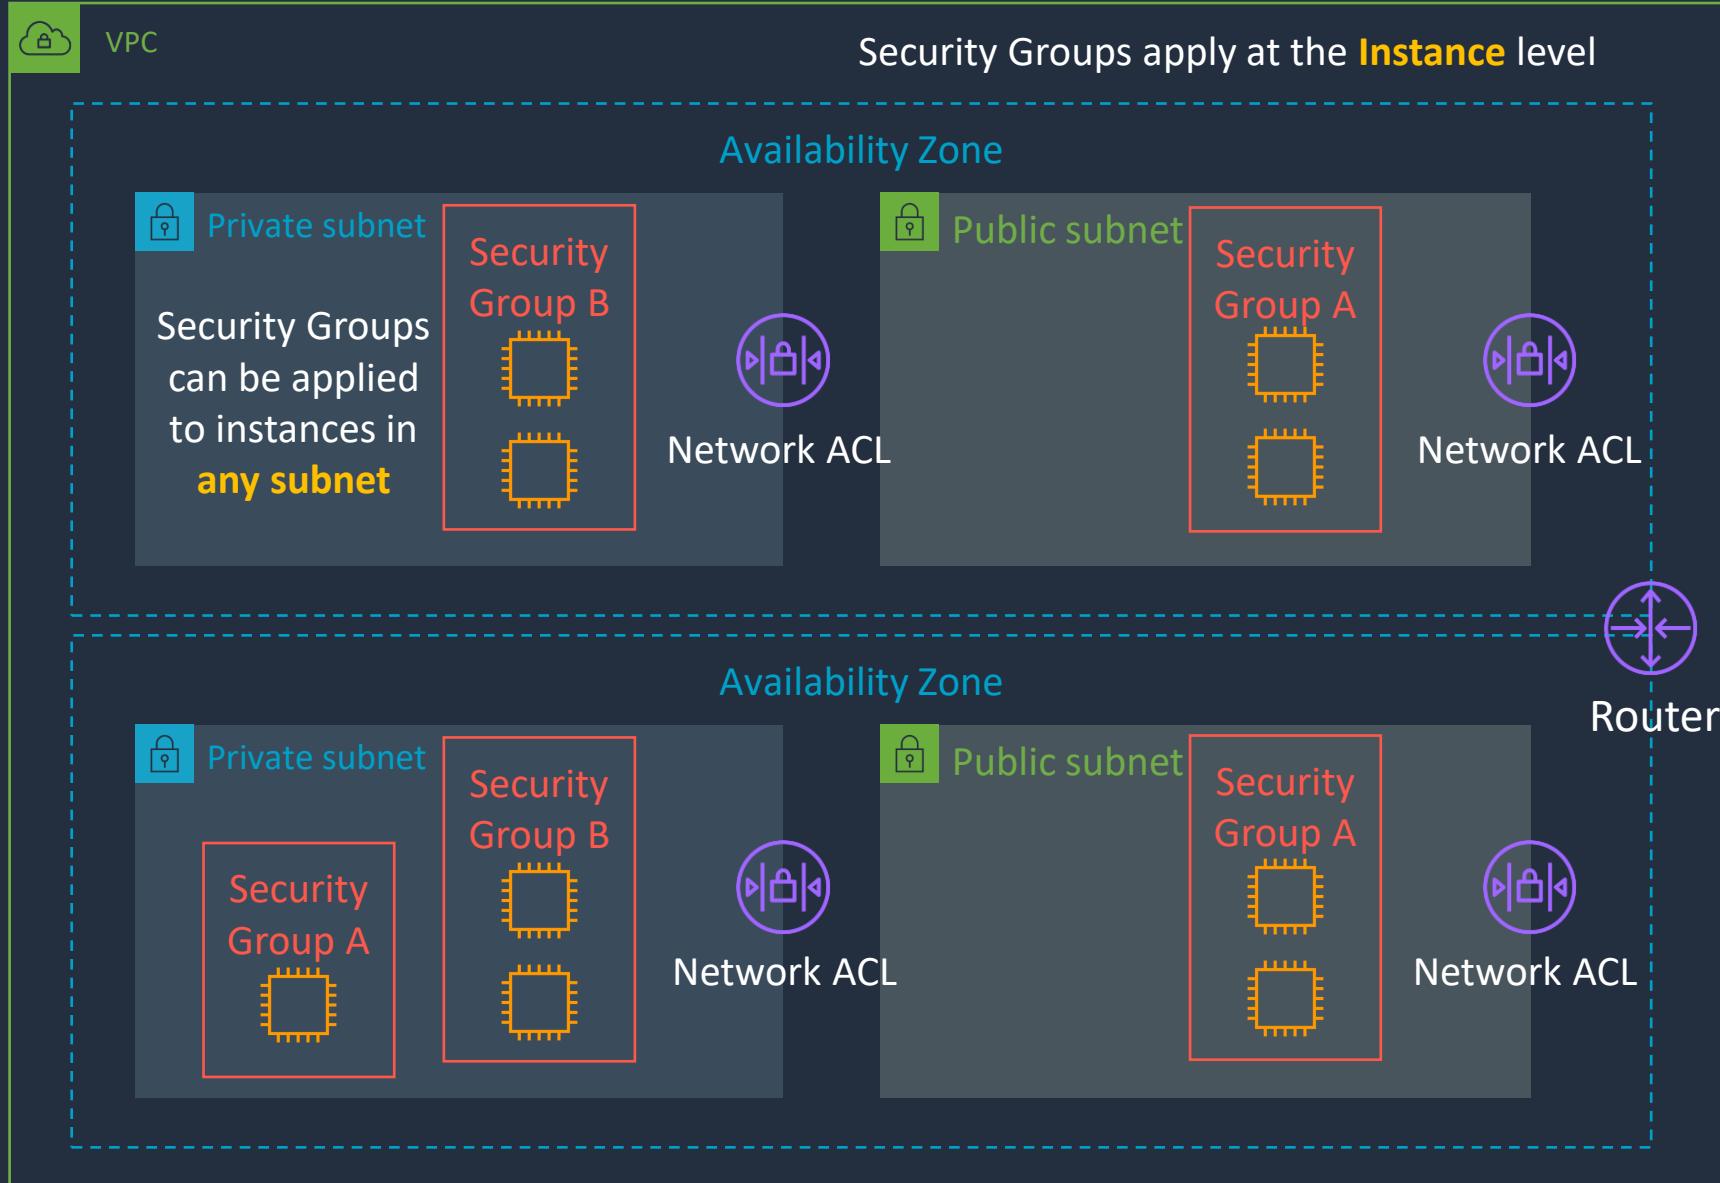


A **stateful** firewall  
allows the return  
traffic automatically

A **stateless** firewall  
checks for an allow  
rule for **both**  
connections



# Security Groups and Network ACLs





# Security Groups & Network Access Control Lists (NACLs)

---

---

Security Group	Network ACL
Operates at the instance (interface) level	Operates at the subnet level
Supports allow rules only	Supports allow and deny rules
Stateful	Stateless
Evaluates all rules	Processes rules in order
Applies to an instance only if associated with a group	Automatically applies to all instances in the subnets its associated with

# Create a Security Group



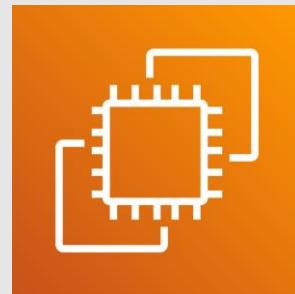
# Configure the AWS CLI



# SECTION 6

## Amazon EC2, Auto Scaling, and Load Balancing

# Amazon EC2 Overview



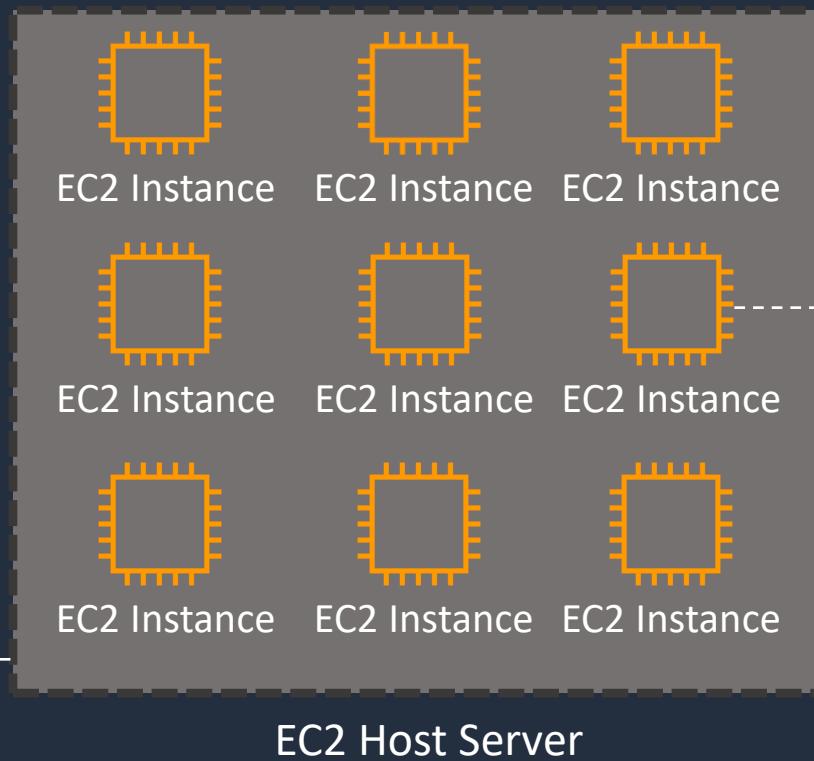
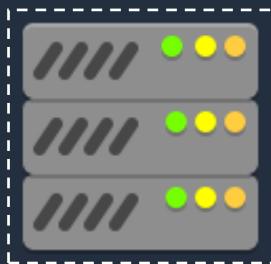


# Amazon Elastic Compute Cloud (EC2)

EC2 instances run  
Windows, Linux, or  
MacOS

An **EC2 instance** is a  
virtual server

EC2 hosts are  
**managed by AWS**



A selection of **instance types**  
come with varying combinations  
of CPU, memory, storage and  
networking



# Public, Private, and Elastic IP addresses

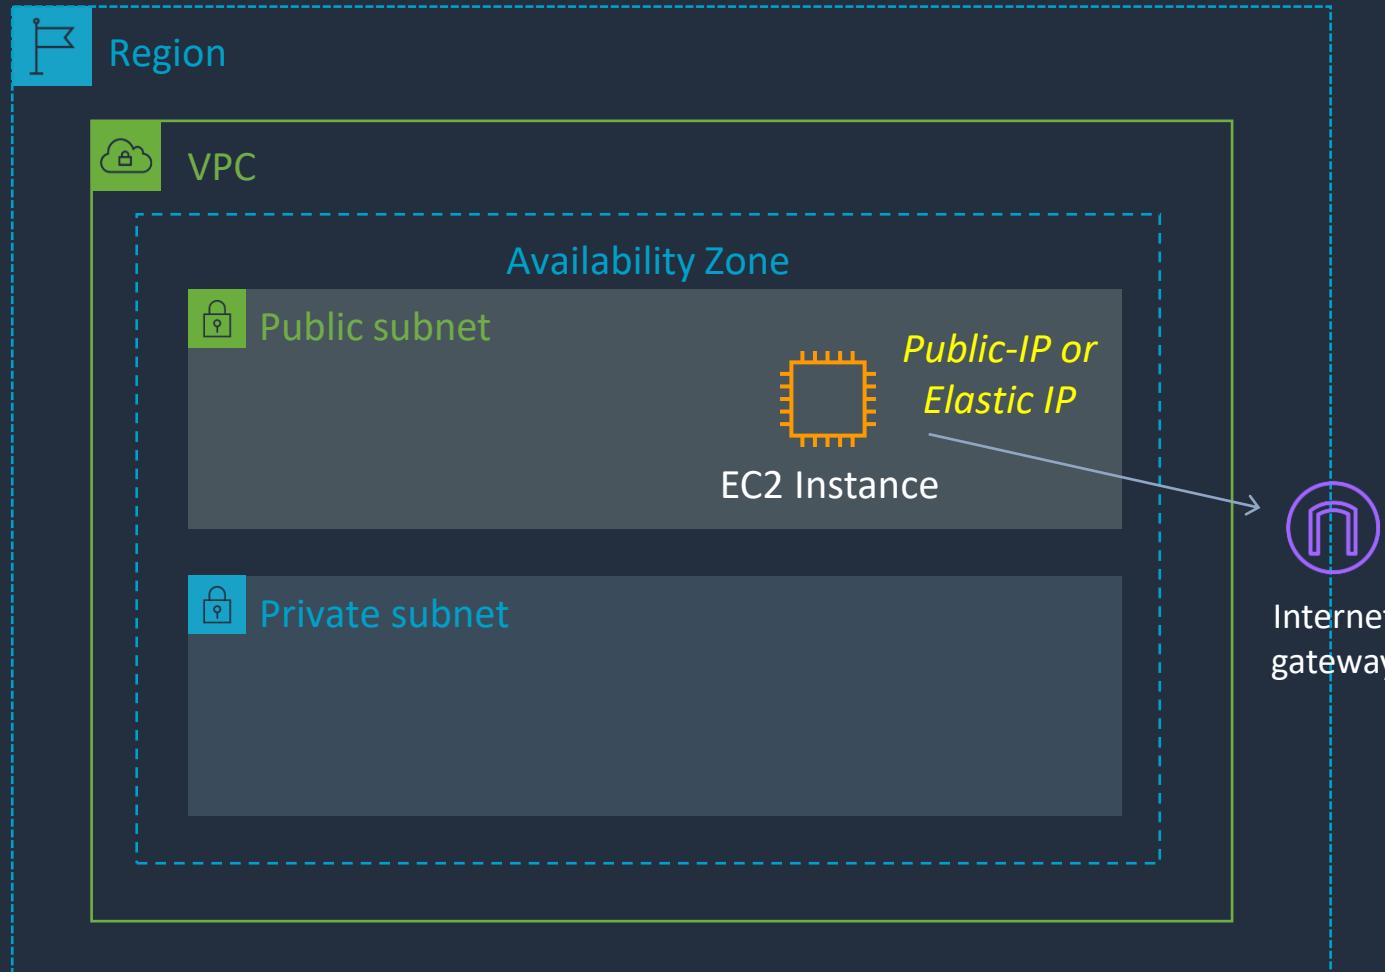
---

---

Type	Description
Public IP address	<p>Lost when the instance is stopped</p> <p>Used in Public Subnets</p> <p>No charge</p> <p>Associated with a private IP address on the instance</p> <p>Cannot be moved between instances</p>
Private IP address	<p>Retained when the instance is stopped</p> <p>Used in Public and Private Subnets</p>
Elastic IP address	<p>Static Public IP address</p> <p>You are charged if not used</p> <p>Associated with a private IP address on the instance</p> <p>Can be moved between instances and Elastic Network Adapters</p>



# Public Subnets



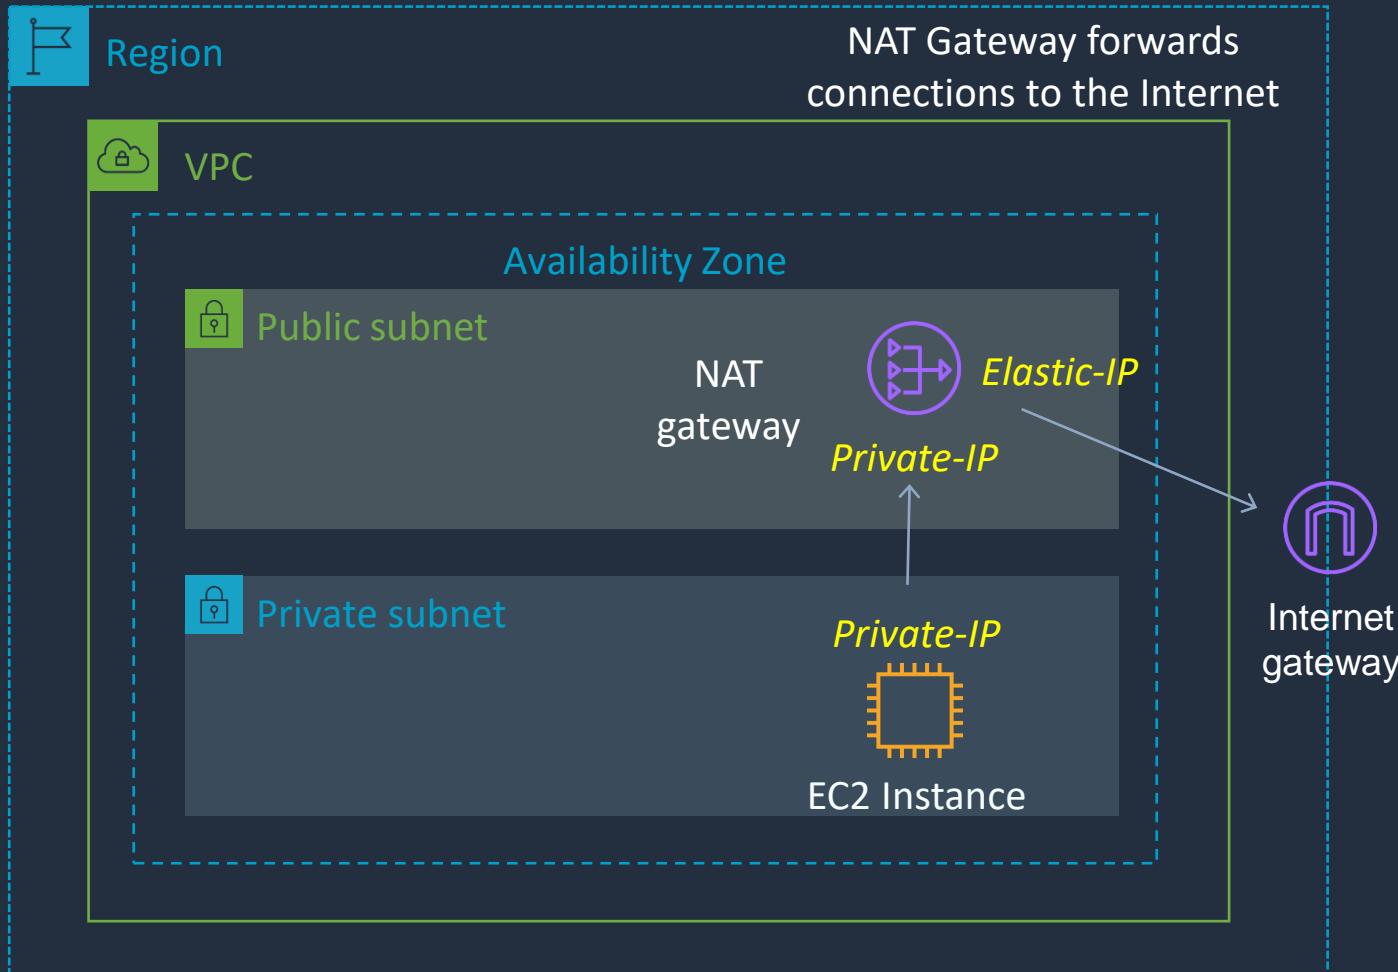
## Public Subnet Route Table

Destination	Target
172.31.0.0/16	Local
0.0.0.0/0	igw-id



# Private Subnets

NAT = Network Address Translation



## Public Subnet Route Table

Destination	Target
172.31.0.0/16	Local
0.0.0.0/0	igw-id

## Private Subnet Route Table

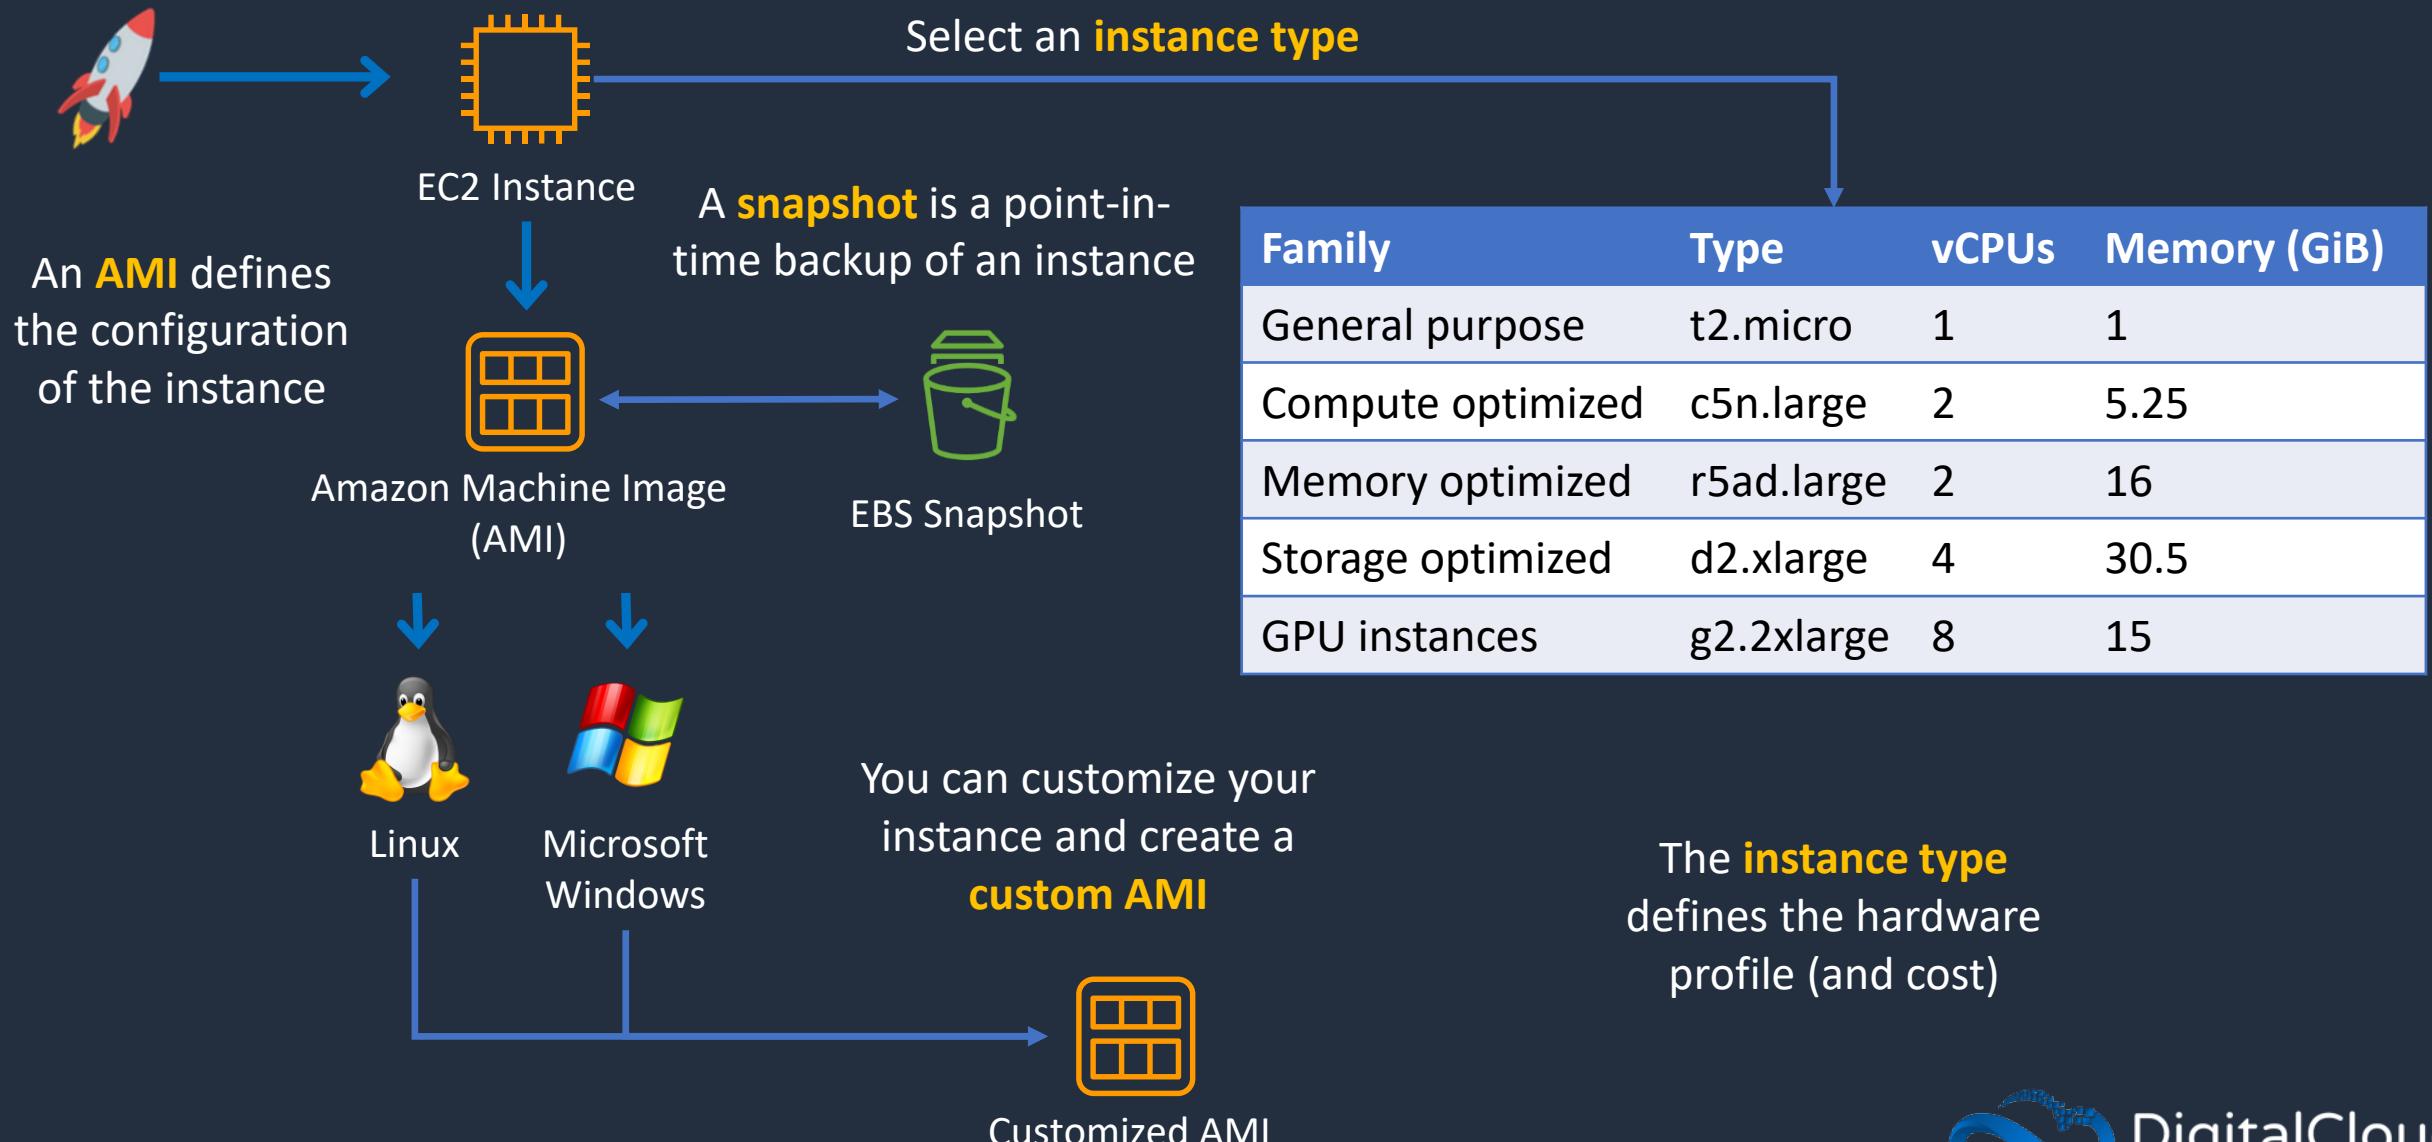
Destination	Target
172.31.0.0/16	Local
0.0.0.0/0	nat-gateway-id

# Launching Amazon EC2 Instances





# Launching an EC2 Instance



# Connecting to Amazon EC2

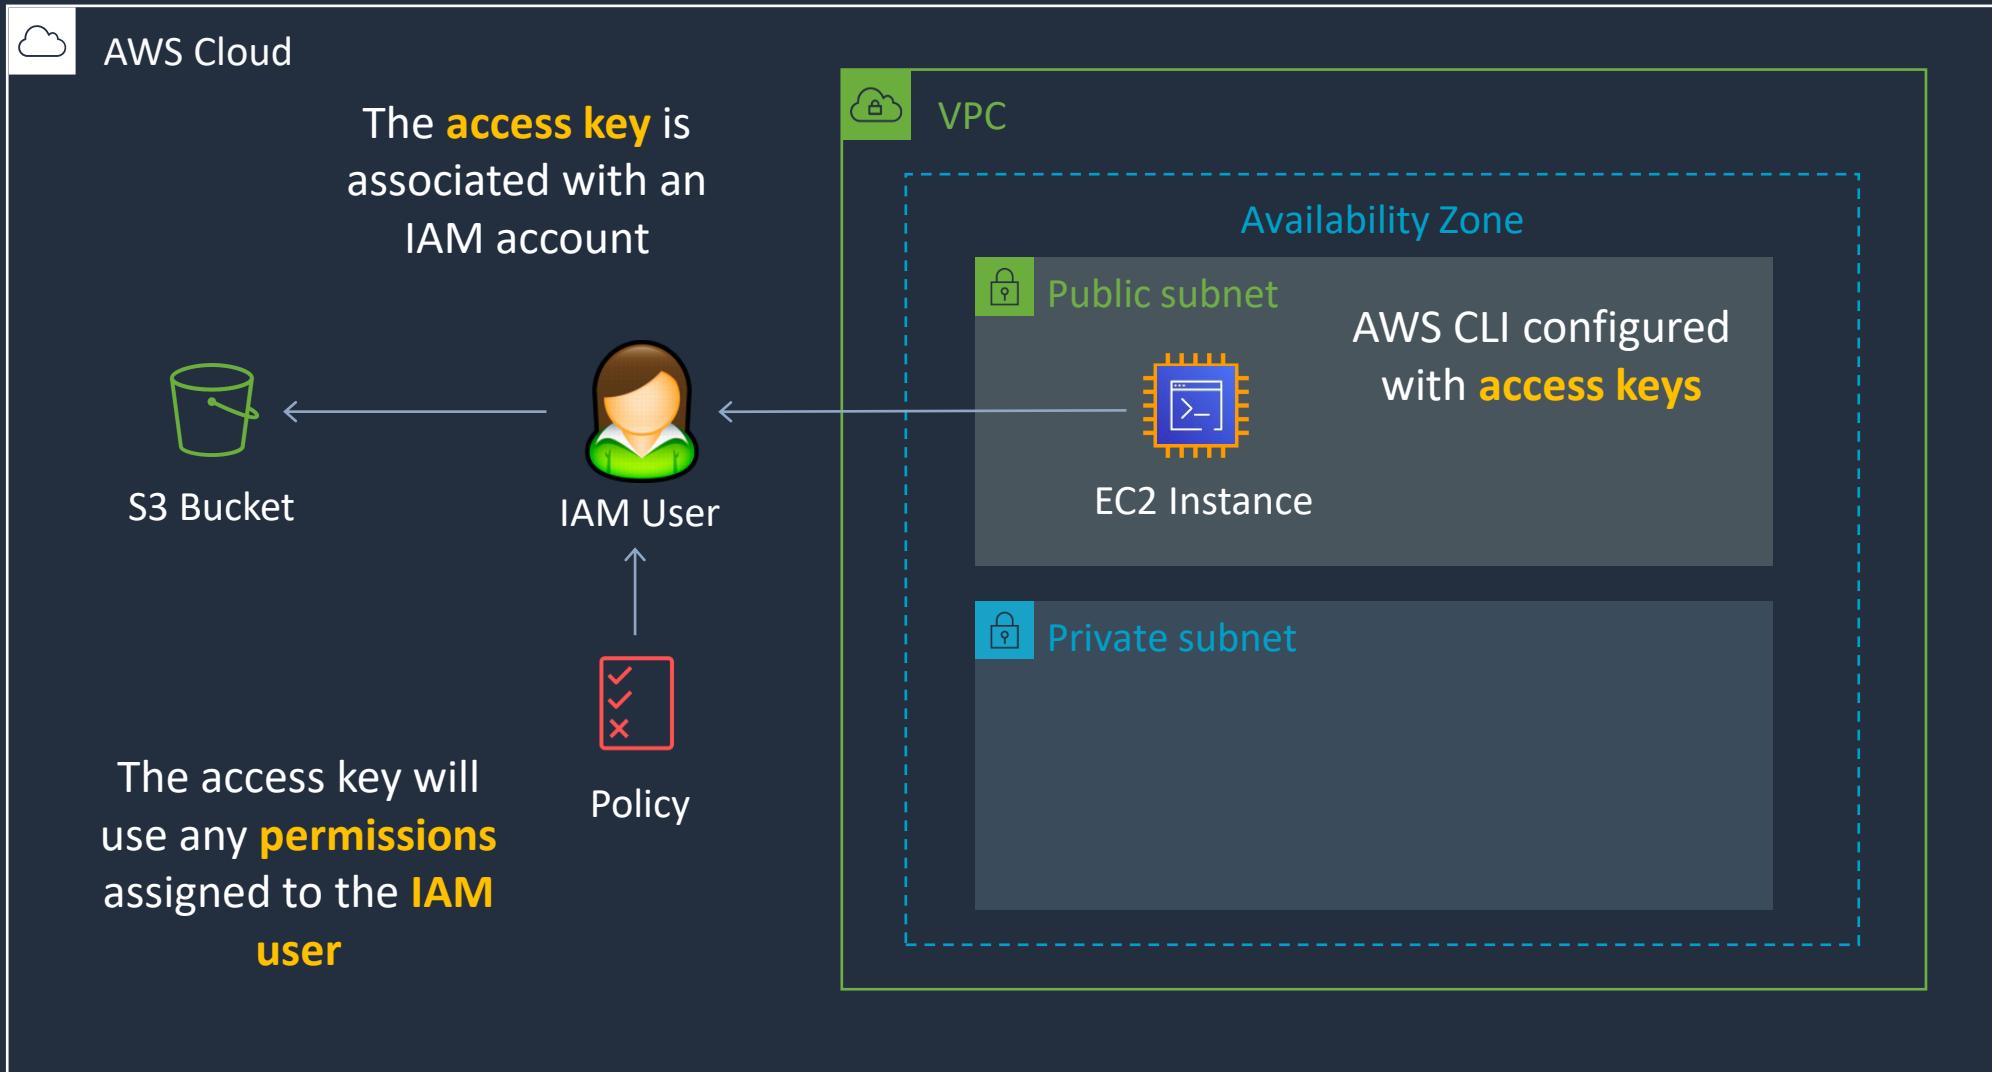


# Access Keys and IAM Roles with EC2



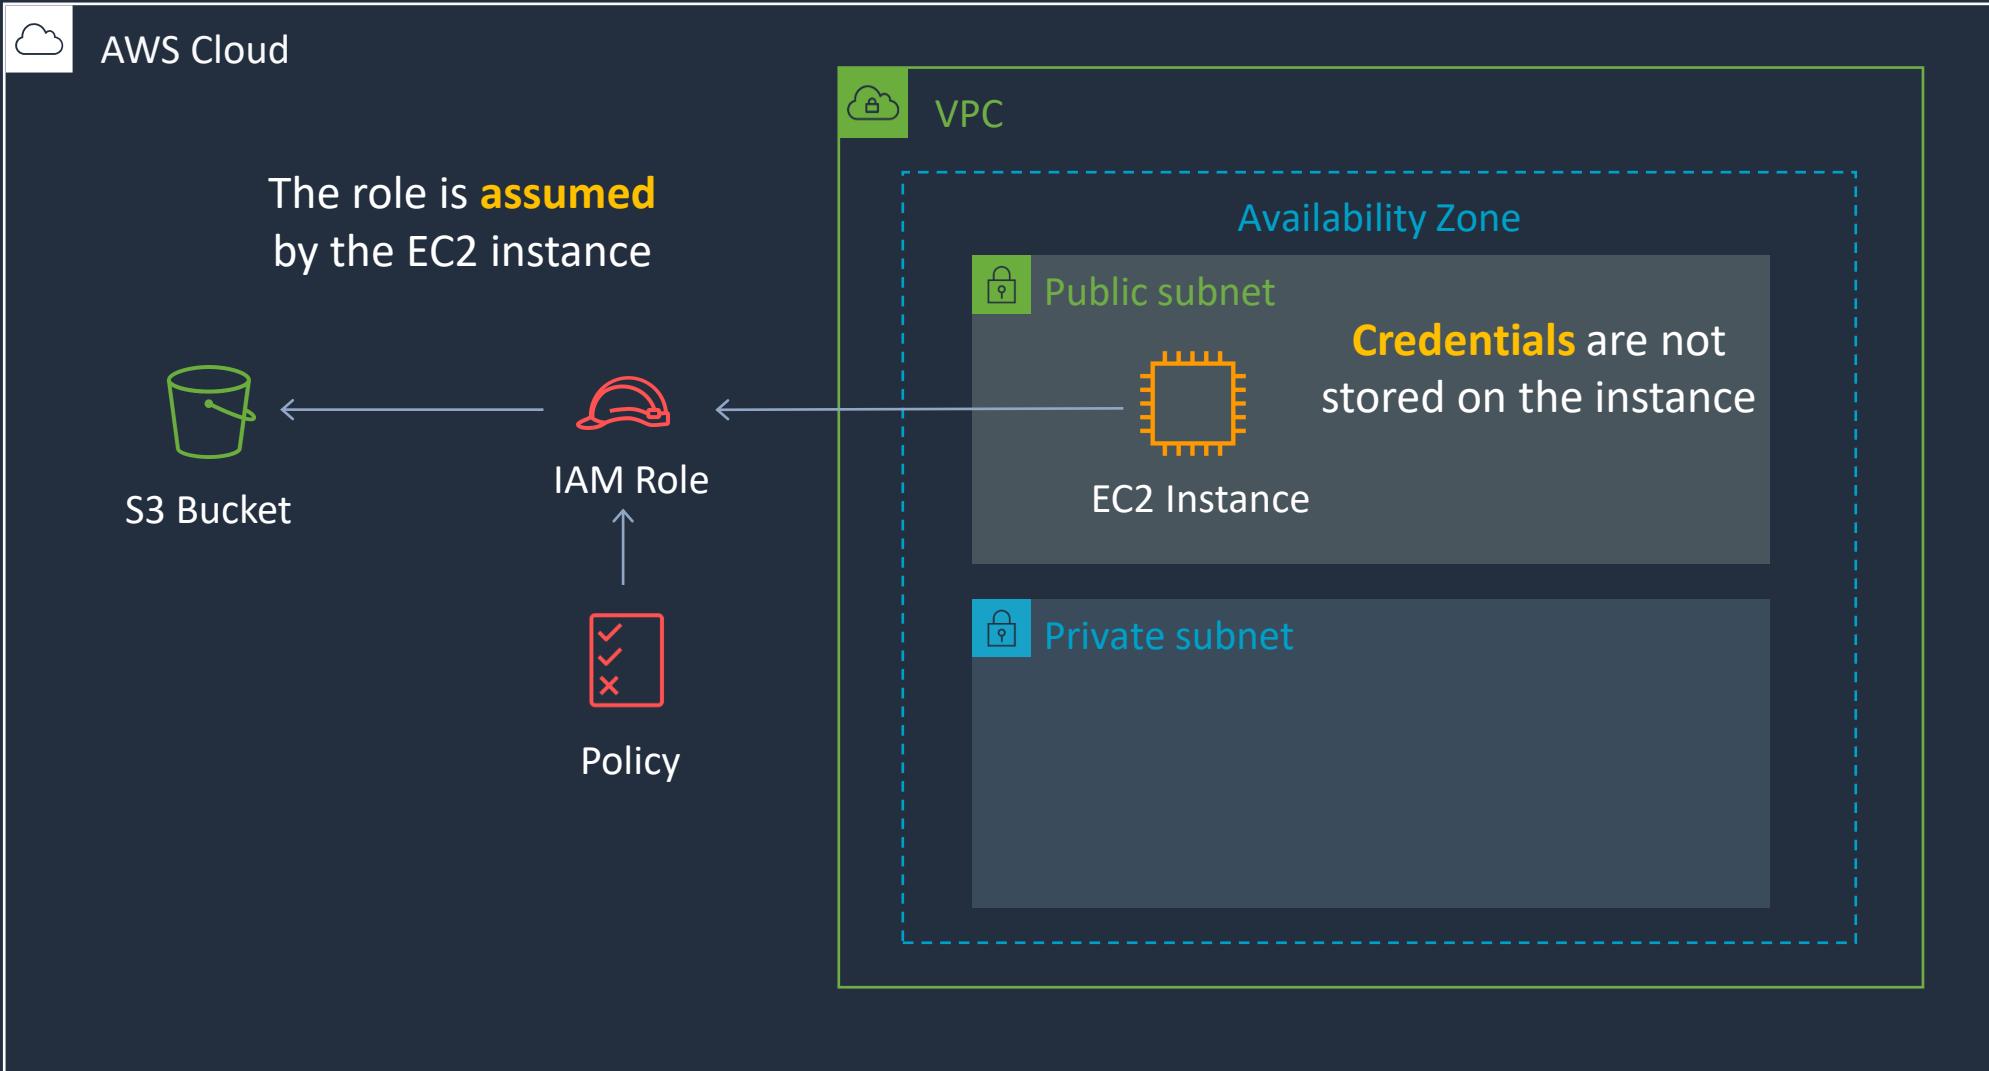


# Using Access Keys with Amazon EC2





# Using Roles with Amazon EC2



# Practice with Access Keys and IAM Roles



# Create a Website with User Data





# Amazon EC2 User Data

The code is run when the instance starts for the **first time**

AWS Management Console



**Batch** and **PowerShell** scripts can be run on Windows

User data (i)  As text  As file  Input is already base64 encoded

```
#!/bin/bash
yum update -y
yum install -y httpd
systemctl start httpd
systemctl enable httpd
```

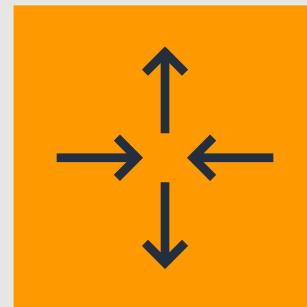
Limited to  
**16 KB**



EC2 Instance

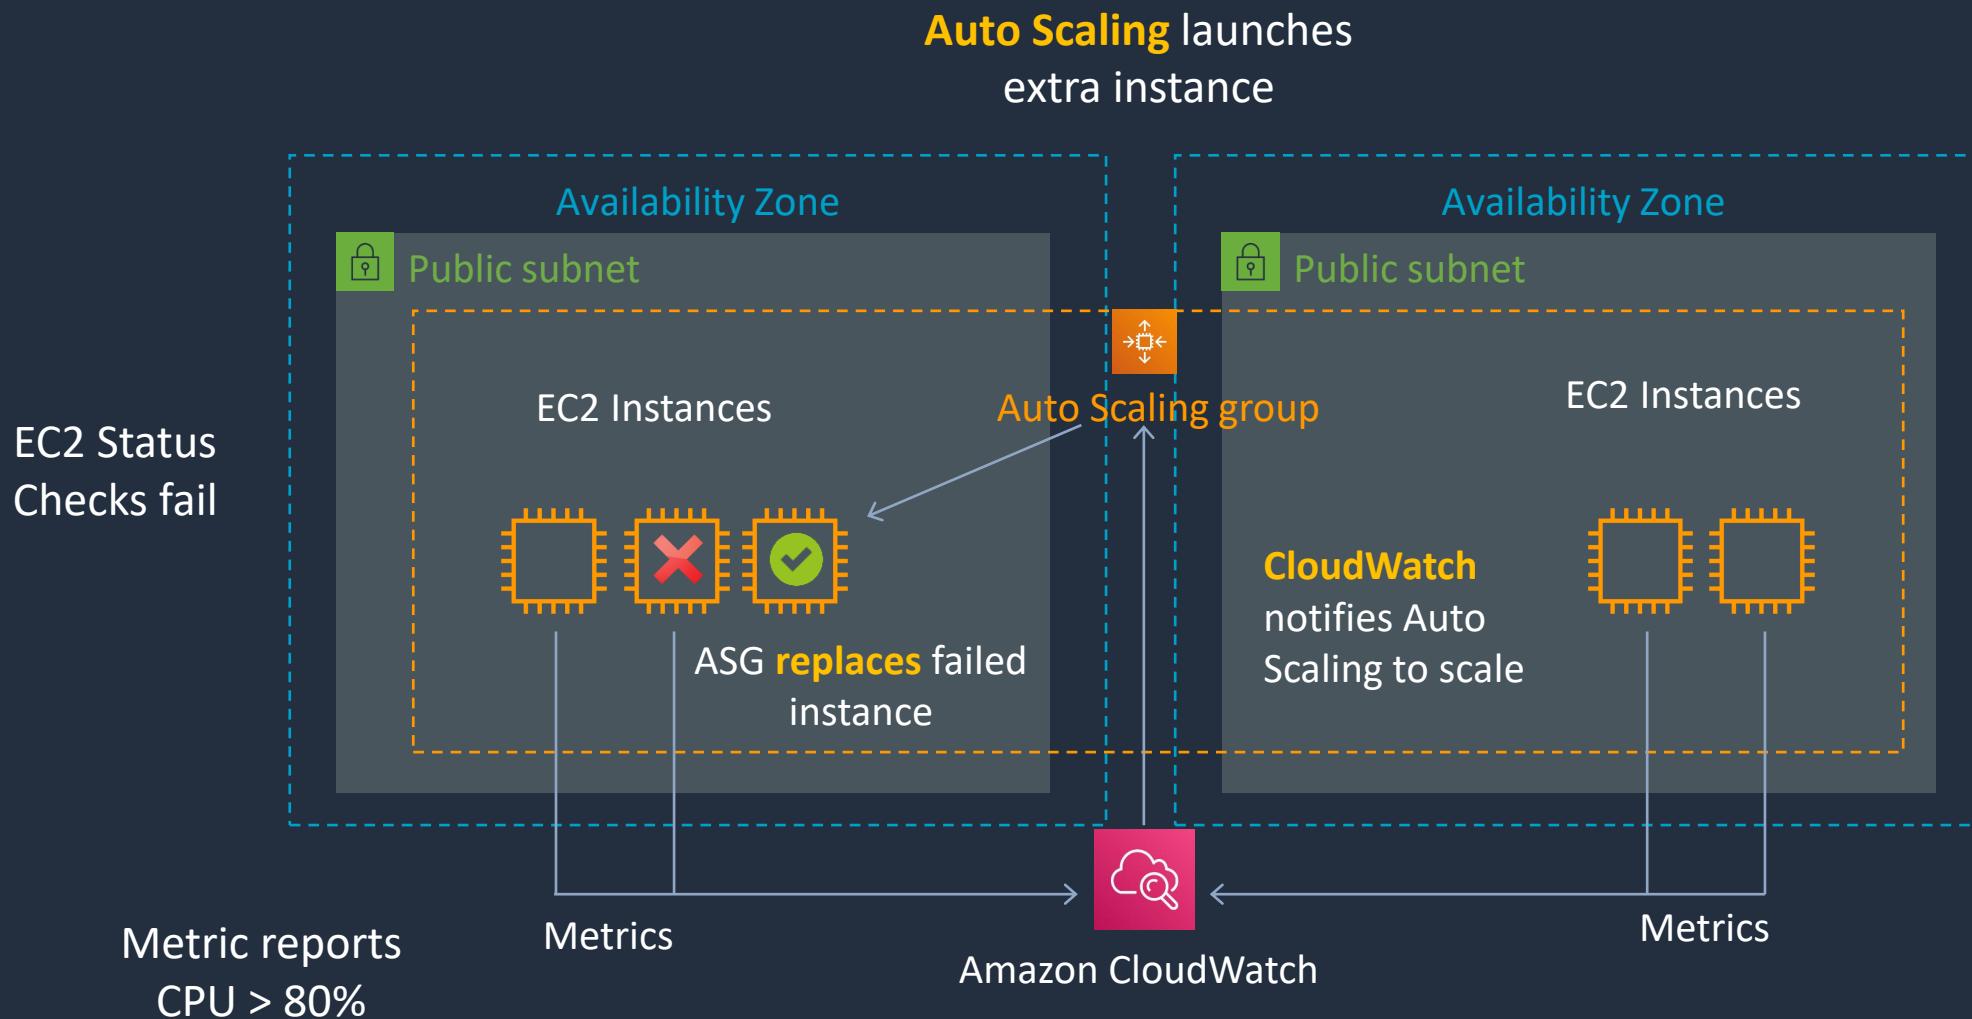
EC2 Instance with a **web service** is launched

# Amazon EC2 Auto Scaling





# Amazon EC2 Auto Scaling



# Create an Auto Scaling Group



# Create a Scaling Policy



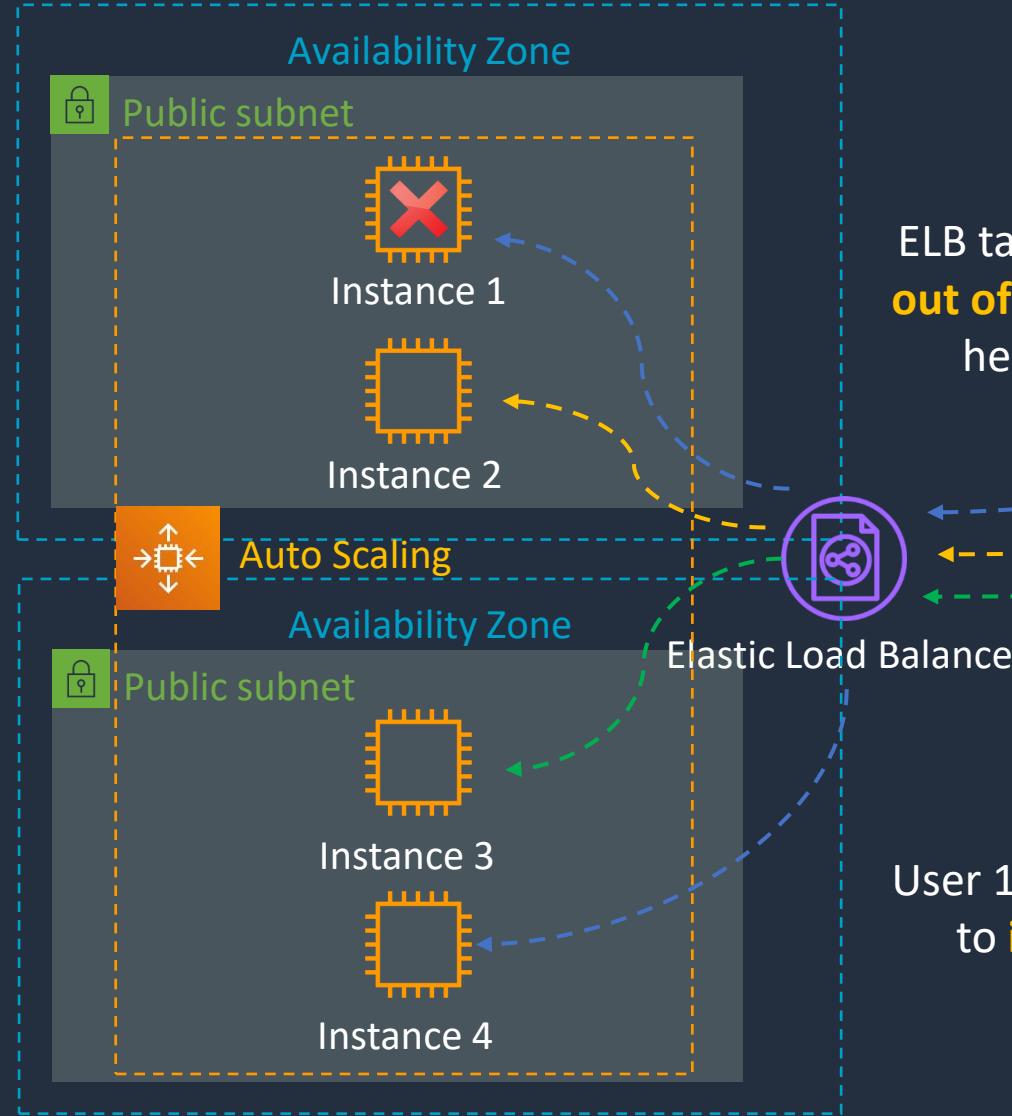
# Amazon Elastic Load Balancing





# Amazon Elastic Load Balancing

EC2 Auto Scaling  
**terminates**  
instance 1



ELB takes instance 1  
**out of service** (failed  
health check)

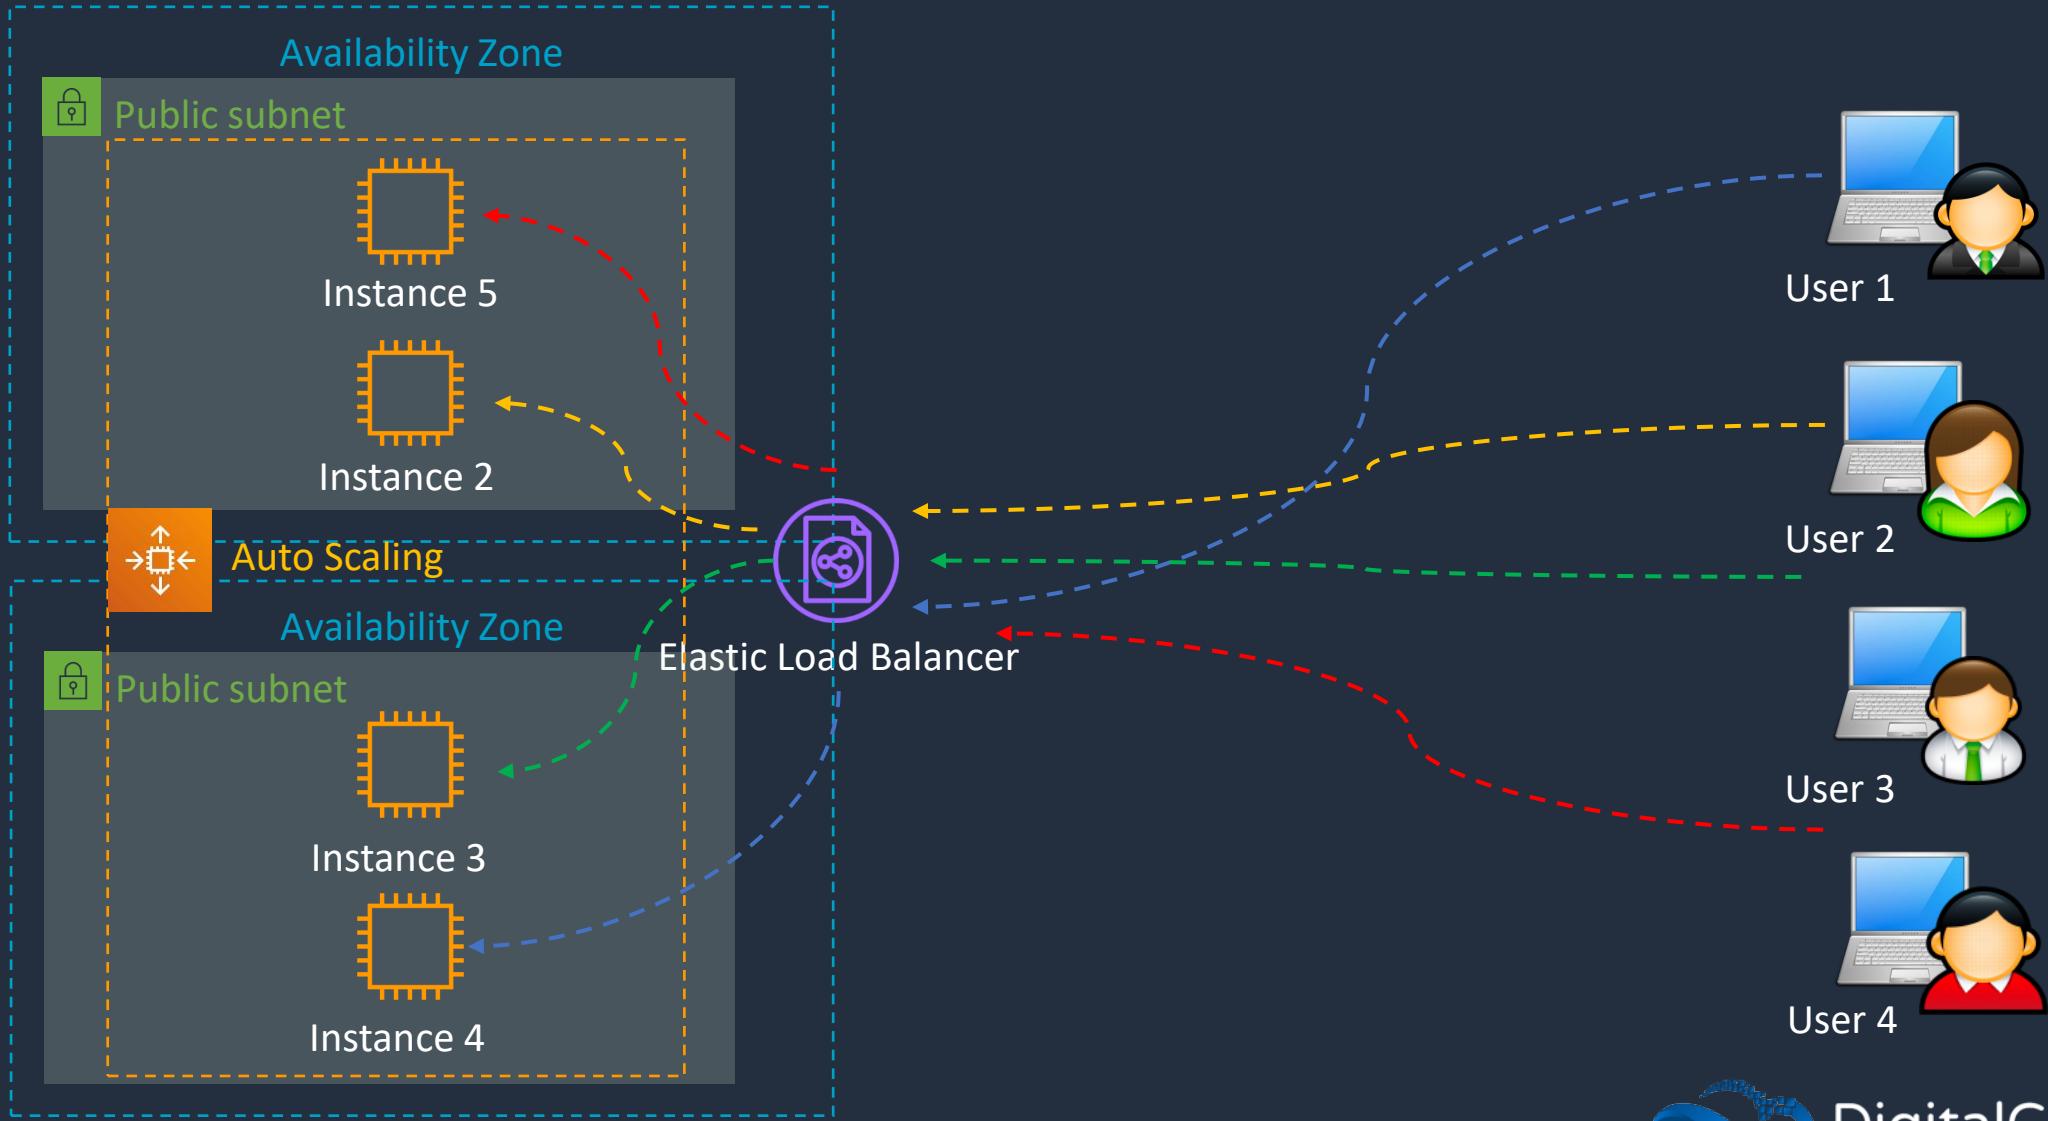
User 1 is connected  
to **instance 4**





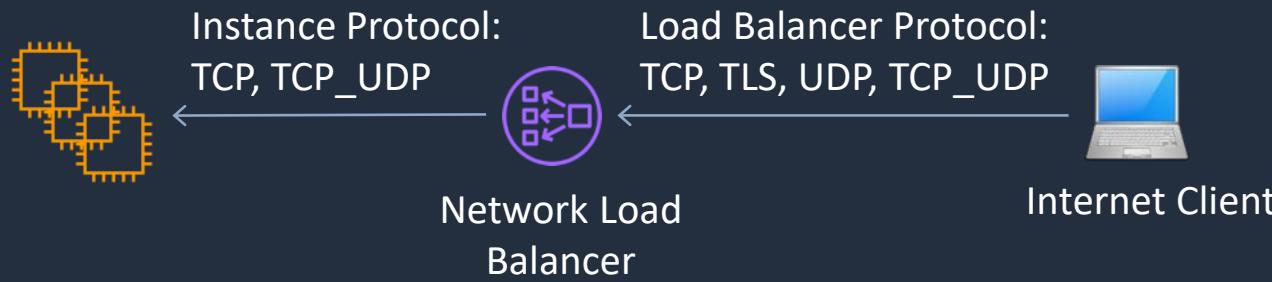
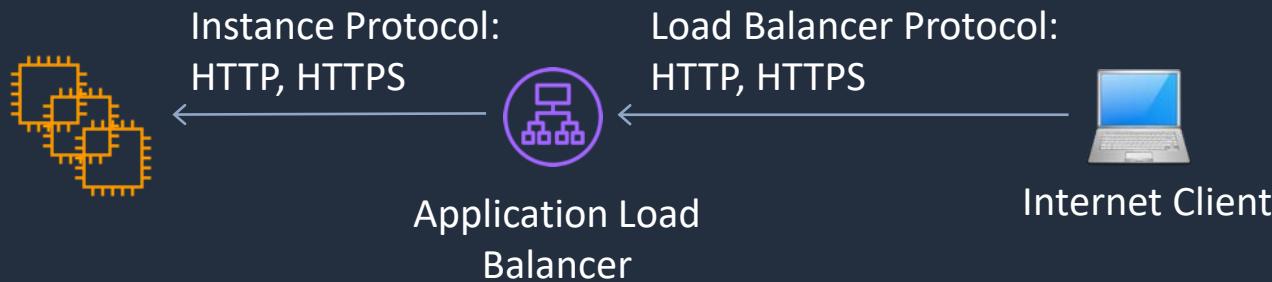
# Amazon Elastic Load Balancing

EC2 Auto Scaling  
**launches**  
instance 5





# Types of Elastic Load Balancer (ELB)



## Application Load Balancer

- Operates at the request level
- Routes based on the content of the request (layer 7)
- Supports path-based routing, host-based routing, query string parameter-based routing, and source IP address-based routing
- Supports instances, IP addresses, Lambda functions and containers as targets

## Network Load Balancer

- Operates at the connection level
- Routes connections based on IP protocol data (layer 4)
- Offers ultra high performance, low latency and TLS offloading at scale
- Can have a static IP / Elastic IP
- Supports UDP and static IP addresses as targets



# ELB Use Cases

---

---

## Application Load Balancer

- Web applications with L7 routing (HTTP/HTTPS)
- Microservices architectures (e.g. Docker containers)
- Lambda targets

## Network Load Balancer

- TCP and UDP based applications
- Ultra-low latency
- Static IP addresses
- VPC endpoint services

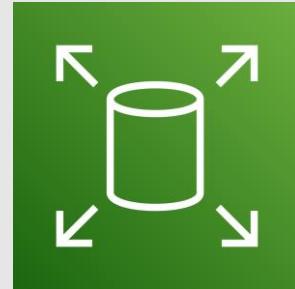
# Create an Application Load Balancer



# SECTION 7

## AWS Storage Services

# Amazon EBS and Instance Stores





# Amazon Elastic Block Store (EBS)

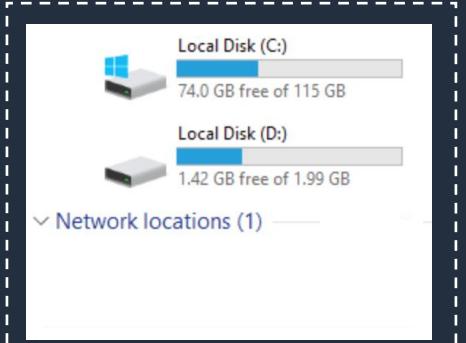
EBS volumes  
exist within an  
**Availability Zone**

Availability Zone

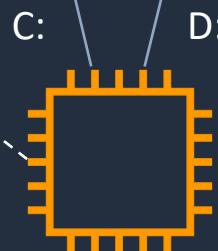


EBS Volume

EBS Volume



The volume is  
**attached** over a  
network



EC2 Instance

The volume is automatically  
replicated **within the AZ**



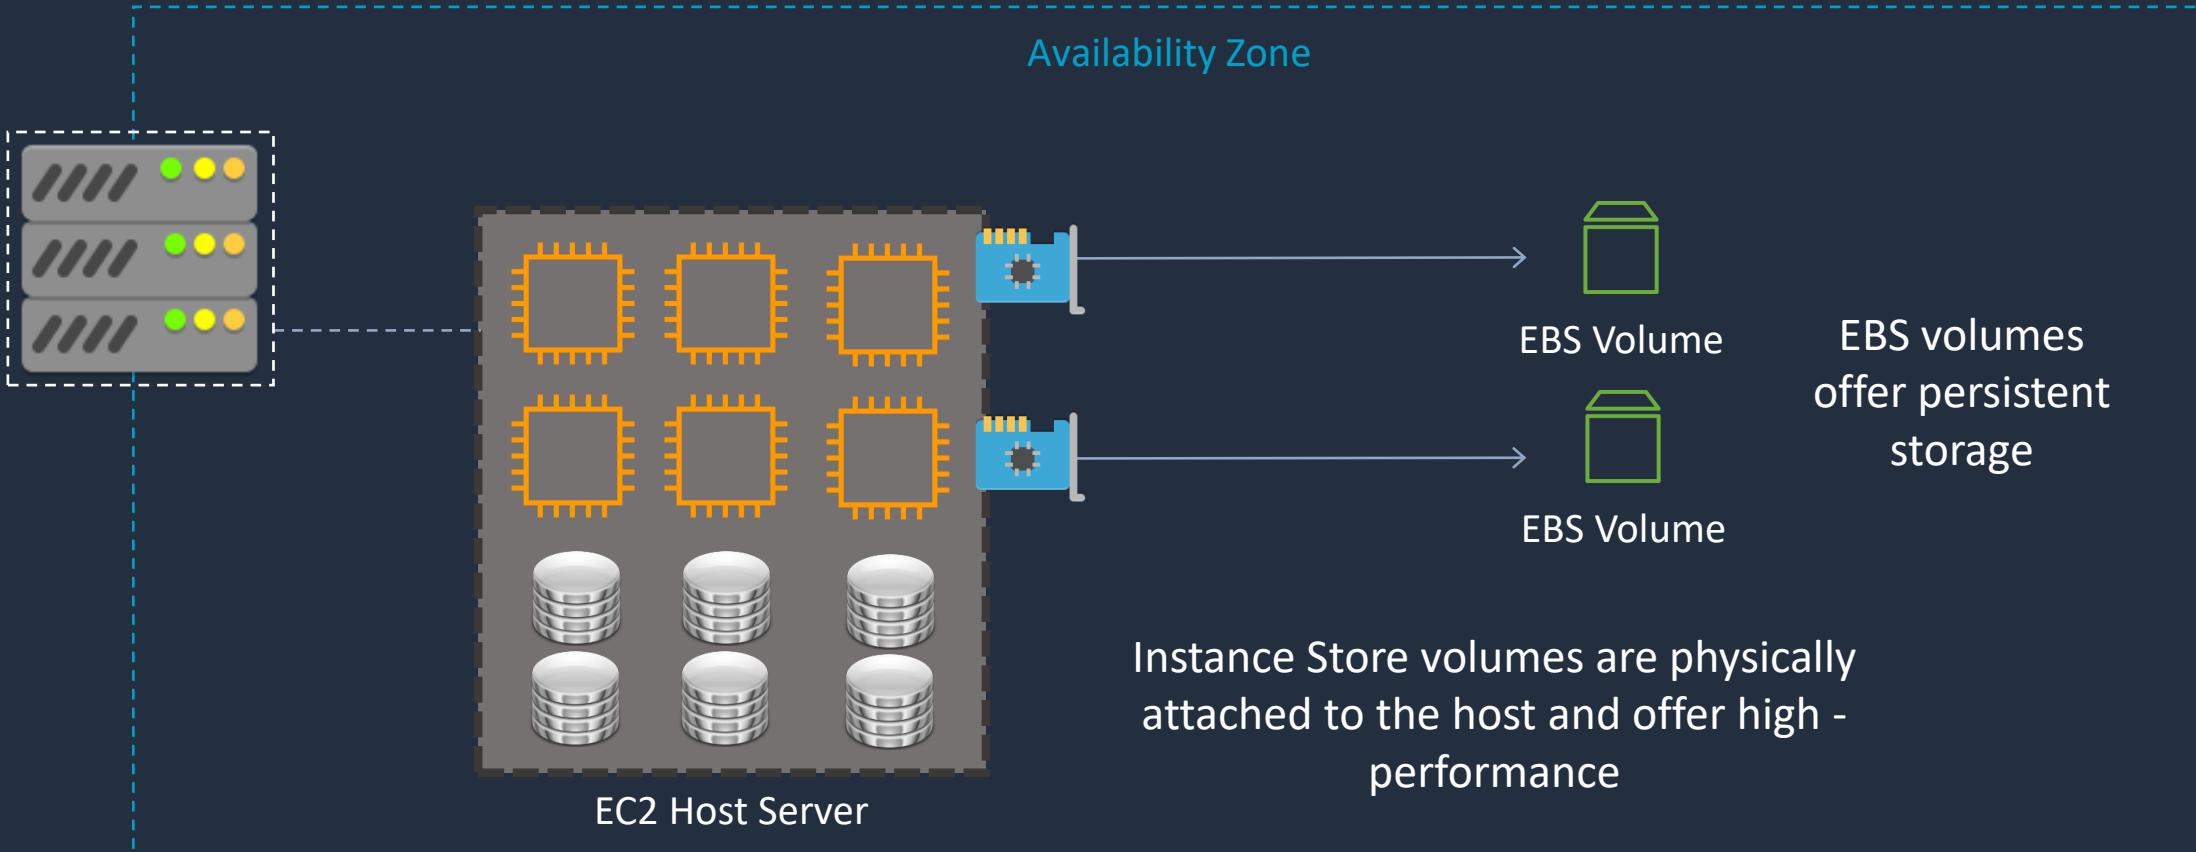
# Amazon EBS Volume Types

	General Purpose SSD		Provisioned IOPS SSD		
Volume type	gp3	gp2	io2 Block Express ‡	io2	io1
Durability	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)	99.999% durability (0.001% annual failure rate)		99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)
Use cases	<ul style="list-style-type: none"><li>Low-latency interactive apps</li><li>Development and test environments</li></ul>		Workloads that require sub-millisecond latency, and sustained IOPS performance or more than 64,000 IOPS or 1,000 MiB/s of throughput		<ul style="list-style-type: none"><li>Workloads that require sustained IOPS performance or more than 16,000 IOPS</li><li>I/O-intensive database workloads</li></ul>
Volume size	1 GiB - 16 TiB		4 GiB - 64 TiB	4 GiB - 16 TiB	
Max IOPS per volume (16 KiB I/O)	16,000		256,000	64,000 †	
Max throughput per volume	1,000 MiB/s	250 MiB/s *	4,000 MiB/s	1,000 MiB/s †	
Amazon EBS Multi-attach	Not supported		Not supported	Supported	
Boot volume	Supported				



# Amazon EBS vs Instance Store

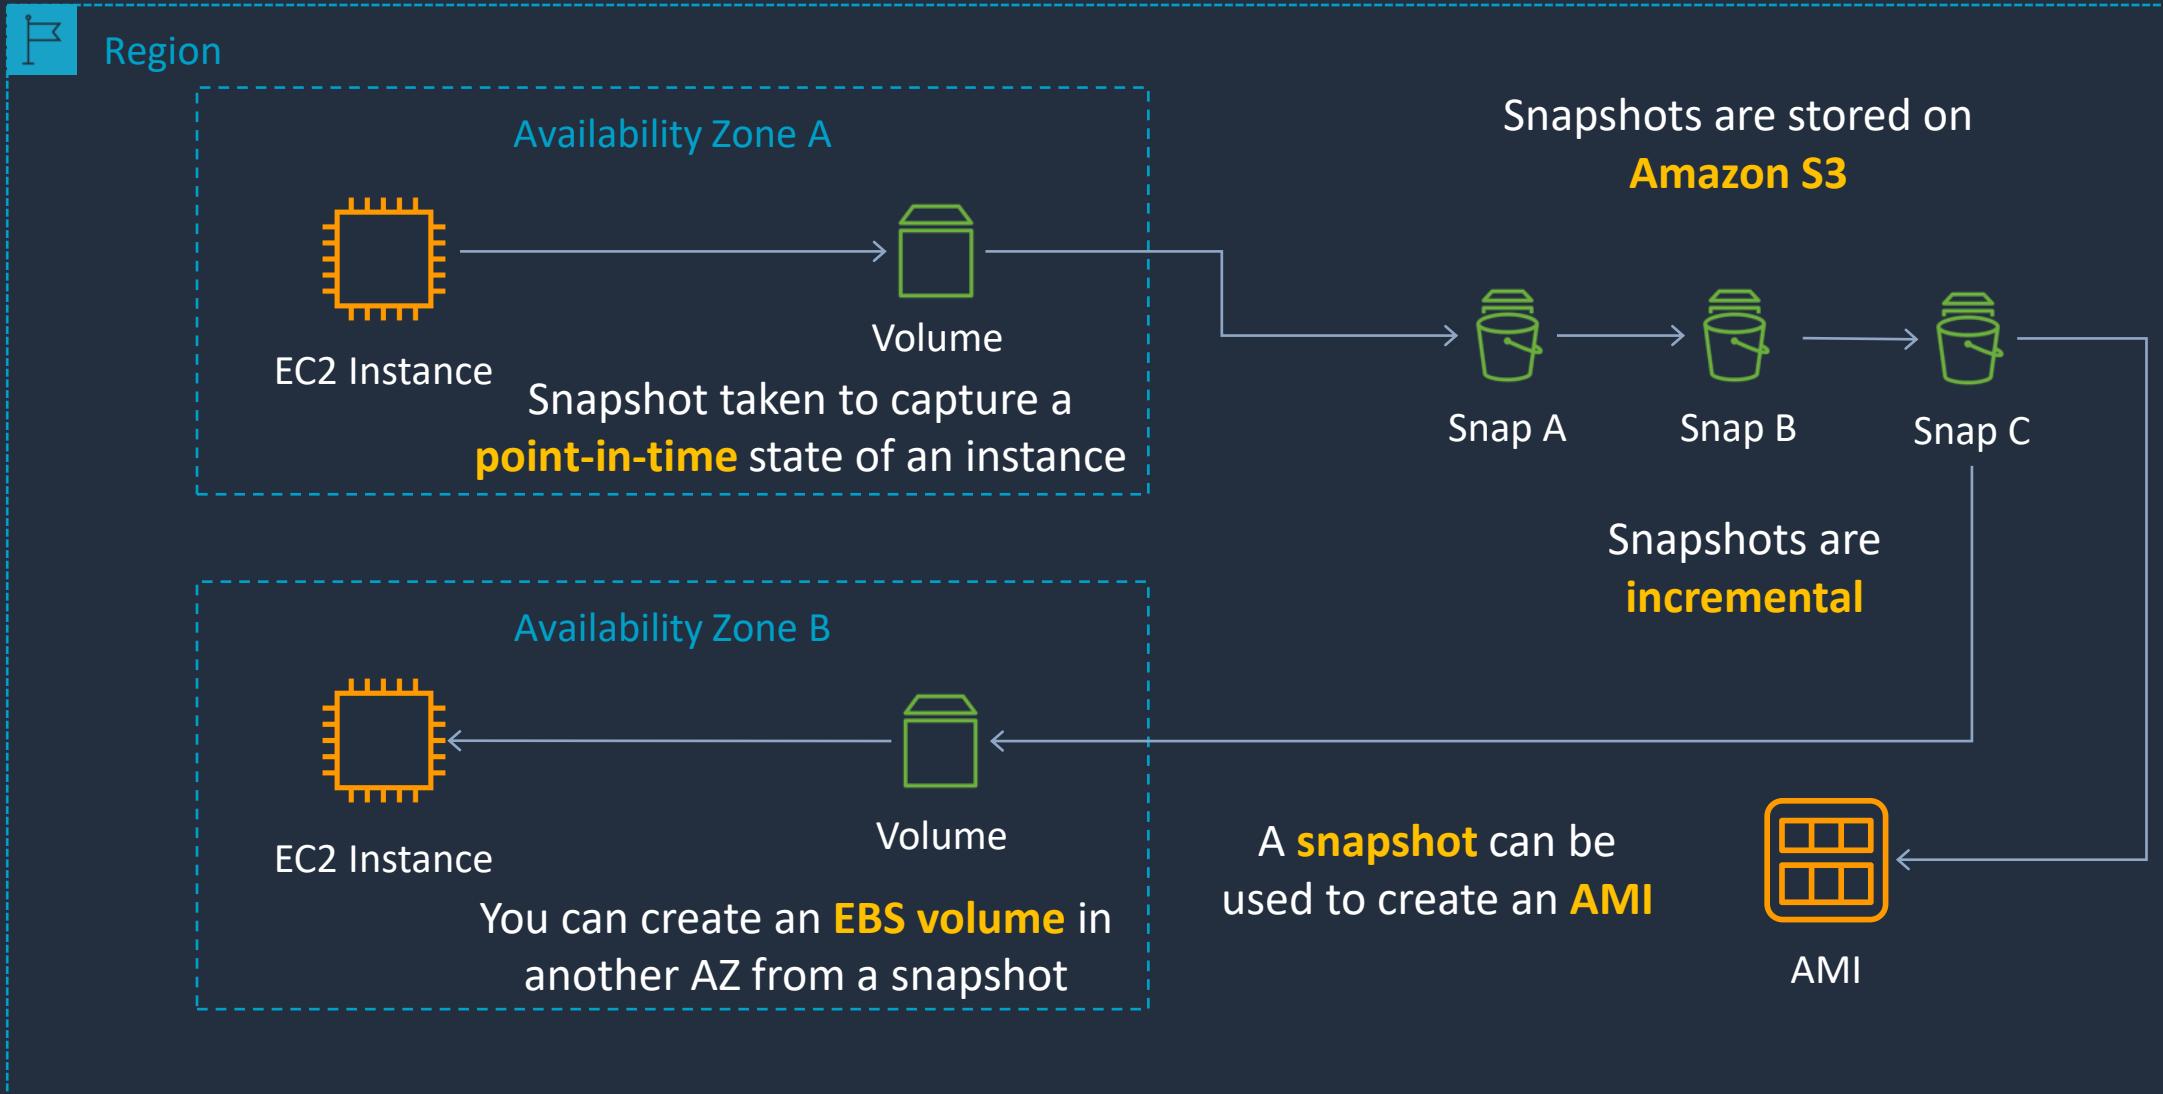
EBS volumes are **attached** over the **network**



Instance Store volumes are **ephemeral** (non-persistent)



# Amazon EBS Snapshots



# Create and Attach an EBS Volume



# EBS Snapshots and AMIs

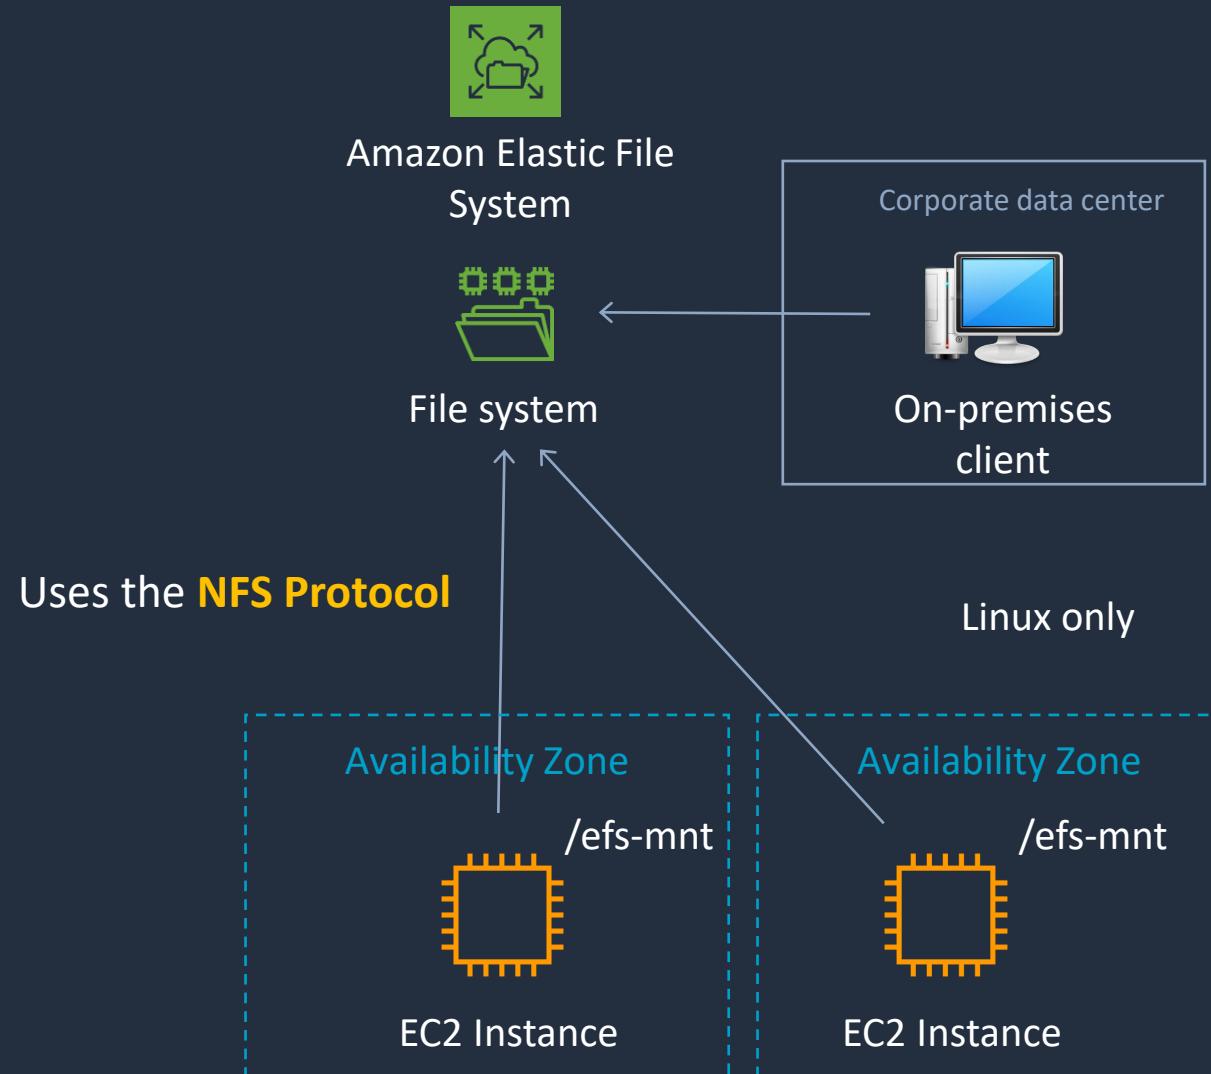


# Amazon Elastic File System (EFS)





# Amazon Elastic File System (EFS)



# Create an Amazon EFS Filesystem





# Mounting EFS File Systems

---

```
mkdir ~/efs-mount-point
```

## # Install NFS client

```
sudo yum -y install nfs-utils
```

## # Mount using the NFS client

```
sudo mount -t nfs4 -o nfsvers=4.1,rsize=1048576,wsize=1048576,hard,timeo=600,retrans=2,noresvport fs-0187058e25d360ce2.efs.us-east-1.amazonaws.com:/ ~/efs-mount-point
```

## # Install EFS utils

```
sudo yum install -y amazon-efs-utils
```

## # Mount using the EFS mount helper:

```
sudo mount -t efs -o tls fs-0187058e25d360ce2.efs.us-east-1.amazonaws.com:/ ~/efs-mount-point
```

# Amazon Simple Storage Service (S3)





# Amazon Simple Storage Service (S3)



S3 Bucket



A **bucket** is a container for objects

An **object** is a file you upload

You can store millions of **objects** in a **bucket**



Accessing objects in a bucket:

`https://bucket.s3.aws-region.amazonaws.com/key`  
`https://s3.aws-region.amazonaws.com/bucket/key`

The **HTTP protocol** is used with a **REST API** (e.g. GET, PUT, POST, SELECT, DELETE)



# Amazon Simple Storage Service (S3)



Bucket

A **bucket** is a container  
for objects

<http://bucket.s3.aws-region.amazonaws.com>

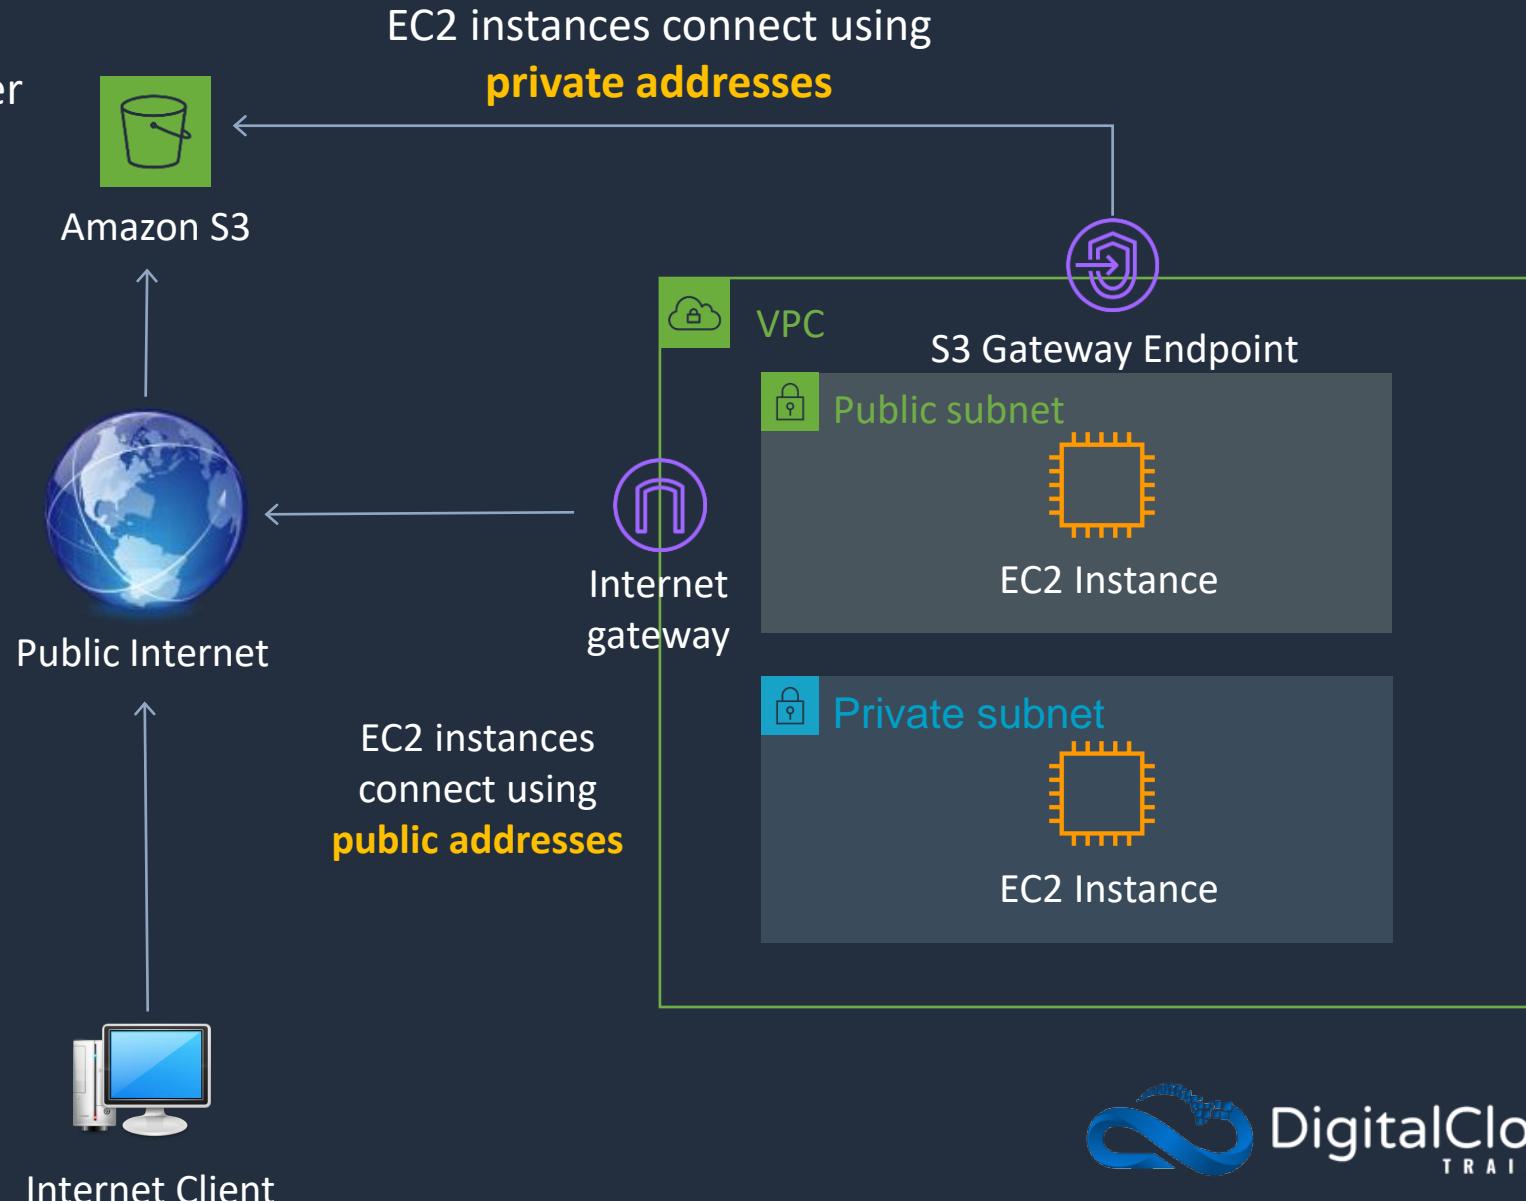
<http://s3.aws-region.amazonaws.com/bucket>



Object

An object consists of:

- Key (name of objects)
- Version ID
- Value (actual data)
- Metadata
- Subresources
- Access control information





# File Storage vs Object Storage

---

File Share



- Data stored in directories
- A hierarchy of directories can be formed
- File systems are mounted to an operating system
- Function like local storage
- Network connection is maintained
- Example is Amazon EFS

Object Store



`http://bucket.s3.aws-region.amazonaws.com`

- Data stored in buckets
- Flat namespace (no hierarchy)
- Hierarchy can be mimicked with prefixes
- Accessed by **REST API** and cannot be mounted
- Network connection is completed after each request
- Example is Amazon S3

# Working with S3 Buckets and Objects



# Create an S3 Static Website



# SECTION 8

## AWS Database Services

# Amazon Relational Database Service (RDS)





# Amazon RDS

---

RDS is a managed,  
**relational database**



RDS runs on EC2  
instances, so you choose  
an **instance type**

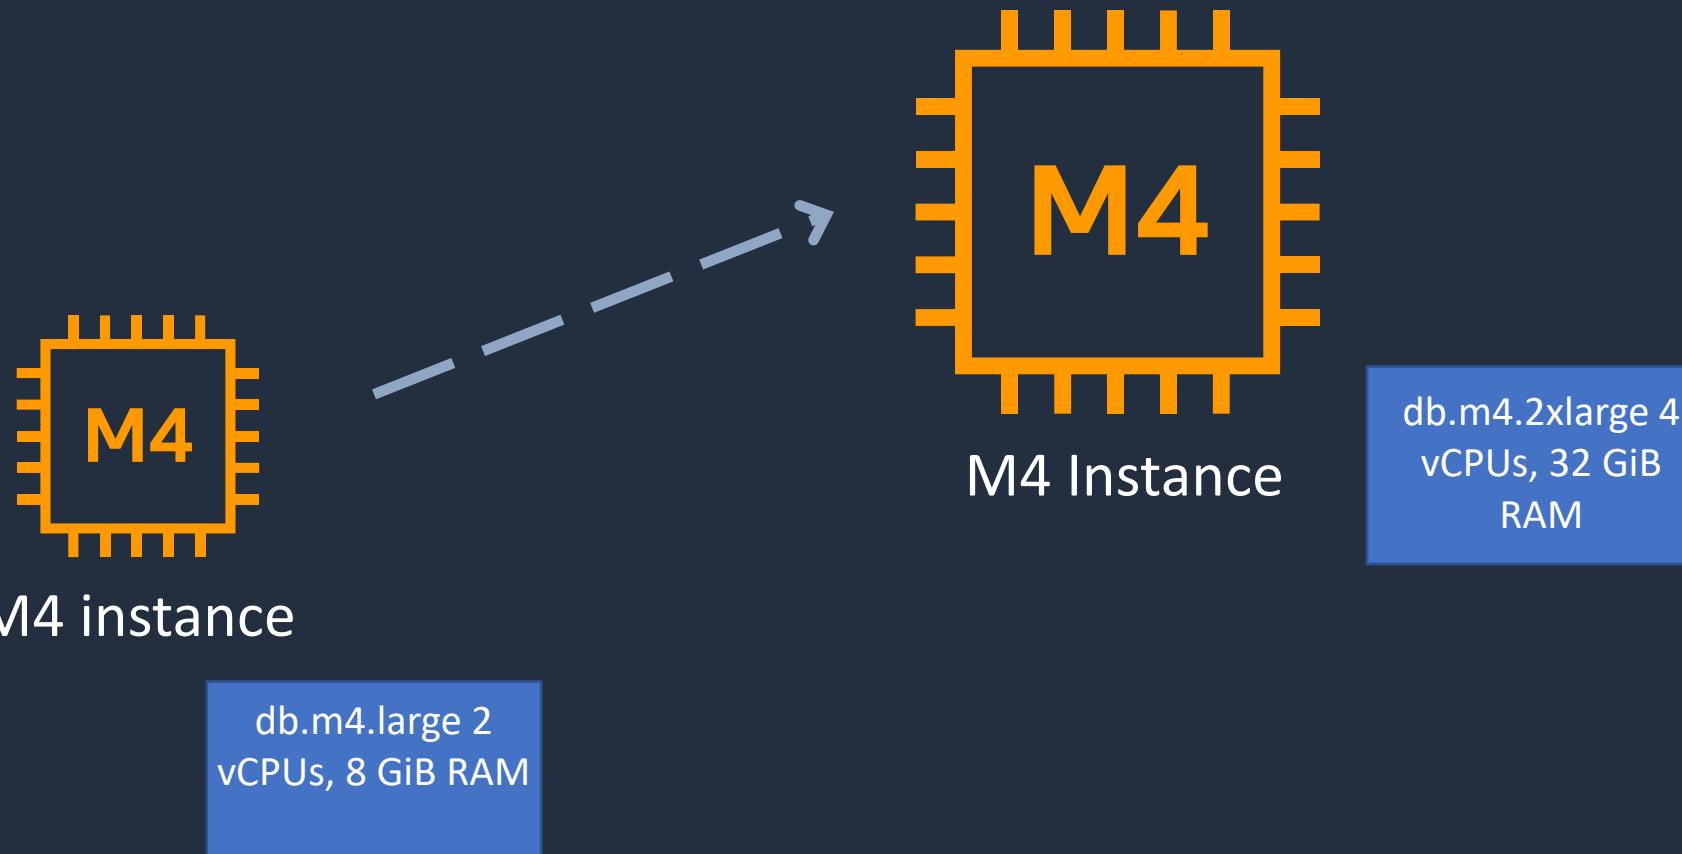
RDS supports the following  
database engines:

- Amazon Aurora
- MySQL
- MariaDB
- Oracle
- Microsoft SQL Server
- PostgreSQL



# Amazon RDS – Scaling up (vertically)

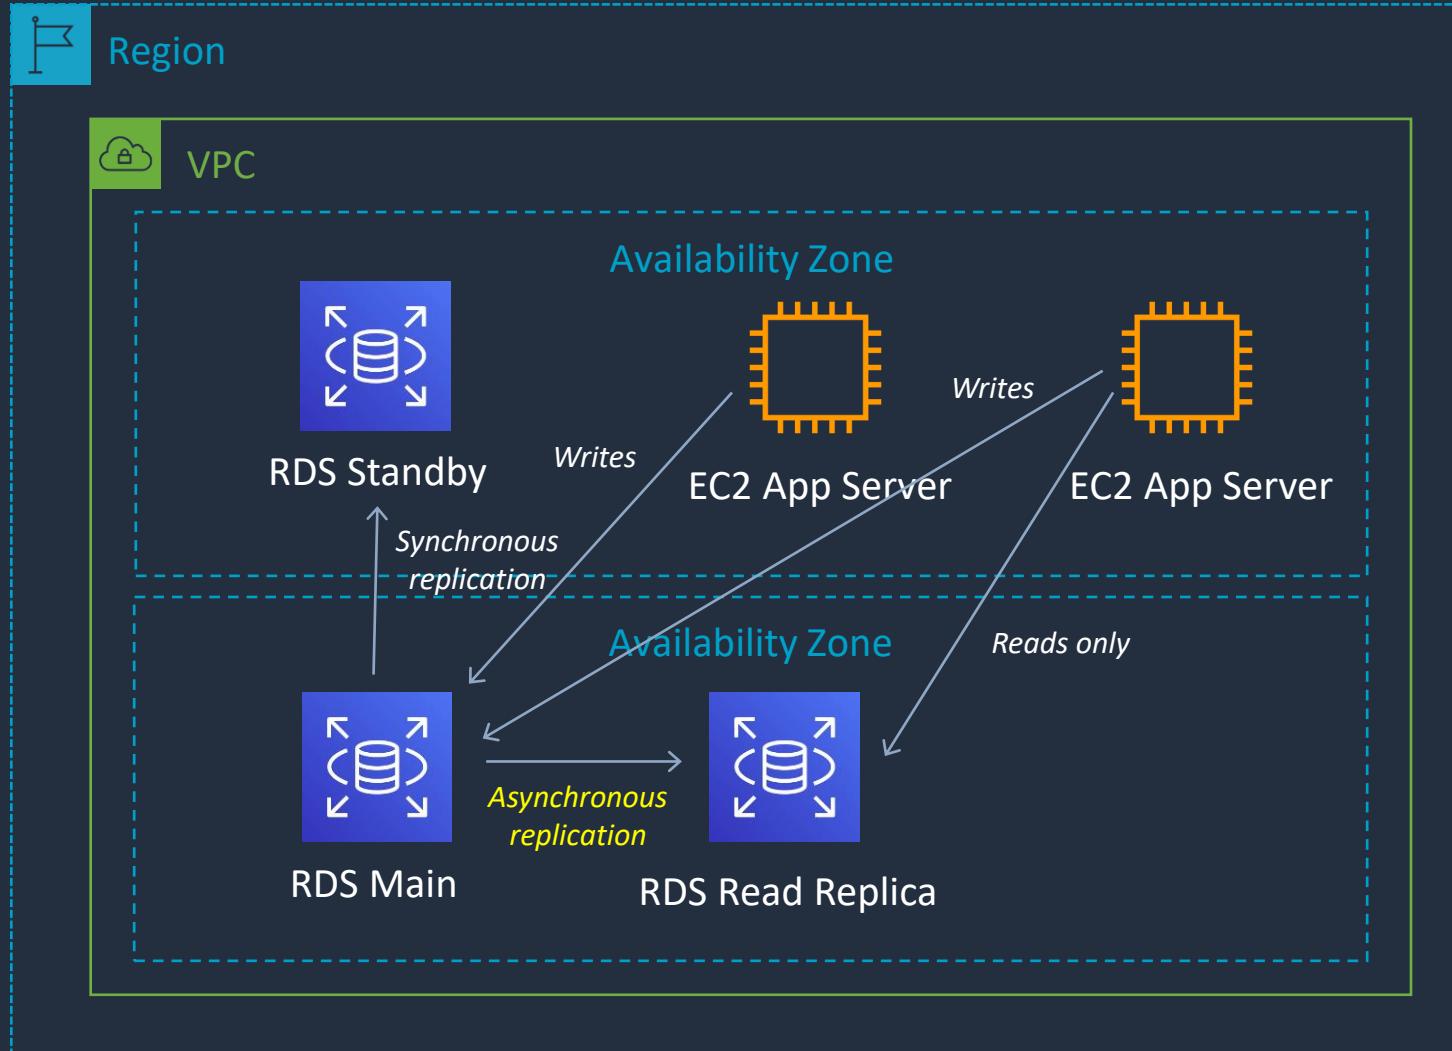
---





# Amazon RDS – Disaster Recovery (DR) and Scaling Out (horizontally)

Multi-AZ creates a passive standby.  
Primarily used for **disaster recovery**



Application servers can read from the read replica and write to the master

Read Replicas are used for scaling database queries (reads)

# Create an Amazon RDS Database



# Add an Amazon RDS Read Replica



# Amazon DynamoDB





# Amazon DynamoDB

Fully managed service. You create tables on an existing database

Offers seamless, horizontal, scaling



DynamoDB Table

DynamoDB is a NoSQL, key-value type of database

Data is replicated across multiple AZs within a region



# Amazon DynamoDB

- DynamoDB is made up of:
  - Tables
  - Items
  - Attributes

userid	orderid	book	price	date
user001	1000092	ISBN100..	9.99	2020.04..
user002	1000102	ISBN100..	24.99	2020.03..
user003	1000168	ISBN2X0..	12.50	2020.04..

# Create an Amazon DynamoDB Table



# SECTION 9

## Automation and DevOps on AWS

# Infrastructure as Code with AWS CloudFormation





# AWS CloudFormation

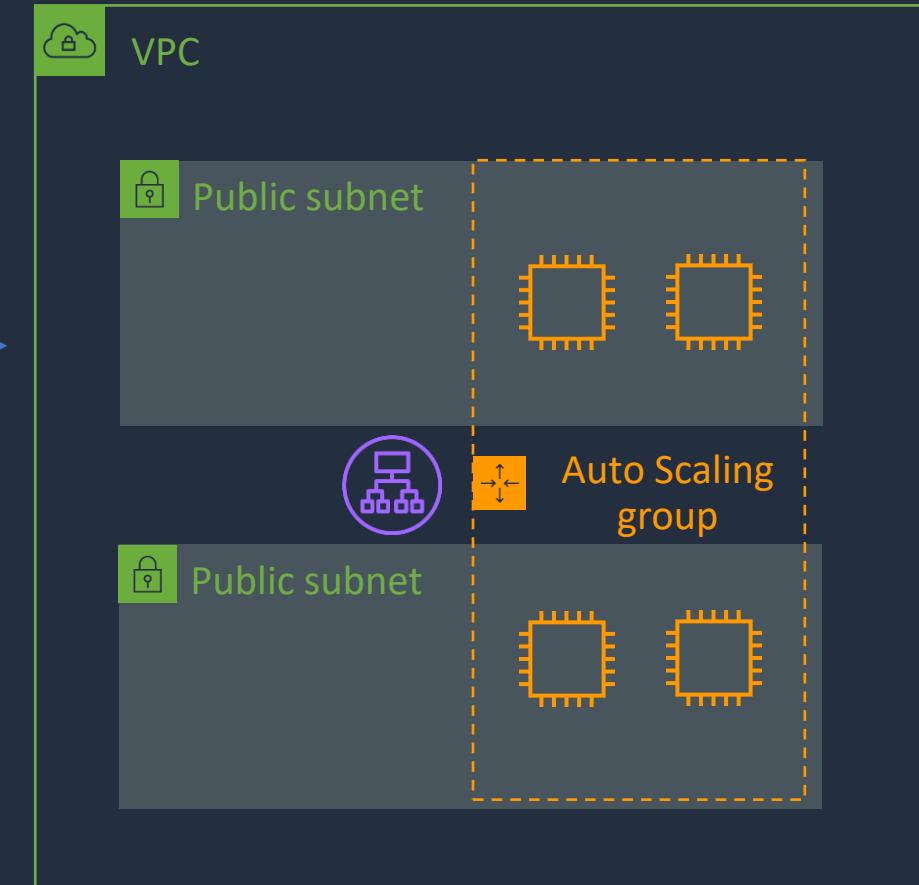
Infrastructure patterns are defined in a **template** file using **code**



AWS CloudFormation

```
1 "AWSTemplateFormatVersion": "2010-09-09",
2
3 "Description": "AWS CloudFormation Sample Template WordPress_Multi_AZ: WordPress is web
4
5 "Parameters": {
6   "VpcId": {
7     "Type": "AWS::EC2::VPC::Id",
8     "Description": "VpcId of your existing Virtual Private Cloud (VPC)",
9     "ConstraintDescription": "must be the VPC Id of an existing Virtual Private Cloud."
10 },
11
12 "Subnets": {
13   "Type": "List<AWS::EC2::Subnet::Id>",
14   "Description": "The list of SubnetIds in your Virtual Private Cloud (VPC)",
15   "ConstraintDescription": "must be a list of at least two existing subnets associated
16 },
```

CloudFormation **builds** your infrastructure according to the **template**





# AWS CloudFormation

Component	Description
Templates	The JSON or YAML text file that contains the instructions for building out the AWS environment
Stacks	The entire environment described by the template and created, updated, and deleted as a single unit
StackSets	AWS CloudFormation StackSets extends the functionality of stacks by enabling you to create, update, or delete stacks across multiple accounts and regions with a single operation
Change Sets	A summary of proposed changes to your stack that will allow you to see how those changes might impact your existing resources before implementing them

# Creating and Updating Stacks



# Deploy a VPC Using CloudFormation



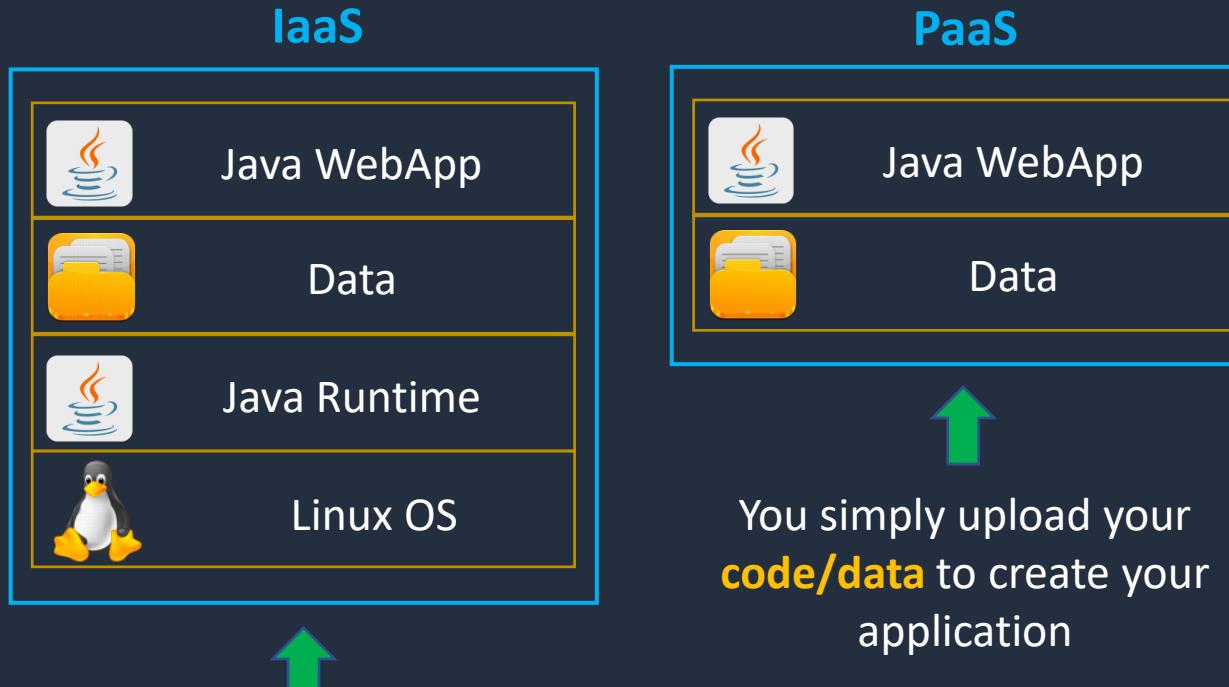
# Platform as a Service with AWS Elastic Beanstalk





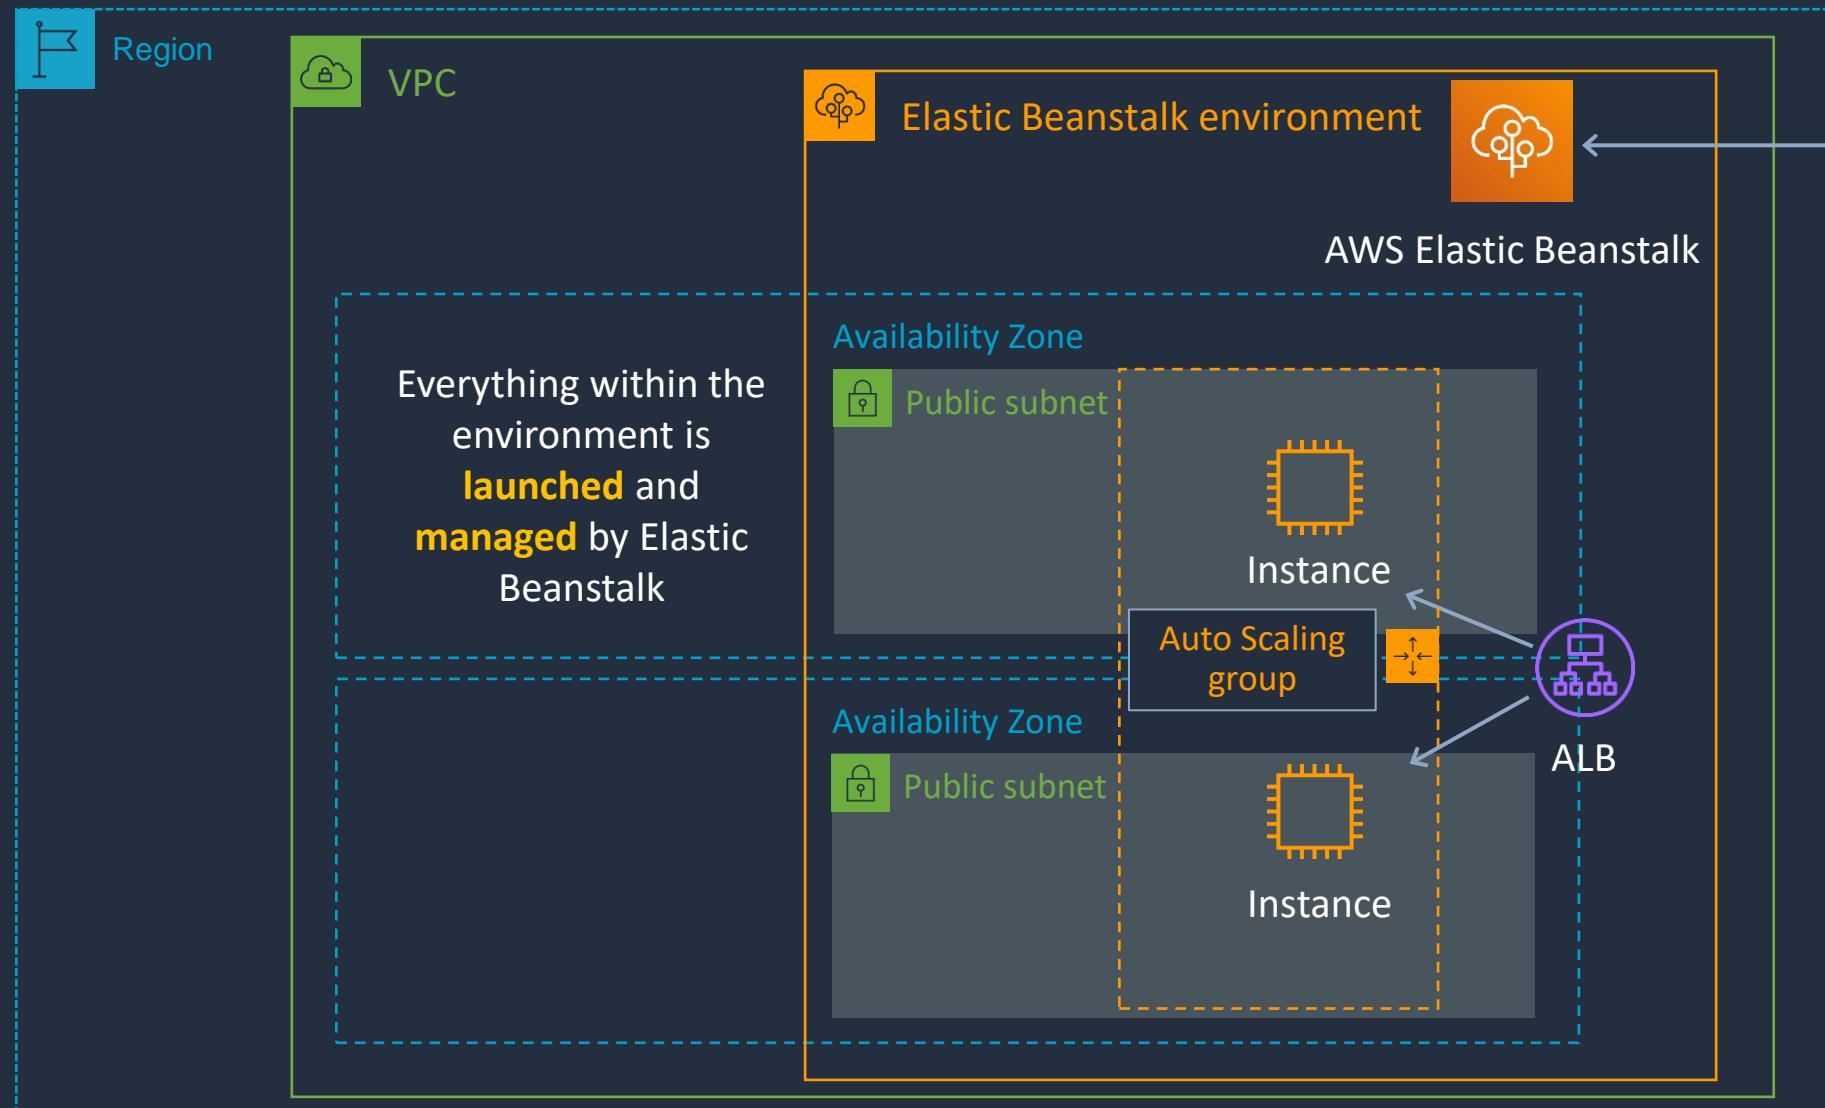
# Cloud Service Models: Comparison

Example is  
**Amazon EC2**





# AWS Elastic Beanstalk



Upload **source code** in ZIP file

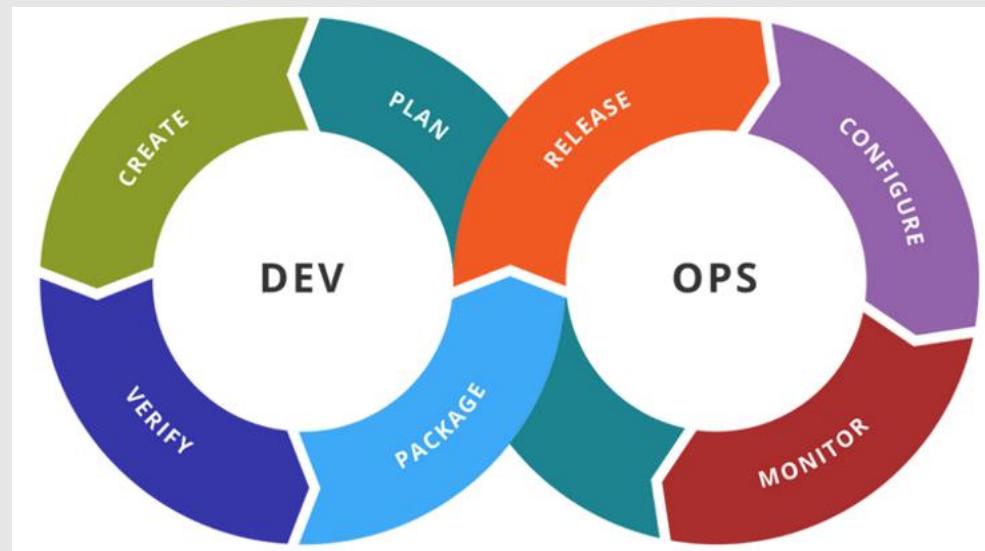


Developer Client

# Create an Elastic Beanstalk Application



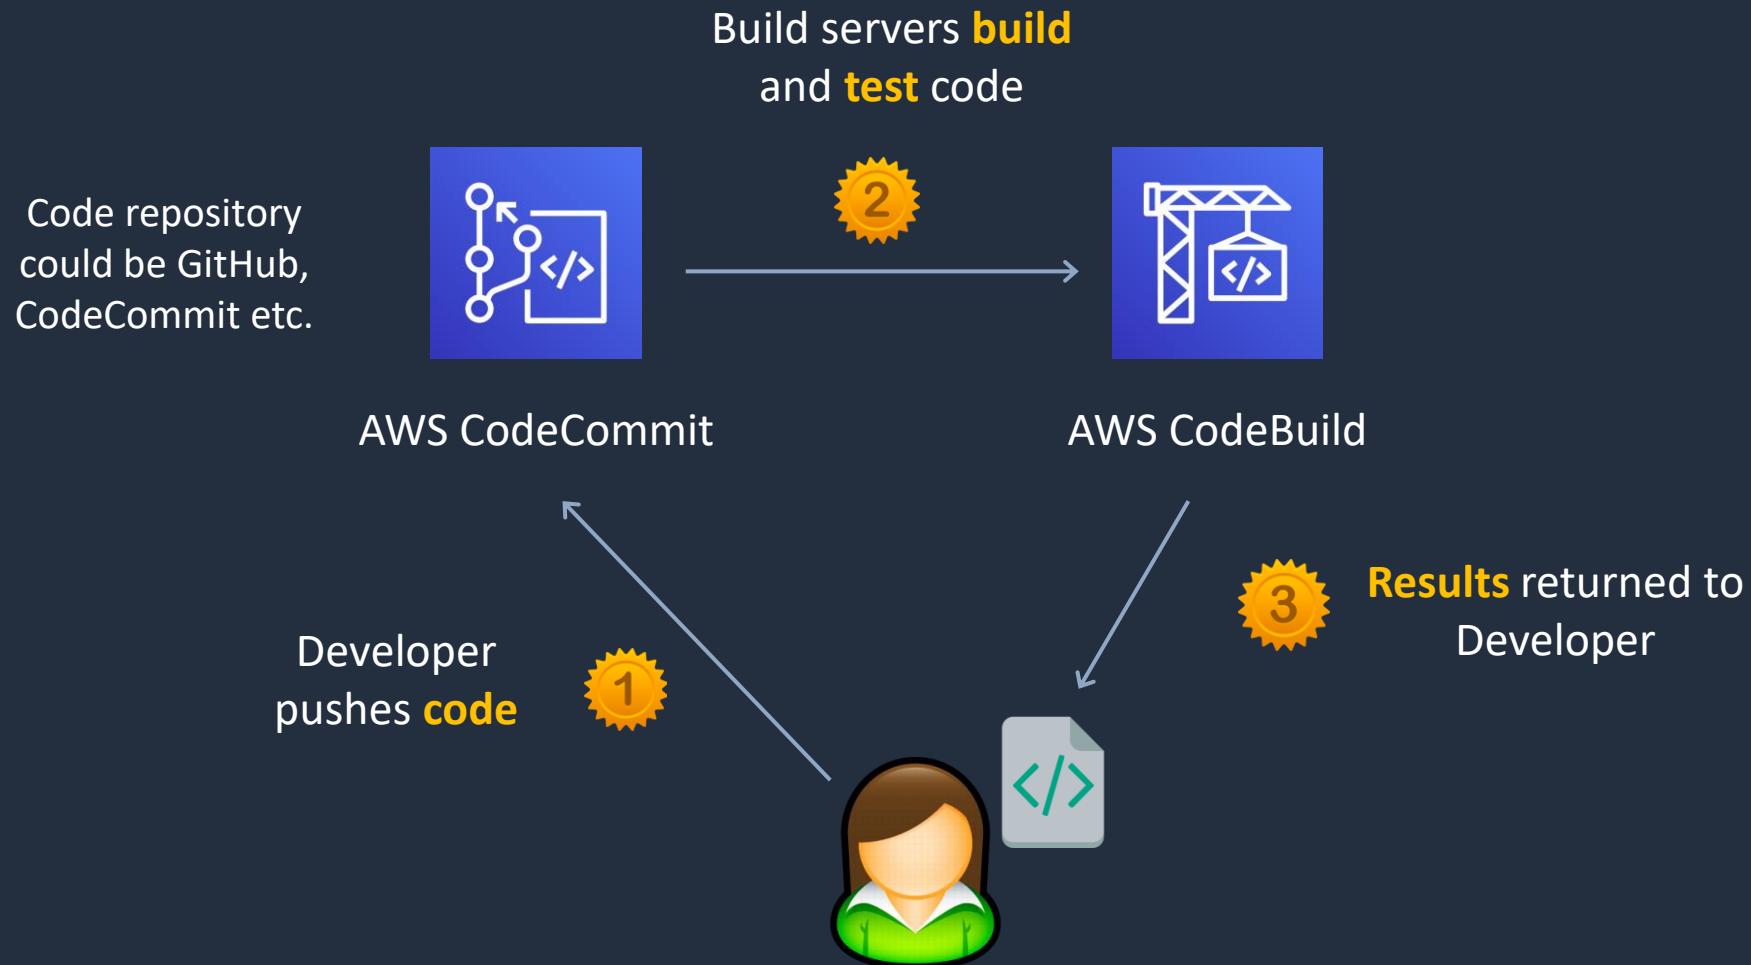
# Continuous Integration and Continuous Delivery (CI/CD)





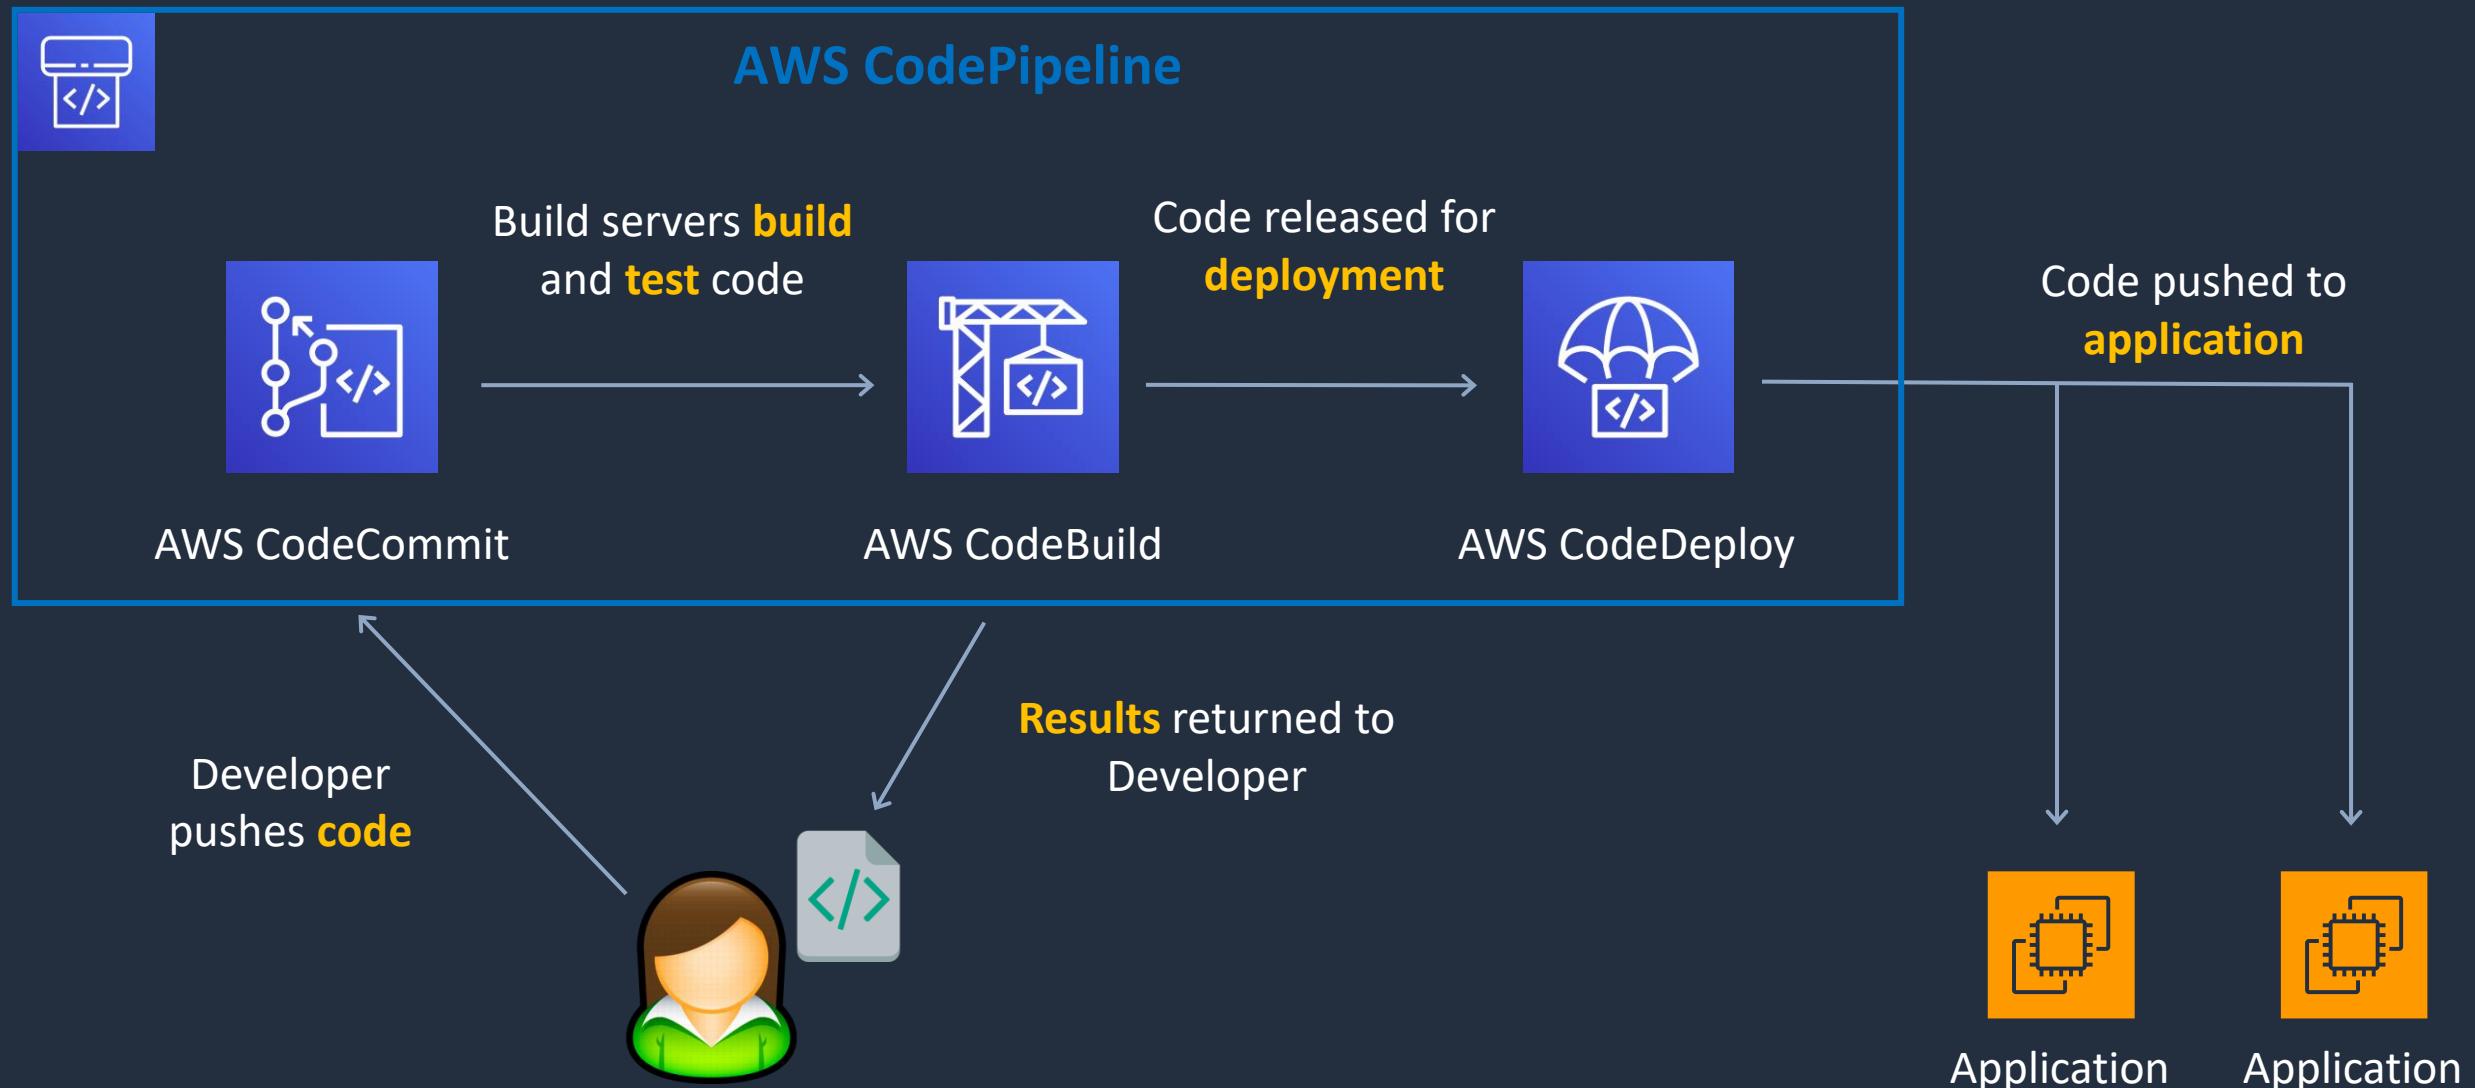
# Continuous Integration

---





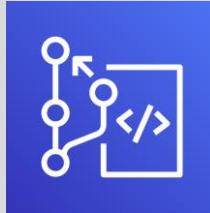
# Continuous Integration & Continuous Delivery





# Continuous Integration and Continuous Delivery

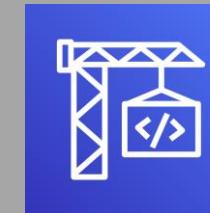
CODE



AWS CodeCommit

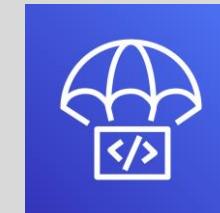
BUILD & TEST

AWS CodePipeline



AWS CodeBuild

DEPLOY



AWS CodeDeploy



# Create an AWS CodeCommit Repository

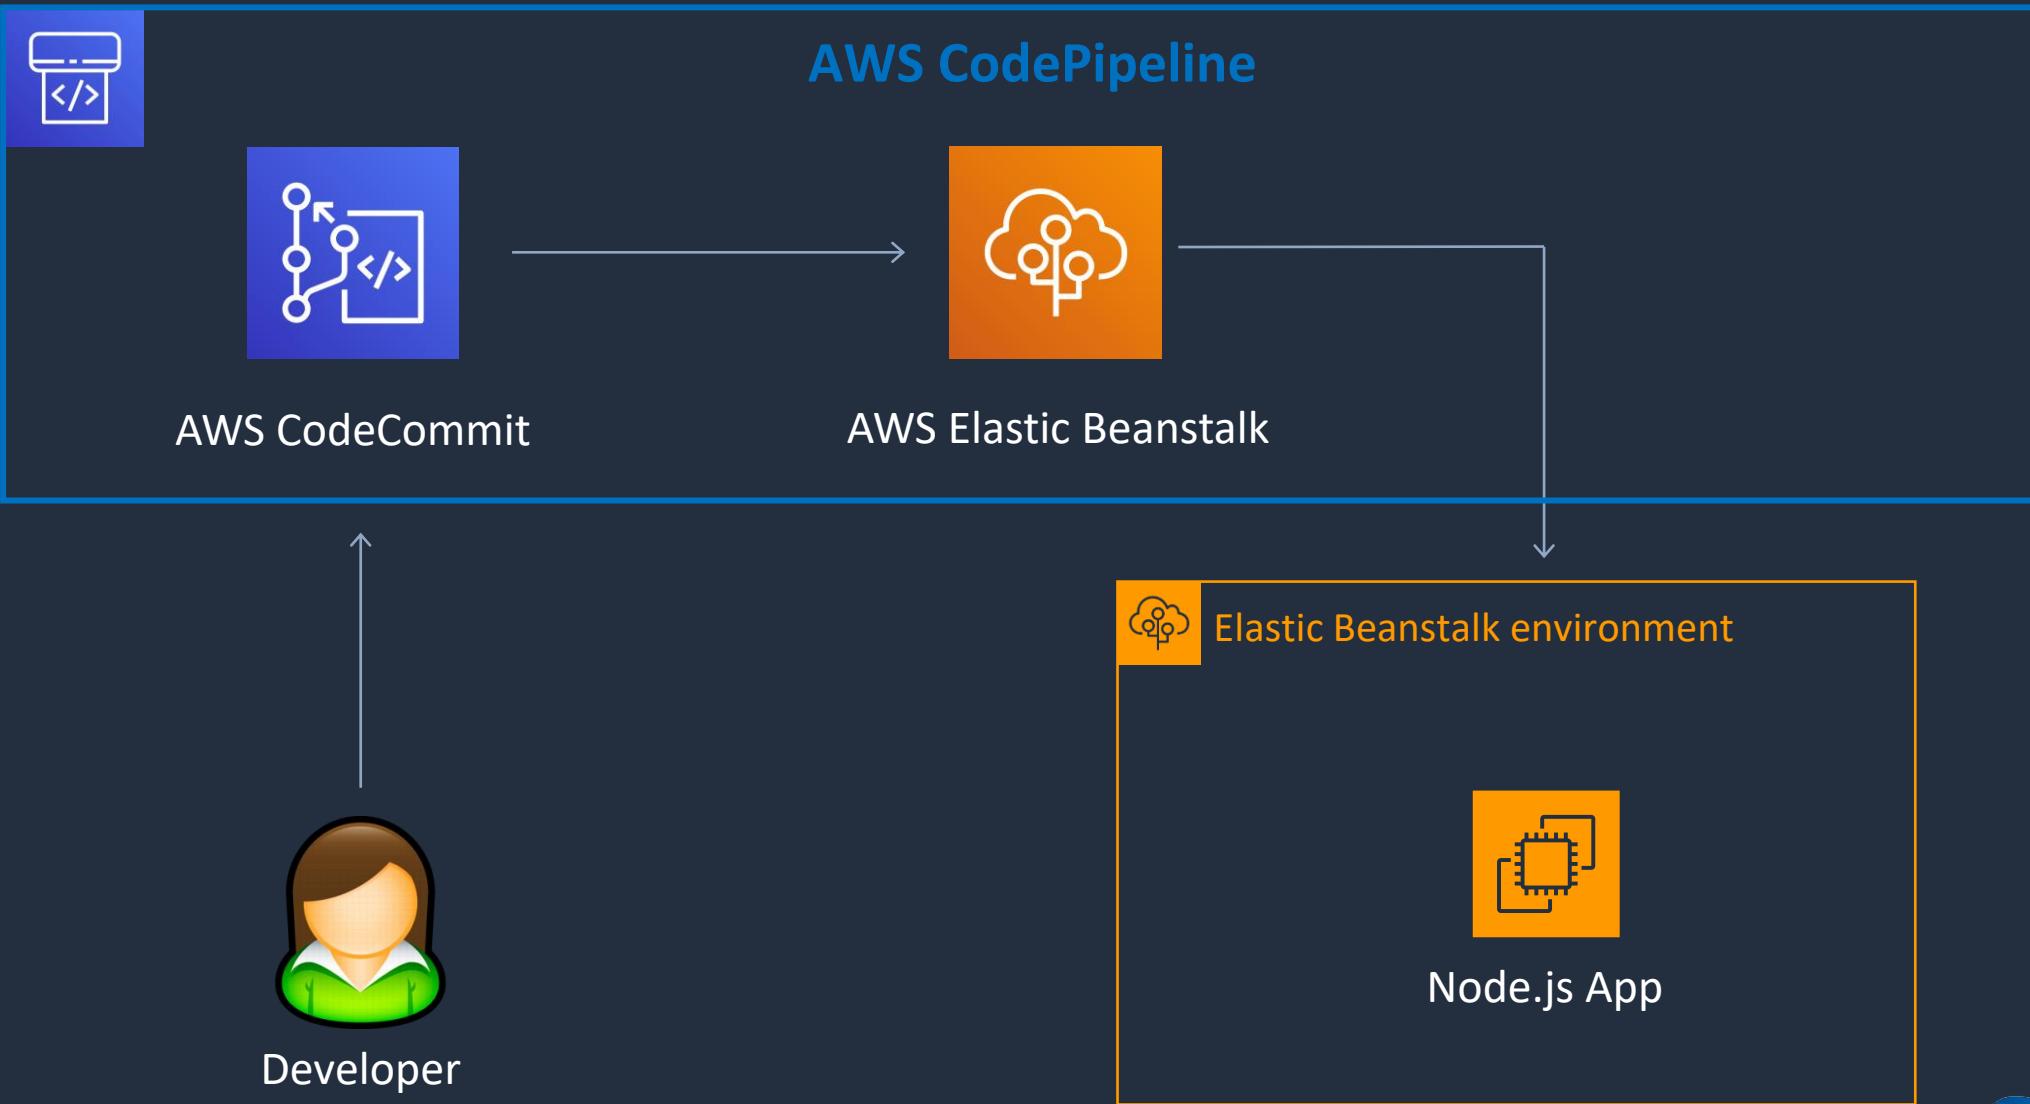


# AWS CodePipeline with AWS Elastic Beanstalk





# AWS CodePipeline with Elastic Beanstalk



# SECTION 10

## DNS, Caching, and Performance Optimization

# Amazon Route 53 DNS

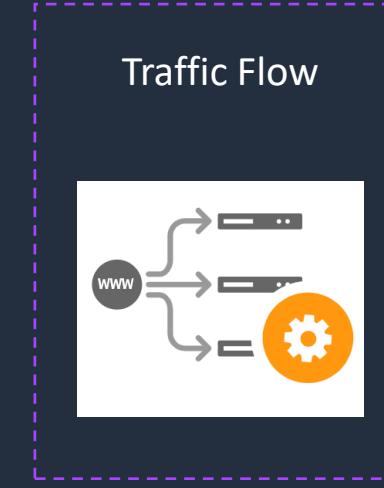
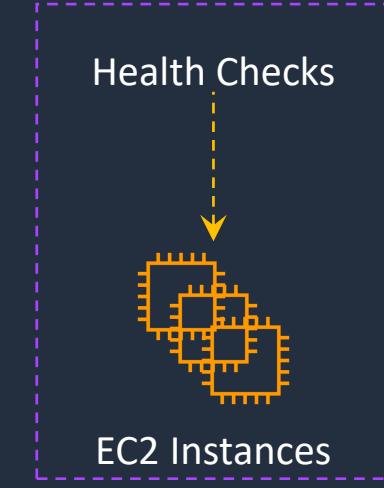




# Amazon Route 53

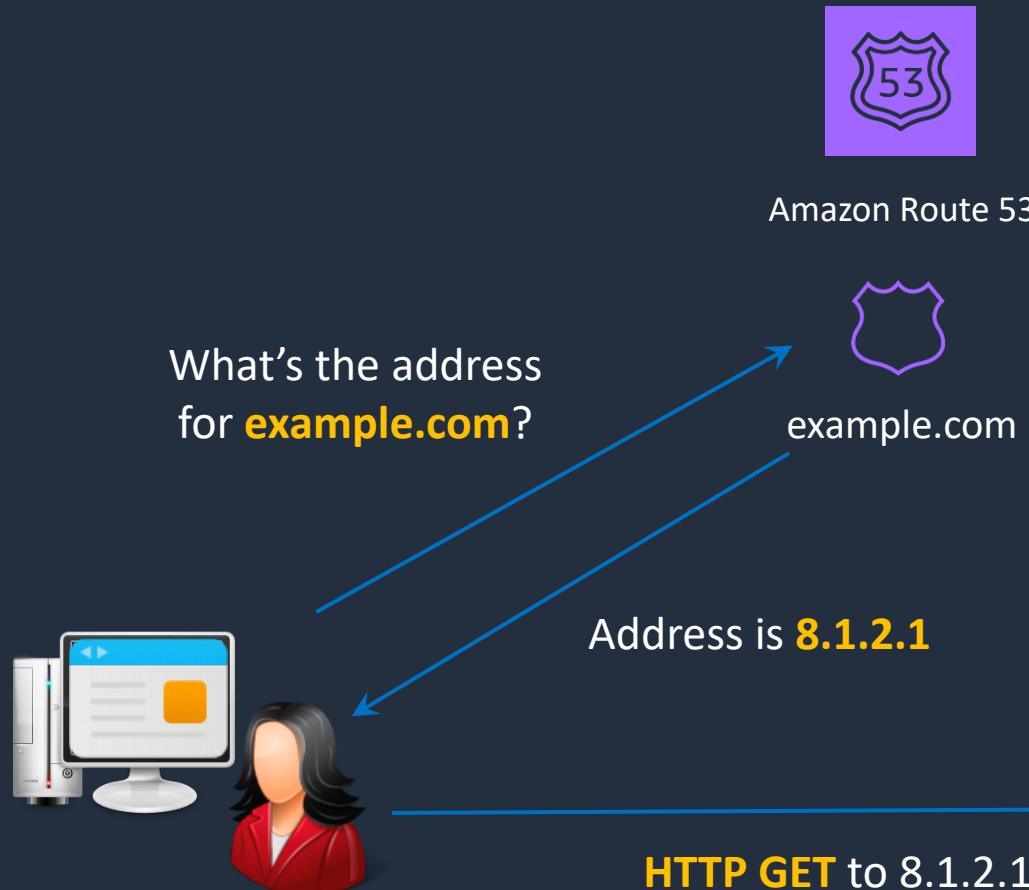


## Amazon Route 53

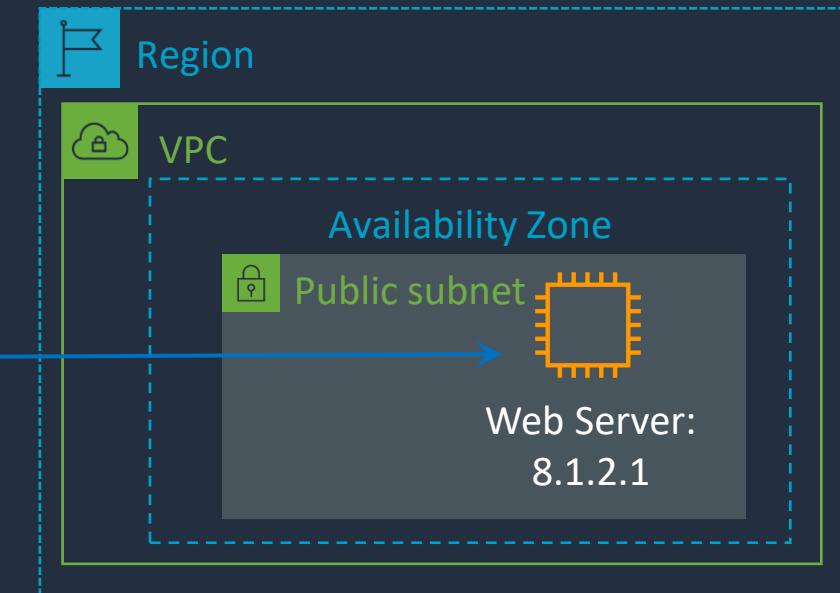




# DNS Resolution with Route 53



A **hosted zone** represents a set of records belonging to a domain



# Route 53 DNS Record Types

Routing Policy	What it does
<b>Simple</b>	Simple DNS response providing the IP address associated with a name
<b>Failover</b>	If primary is down (based on health checks), routes to secondary destination
<b>Geolocation</b>	Uses geographic location you're in (e.g. Europe) to route you to the closest region
<b>Geoproximity</b>	Routes you to the closest region within a geographic area
<b>Latency</b>	Directs you based on the lowest latency route to resources
<b>Multivalue answer</b>	Returns several IP addresses and functions as a basic load balancer
<b>Weighted</b>	Uses the relative weights assigned to resources to determine which to route to

# Register a Domain using Route 53

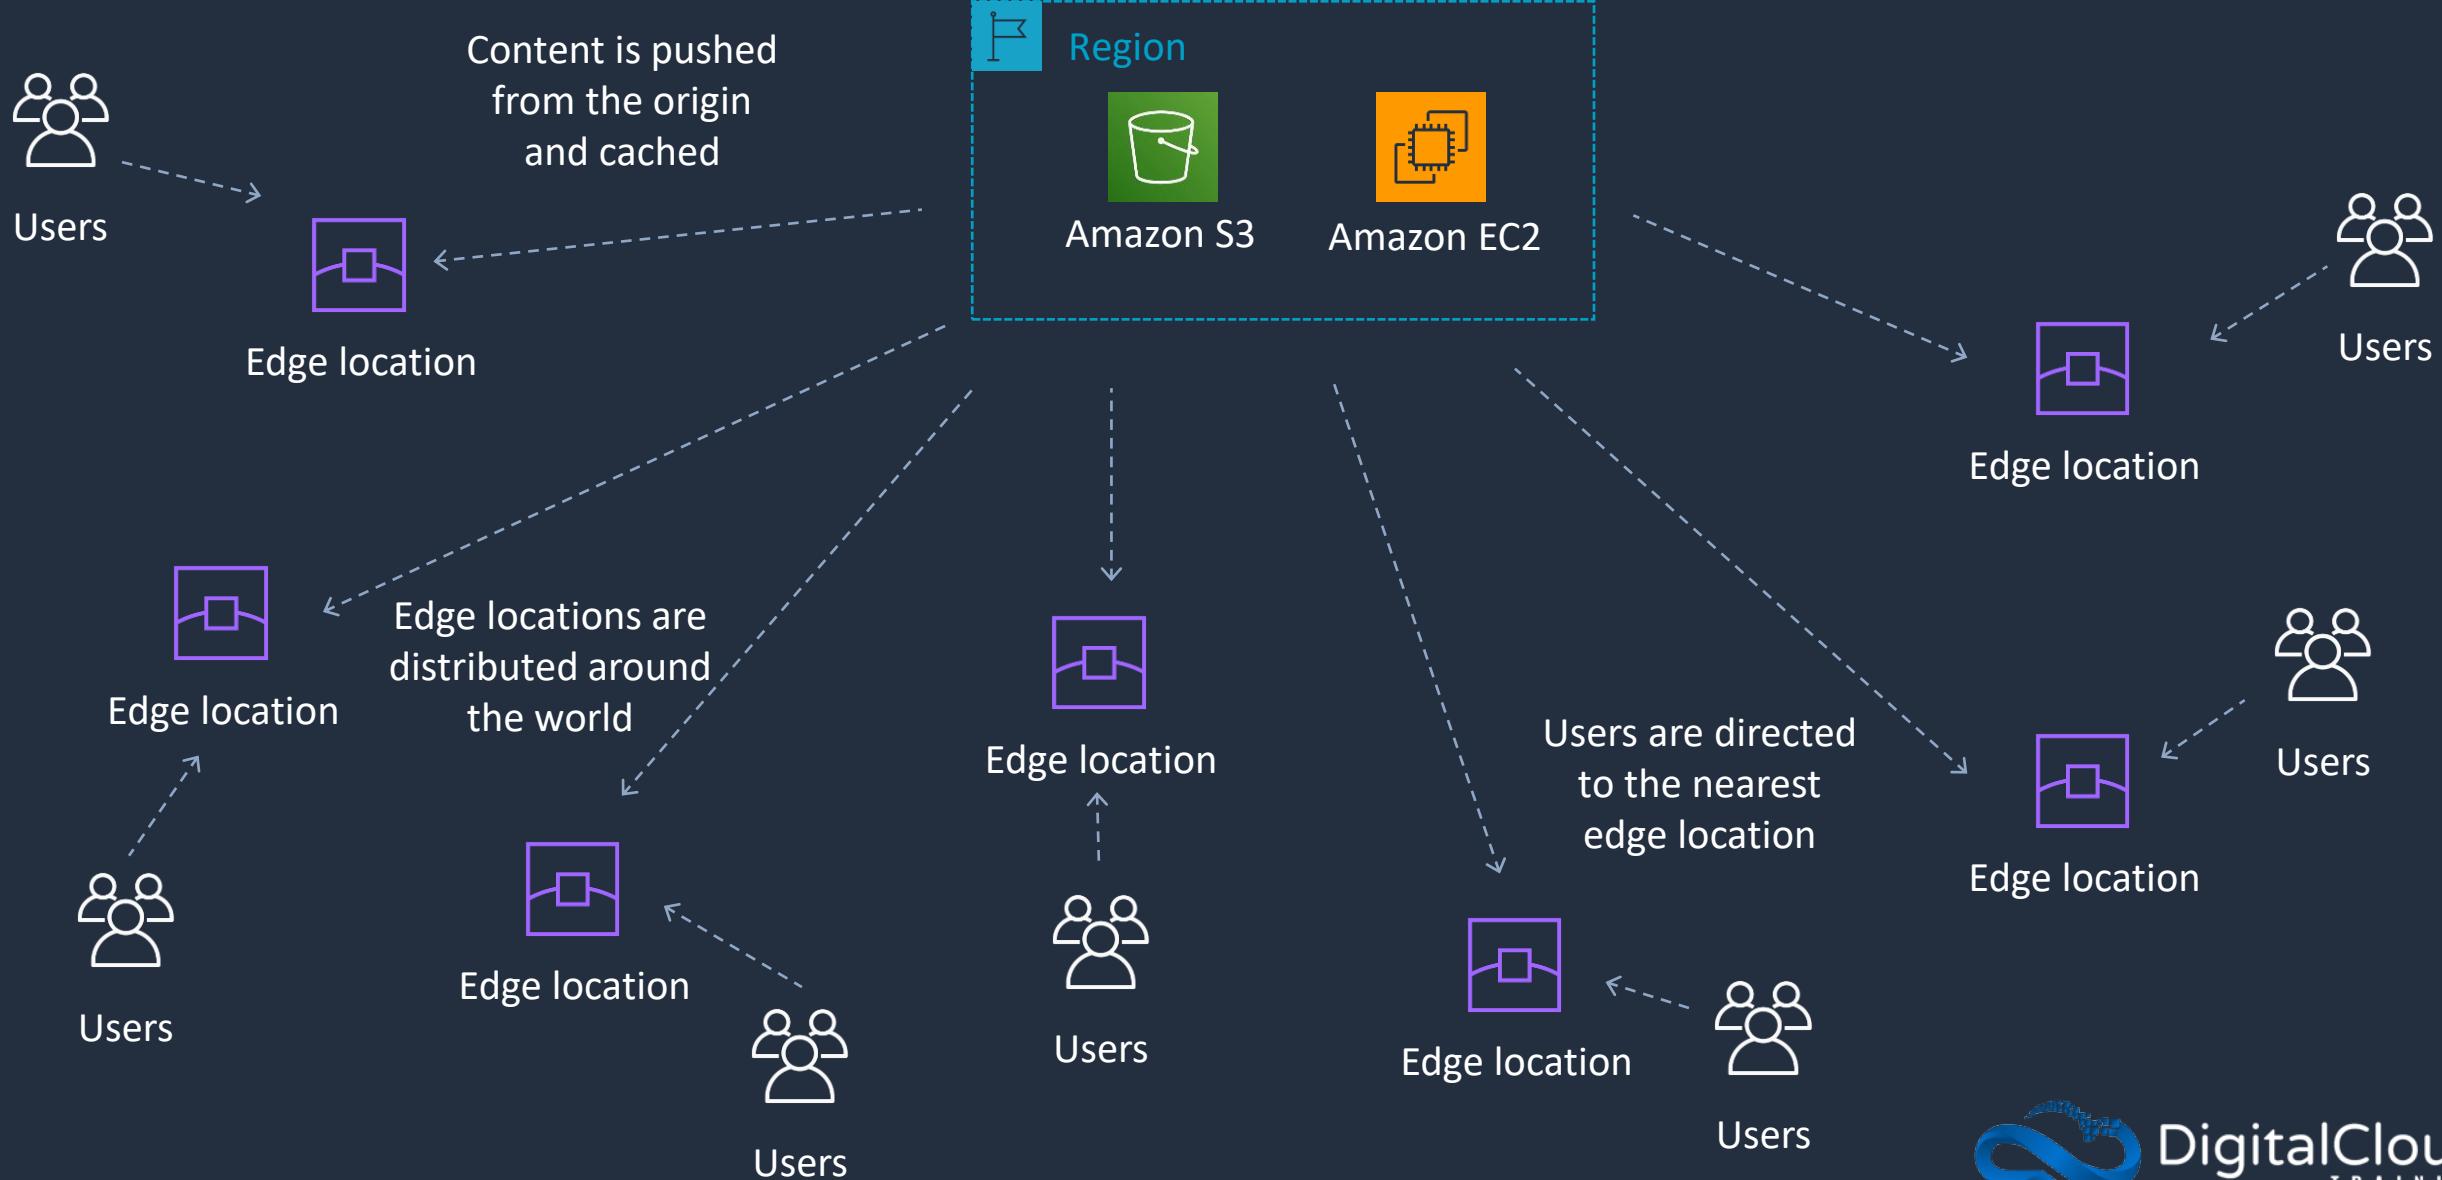


# Amazon CloudFront





# Amazon CloudFront



# Create an Amazon CloudFront Distribution

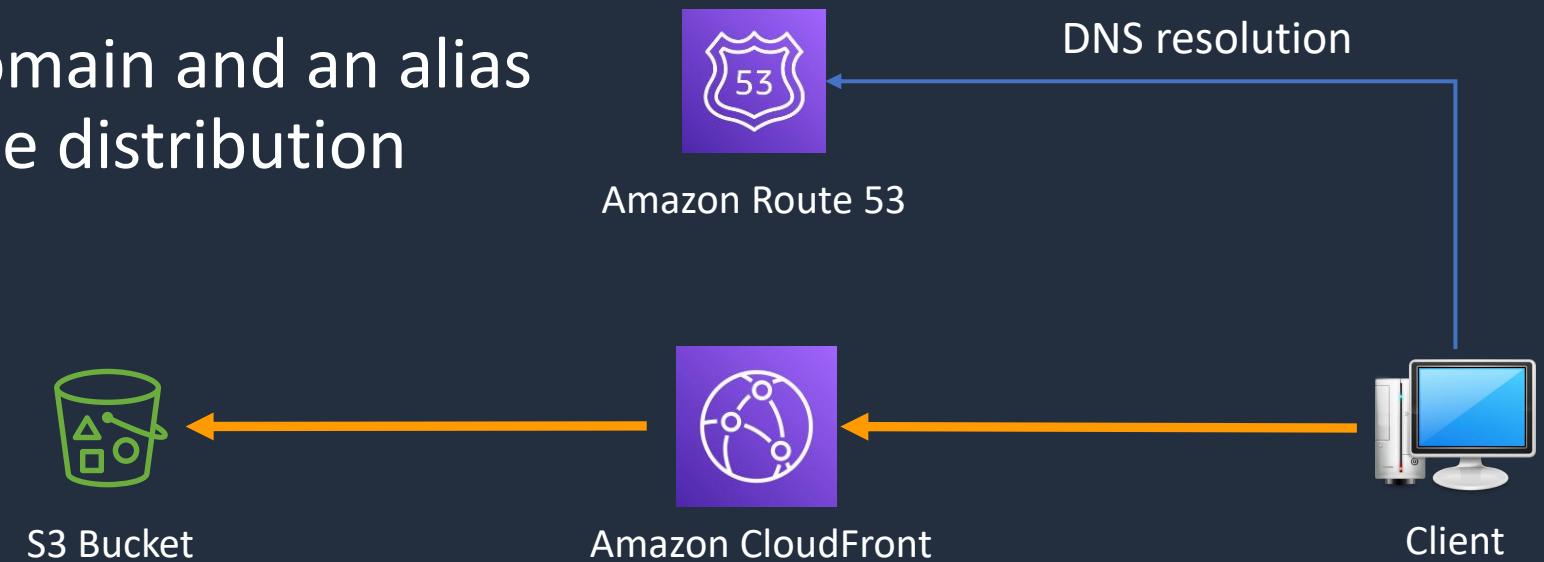




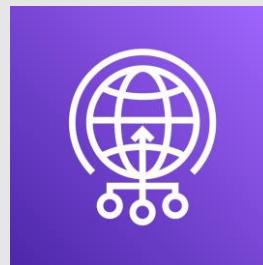
# Amazon CloudFront Distribution

We will create:

- Amazon S3 bucket
- AWS Certificate Manager SSL/TLS certificate
- Amazon CloudFront distribution
- Amazon Route 53 domain and an alias record pointing to the distribution



# AWS Global Accelerator



# AWS Global Accelerator



Users in US

Resolve dctlabs.com



Answer:  
51.45.2.12  
53.58.31.89

Amazon Route 53



Edge location



Global Accelerator

Addresses:  
51.45.2.12  
53.58.31.89

Static **anycast**  
IP addresses

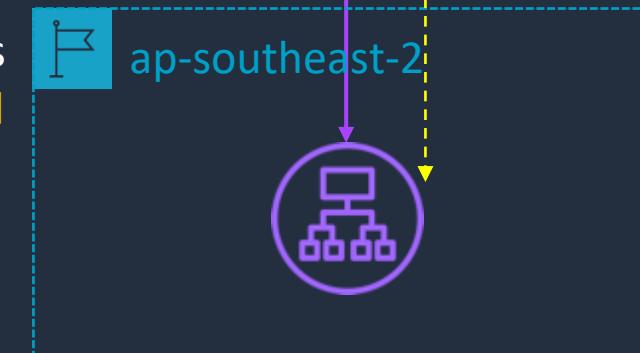
Requests are  
routed to the  
**optimal endpoint**

AWS Global Network

Users are **redirected**  
to another **endpoint**



Traffic traverses  
the **AWS global  
network**



DigitalCloud  
TRAINING

# SECTION 11

## Containers and Serverless Computing

# Docker Containers on Amazon ECS





# Amazon ECS

ECS **Services** are used to maintain a **desired count** of tasks

An ECS **Task** is created from a **Task Definition**

Task Definition

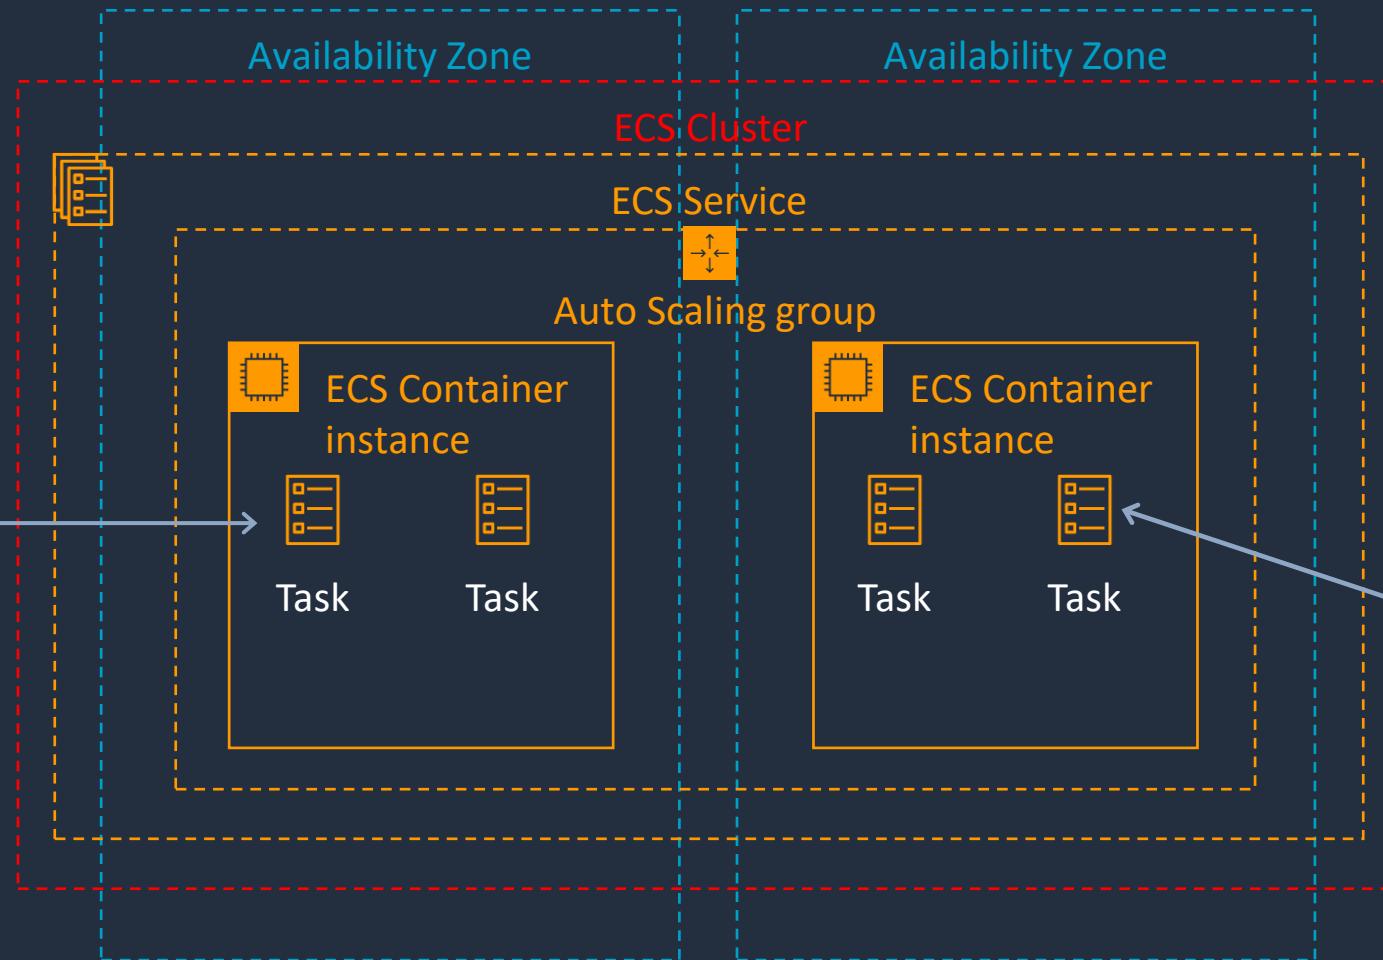
```
{
  "containerDefinitions": [
    {
      "name": "wordpress",
      "links": [
        "mysql"
      ],
      "image": "wordpress",
      "essential": true,
      "portMappings": [
        {
          "containerPort": 80,
          "hostPort": 80
        }
      ],
      "memory": 500,
      "cpu": 10
    }
  ]
}
```

An ECS **Task** is a running Docker container



Amazon Elastic Container Service

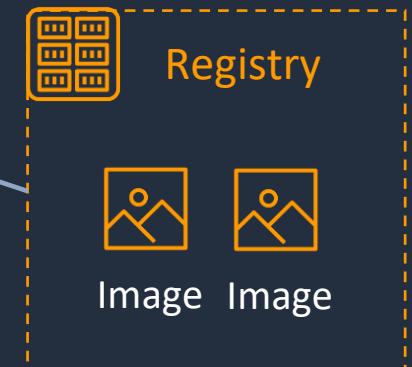
An Amazon **ECS Cluster** is a logical grouping of **tasks or services**



Docker **images** can be stored in **Amazon ECR**



Amazon Elastic Container Registry





# Amazon ECS Components

Elastic Container Service (ECS)		Description
Cluster		Logical grouping of tasks or services
Container instance		EC2 instance running the the ECS agent
Task Definition		Blueprint that describes how a docker container should launch
Task		A running container using settings in a Task Definition
Service		Defines long running tasks – can control task count with Auto Scaling and attach an ELB

# Create an AWS Fargate Cluster

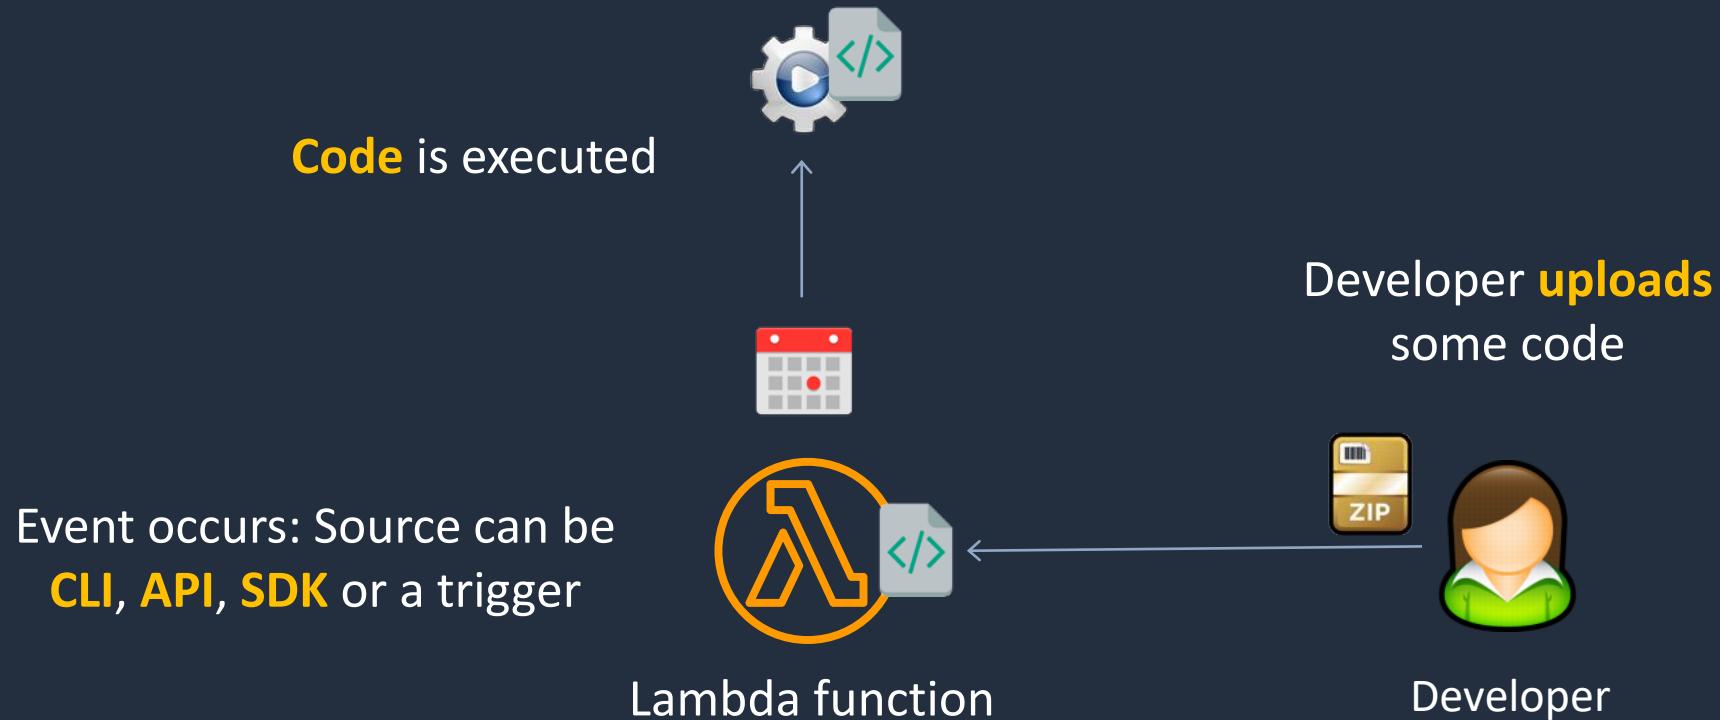


# Serverless with AWS Lambda





# AWS Lambda



# Create an AWS Lambda Function

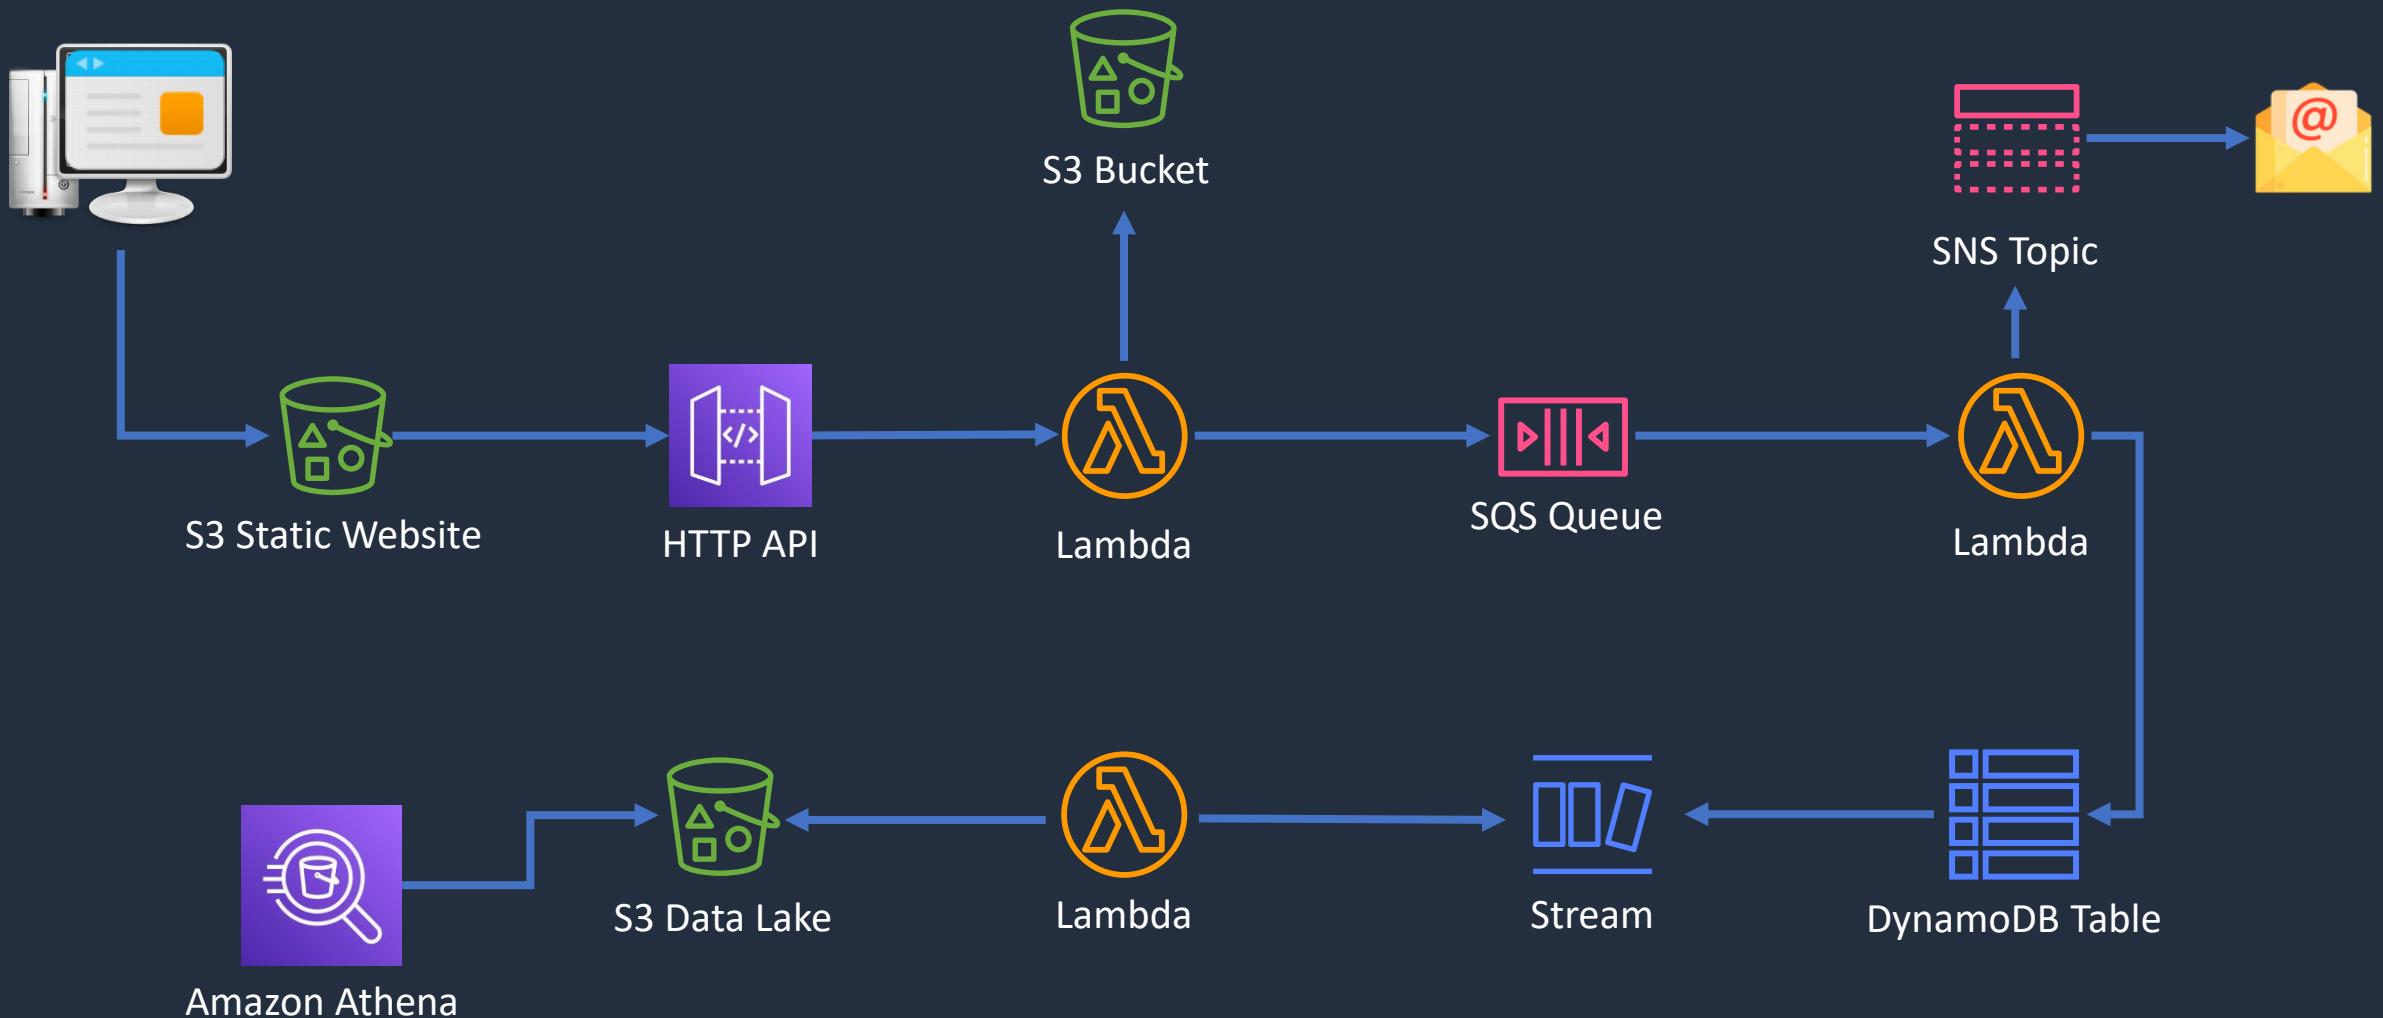


# Application Integration Services



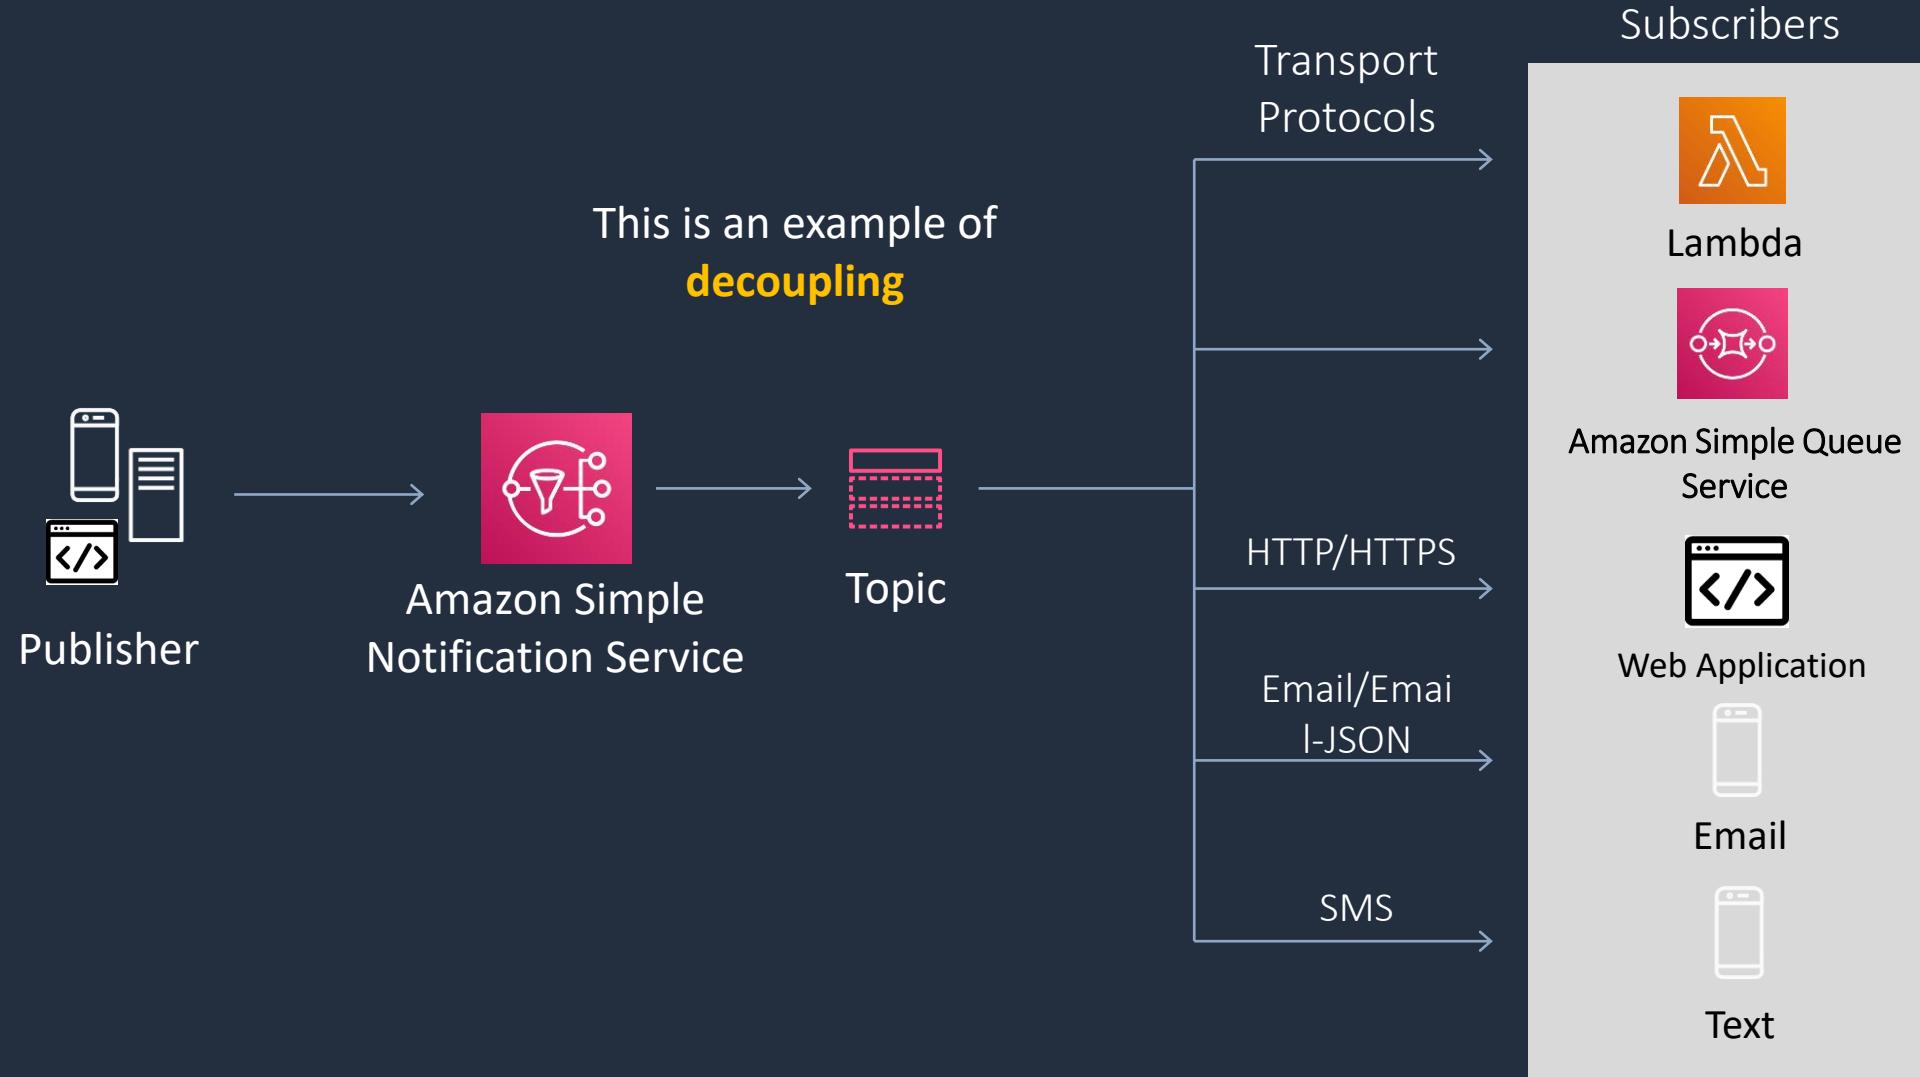


# Serverless Services and Event-Driven Architecture



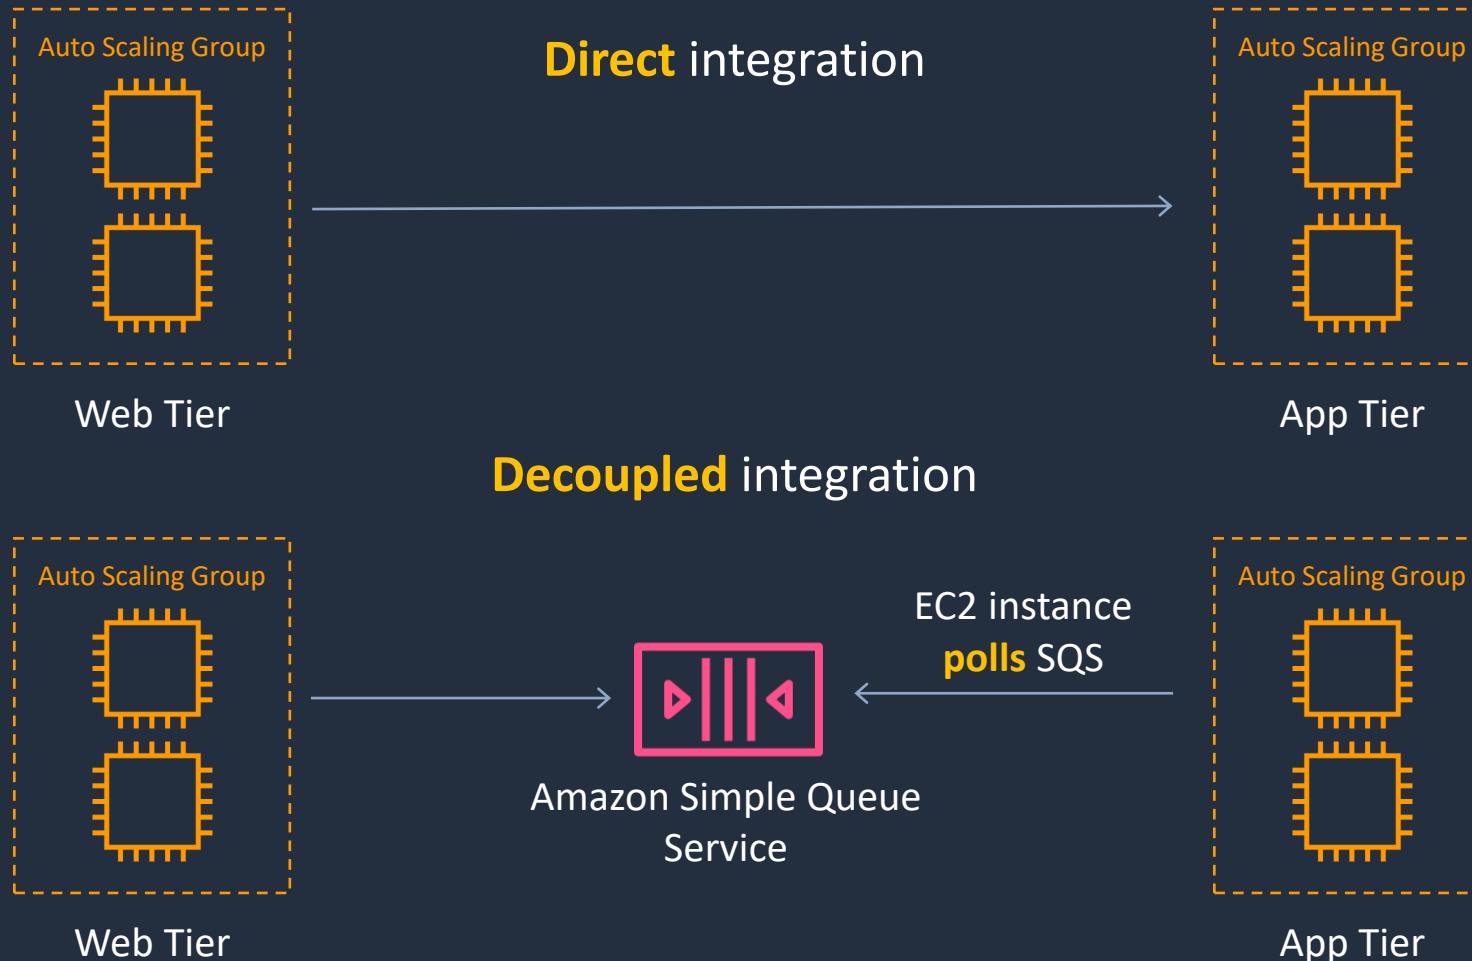


# Amazon Simple Notification Service (SNS)





# Amazon Simple Queue Service (SQS)

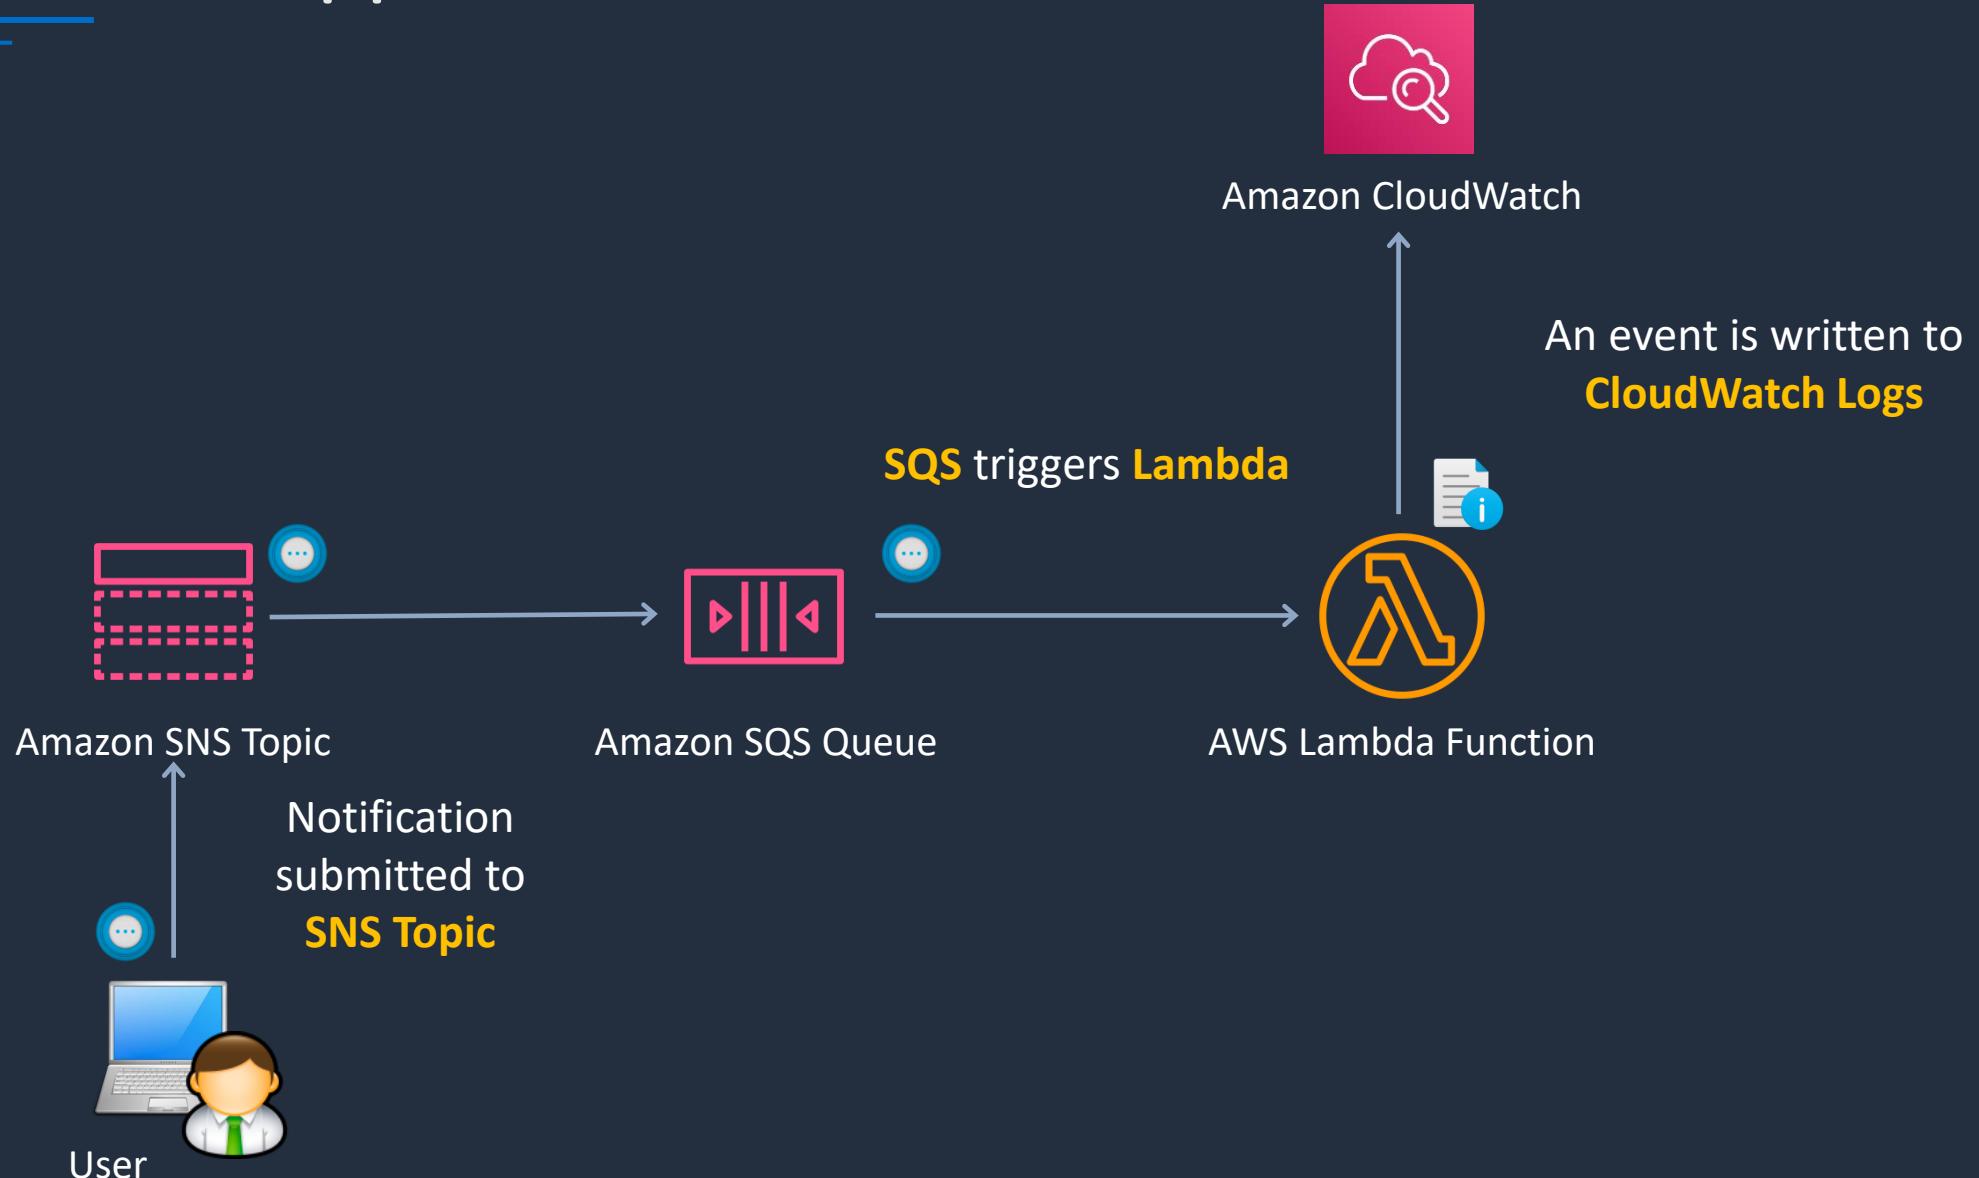


# Create a Serverless Application





# Serverless Application



# Amazon EventBridge

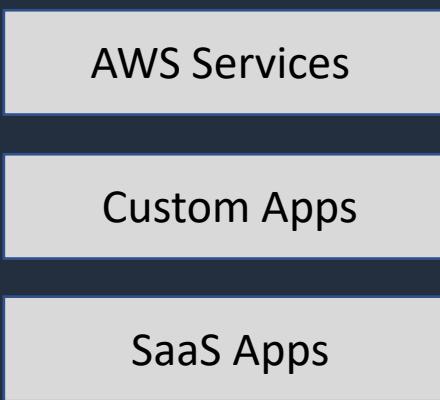




# Amazon EventBridge

EventBridge used to be known as **CloudWatch Events**

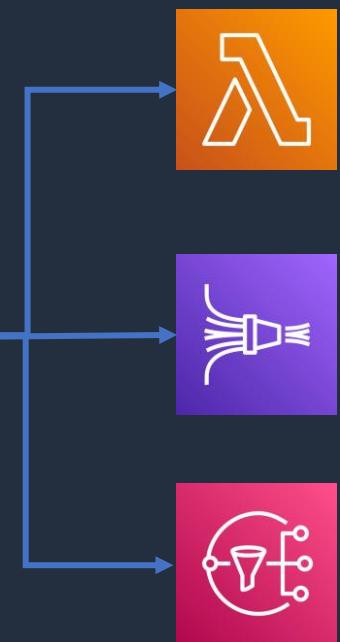
## Event Sources



## Events

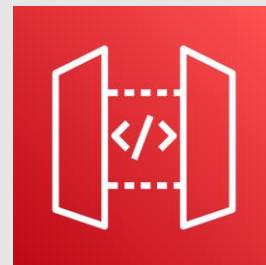
EventBridge  
event bus

## Rules



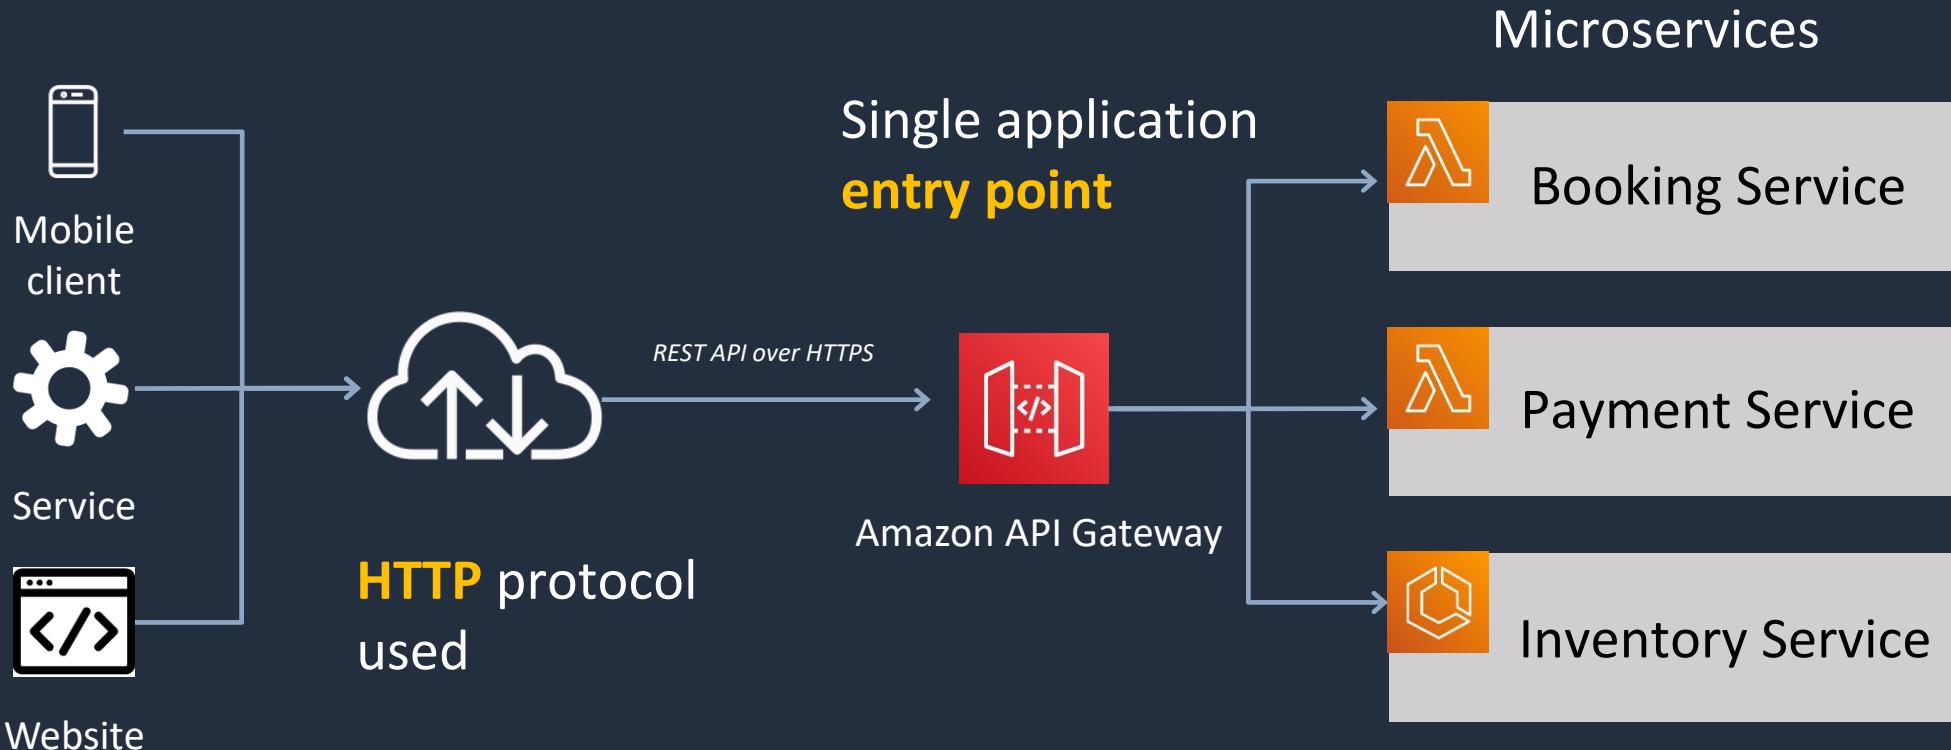
## Targets

# Amazon API Gateway



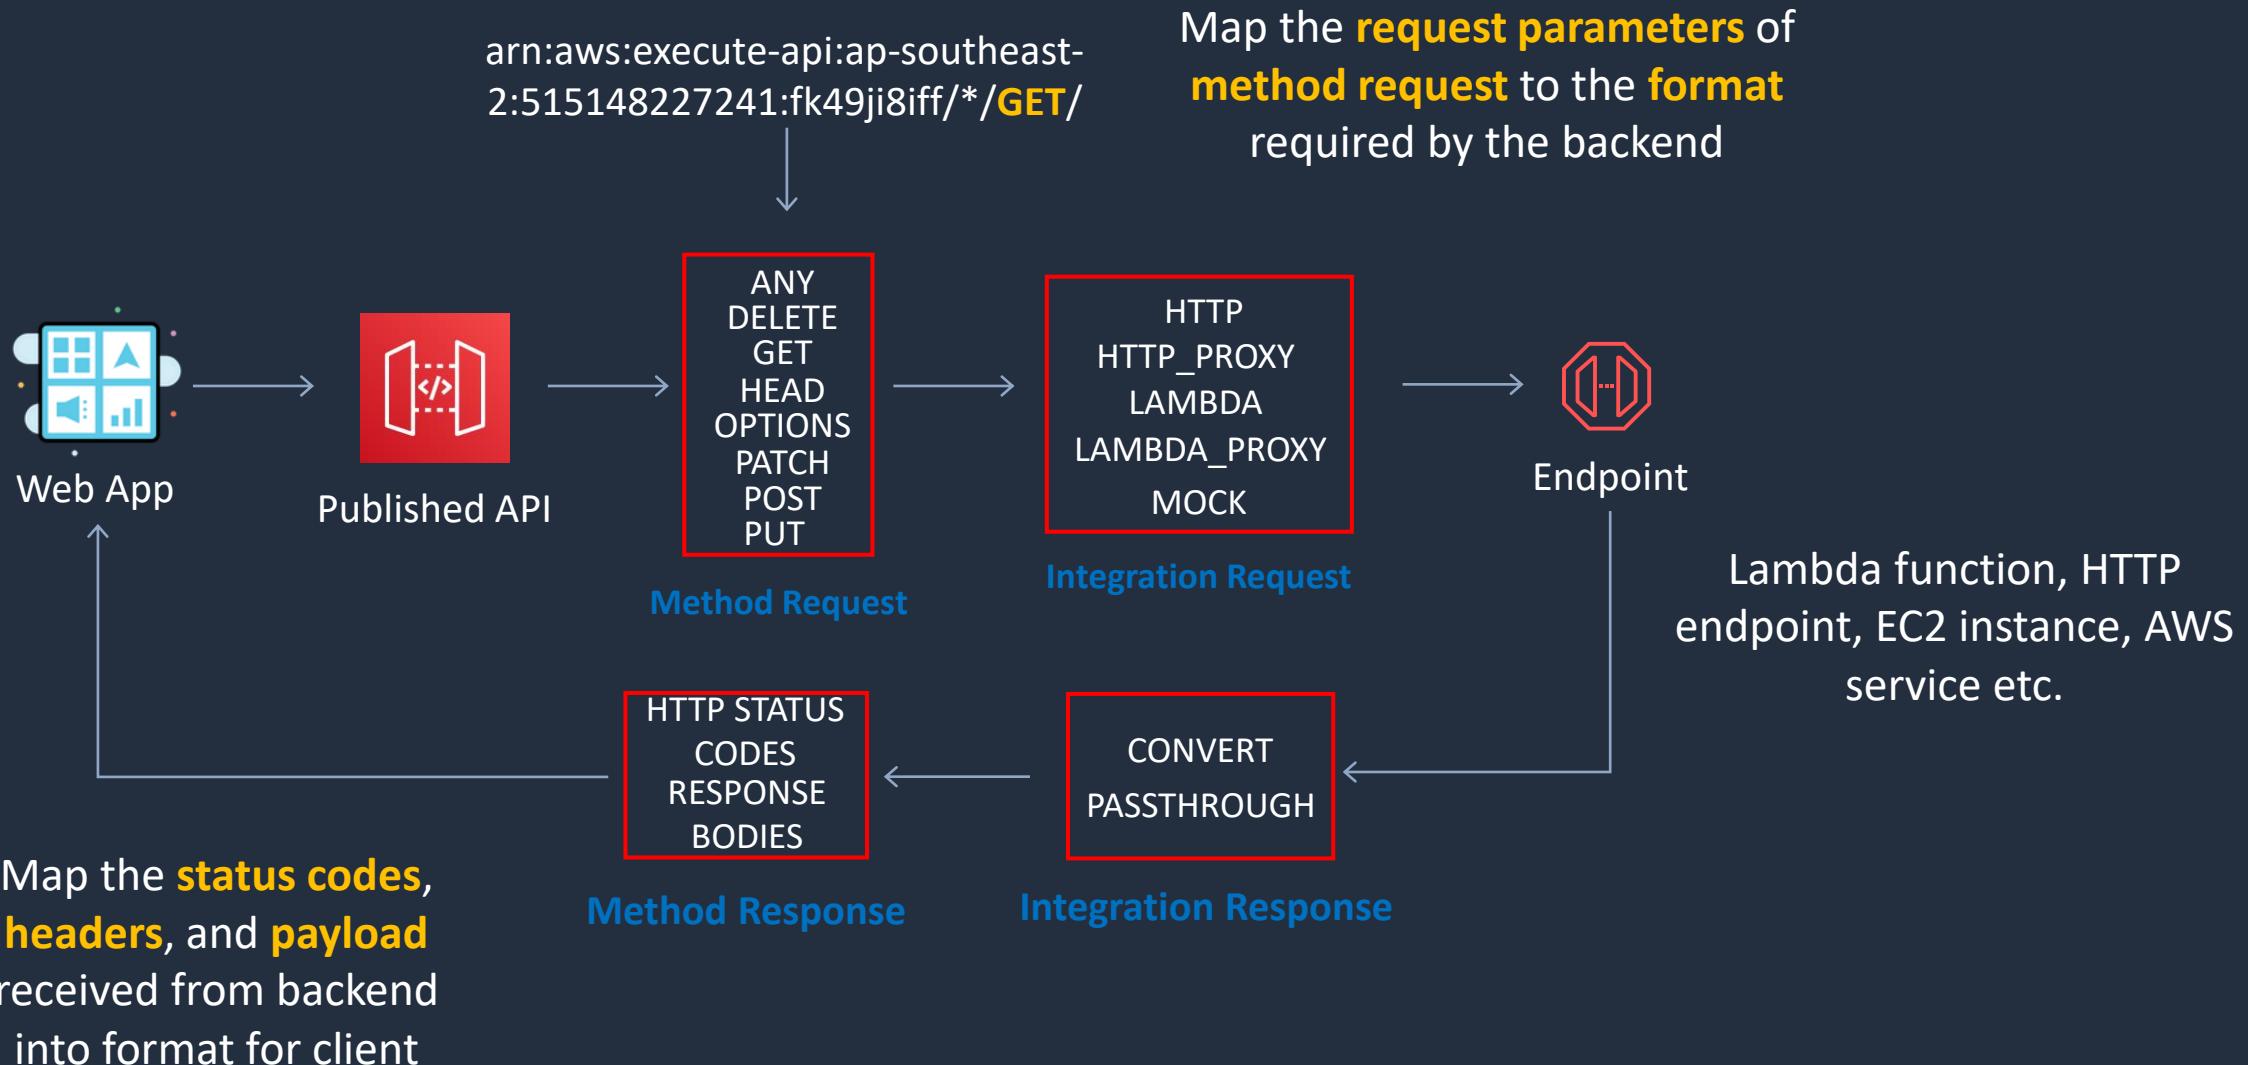


# REST API with Amazon API Gateway





# Structure of a REST API



# Create a REST API with Lambda Proxy Integration



# SECTION 13

Get Certified on AWS

# Get Certified on AWS



## Why work in cloud computing?

1. Job demand
2. Globally relevant skills
3. Rewarding career paths
4. Great salaries

## Why get AWS certified?

1. Demonstrate skills to employers
2. Differentiate yourself
3. Gain knowledge
4. Develop practical skills



# AWS Certification

New to the industry

## FOUNDATIONAL

**Six months** of fundamental AWS Cloud and industry knowledge



## ASSOCIATE

**One year** of experience solving problems and implementing solutions using the AWS Cloud



Some experience

## PROFESSIONAL

**Two years** of experience designing, operating, and troubleshooting solutions using the AWS Cloud



Highly experienced

## SPECIALTY

Technical AWS Cloud experience in the Specialty domain as specified in the exam guide



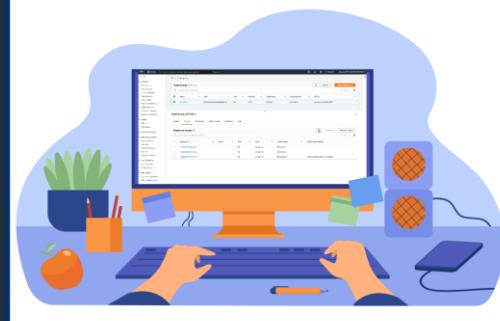
Specialized in an IT role

## ON-DEMAND TRAINING COURSES



Prepare for your next AWS certification with video courses & practice exams

## HANDS-ON CHALLENGE LABS



Build hands-on cloud skills in a secure sandbox environment

## LIVE AWS BOOTCAMPS



Get ready for your next cloud job with real-world projects in a virtual classroom